

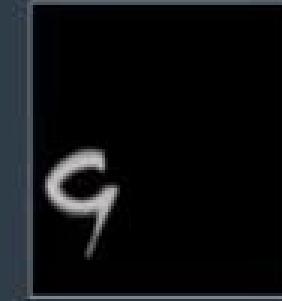
MNIST Database:

Centered, heavily
pre-processed images

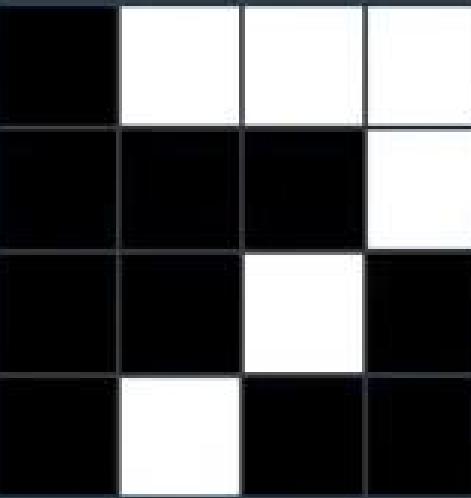
6	8	3	6	8	4	4	8	1	2
9	1	0	2	7	1	3	4	5	0
4	2	6	9	1	2	0	7	3	5
2	0	7	8	6	3	4	1	9	6
7	5	9	3	9	5	2	0	8	4



VS



Flattening



row 1

row 2

row 3

row 4

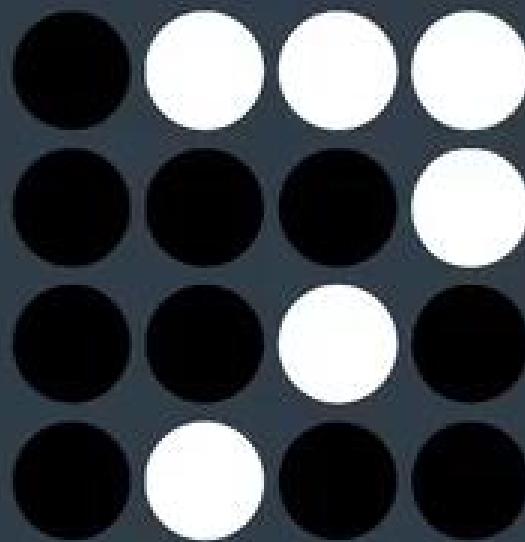


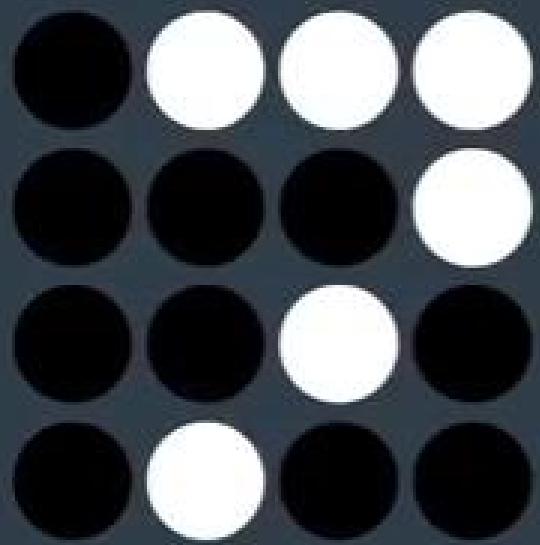
MLPs

- Only use **fully** connected layers
- Only accept **vectors** as input

CNNs

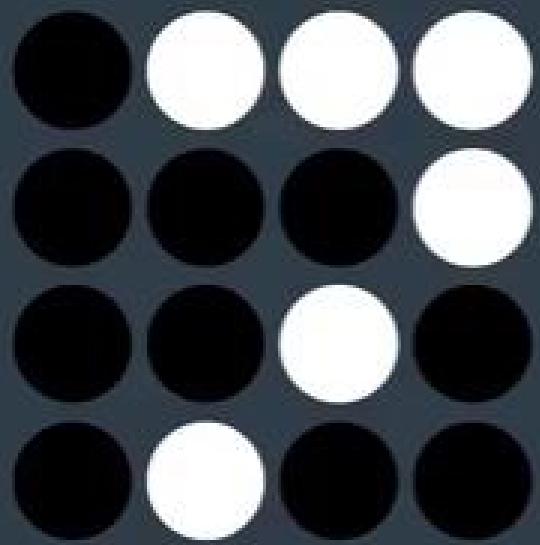
- Also use **sparsely** connected layers
- Also accept **matrices** as input



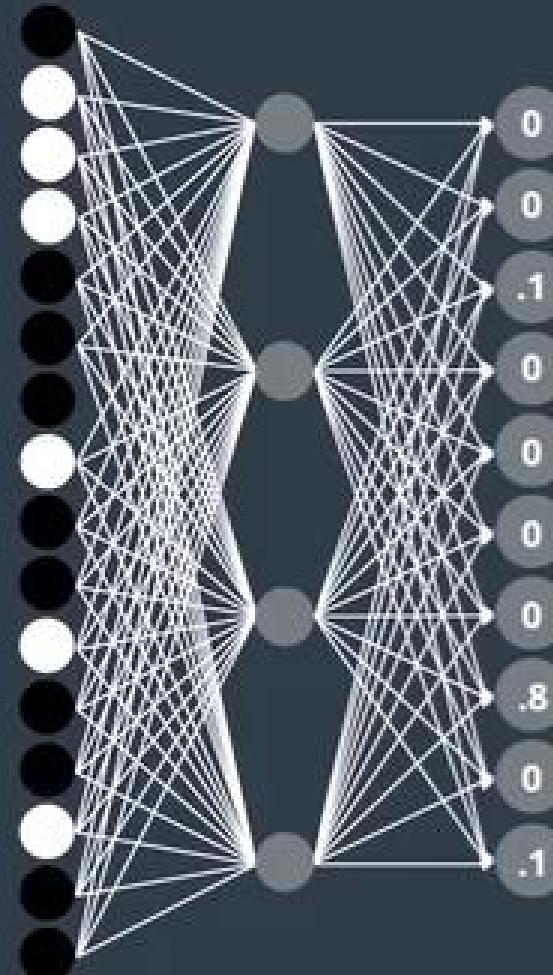


flatten
→





flatten
→





Input



Input

Neuron

→ 6

Output

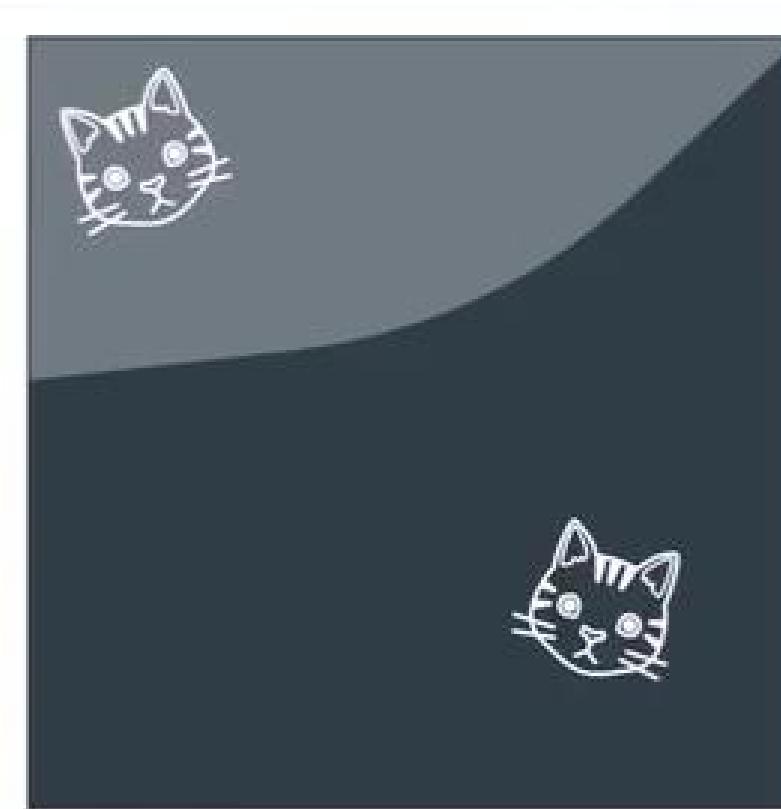




Cat



Cat



Cat

HIGH-PASS FILTERS

FILTERS



1. Filter out unwanted information
2. Amplify features of interest

HIGH-PASS FILTERS



HIGH-PASS FILTERS

- Sharpen an image
- Enhance ***high-frequency*** parts of an image



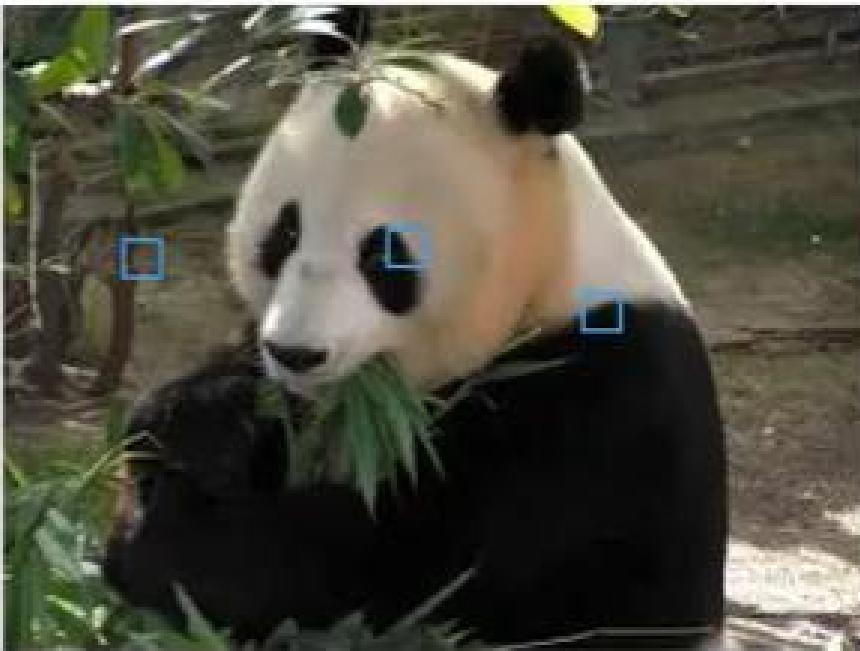
HIGH-PASS FILTERS

- Sharpen an image
- Enhance ***high-frequency*** parts of an image

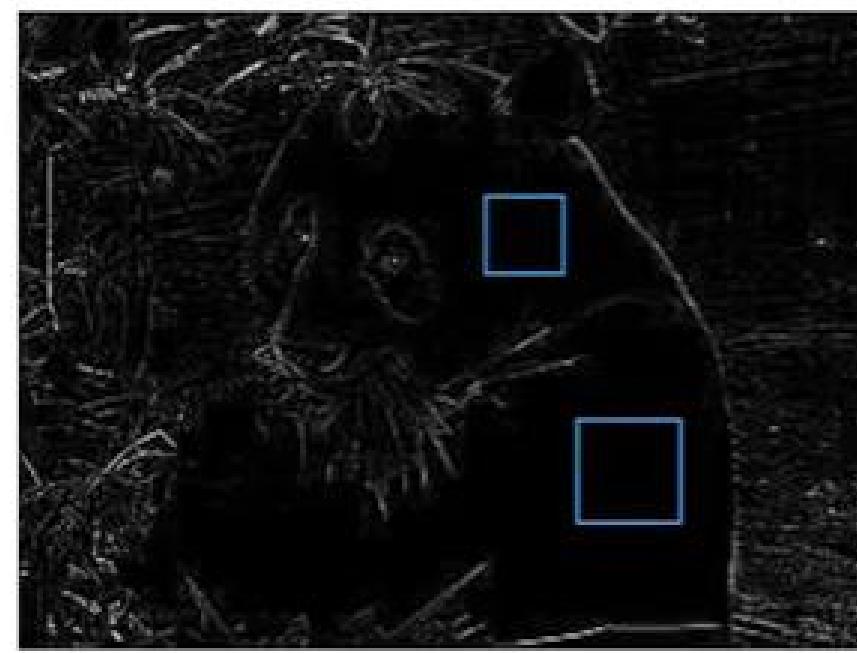
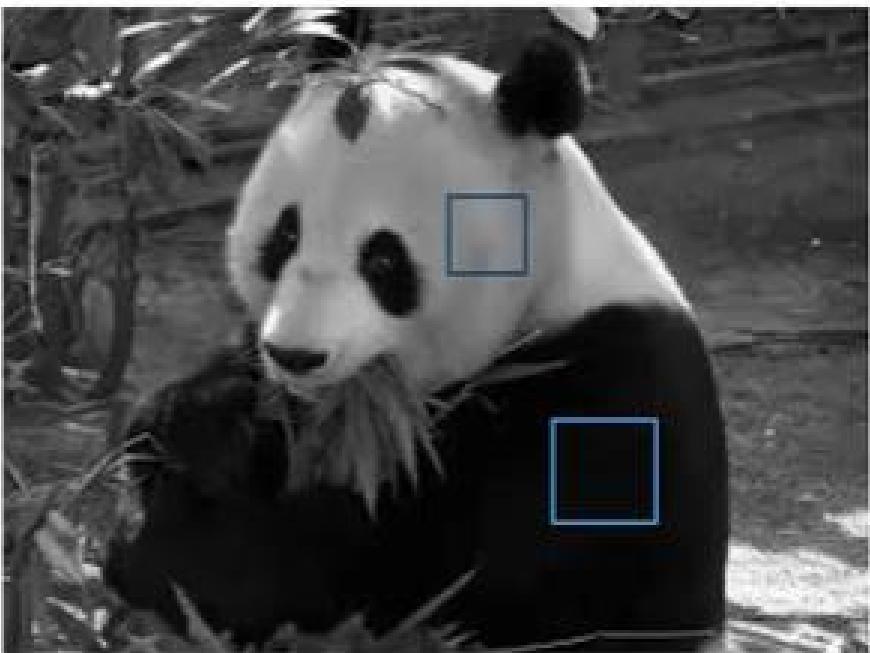


HIGH-PASS FILTERS

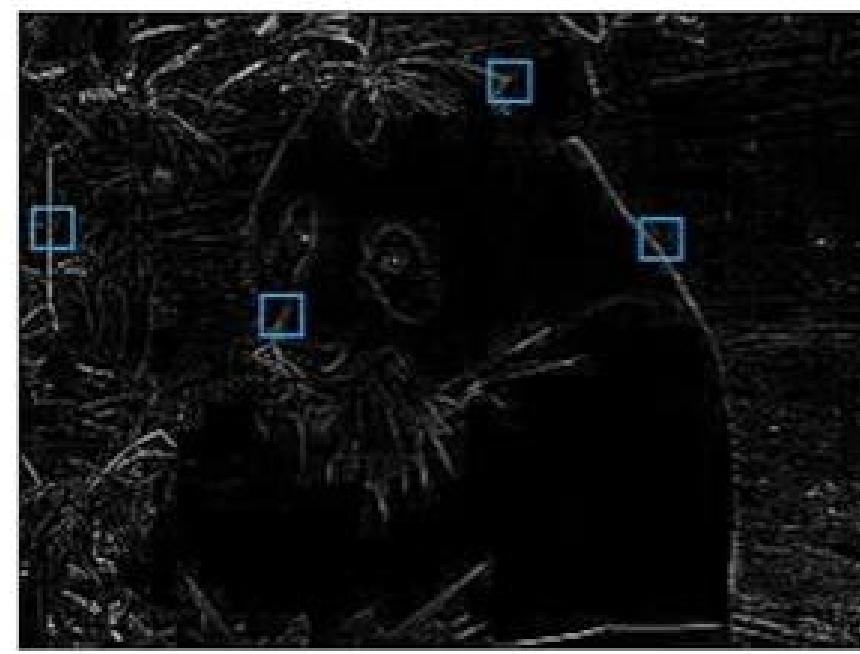
- Sharpen an image
- Enhance ***high-frequency*** parts of an image



HIGH-PASS FILTERS

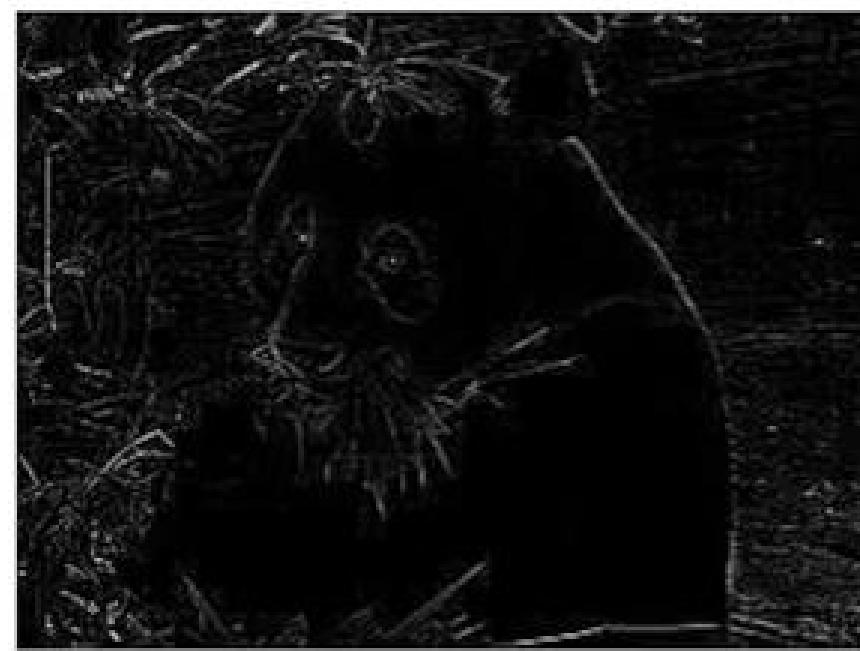


HIGH-PASS FILTERS



EDGE DETECTION

Emphasize Edges



Edges are areas in an image where the intensity changes very quickly,
and they often indicate object boundaries

CONVOLUTION KERNELS

A kernel is a matrix of numbers that modifies an image

0	-1	0
-1	4	-1
0	-1	0

edge detection filter

CONVOLUTION KERNELS

0	-1	0
-1	4	-1
0	-1	0

edge detection filter

$$0 + -1 + 0 + -1 + 4 + -1 + 0 + -1 + 0 = \mathbf{0}$$

CONVOLUTION KERNELS

0	-1	0
-1	4	-1
0	-1	0

edge detection filter

$$0 + -1 + 0 + -1 + 4 + -1 + 0 + -1 + 0 = \mathbf{0}$$

CONVOLUTION KERNELS

0	-1	0
-1	4	-1
0	-1	0

edge detection filter

$$0 + -1 + 0 + -1 + 4 + -1 + 0 + -1 + 0 = \mathbf{0}$$

Operation

Filter

Convolved
Image**Identity**

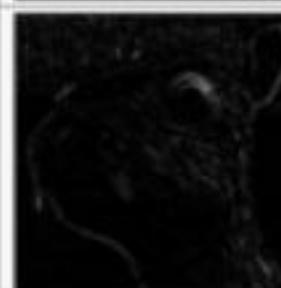
$$\begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix}$$



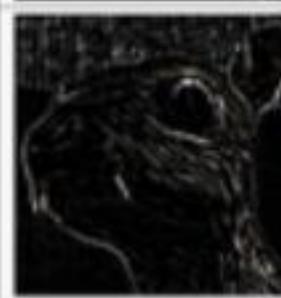
$$\begin{bmatrix} 1 & 0 & -1 \\ 0 & 0 & 0 \\ -1 & 0 & 1 \end{bmatrix}$$

**Edge detection**

$$\begin{bmatrix} 0 & 1 & 0 \\ 1 & -4 & 1 \\ 0 & 1 & 0 \end{bmatrix}$$



$$\begin{bmatrix} -1 & -1 & -1 \\ -1 & 8 & -1 \\ -1 & -1 & -1 \end{bmatrix}$$

**Sharpen**

$$\begin{bmatrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{bmatrix}$$

**Box blur**
(normalized)

$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

**Gaussian blur**
(approximation)

$$\frac{1}{16} \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix}$$



CONVOLUTION KERNELS

0	-1	0
-1	8	-1
0	-1	0



$$0 + -1 + 0 + -1 + \mathbf{8} + -1 + 0 + -1 + 0 = \boxed{4}$$

CONVOLUTION KERNELS

0	-1	0
-1	4	-2
0	-1	0



$$0 + -1 + 0 + -1 + 4 + \textcolor{yellow}{-2} + 0 + -1 + 0 = \textcolor{blue}{-1}$$

CONVOLUTION

$$\begin{bmatrix} 0 & -1 & 0 \\ -1 & 4 & -1 \\ 0 & -1 & 0 \end{bmatrix}$$

*



K

F(x,y)

$K * F(x,y) = \text{output image}$

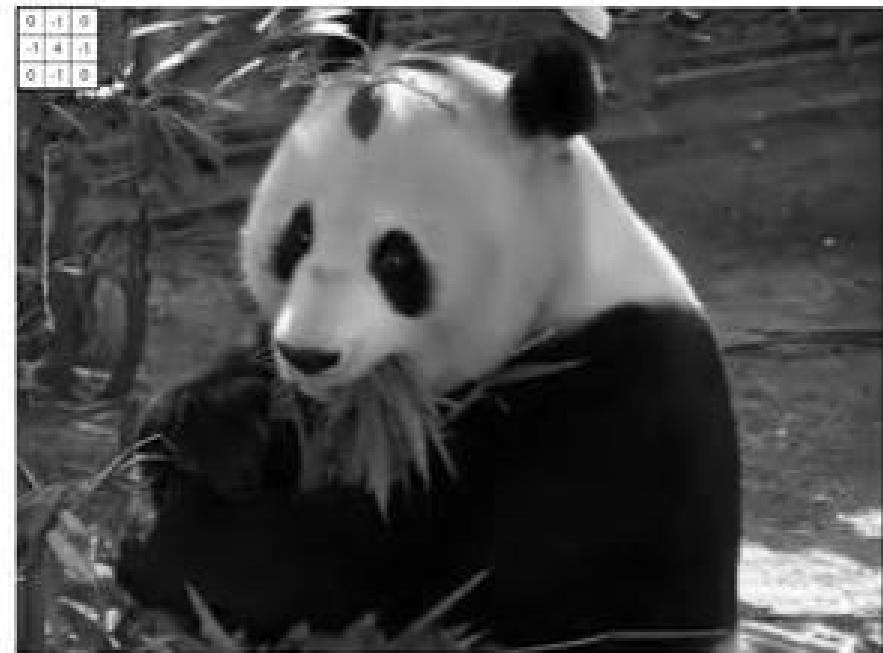
CONVOLUTION

0	-1	0
-1	4	-1
0	-1	0



$$K * F(x,y) = \text{output image}$$

CONVOLUTION



$$K * F(x,y) = \text{output image}$$

CONVOLUTION



$$K * F(x,y) = \text{output image}$$

CONVOLUTION



$$K * F(x,y) = \text{output image}$$

CONVOLUTION

0	-1	0
-1	4	-1
0	-1	0



CONVOLUTION

0	-1	0
-1	4	-1
0	-1	0



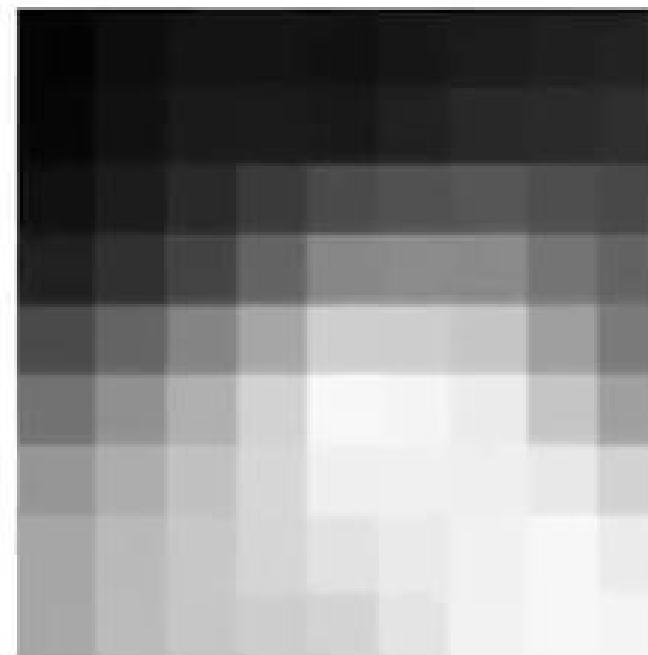
CONVOLUTION

0	-1	0
-1	4	-1
0	-1	0



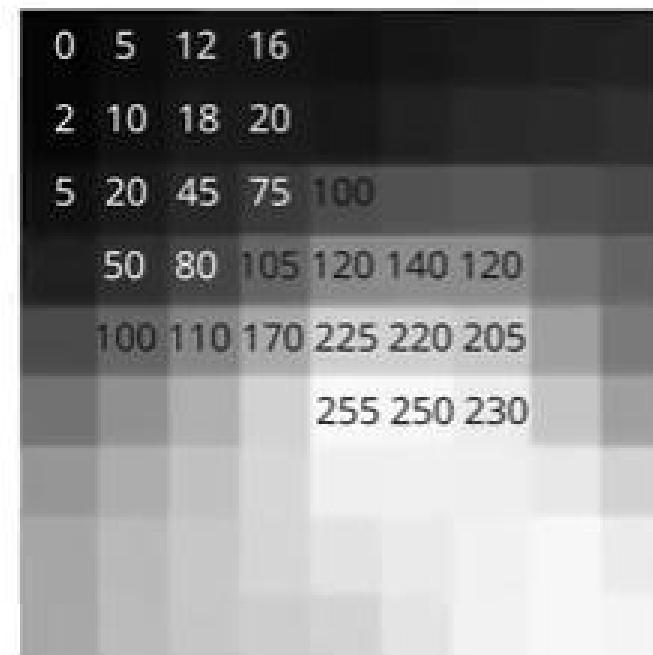
CONVOLUTION

0	-1	0
-1	4	-1
0	-1	0



CONVOLUTION

0	-1	0
-1	4	-1
0	-1	0



CONVOLUTION

0	-1	0
-1	4	-1
0	-1	0



CONVOLUTION

0	-1	0
-1	4	-1
0	-1	0



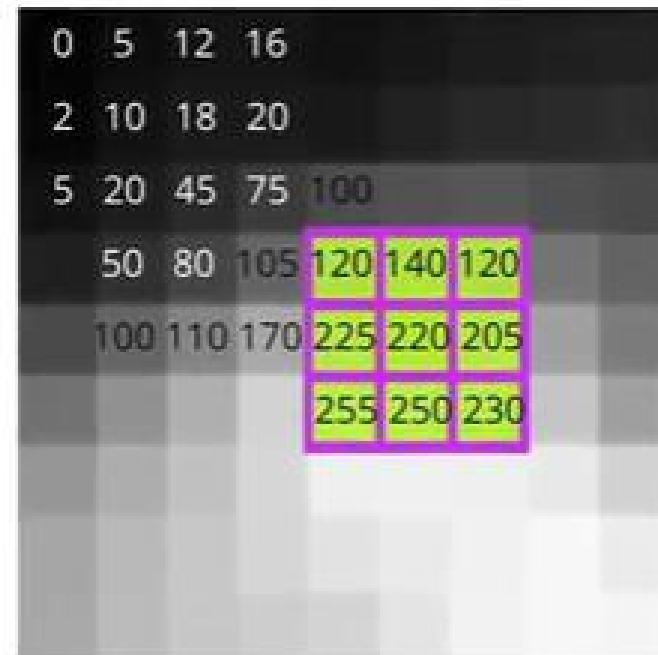
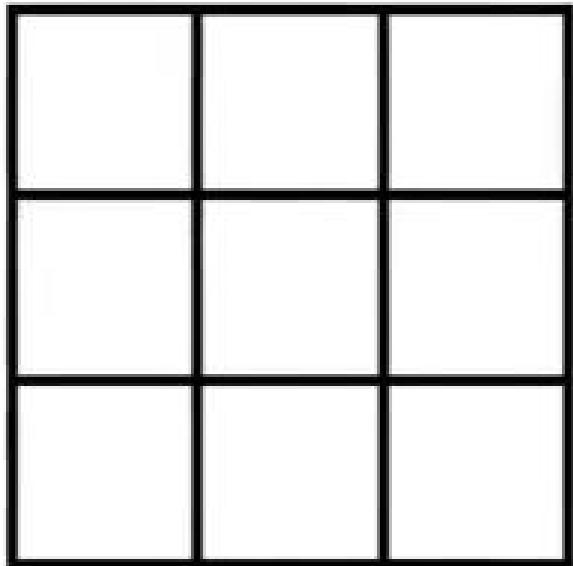
CONVOLUTION

0	-1	0
-1	4	-1
0	-1	0



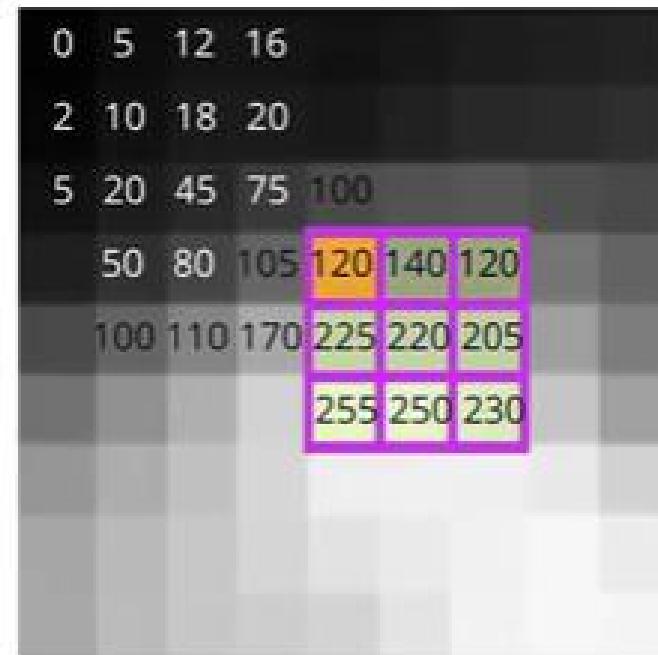
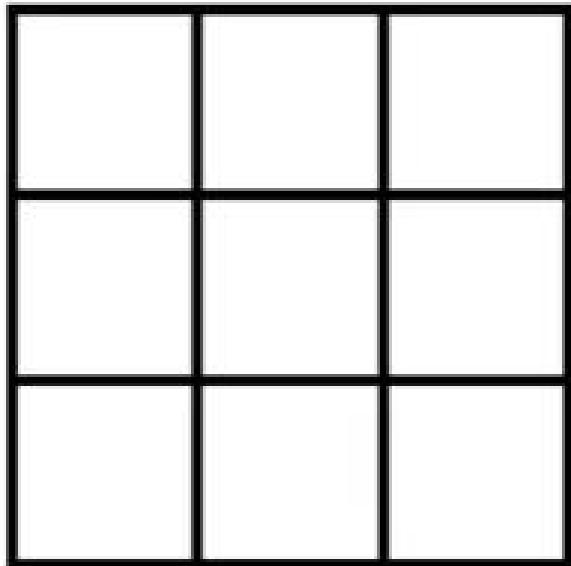
CONVOLUTION

0	-1	0
-1	4	-1
0	-1	0



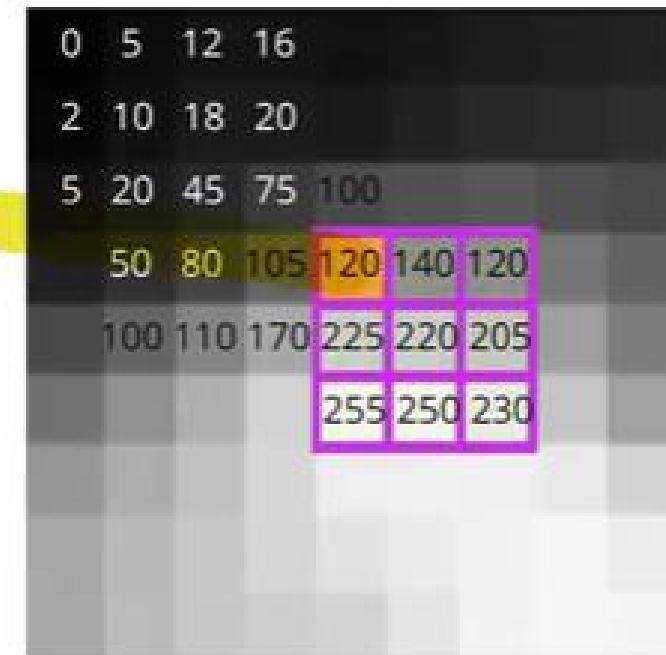
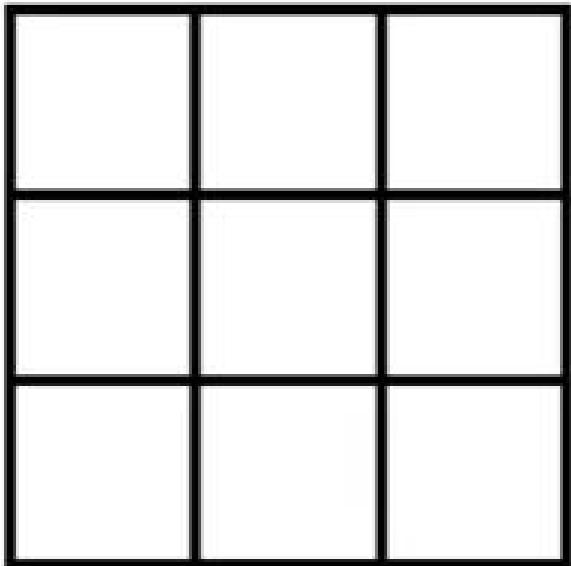
CONVOLUTION

0	-1	0
-1	4	-1
0	-1	0



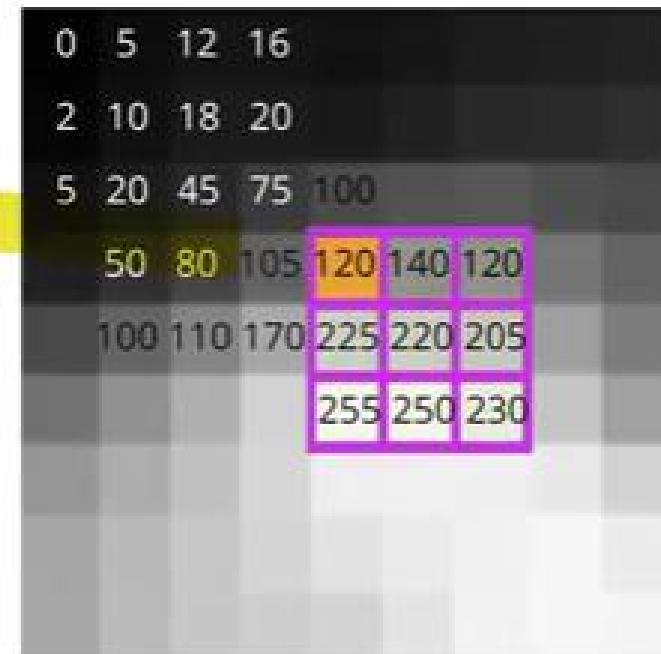
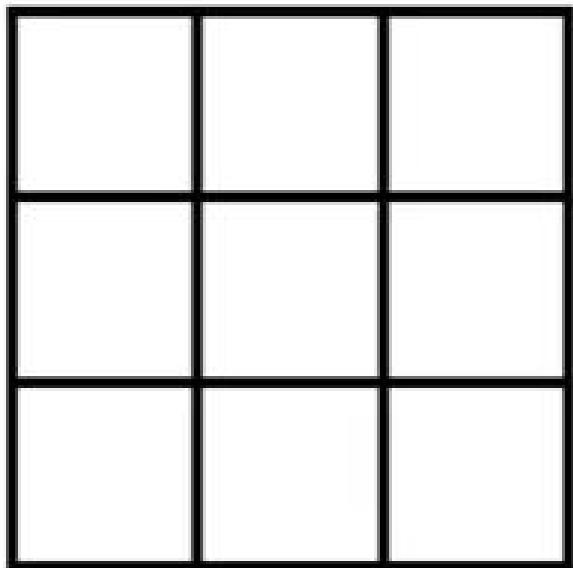
CONVOLUTION

0	-1	0
-1	4	-1
0	-1	0

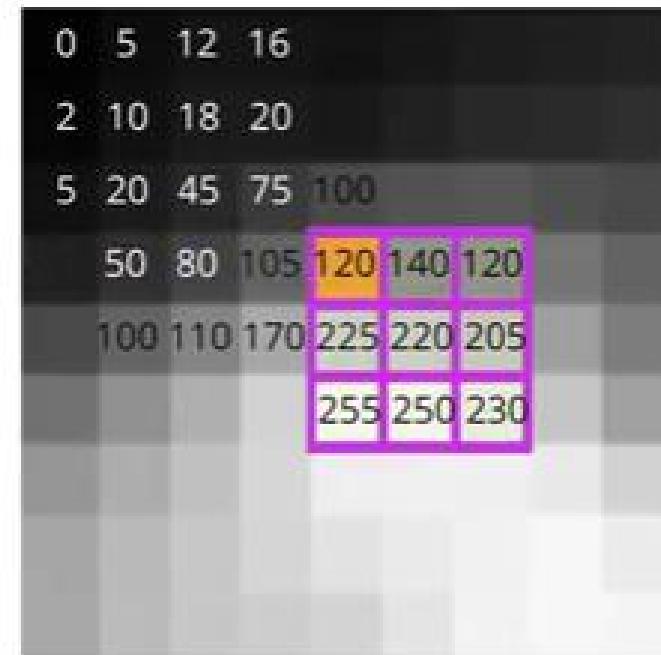
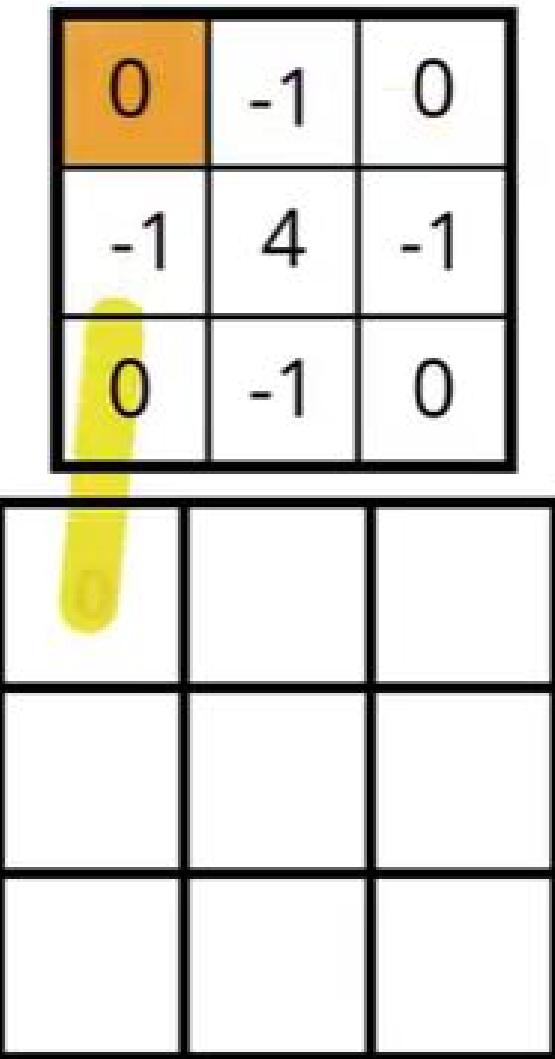


CONVOLUTION

0	-1	0
-1	4	-1
0	-1	0



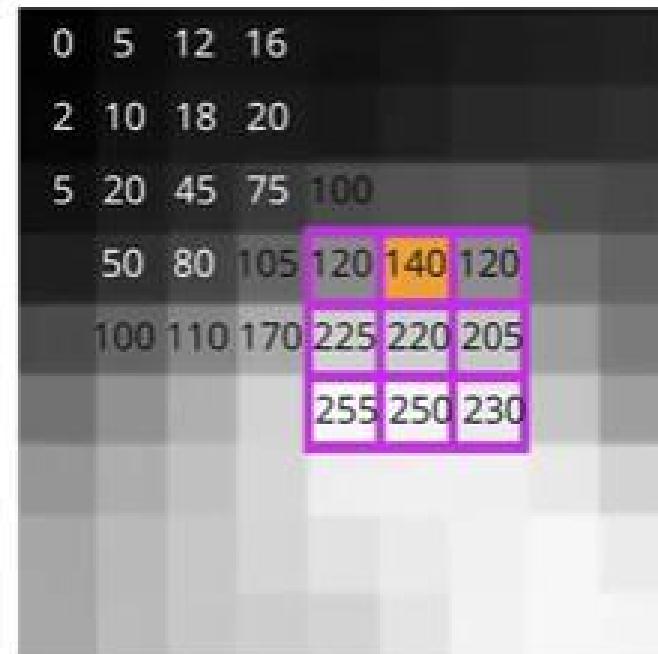
CONVOLUTION



CONVOLUTION

0	-1	0
-1	4	-1
0	-1	0

0		



CONVOLUTION

0	-1	0
-1	4	-1
0	-1	0

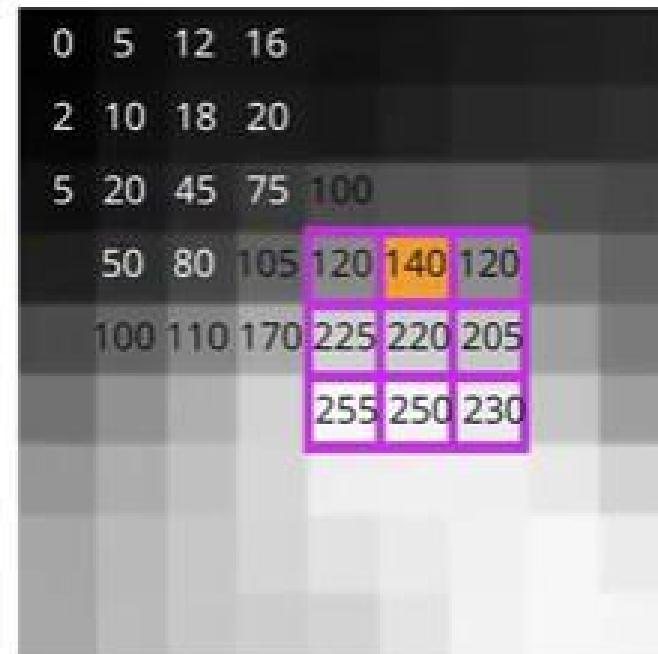
0		

0	5	12	16
2	10	18	20
5	20	45	75
100	50	80	105
120	120	140	120
205	225	220	205
230	250	255	230

CONVOLUTION

0	-1	0
-1	4	-1
0	-1	0

0	-140	



CONVOLUTION

0	-1	0
-1	4	-1
0	-1	0

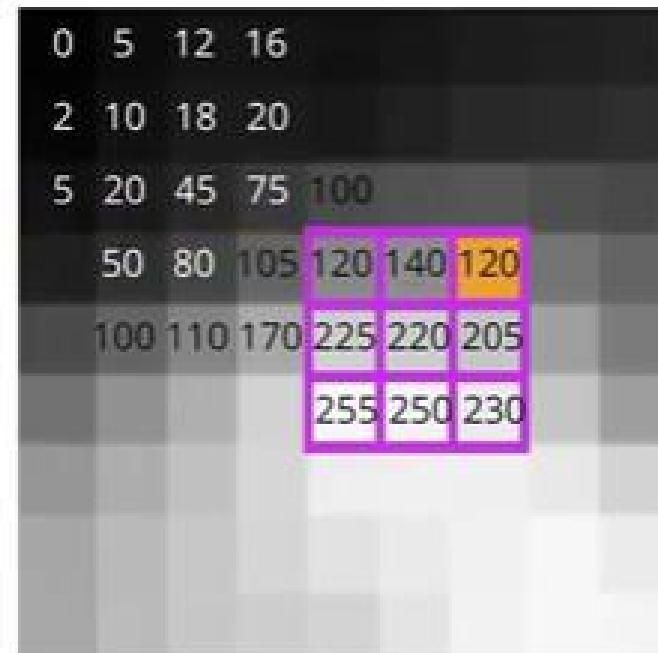
0	-140	

0	5	12	16	
2	10	18	20	
5	20	45	75	100
50	80	105	120	140
100	110	170	225	220
			205	
			255	250
			230	

CONVOLUTION

0	-1	0
-1	4	-1
0	-1	0

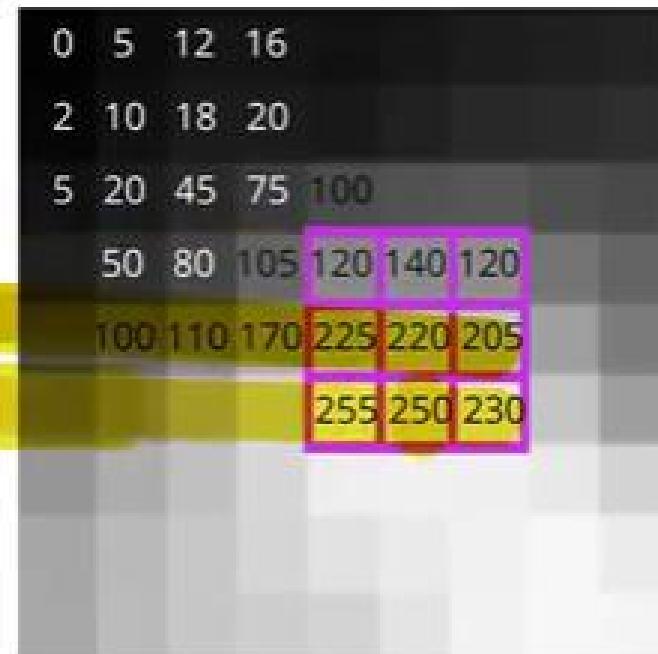
0	-140	



CONVOLUTION

0	-1	0
-1	4	-1
0	-1	0

0	-140	0



CONVOLUTION

0	-1	0
-1	4	-1
0	-1	0

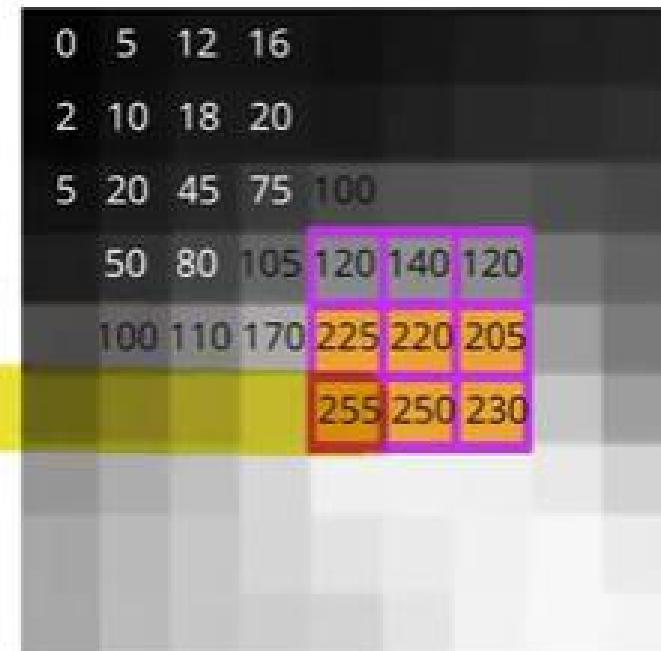
0	-140	0



CONVOLUTION

0	-1	0
-1	4	-1
0	-1	0

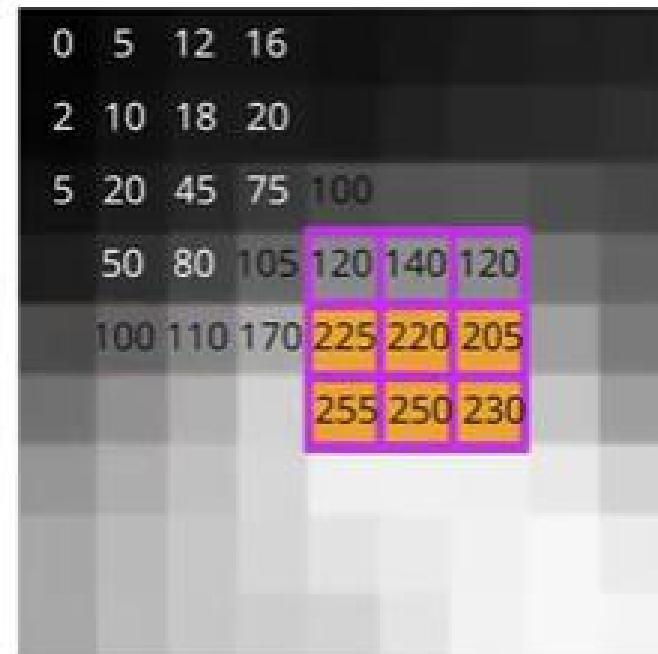
0	-140	0



CONVOLUTION

0	-1	0
-1	4	-1
0	-1	0

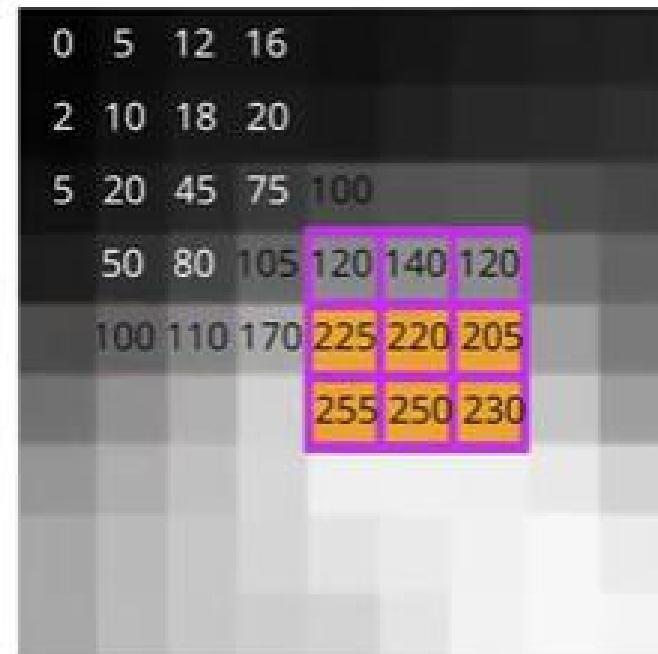
0	-140	0



CONVOLUTION

0	-1	0
-1	4	-1
0	-1	0

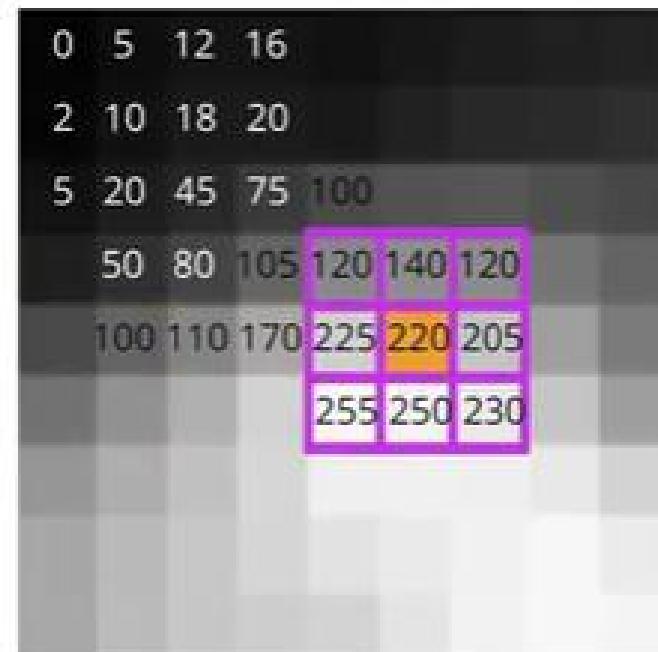
0	-140	0
-225		-205
0	0	0



CONVOLUTION

0	-1	0
-1	4	-1
0	-1	0

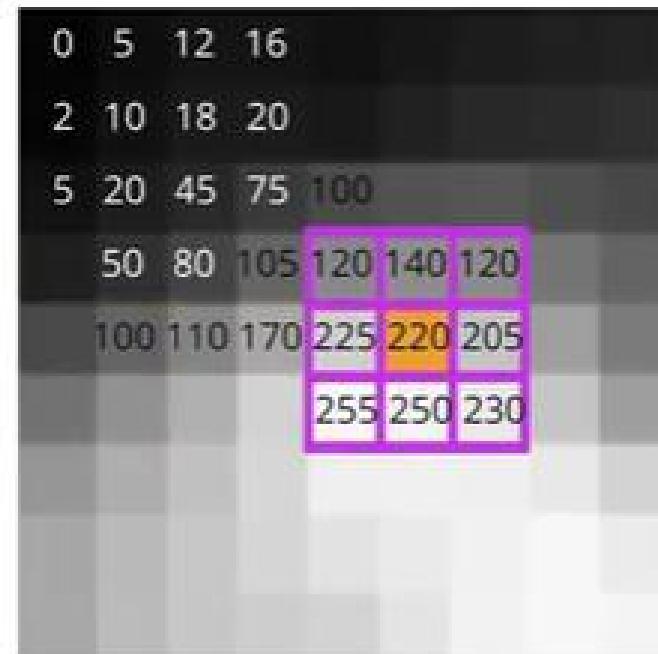
0	-140	0
-225		-205
0	-250	0



CONVOLUTION

0	-1	0
-1	4	-1
0	-1	0

0	-140	0
-225		-205
0	-250	0



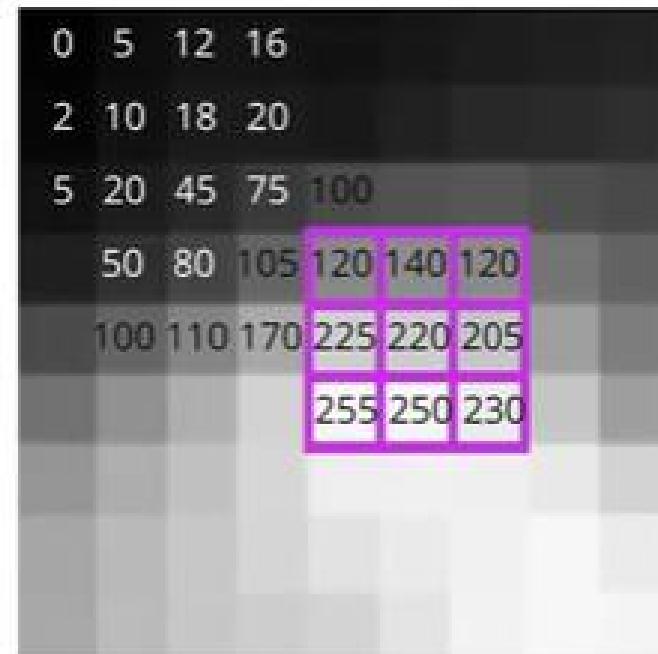
CONVOLUTION

0	-1	0
-1	4	-1
0	-1	0

0	-140	0
-225	880	-205
0	-250	0

+

=

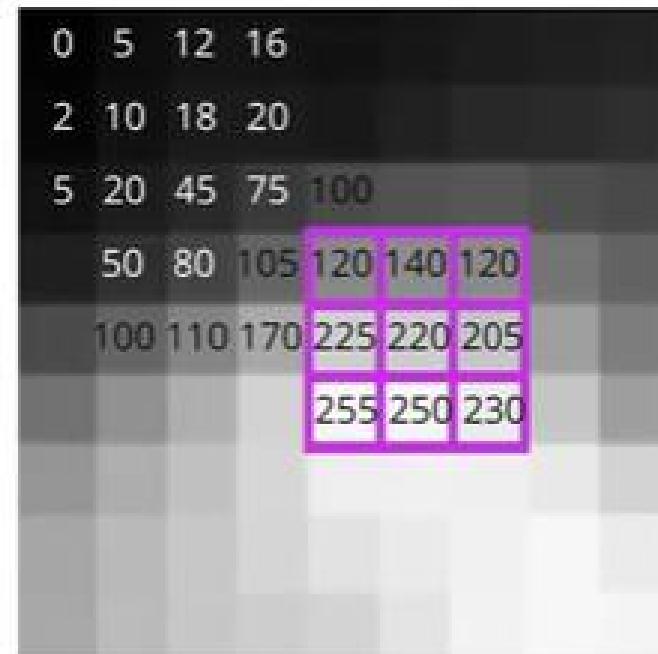


CONVOLUTION

0	-1	0
-1	4	-1
0	-1	0

0	-140	0
-225	880	-205
0	-250	0

= 60

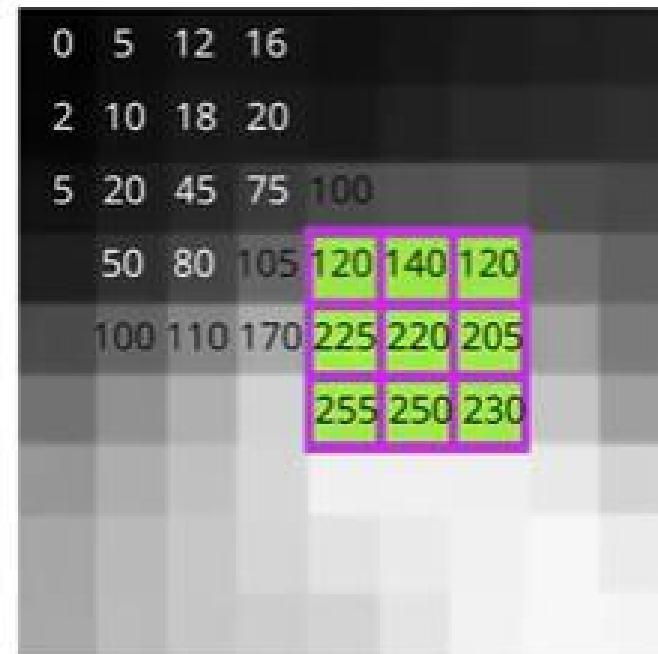


CONVOLUTION

0	-1	0
-1	4	-1
0	-1	0

0	-140	0
-225	880	-205
0	-250	0

= 60

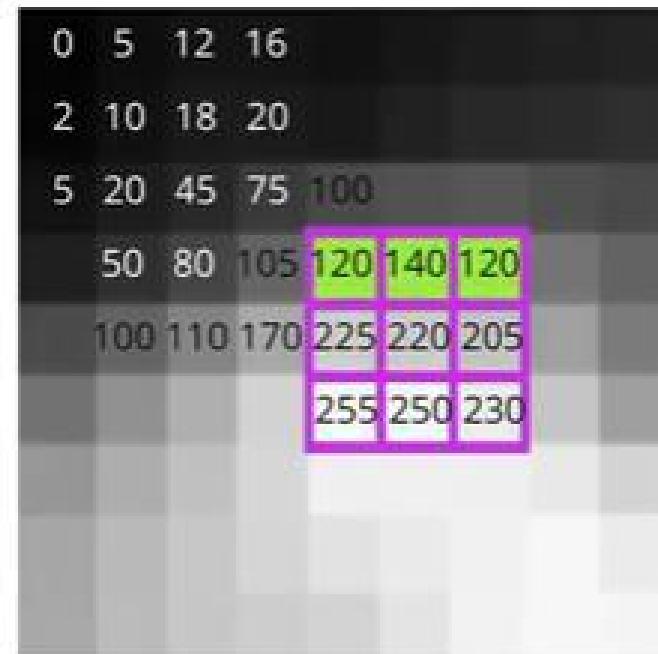


CONVOLUTION

0	-1	0
-1	4	-1
0	-1	0

0	-140	0
-225	880	-205
0	-250	0

= 60



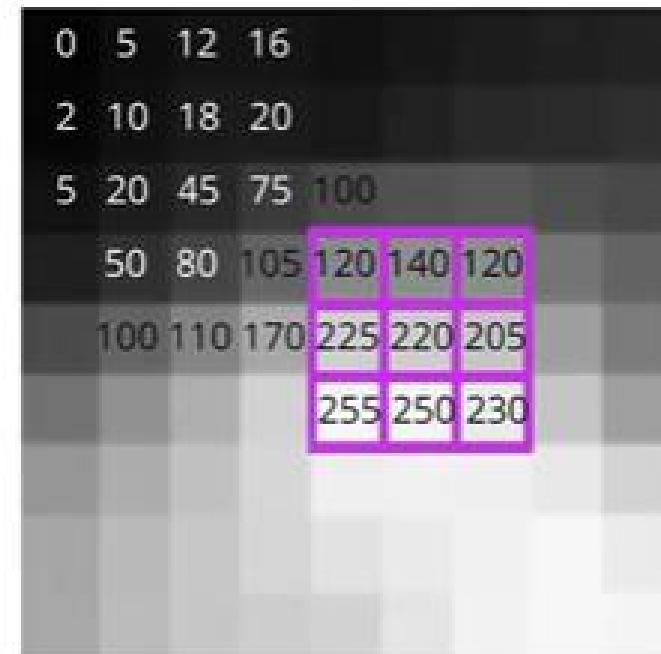
CONVOLUTION

Weights

0	-1	0
-1	4	-1
0	-1	0

0	-140	0
-225	880	-205
0	-250	0

$$= 60$$



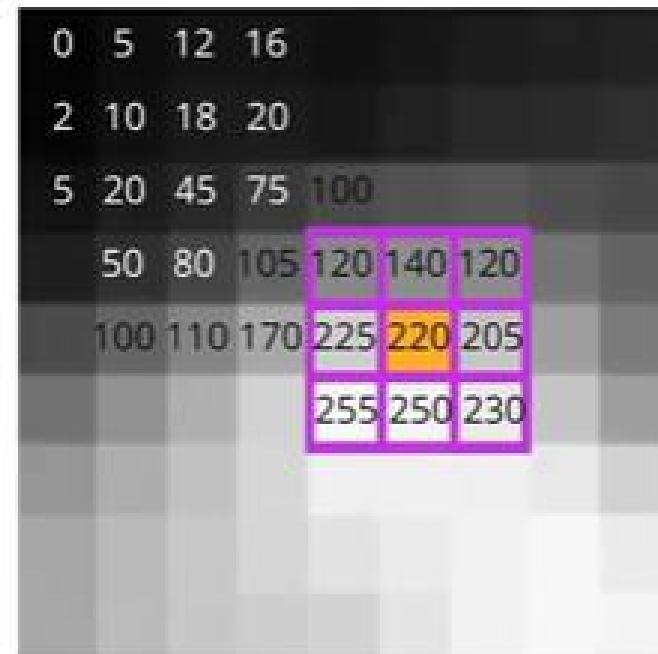
CONVOLUTION

Weights

0	-1	0
-1	4	-1
0	-1	0

0	-140	0
-225	880	-205
0	-250	0

$$= 60$$



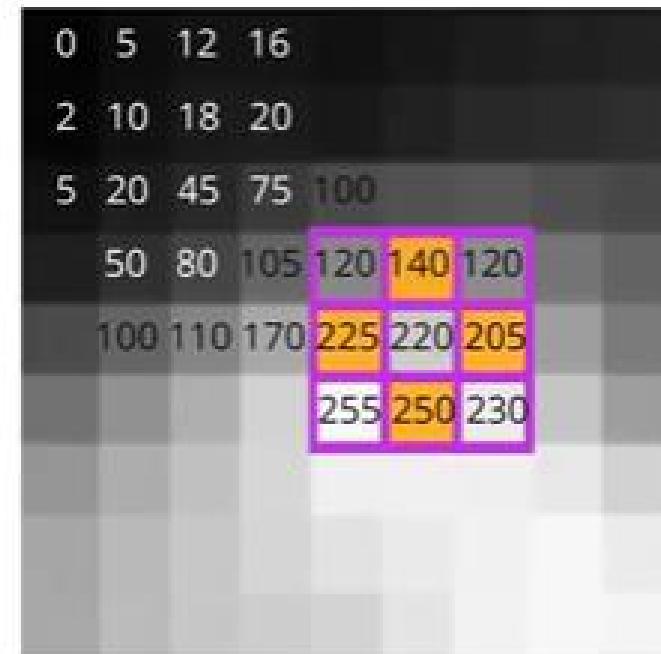
CONVOLUTION

Weights

0	-1	0
-1	4	-1
0	-1	0

0	-140	0
-225	880	-205
0	-250	0

$$= 60$$

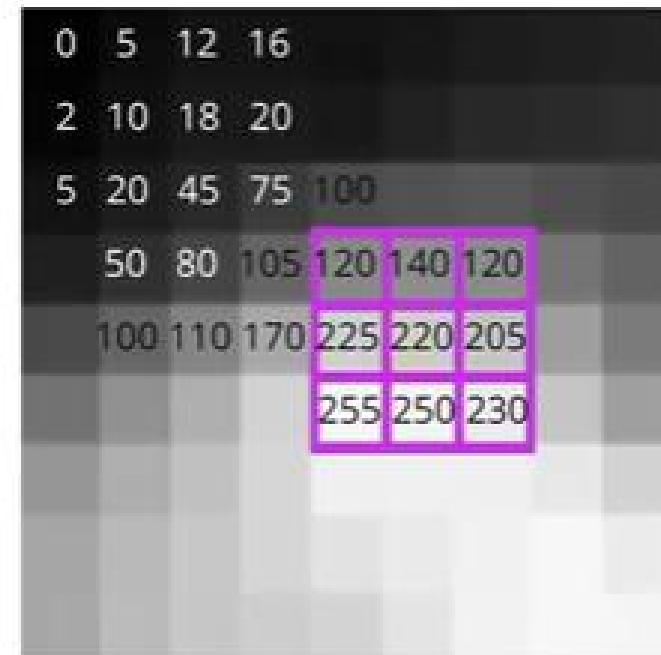


CONVOLUTION

Weights

0	-1	0
-1	4	-1
0	-1	0

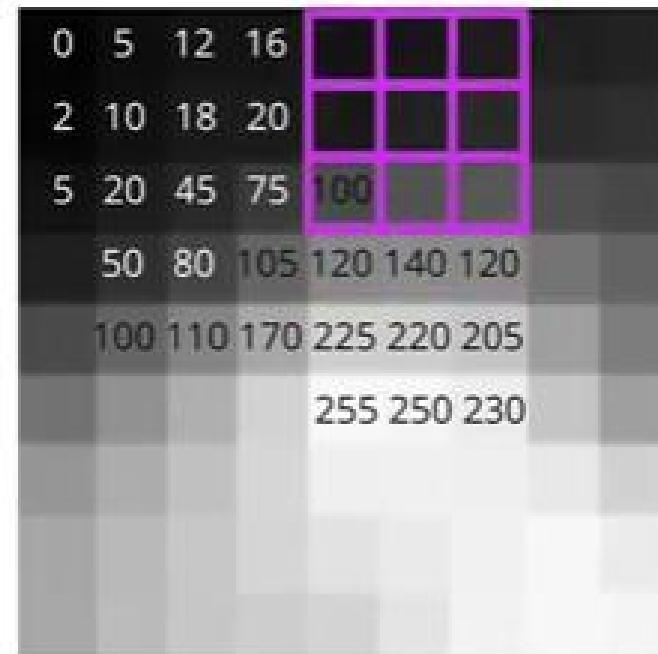
0	-140	0
-225	880	-205
0	-250	0



= **60** Pixel value in the output image

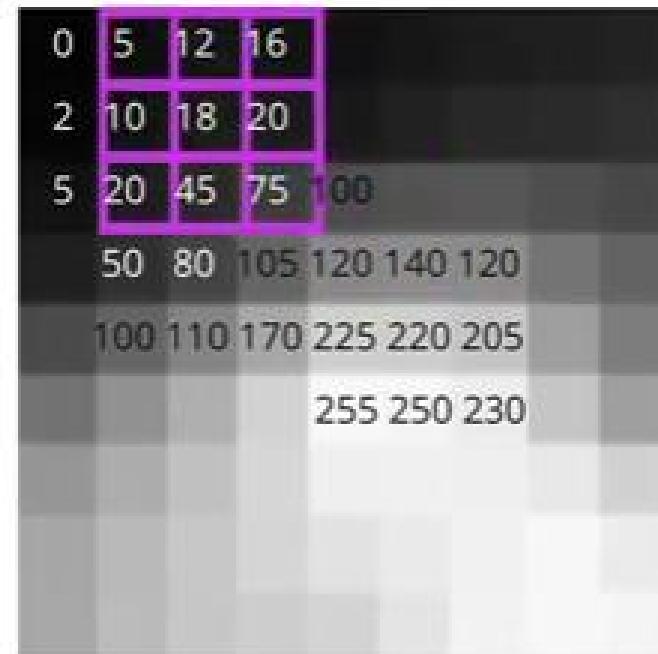
CONVOLUTION

0	-1	0
-1	4	-1
0	-1	0



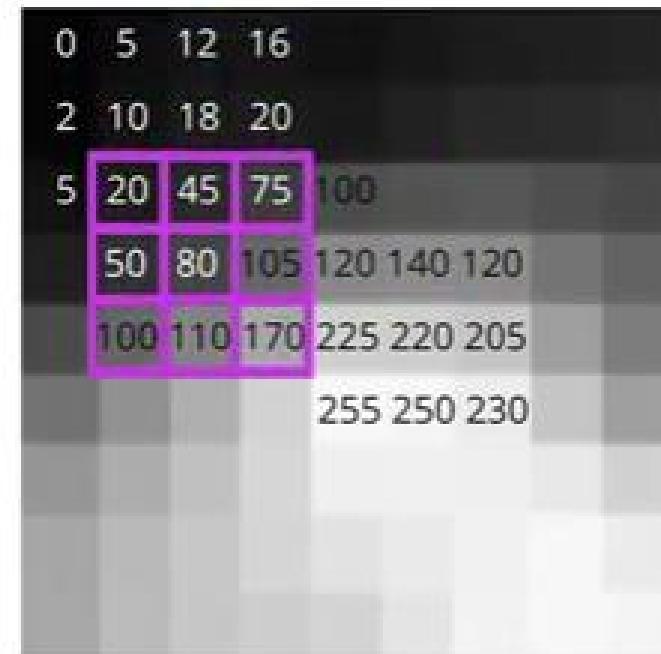
CONVOLUTION

0	-1	0
-1	4	-1
0	-1	0



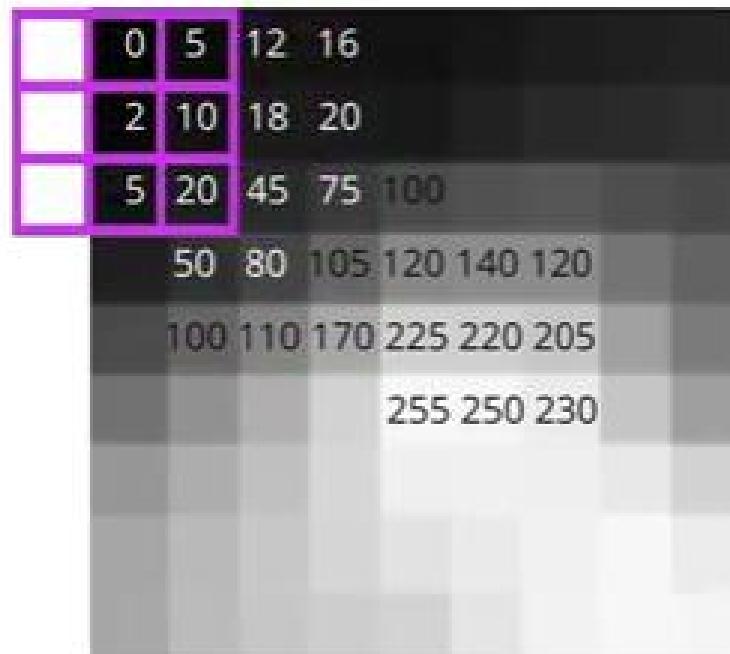
CONVOLUTION

0	-1	0
-1	4	-1
0	-1	0



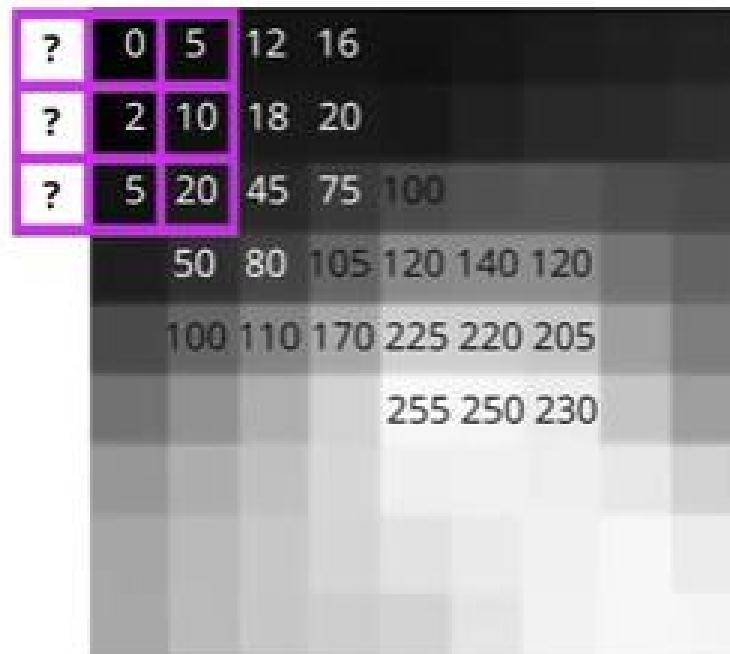
CONVOLUTION

0	-1	0
-1	4	-1
0	-1	0



CONVOLUTION

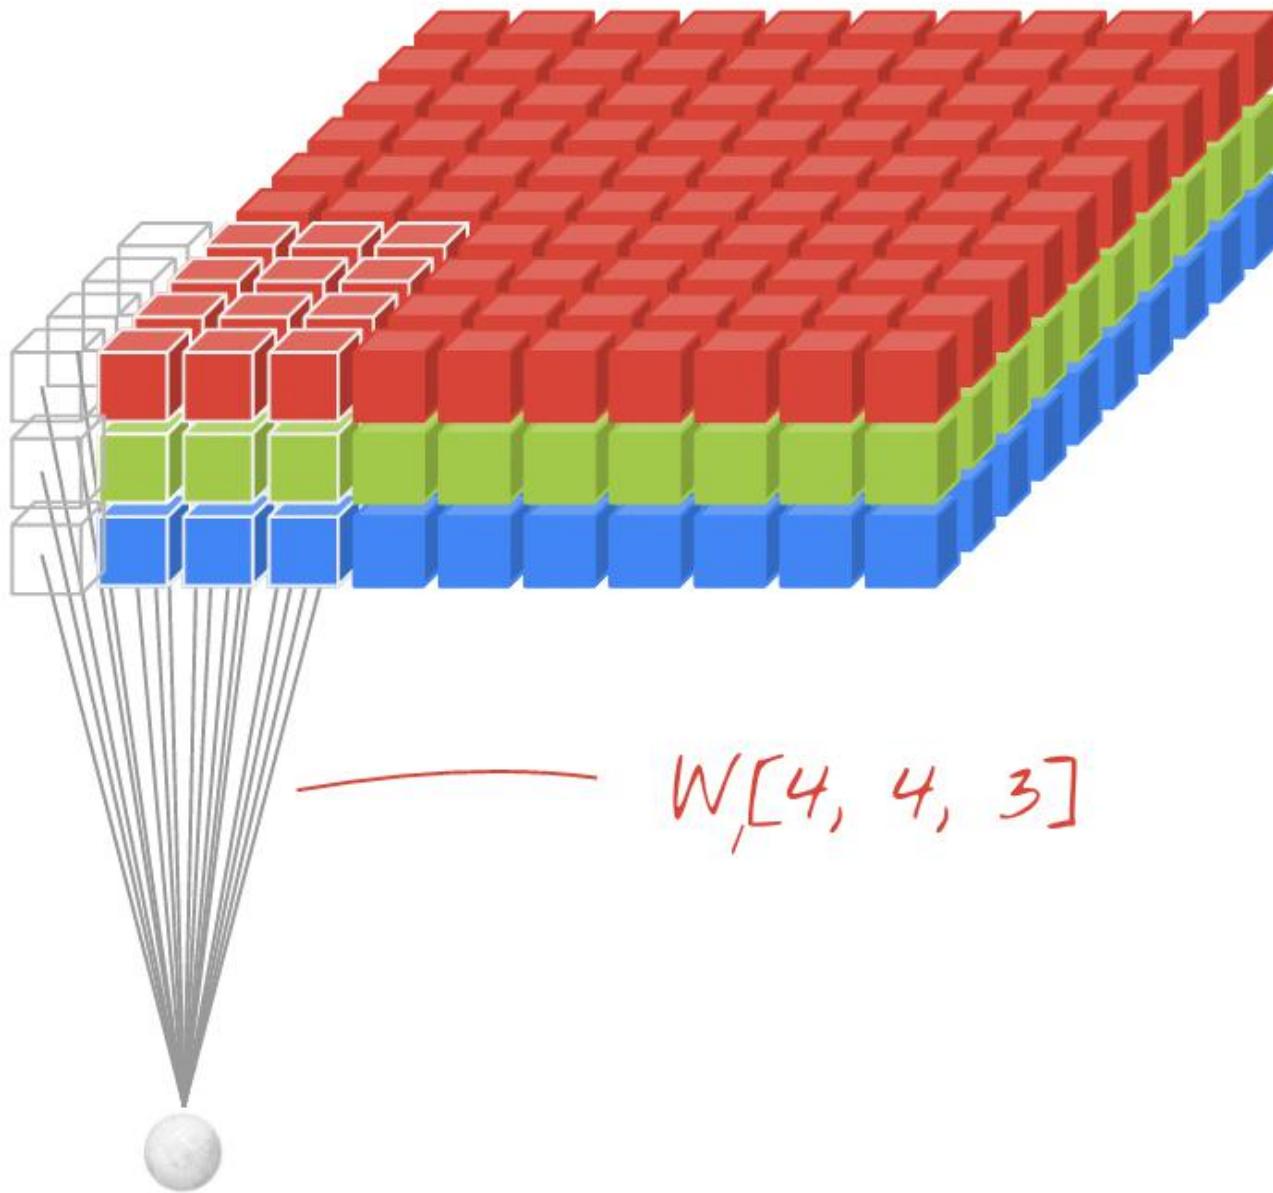
0	-1	0
-1	4	-1
0	-1	0



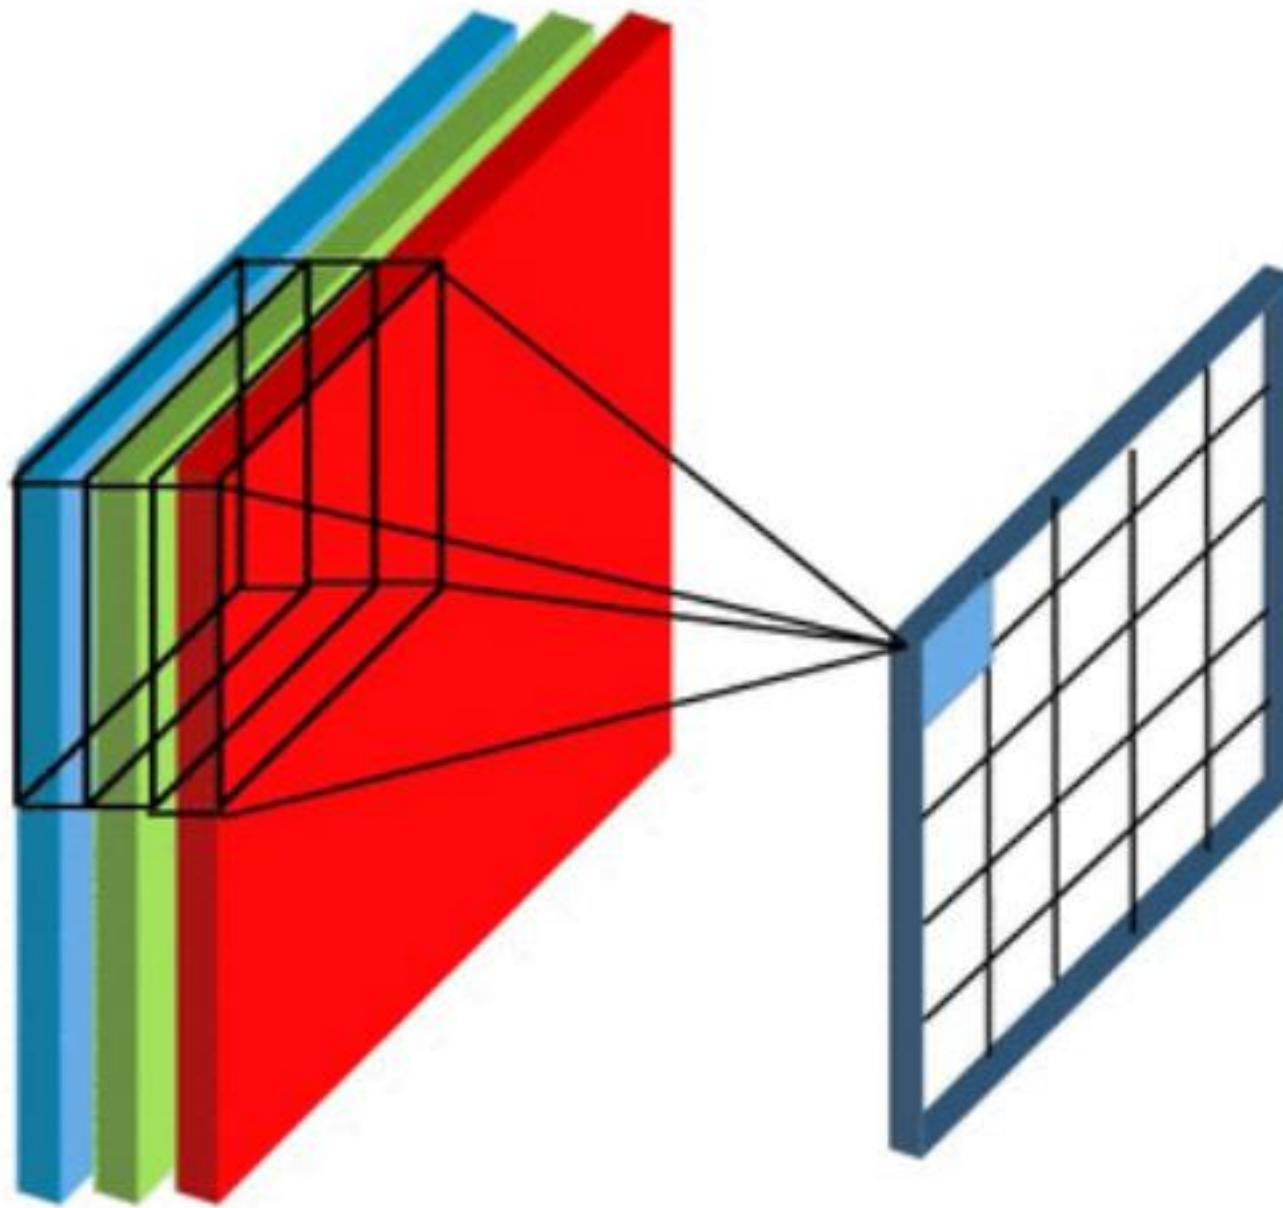
Padding

Sometimes filter does not perfectly fit the input image. We have two options:

- Pad the picture with zeros (zero-padding) so that it fits
- Drop the part of the image where the filter did not fit. This is called valid padding which keeps only valid part of the image.



$W[4, 4, 3]$



1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

*

1	0	1
0	1	0
1	0	1

5 x 5 – Image Matrix

3 x 3 – Filter Matrix

1 _{*1}	1 _{*0}	1 _{*1}	0	0
0 _{*0}	1 _{*1}	1 _{*0}	1	0
0 _{*1}	0 _{*0}	1 _{*1}	1	1
0	0	1	1	0
0	1	1	0	0

Image

4		

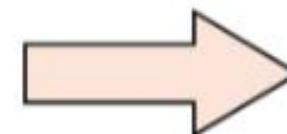
Convolved
Feature

Strides

Stride is the number of pixels shifts over the input matrix. When the stride is 1 then we move the filters to 1 pixel at a time. When the stride is 2 then we move the filters to 2 pixels at a time and so on. The below figure shows convolution would work with a stride of 2.

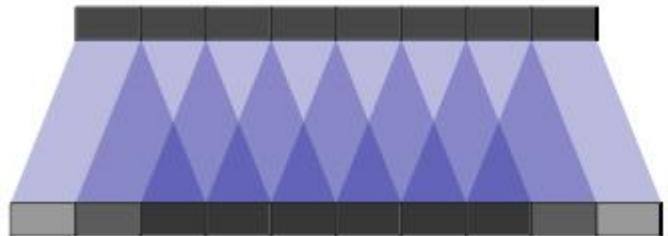
1	2	3	4	5	6	7
11	12	13	14	15	16	17
21	22	23	24	25	26	27
31	32	33	34	35	36	37
41	42	43	44	45	46	47
51	52	53	54	55	56	57
61	62	63	64	65	66	67
71	72	73	74	75	76	77

Convolve with 3x3
filters filled with ones

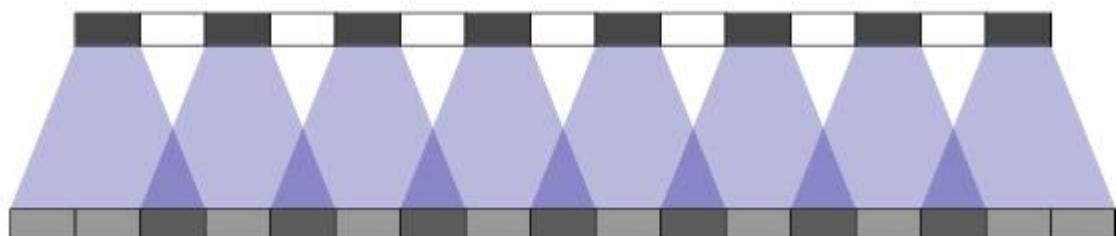


108	126	
288	306	

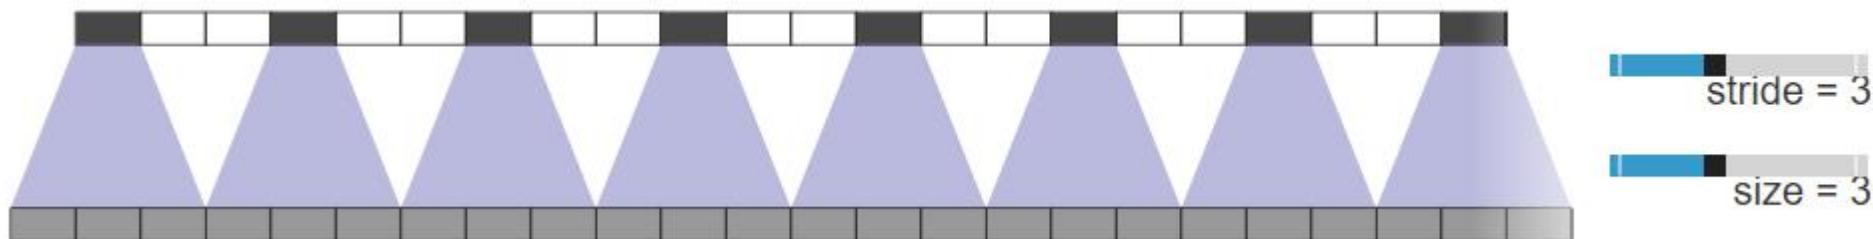
Figure 6 : Stride of 2 pixels



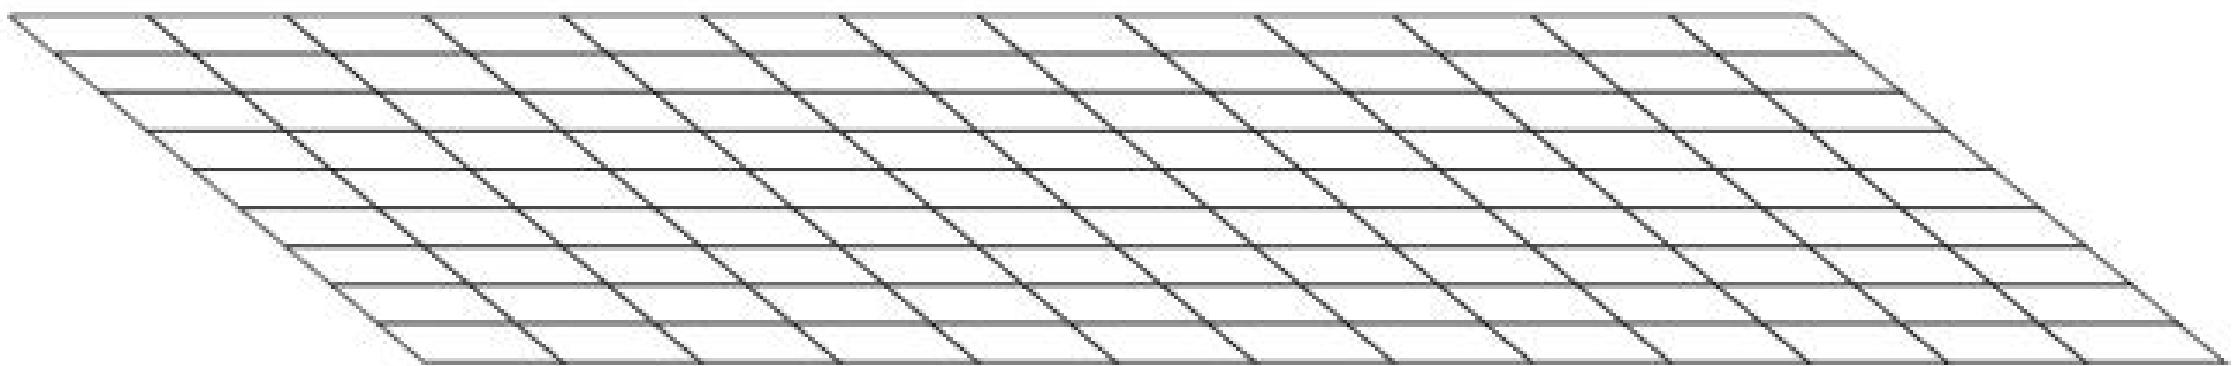
■ stride = 1
■ size = 3



■ stride = 2
■ size = 3



■ stride = 3
■ size = 3



How to compute the output volume [W2xH2xD2]?

Answer:

- $W2 = \lfloor (W1 - F + 2P) / S \rfloor + 1$
- $H2 = \lfloor (H1 - F + 2P) / S \rfloor + 1$
- $D2 = K$

where:

- $[W1 \times H1 \times D1]$: input volume size
- F: receptive field size
- S: stride
- P: amount of zero padding used on the border.
- K: depth

Kernel Size

Example:

What is the output volume of the first Convolutional Layer of Krizhevsky et al. architecture that won the ImageNet challenge in 2012?

Input size: $[227 \times 227 \times 3]$, W=227, F=11, S=4, P=0, and K=96.

Answer:

- $(227 - 11) / 4 + 1 = 55$
- The size of the Conv layer output volume is $[55 \times 55 \times 96]$.

Input Volume (+pad 1) (7x7x3)

 $x[:, :, 0]$

0	0	0	0	0	0	0
0	0	0	1	0	2	0
0	1	0	2	0	1	0

0	1	0	2	2	0	0
0	2	0	0	2	0	0
0	2	1	2	2	0	0
0	0	0	0	0	0	0

0	0	0	0	0	0	0
0	2	1	2	1	1	0
0	2	1	2	0	1	0

0	0	2	1	0	1	0
0	1	2	2	2	2	0
0	0	1	2	0	1	0
0	0	0	0	0	0	0

0	0	0	0	0	0	0
0	2	1	1	2	0	0
0	1	0	0	1	0	0

0	0	1	0	0	0	0
0	1	0	2	1	0	0
0	2	2	1	1	1	0
0	0	0	0	0	0	0

Filter W0 (3x3x3)

 $w0[:, :, 0]$

-1	0	1
0	0	1
1	-1	1

-1	0	1
1	-1	1
0	1	0

-1	1	1
1	1	0
0	-1	0

1
b0[:, :, 0]
0

Filter W1 (3x3x3)

 $w1[:, :, 0]$

0	1	-1
0	-1	0
0	-1	1

-1	0	0
1	-1	0
1	-1	0

-1	1	-1
0	-1	-1
1	0	0

0
b1[:, :, 0]
0

Output Volume (3x3x2)

 $o[:, :, 0]$

2	3	3
3	7	3
8	10	-3

o[:, :, 1]		
-8	-8	-3
-3	1	0
-3	-8	-5

toggle movement

$$W2 = [(W1 - F + 2P)/S] + 1$$

$$H2 = [(H1 - F + 2P)/S] + 1$$

$$D2 = K$$

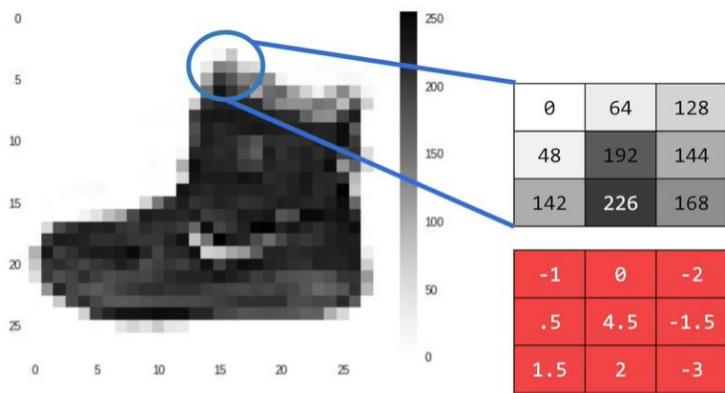
$$W = [5 - 3 + 2 * 1] / 2 + 1 = 3$$

$$H = [5 - 3 + 2 * 1] / 2 + 1 = 3$$

$$D = 2$$

$$\text{Output} = 3 \times 3 \times 2$$

output shape is (stacked 2D) **3D = 2D x N** matrix

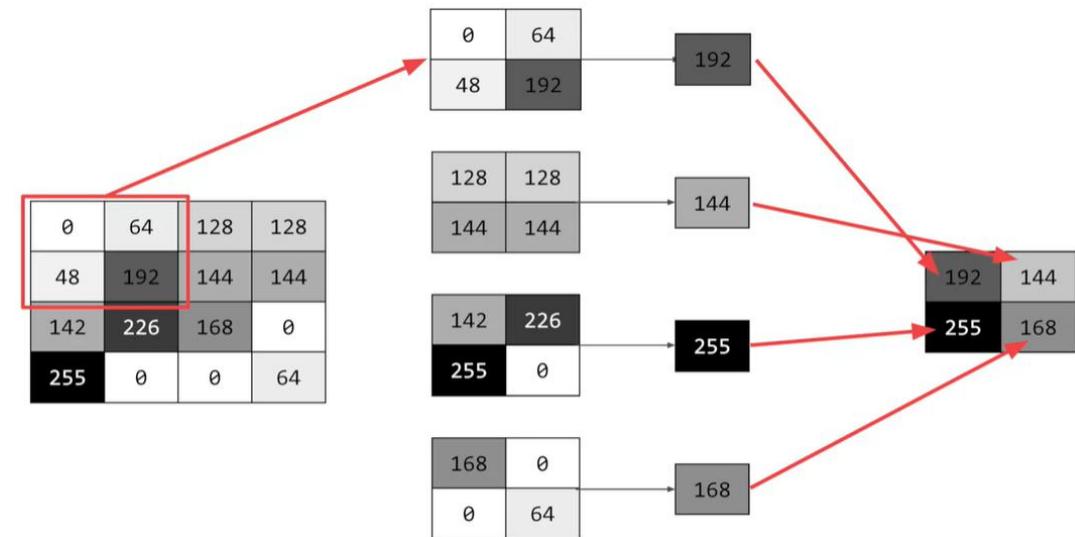


```
CURRENT_PIXEL_VALUE = 192
NEW_PIXEL_VALUE = (-1 * 0) + (0 * 64) + (-2 * 128) +
(.5 * 48) + (4.5 * 192) + (-1.5 * 144) +
(1.5 * 42) + (2 * 226) + (-3 * 168)
```

Current Pixel Value is 192

Consider neighbor Values

Filter Definition



Non Linearity (ReLU)

ReLU stands for Rectified Linear Unit for a non-linear operation. The output is $f(x) = \max(0, x)$.

Why ReLU is important : ReLU's purpose is to introduce non-linearity in our ConvNet. Since, the real world data would want our ConvNet to learn would be non-negative linear values.

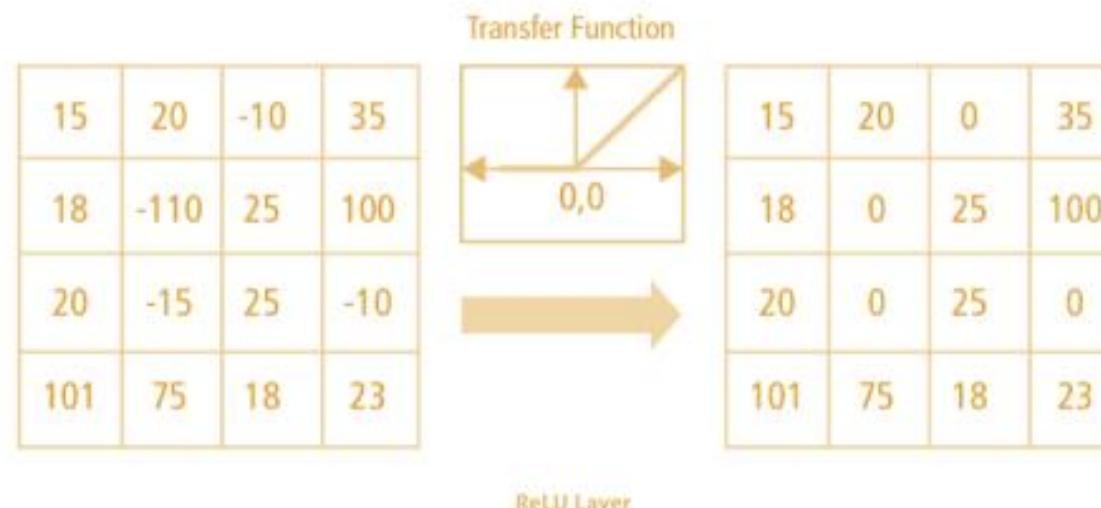
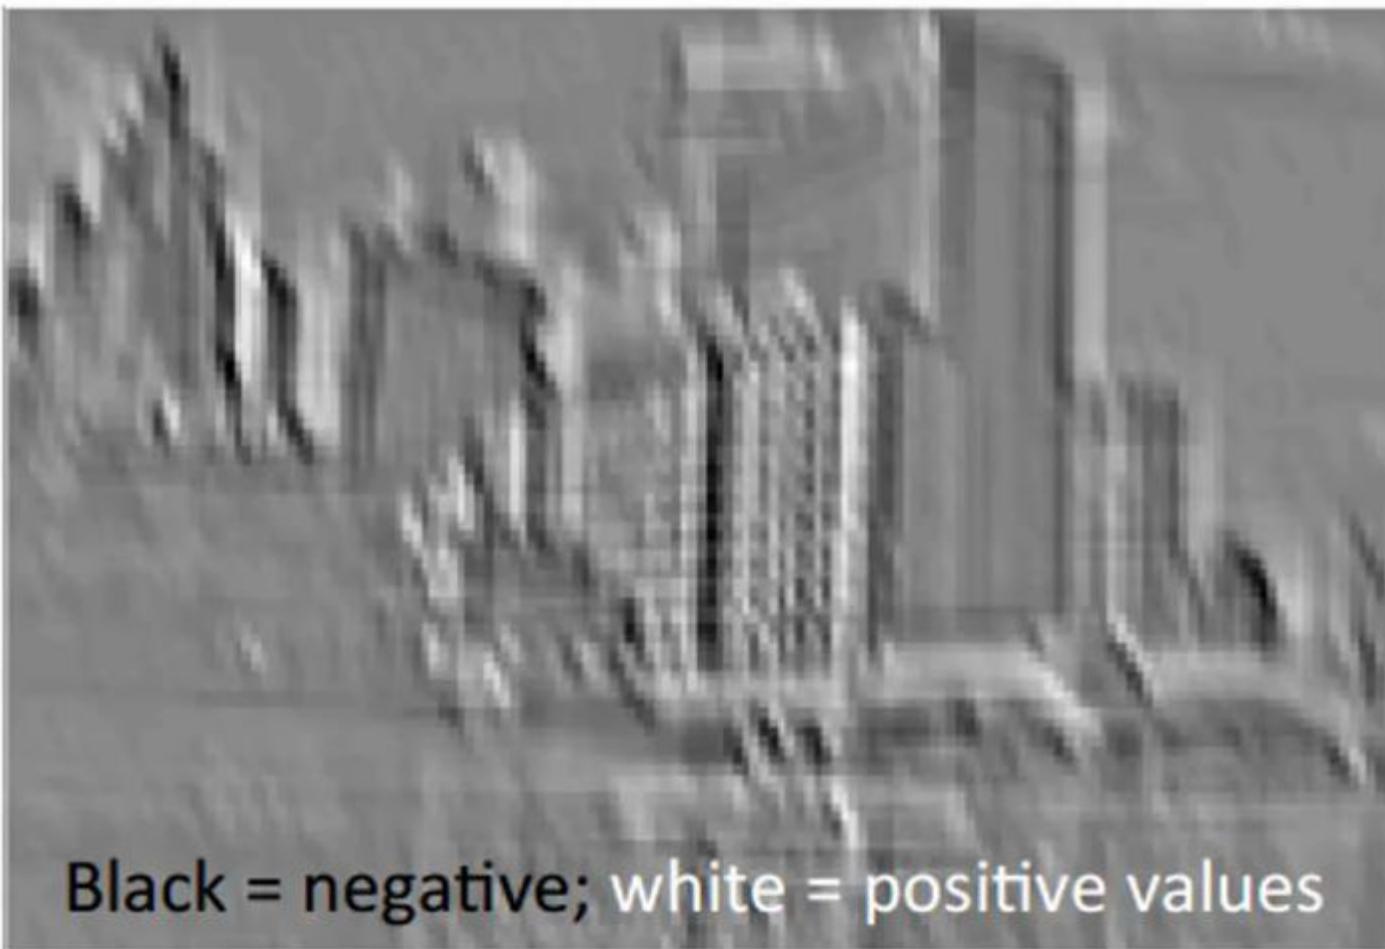


Figure 7 : ReLU operation

ReLU Effect



Kernel Effect

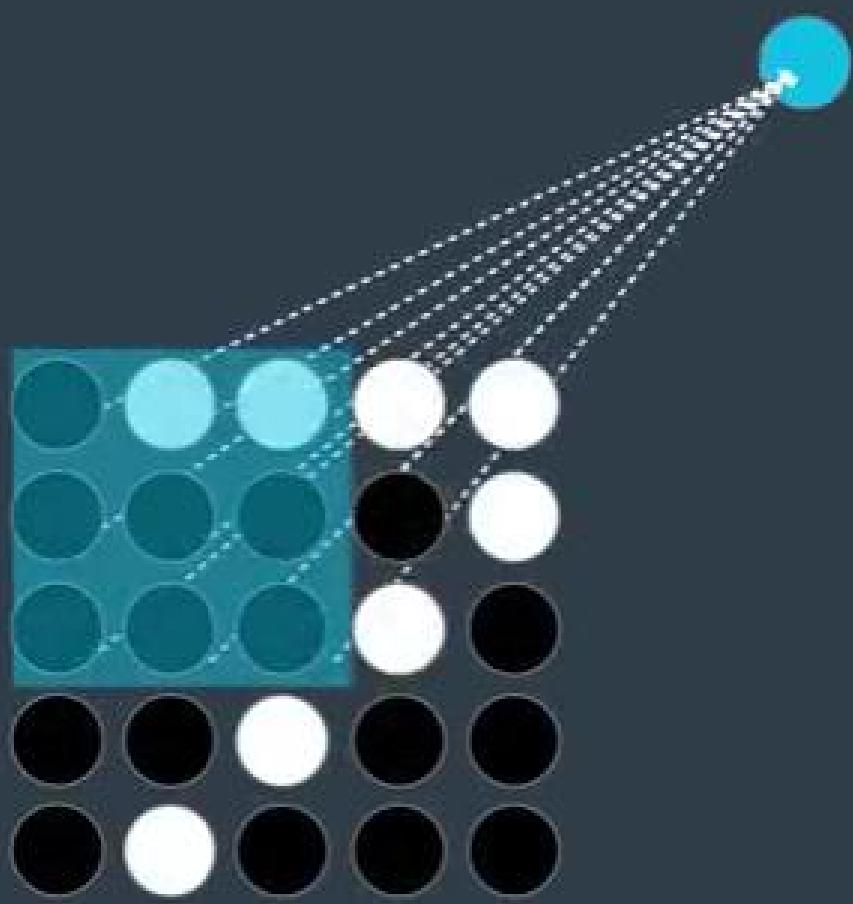


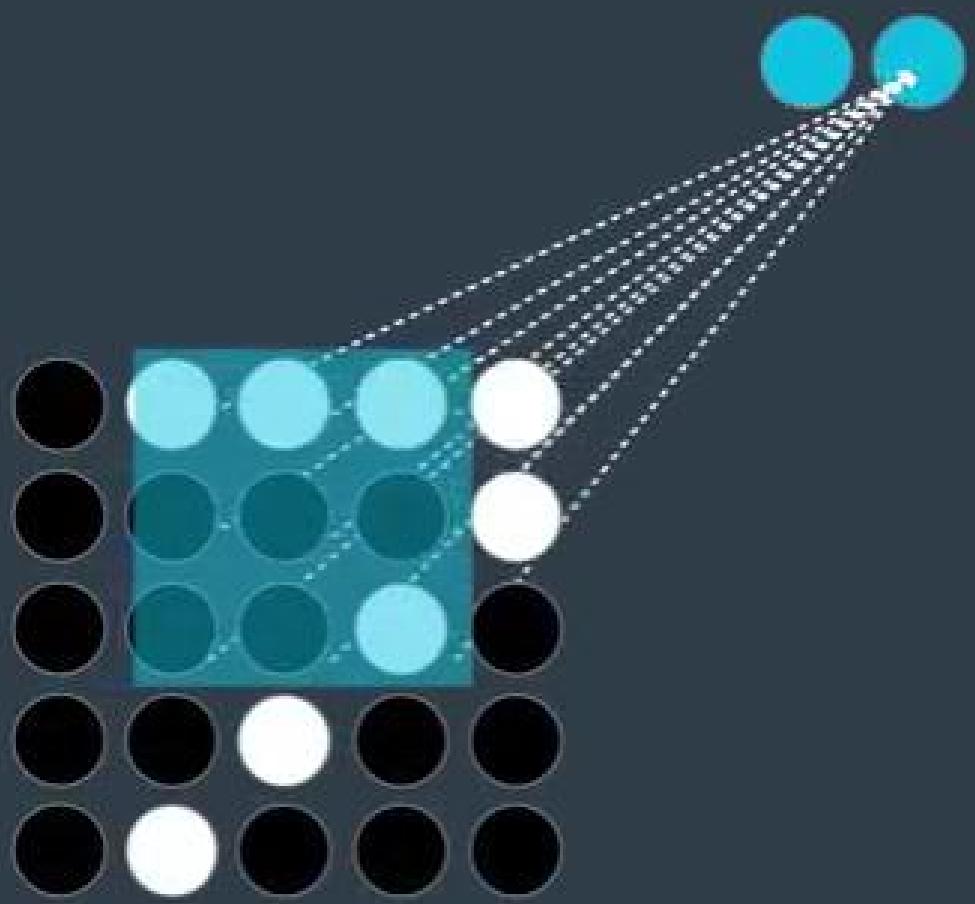
Black = negative; white = positive values

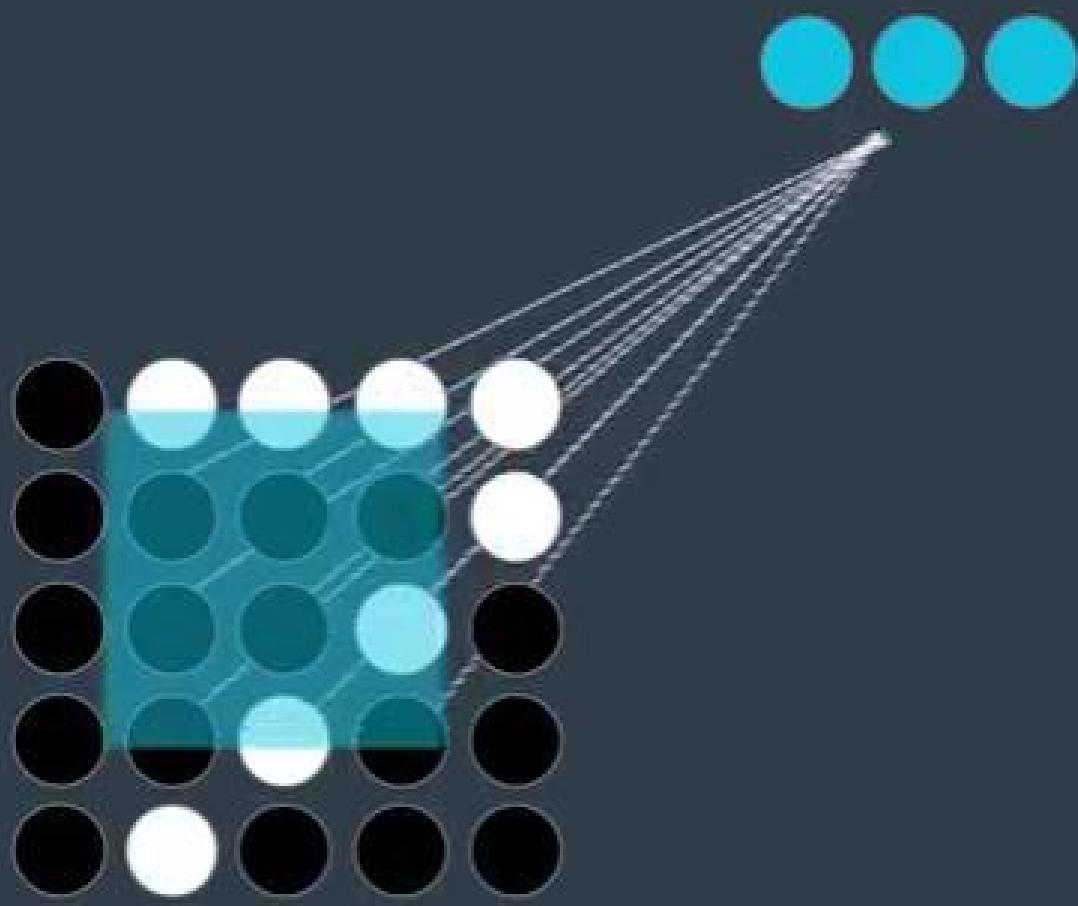
ReLU Effect

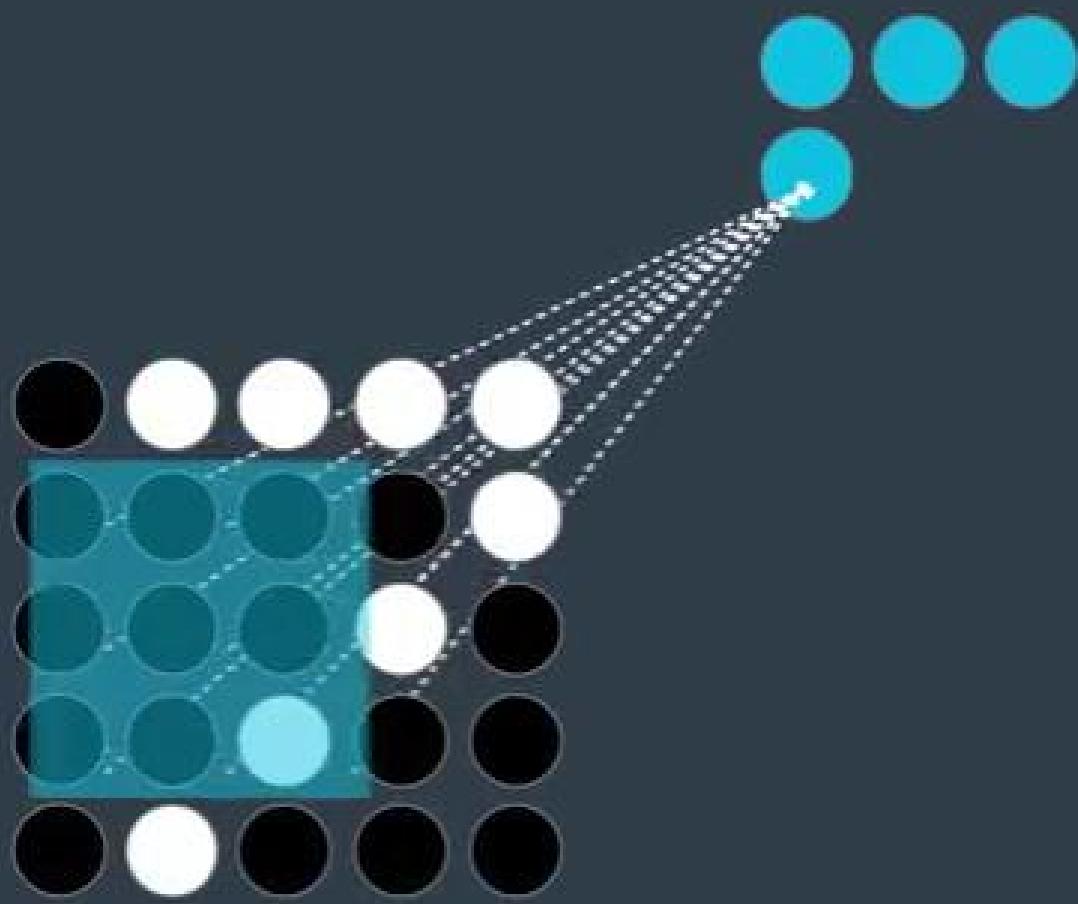


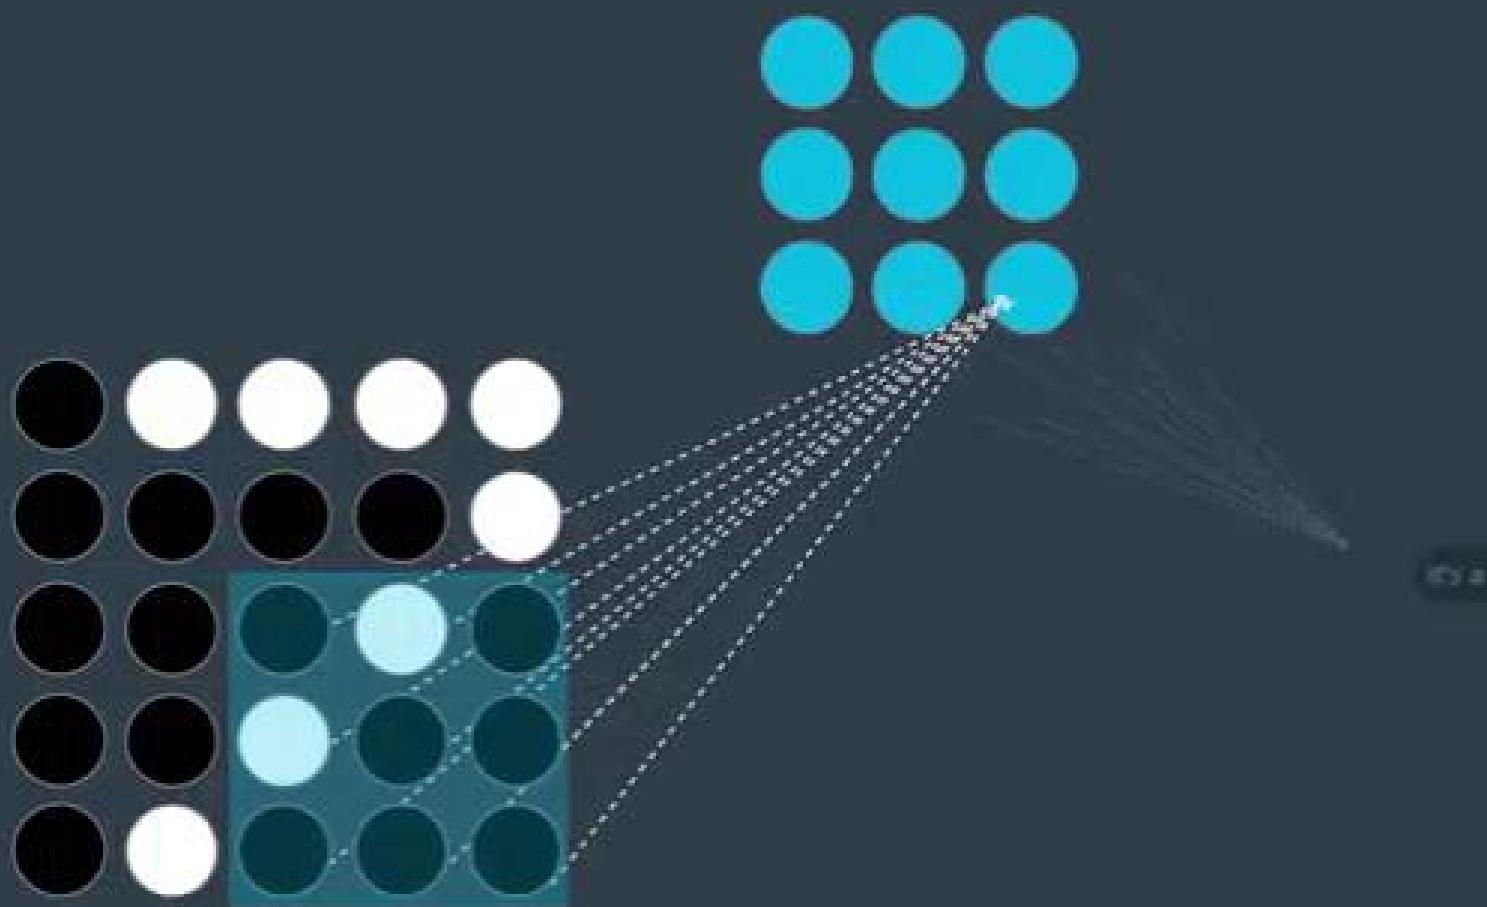
Only non-negative values



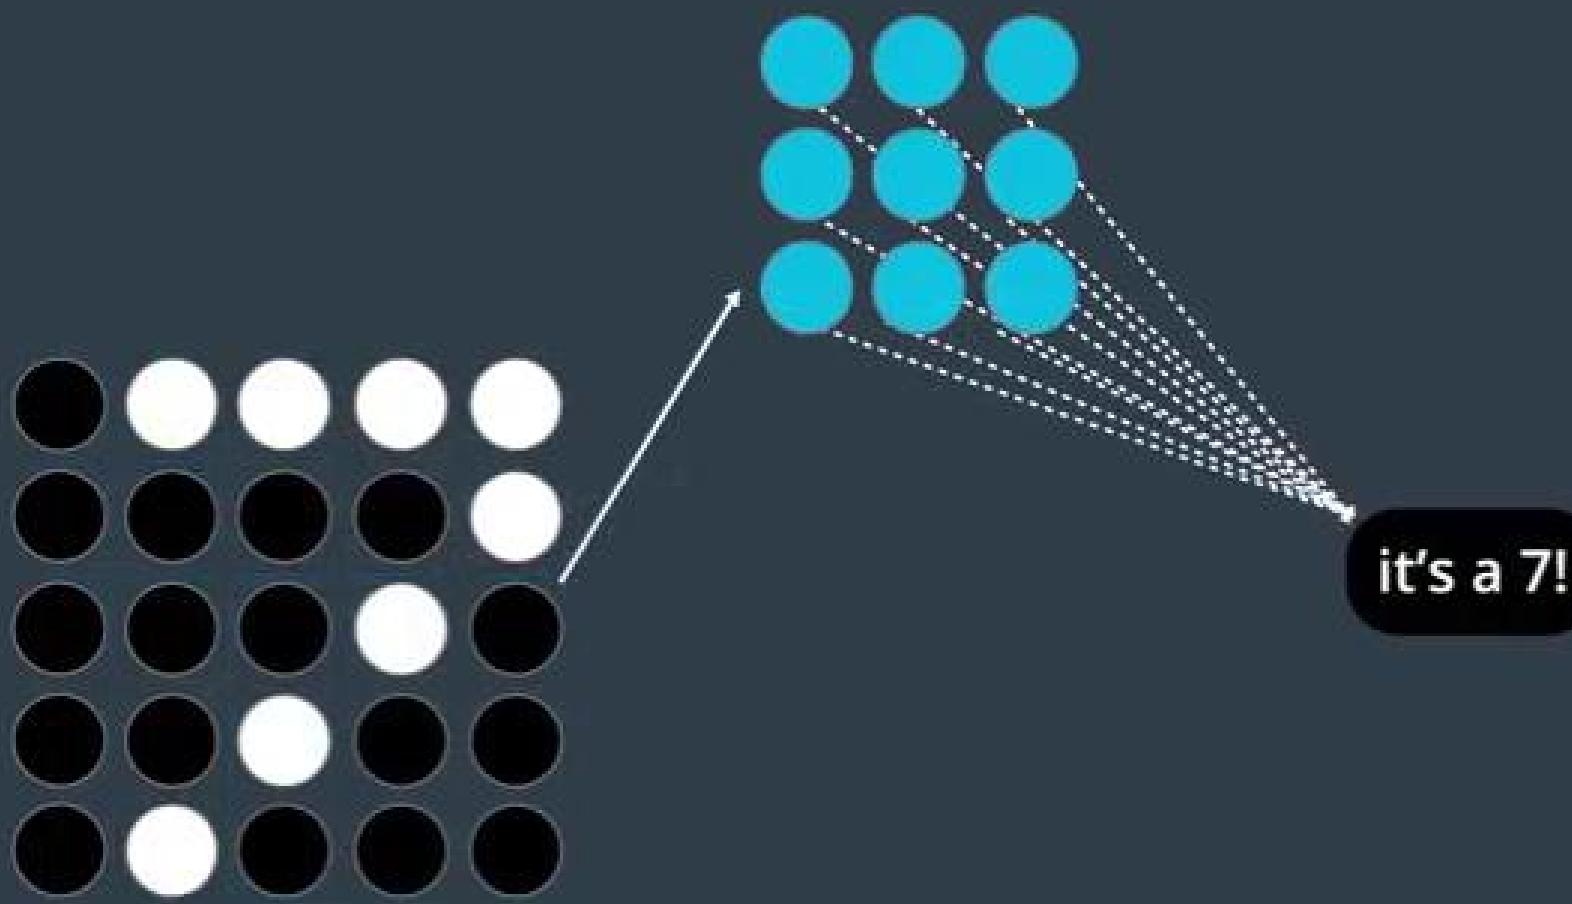


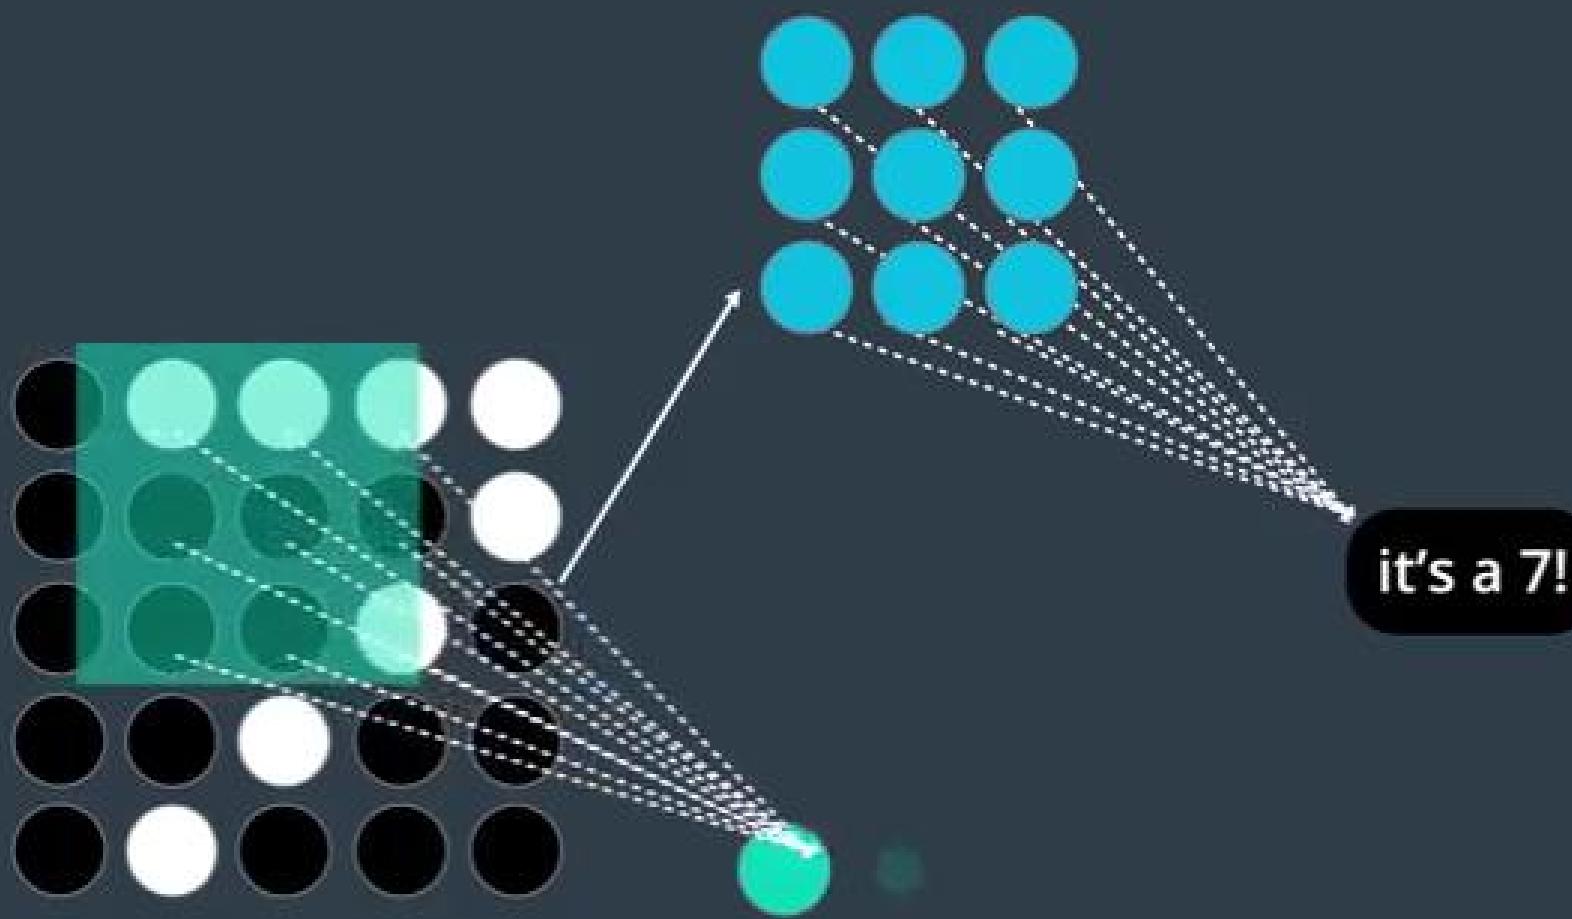




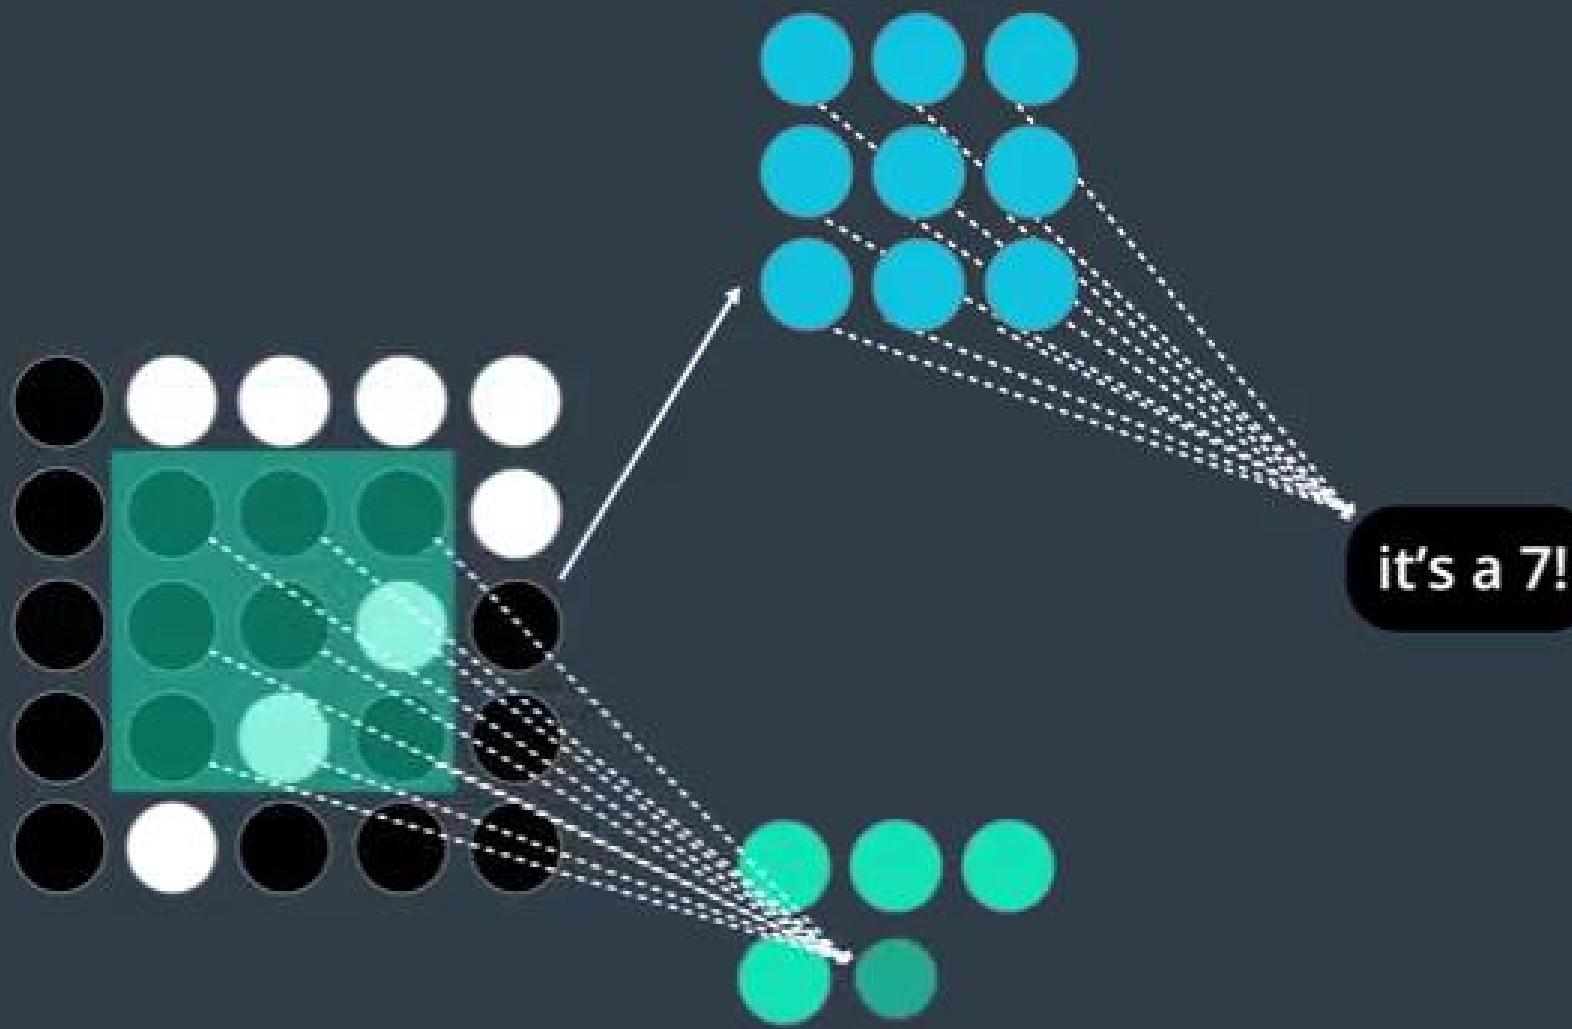


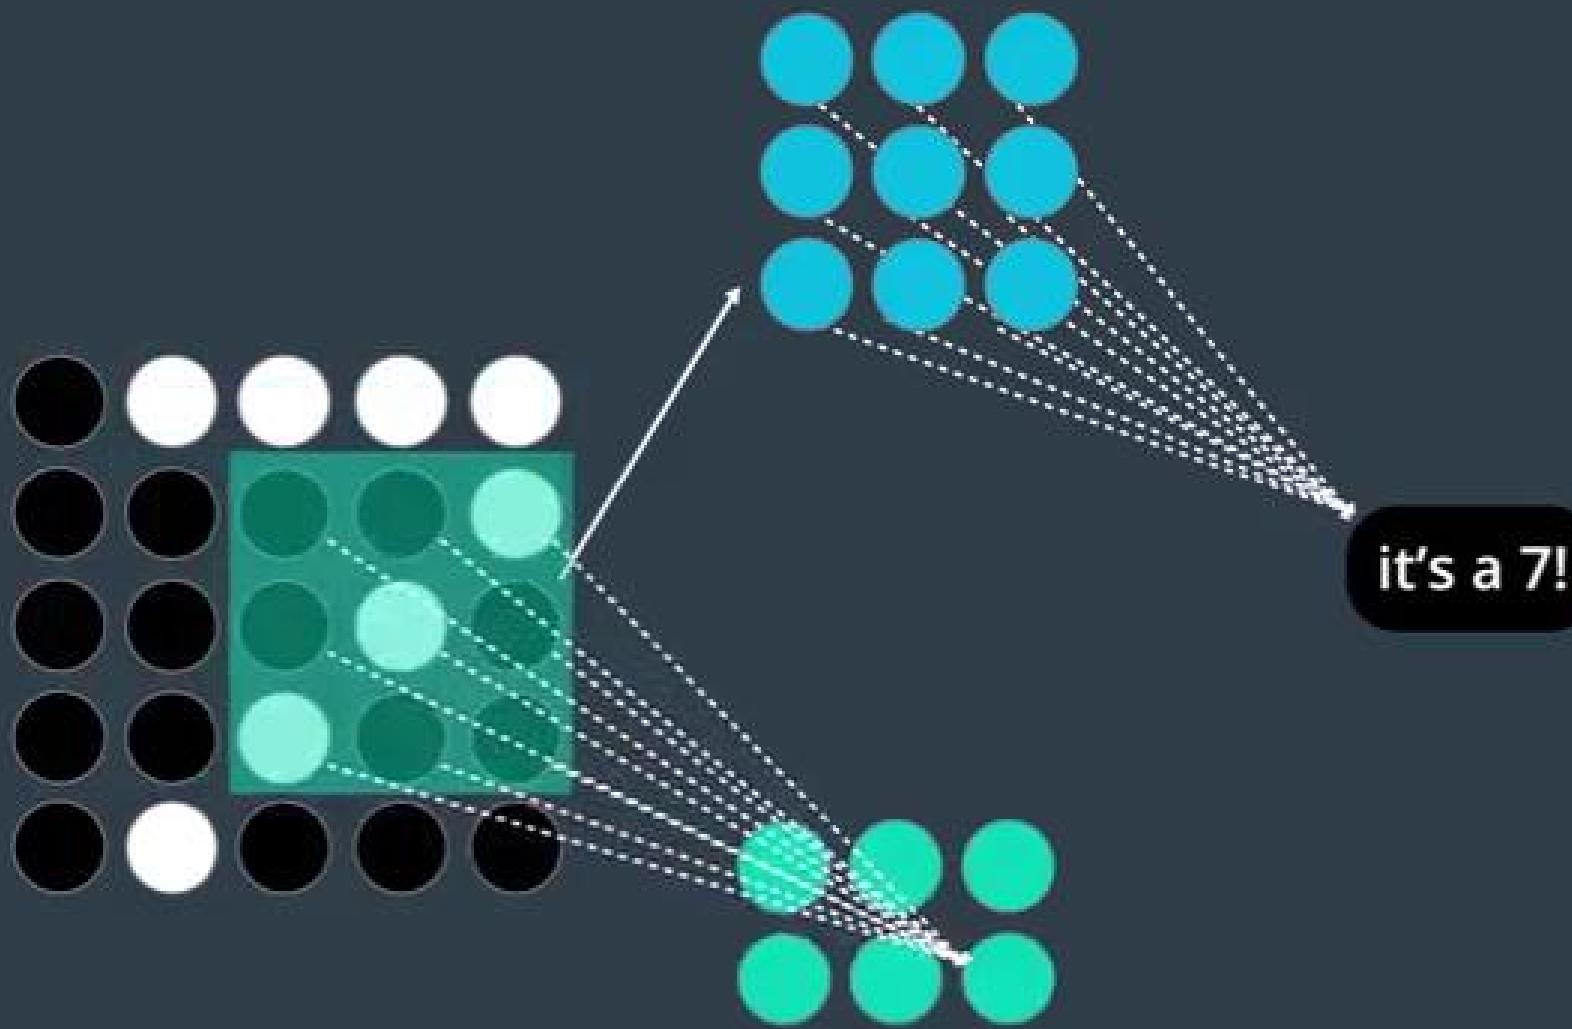
6/27

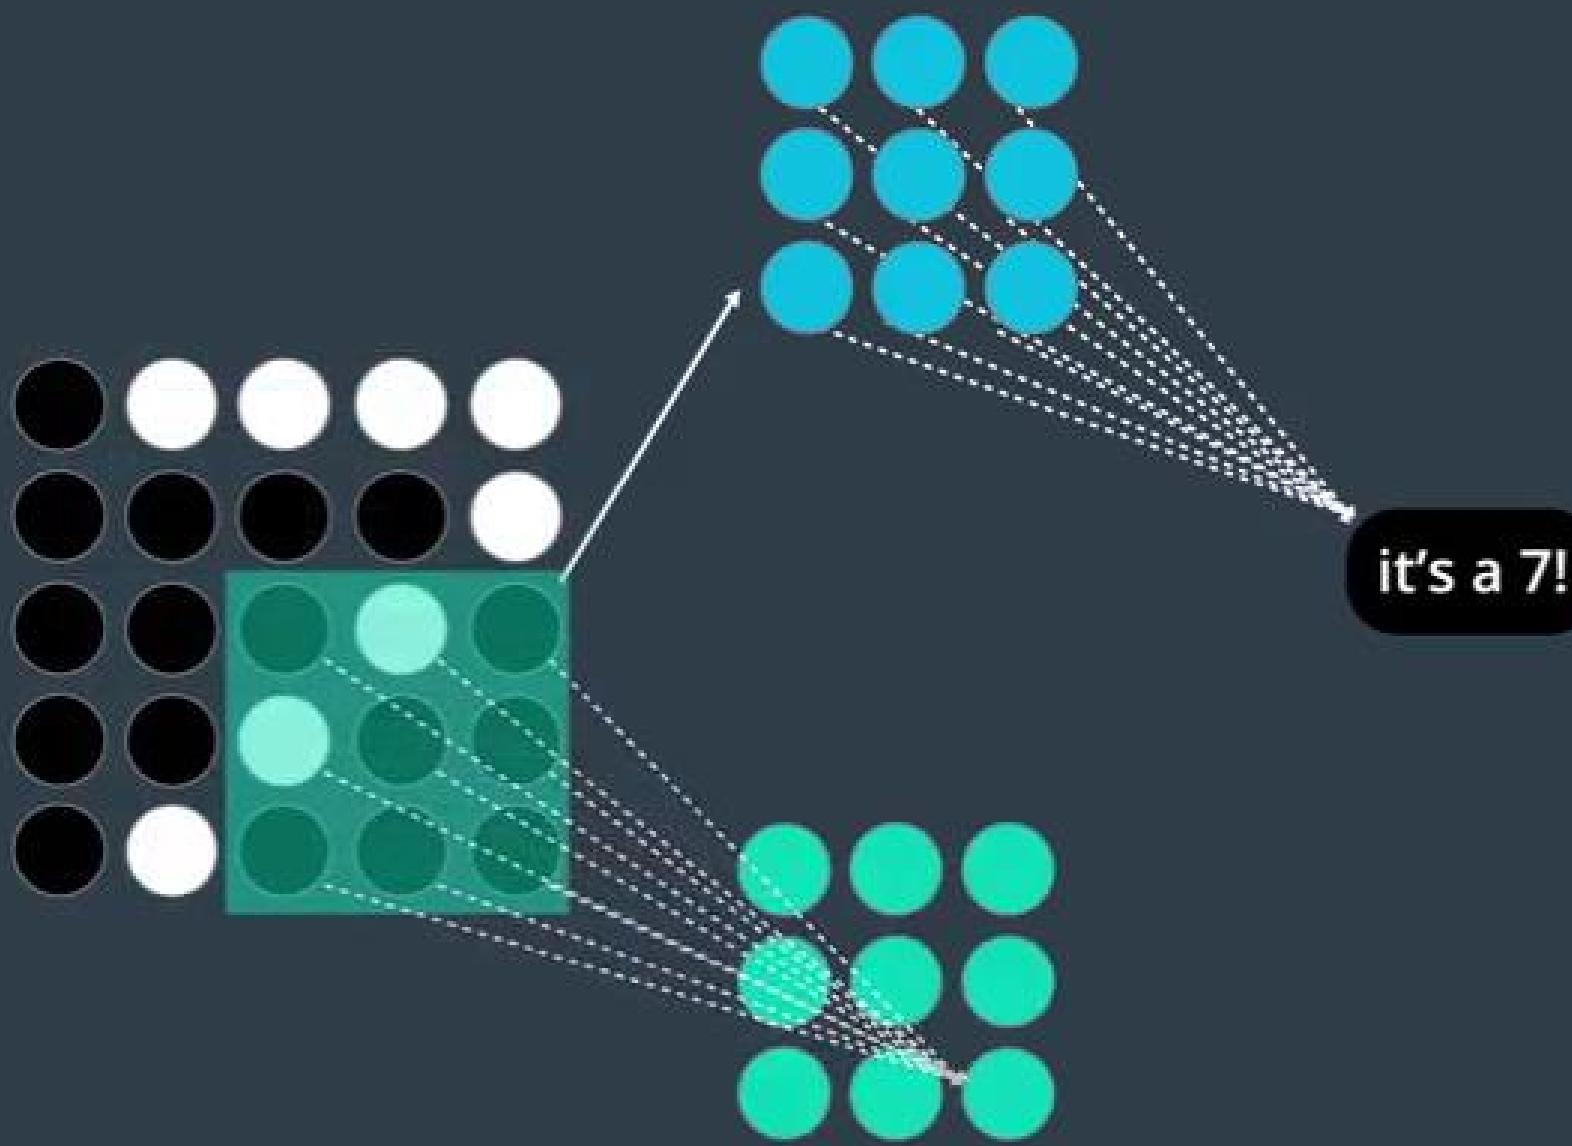


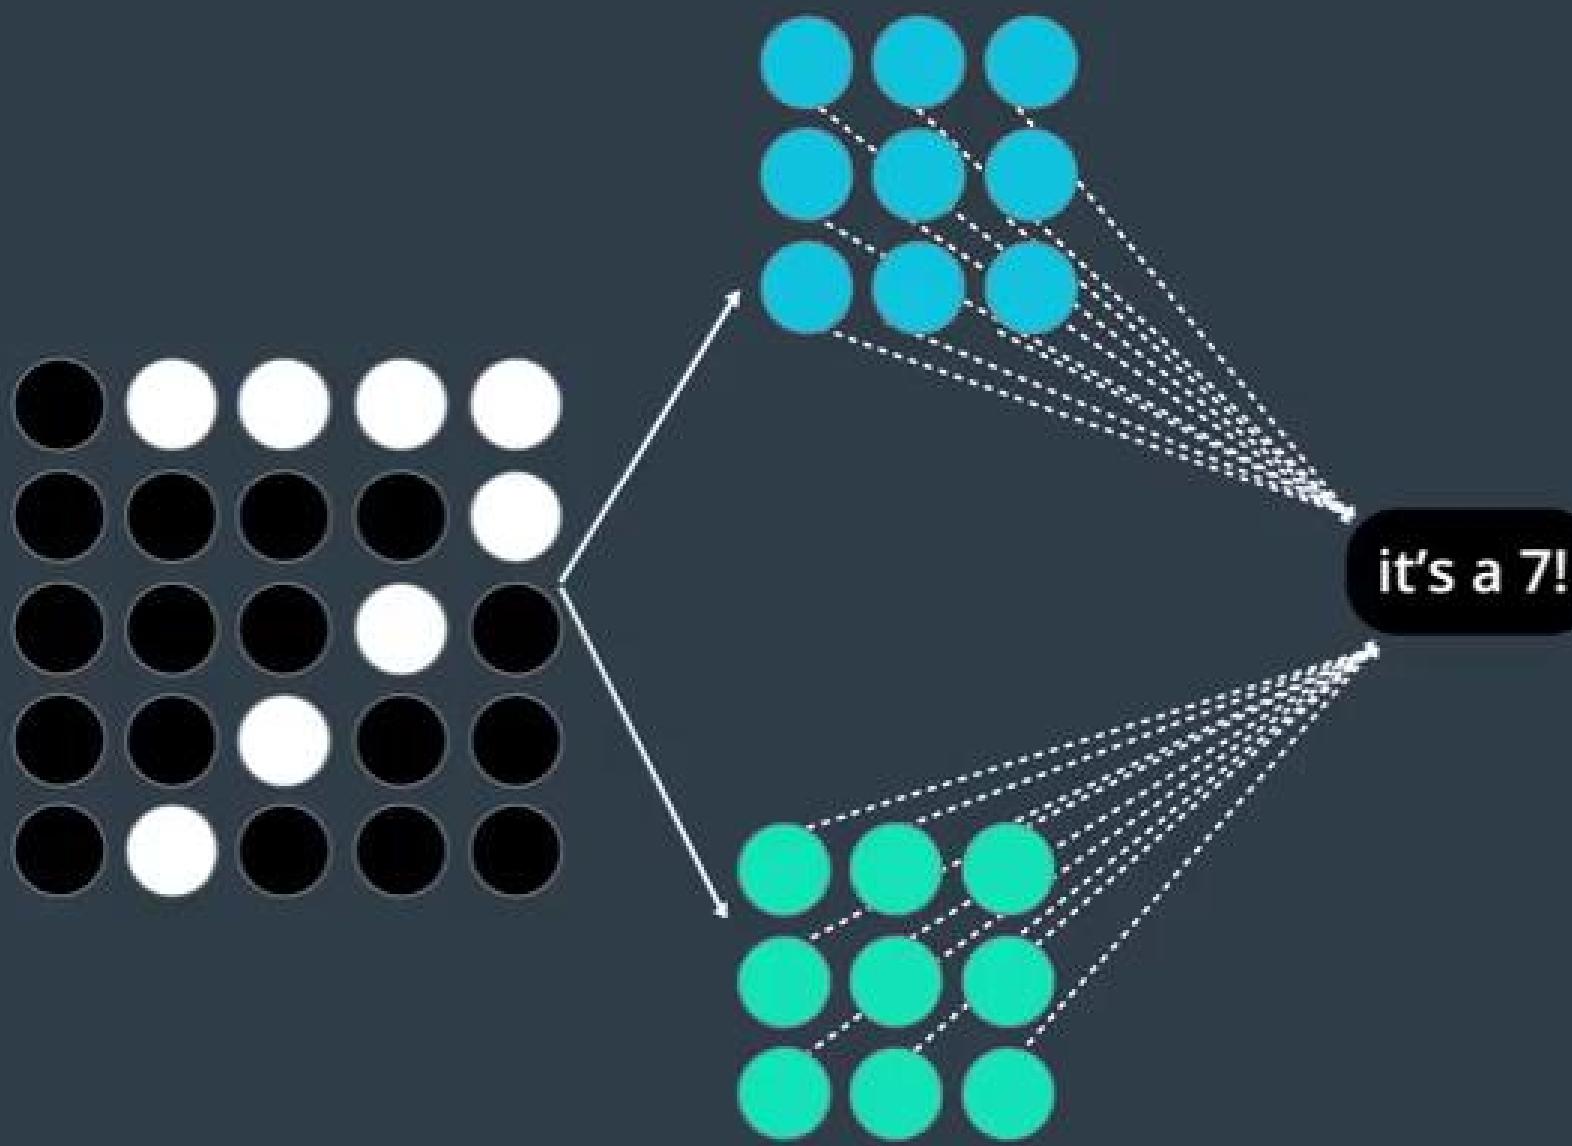




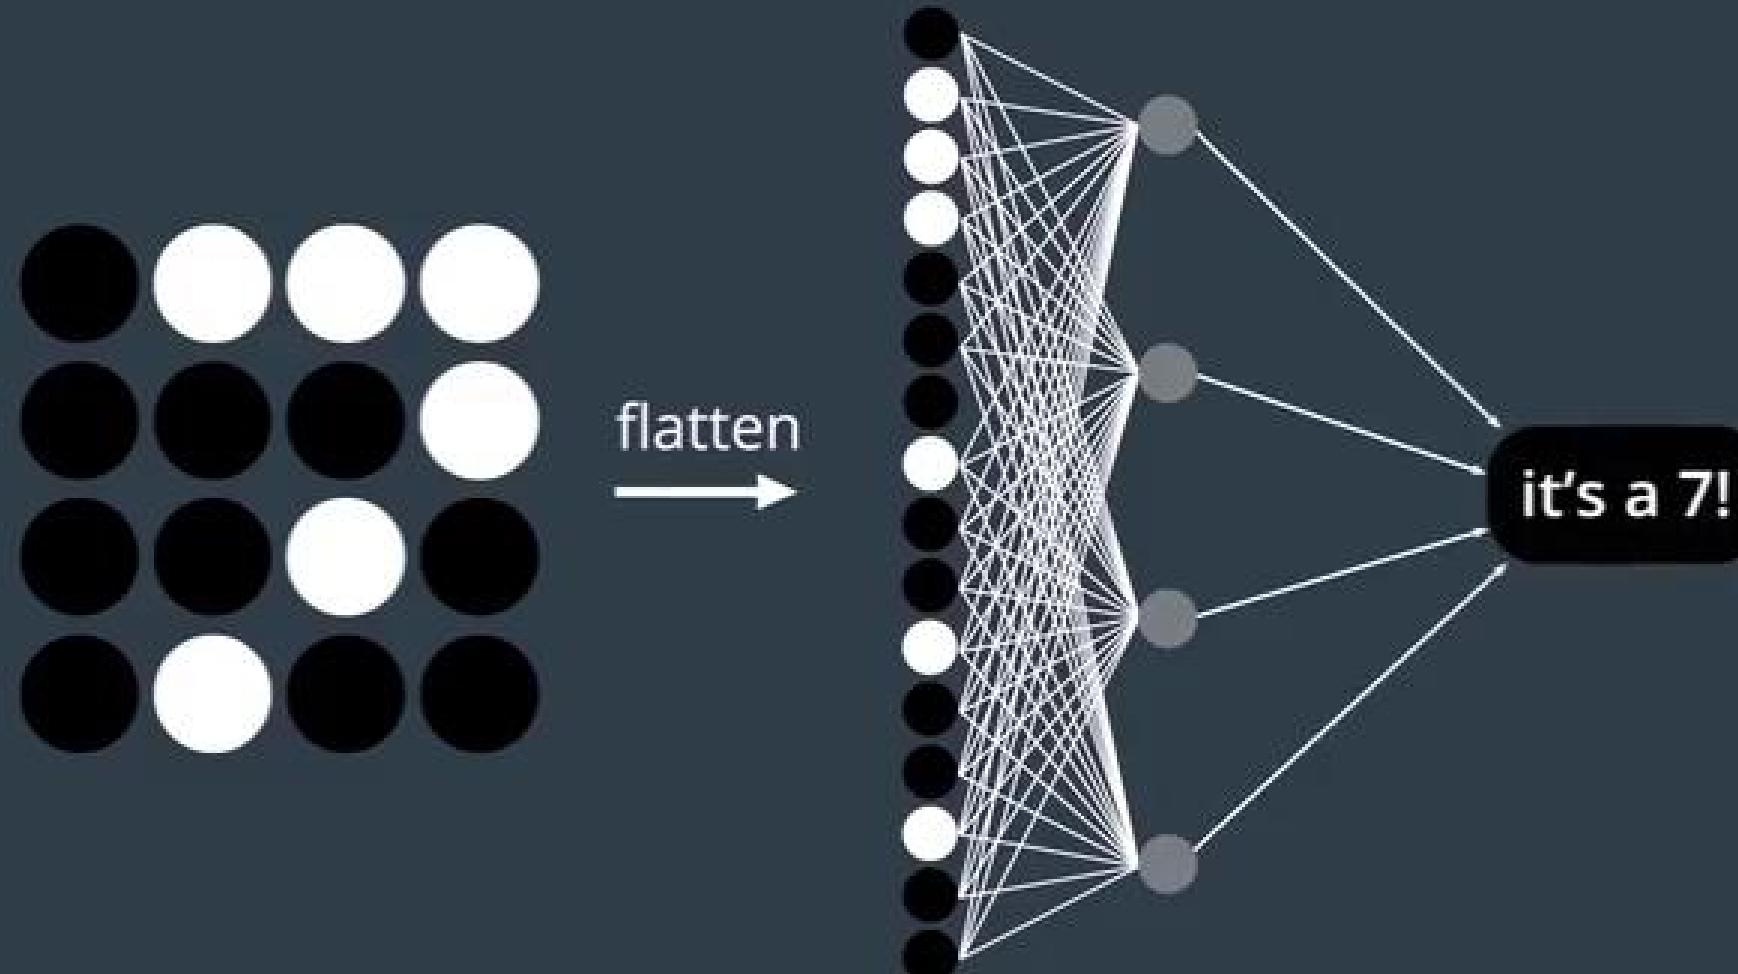




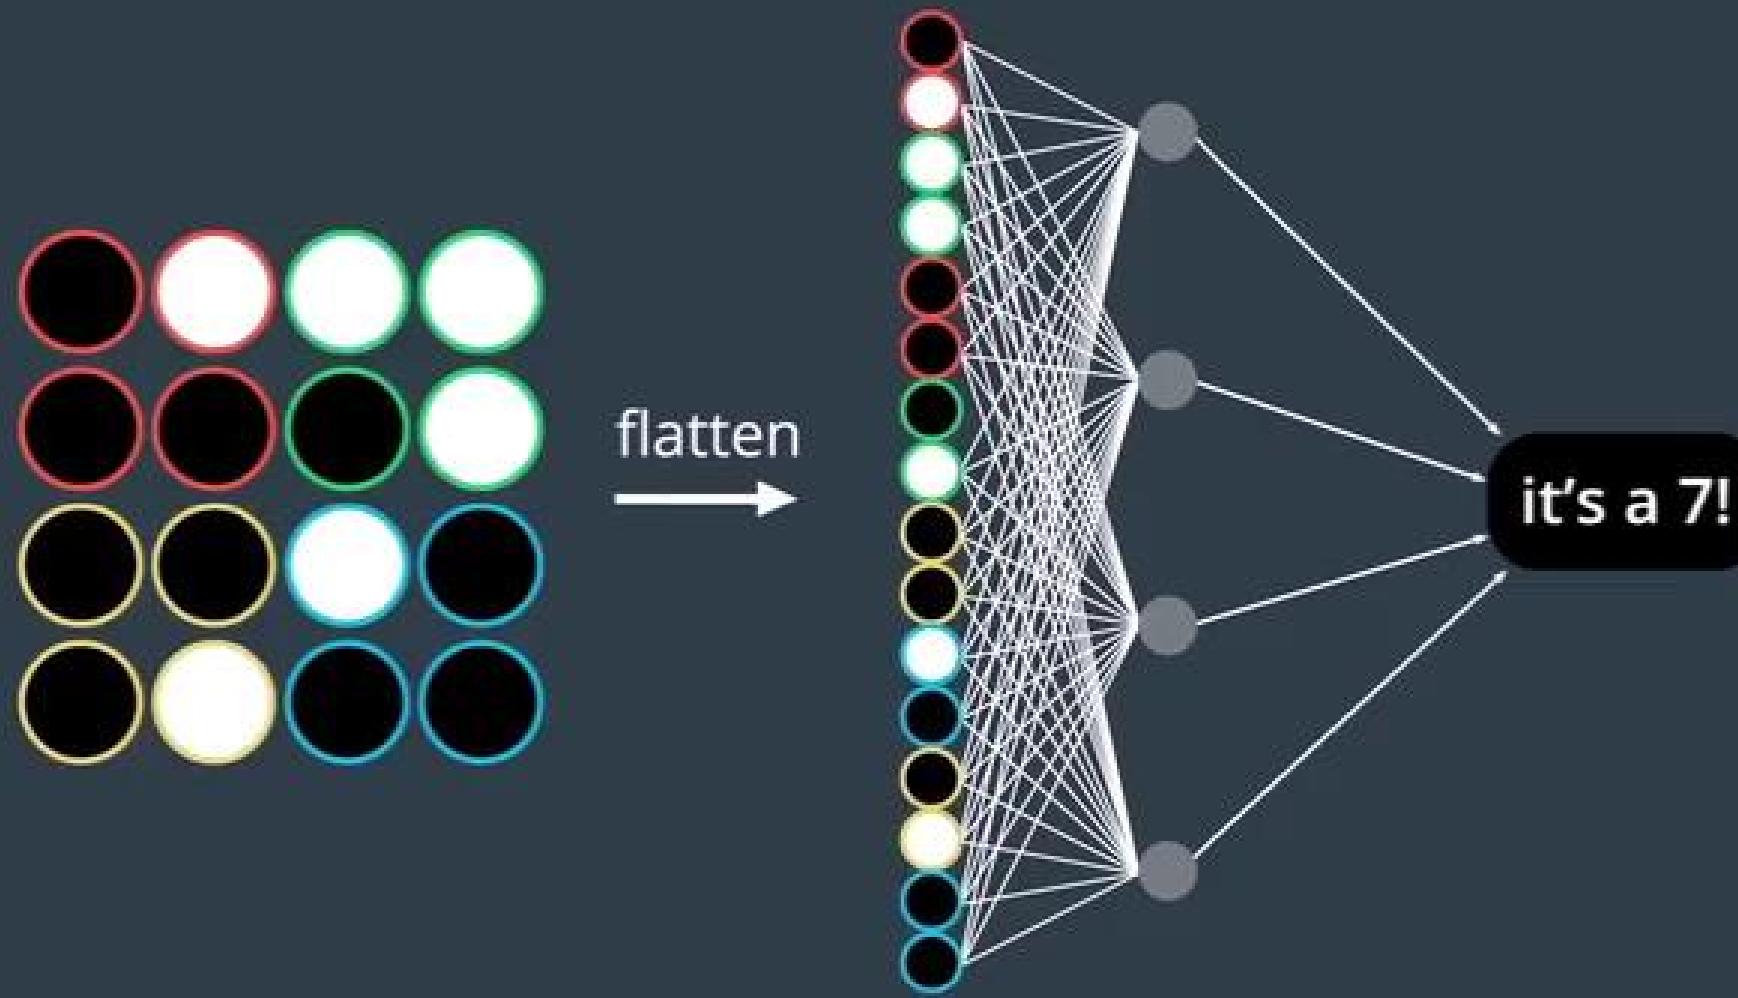




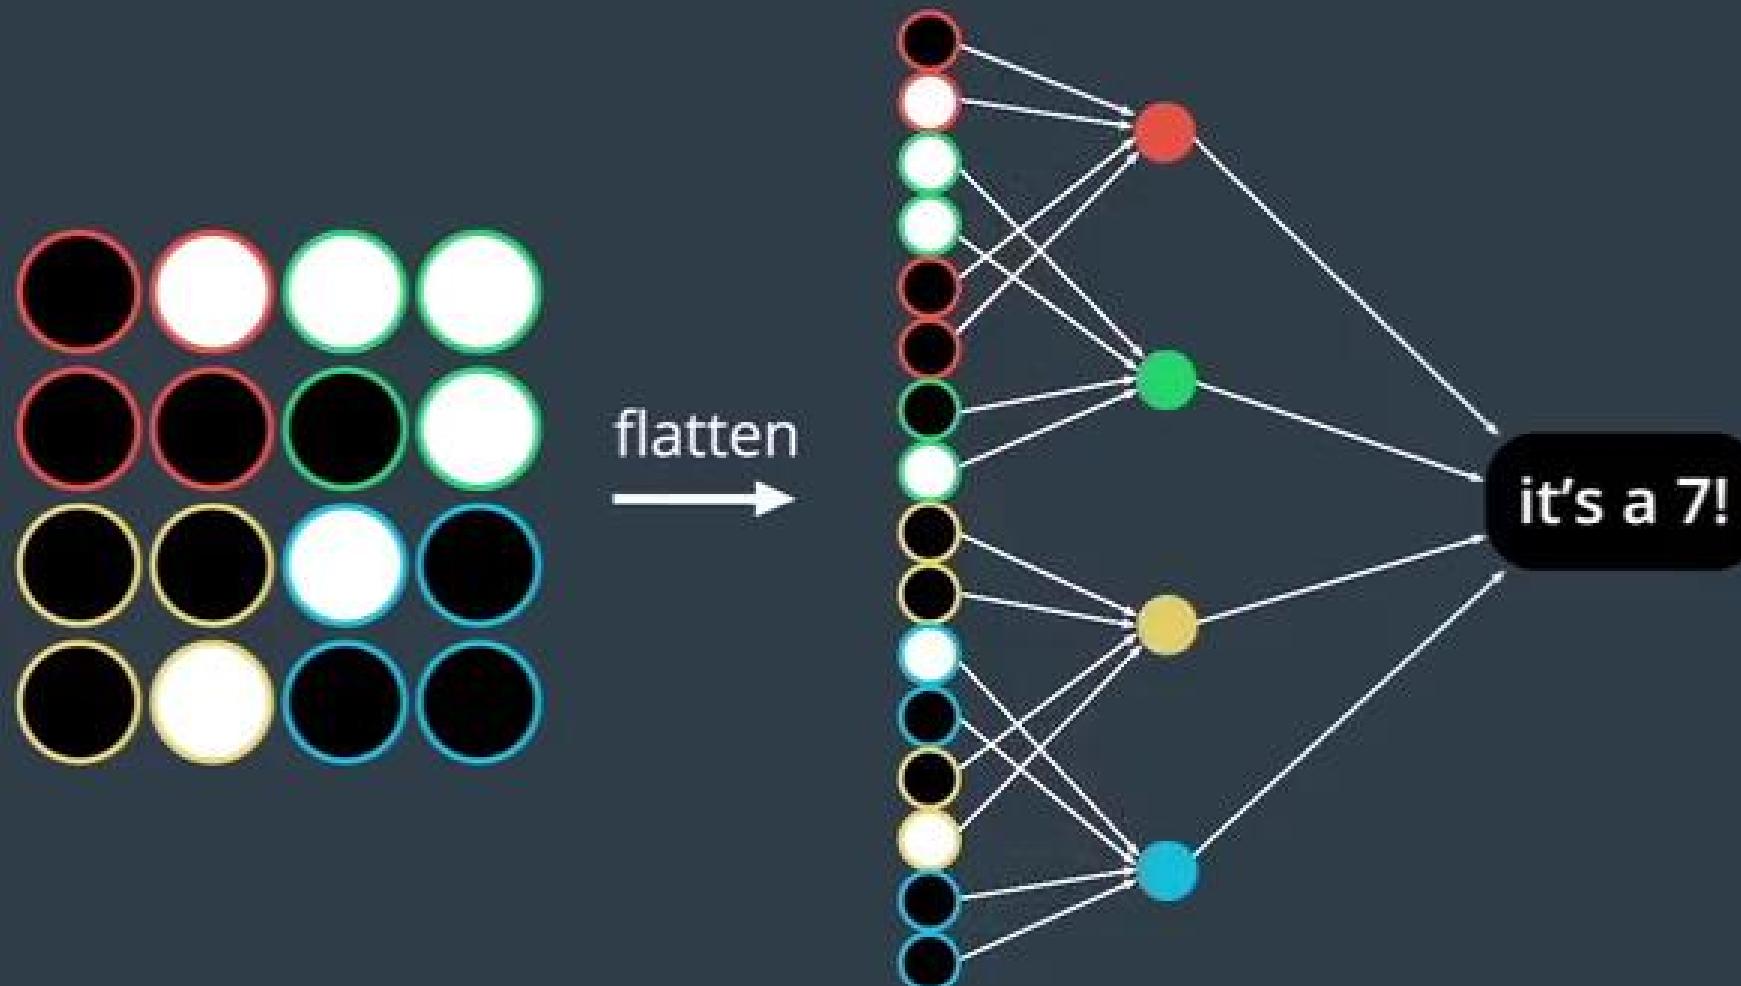
Densely Connected Layers



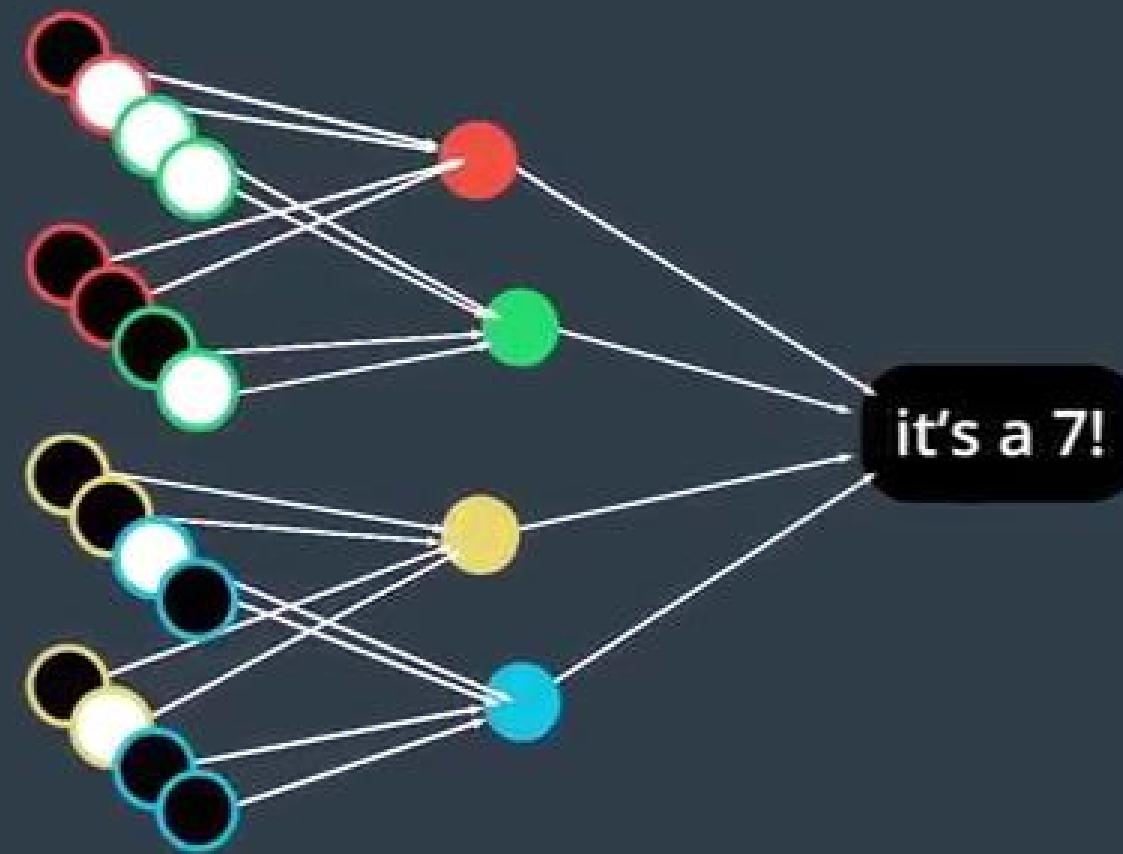
Densely Connected Layers



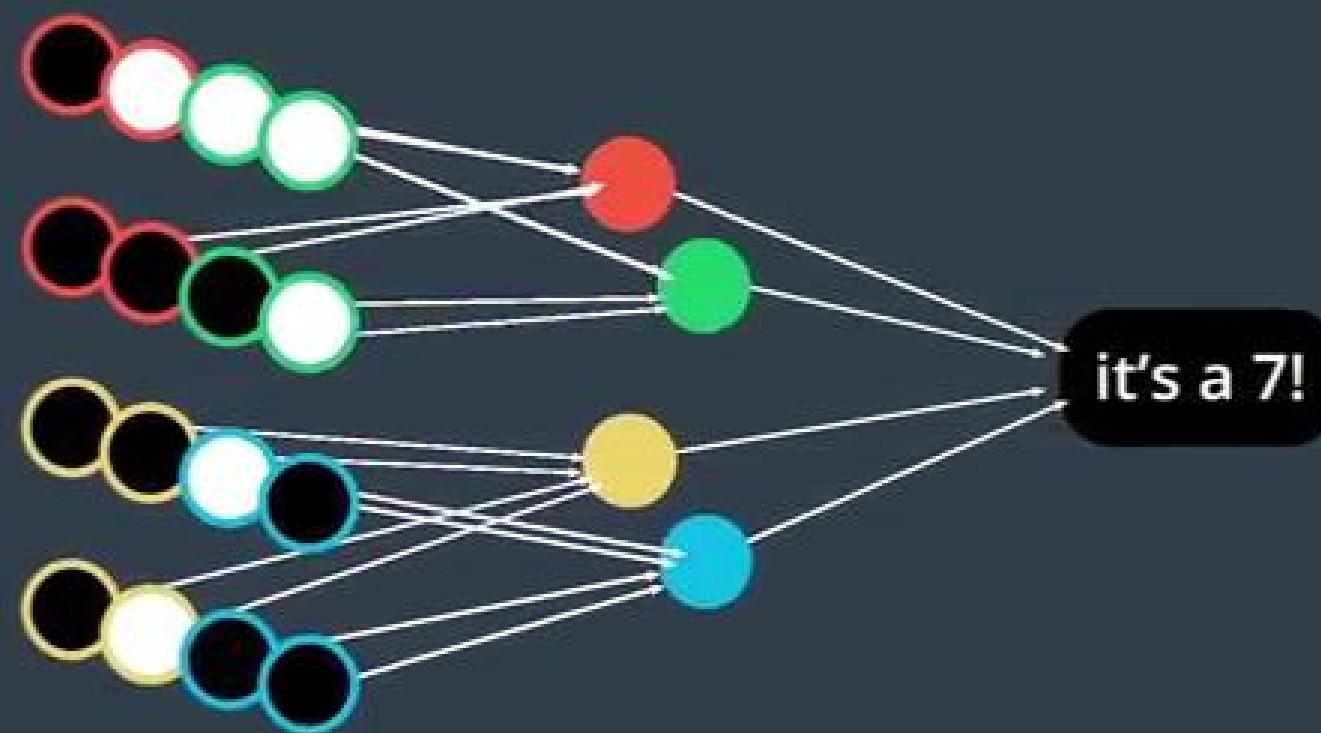
Locally Connected Layers



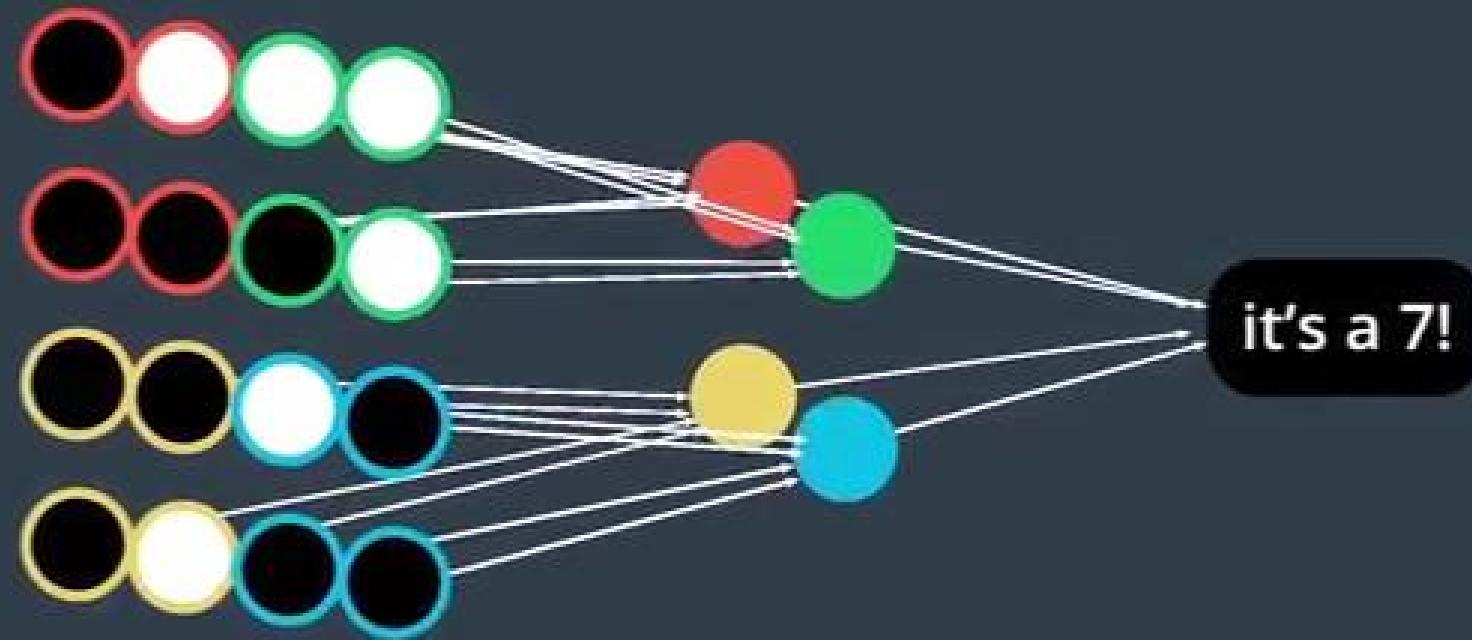
Locally Connected Layers



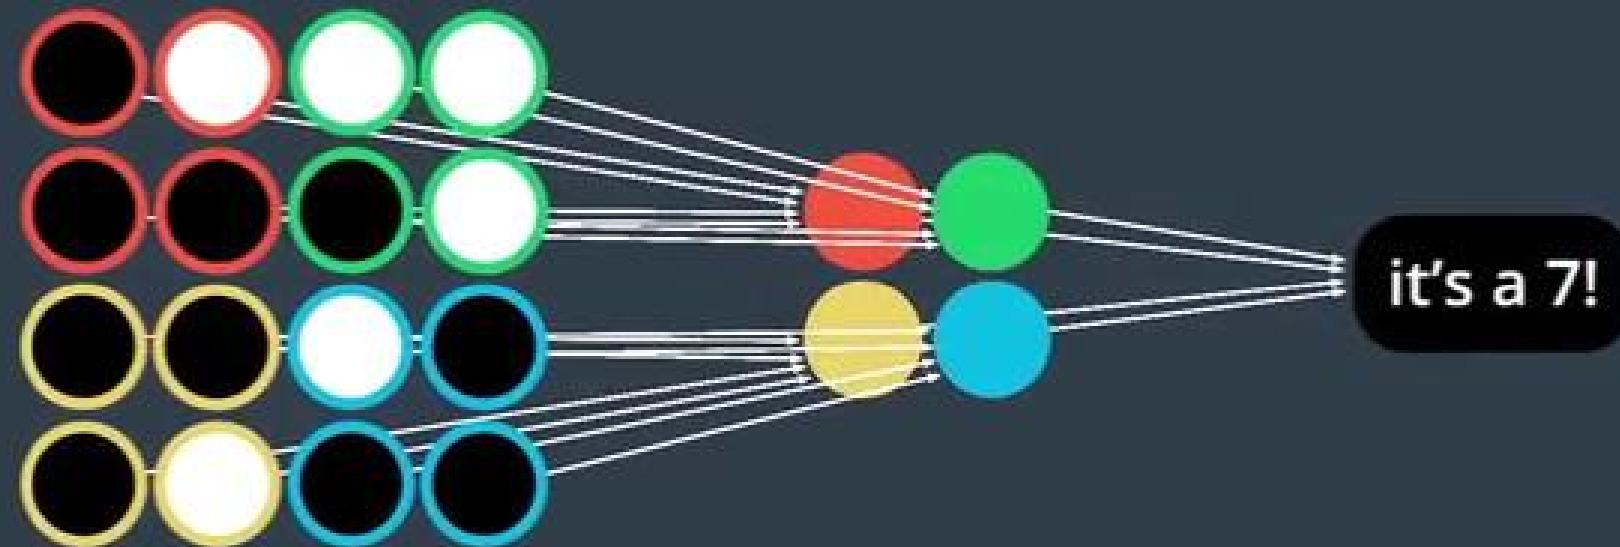
Locally Connected Layers



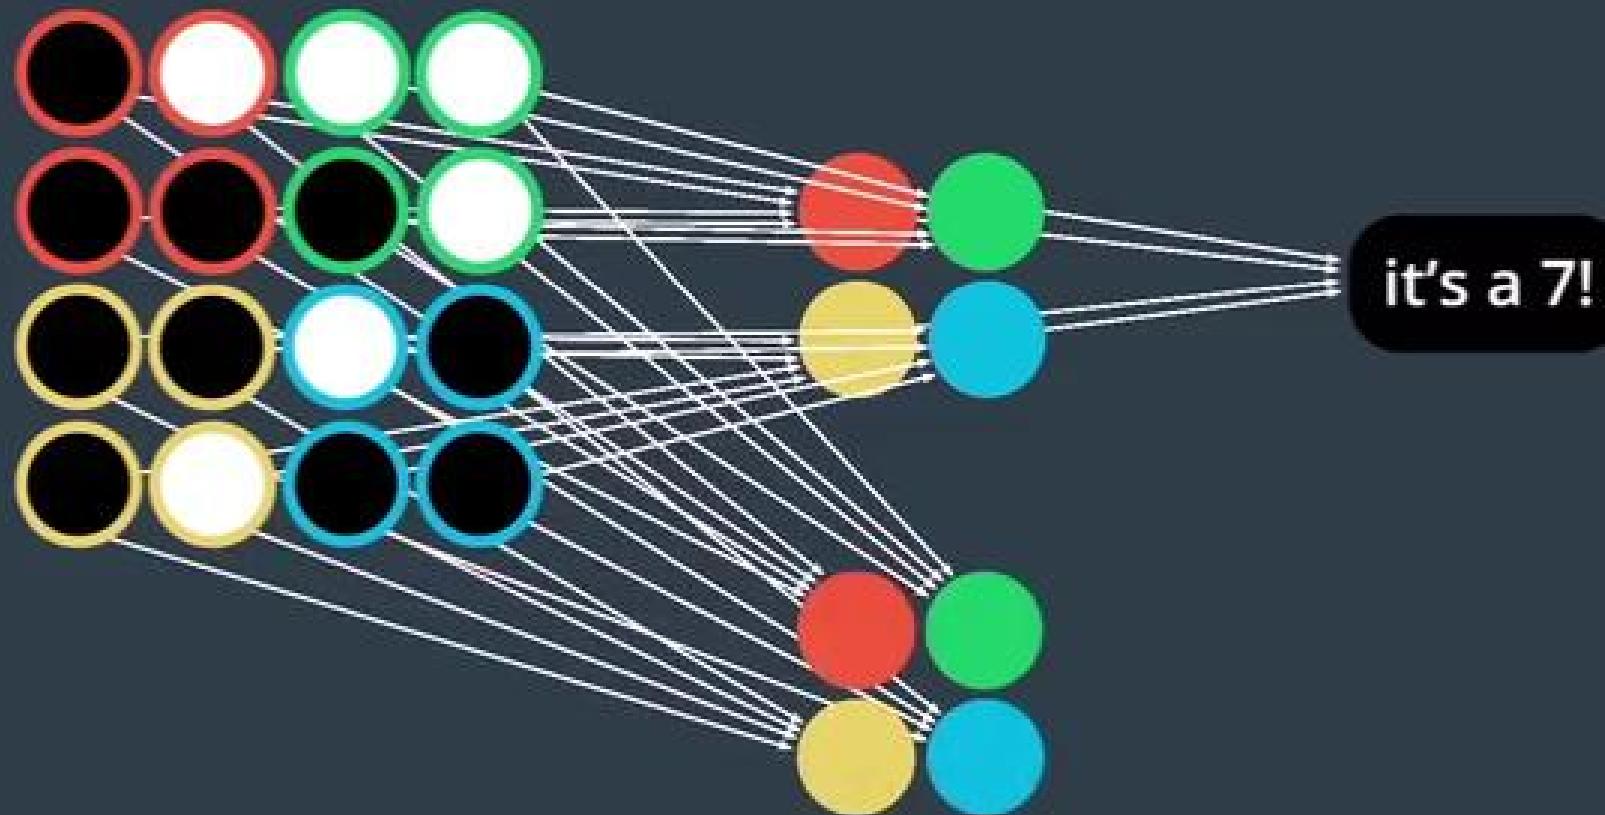
Locally Connected Layers



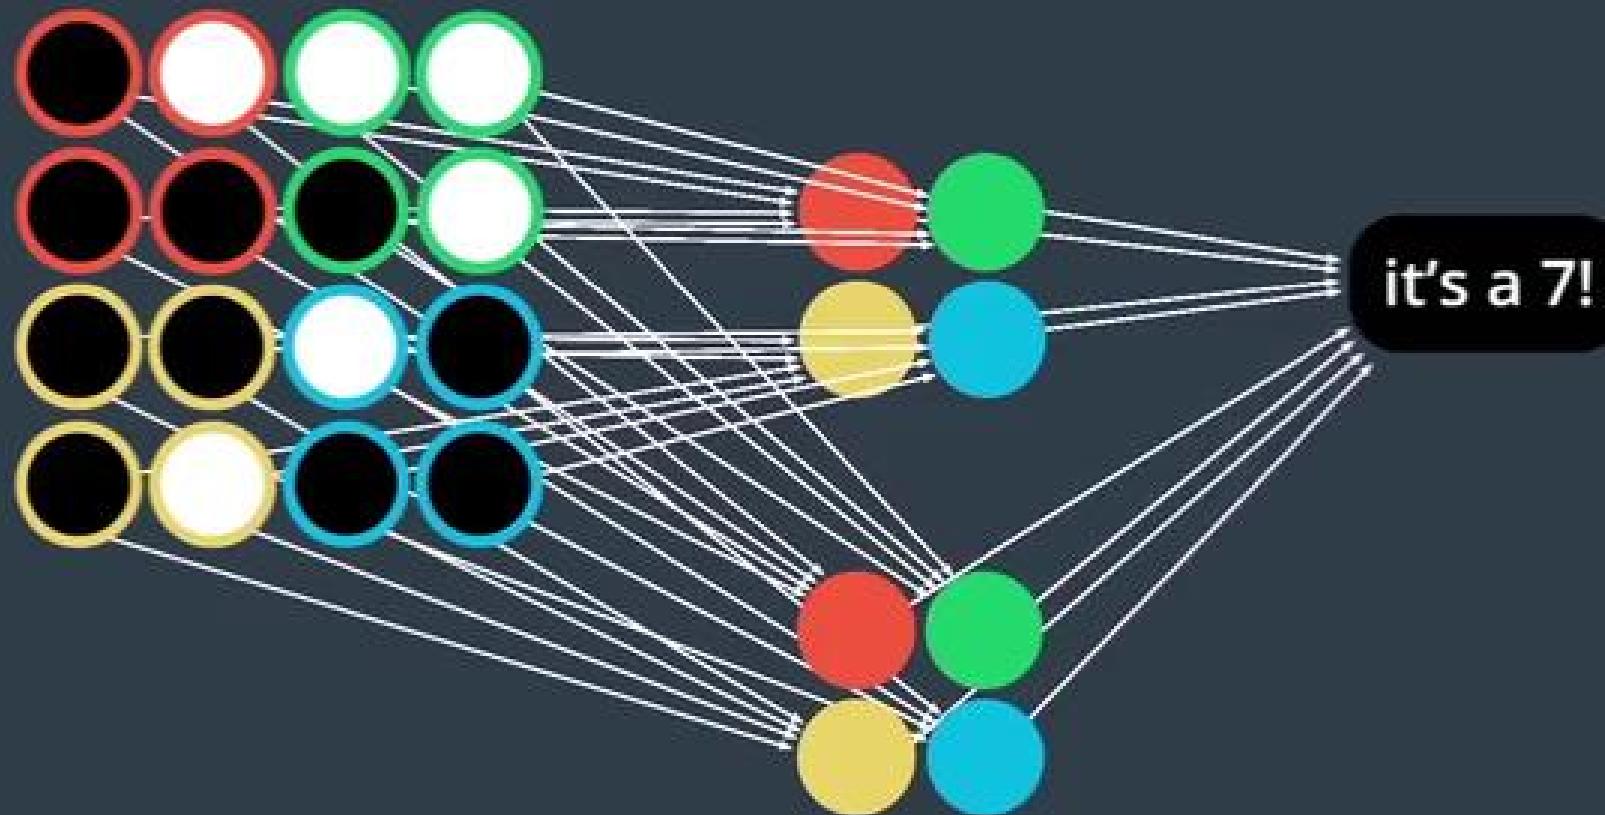
Locally Connected Layers

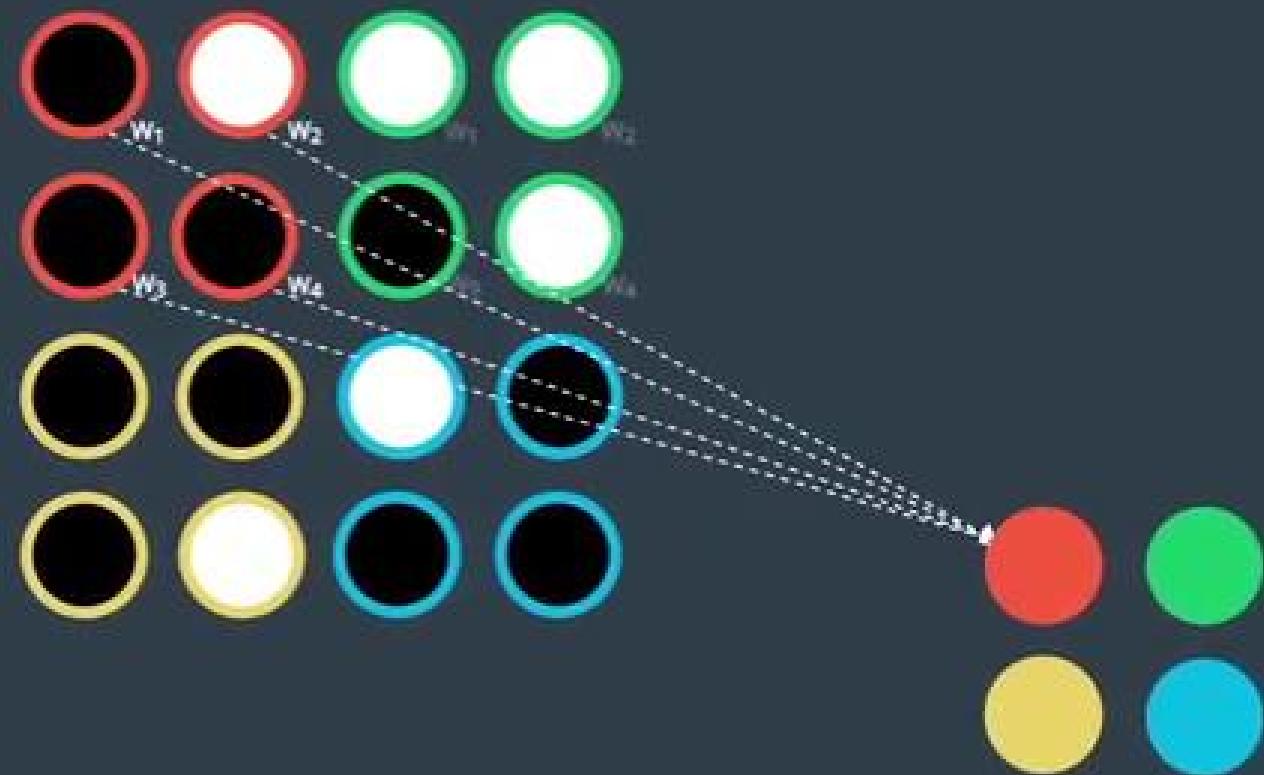


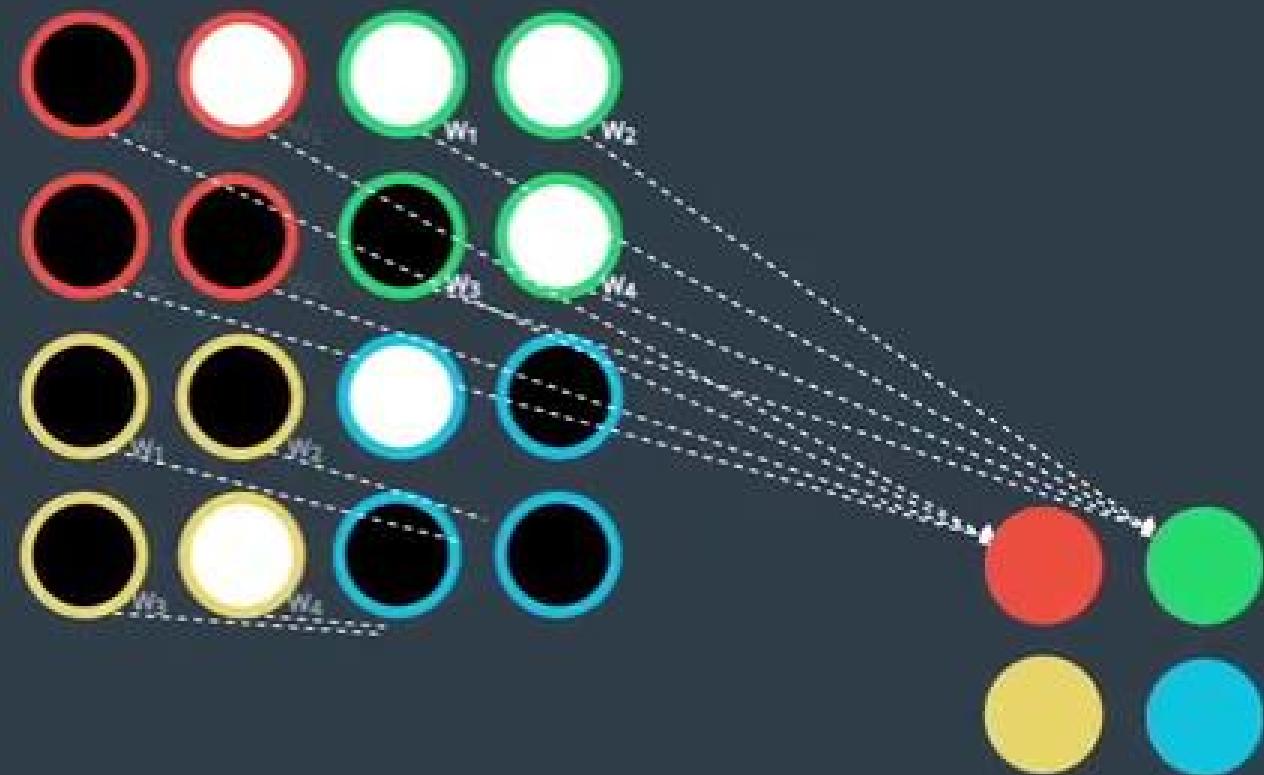
Locally Connected Layers

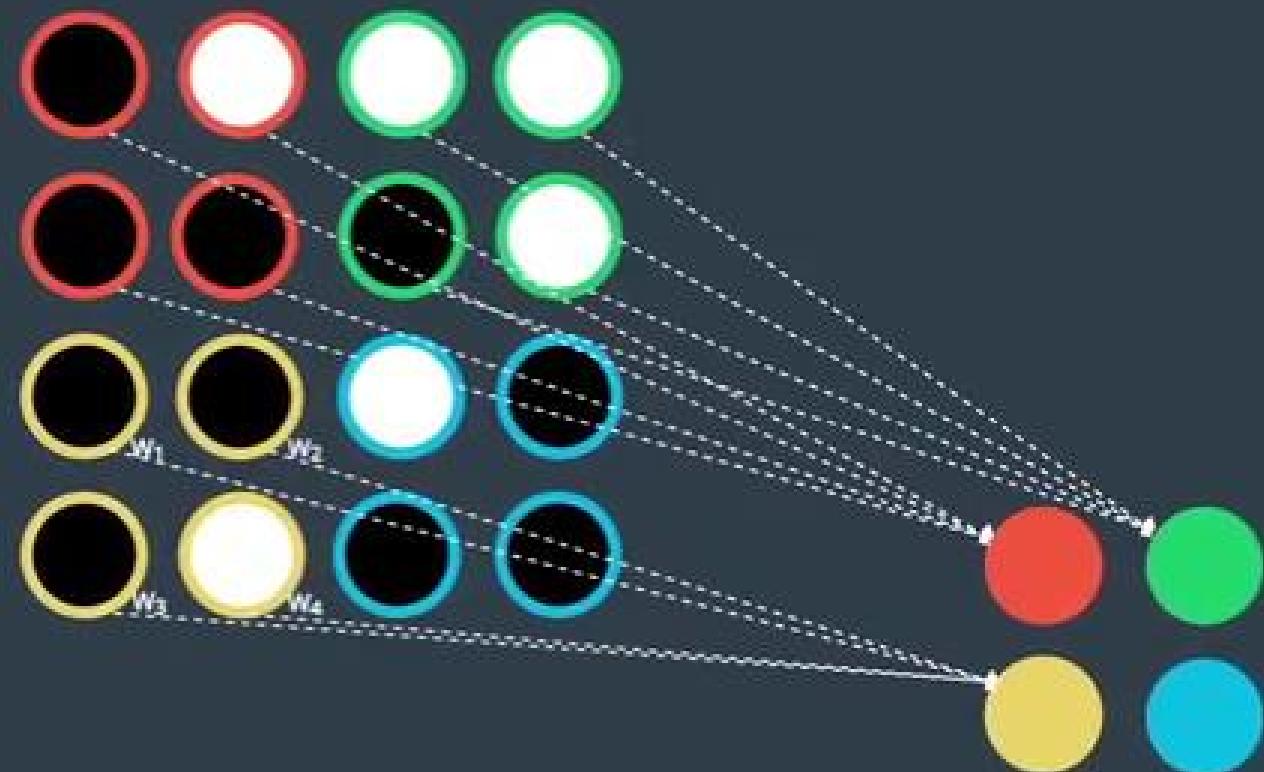


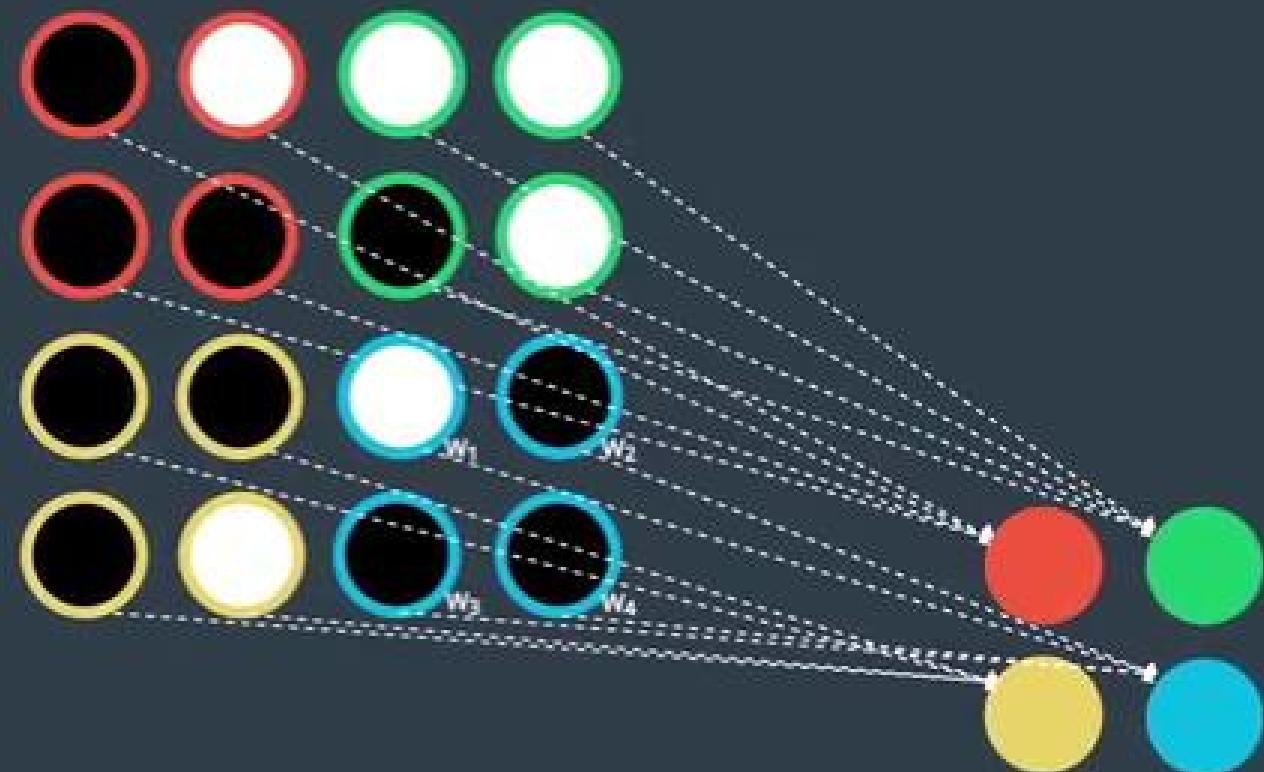
Locally Connected Layers



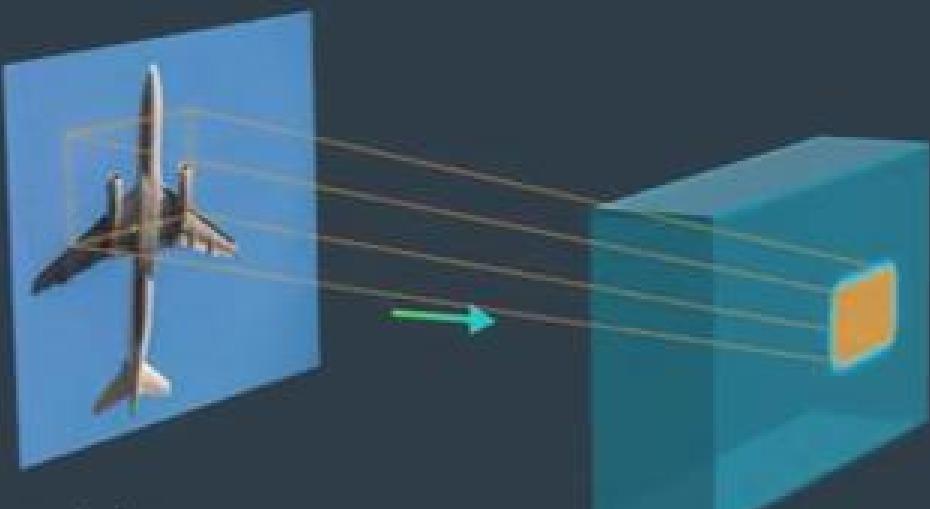








Convolutional Kernels

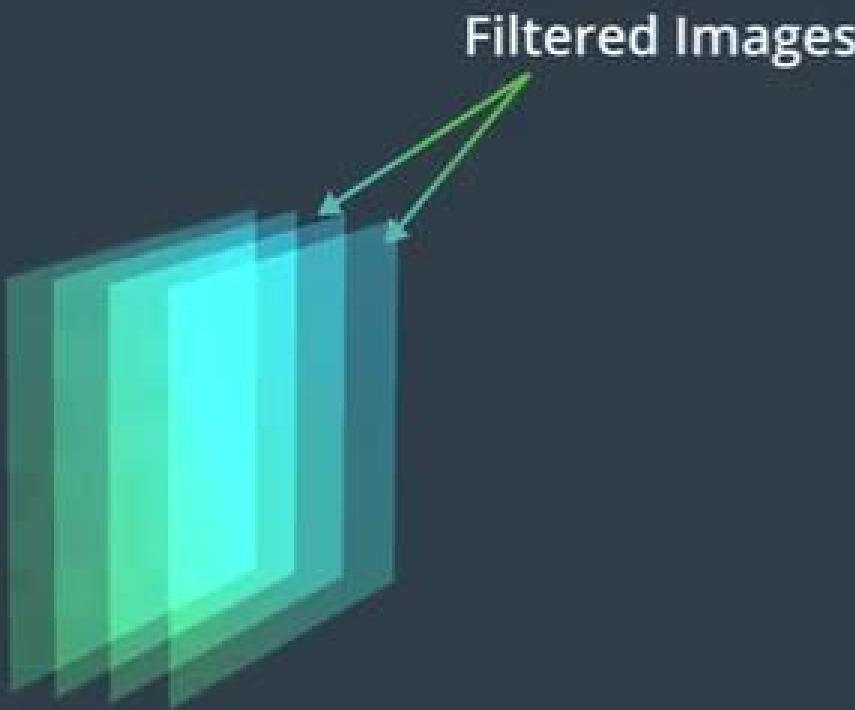


Input Image

Convolutional Layer



Input Image

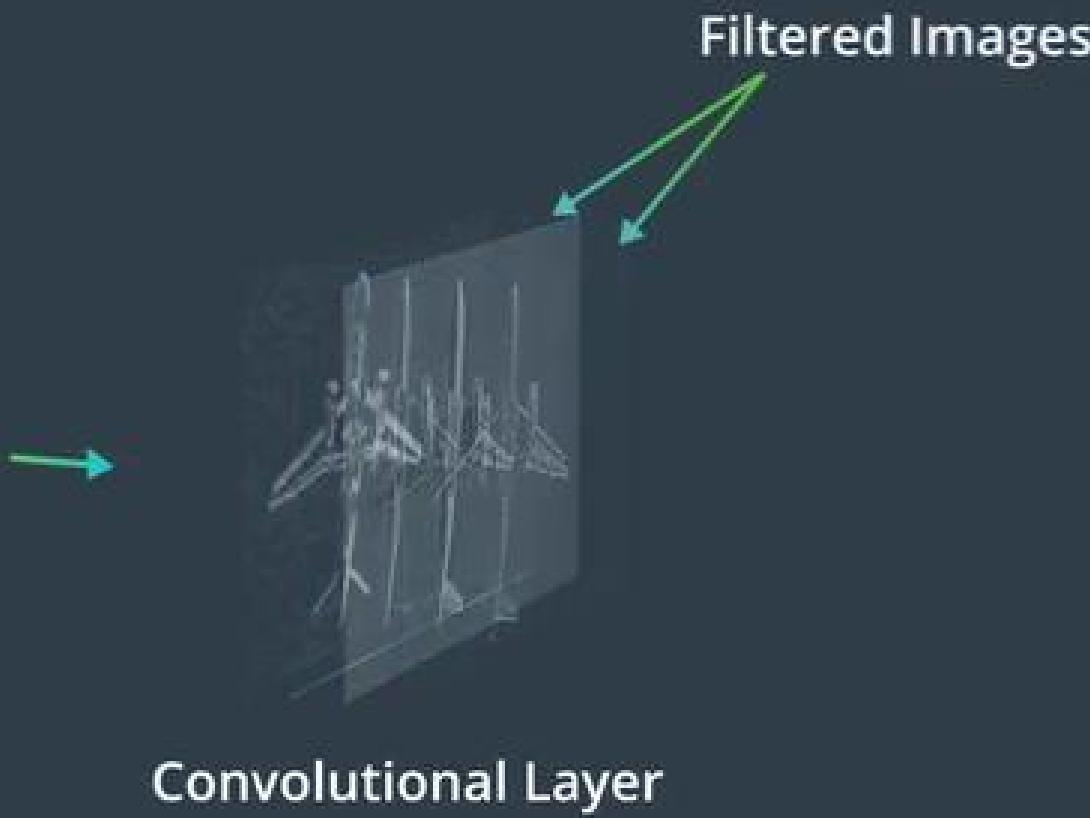


Convolutional Layer

Filtered Images



Input Image

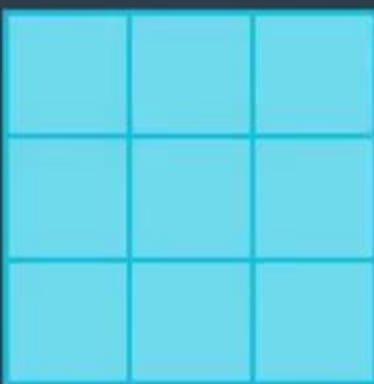


Convolutional Layer

Window

width = 3

height = 3

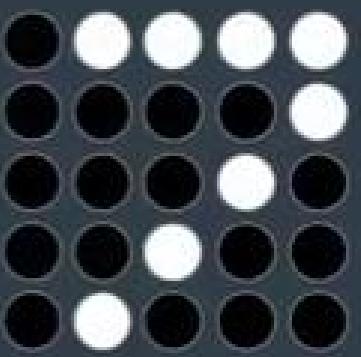


Filter

width = 3

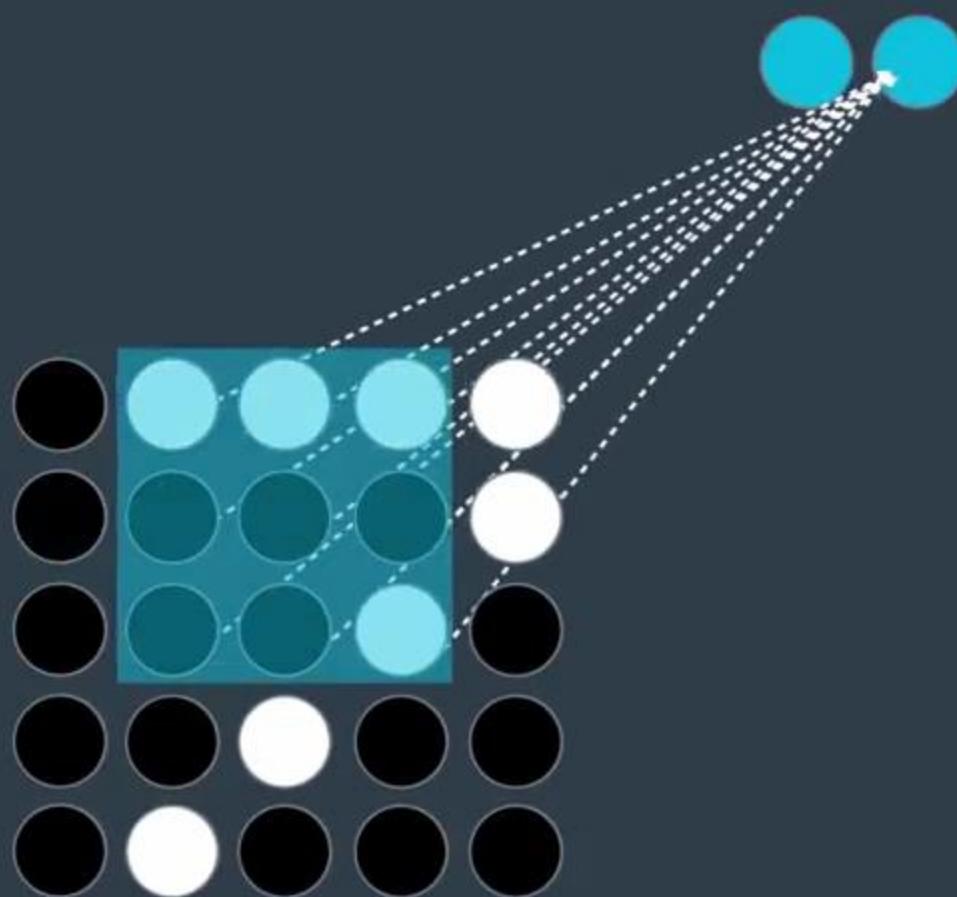
height = 3

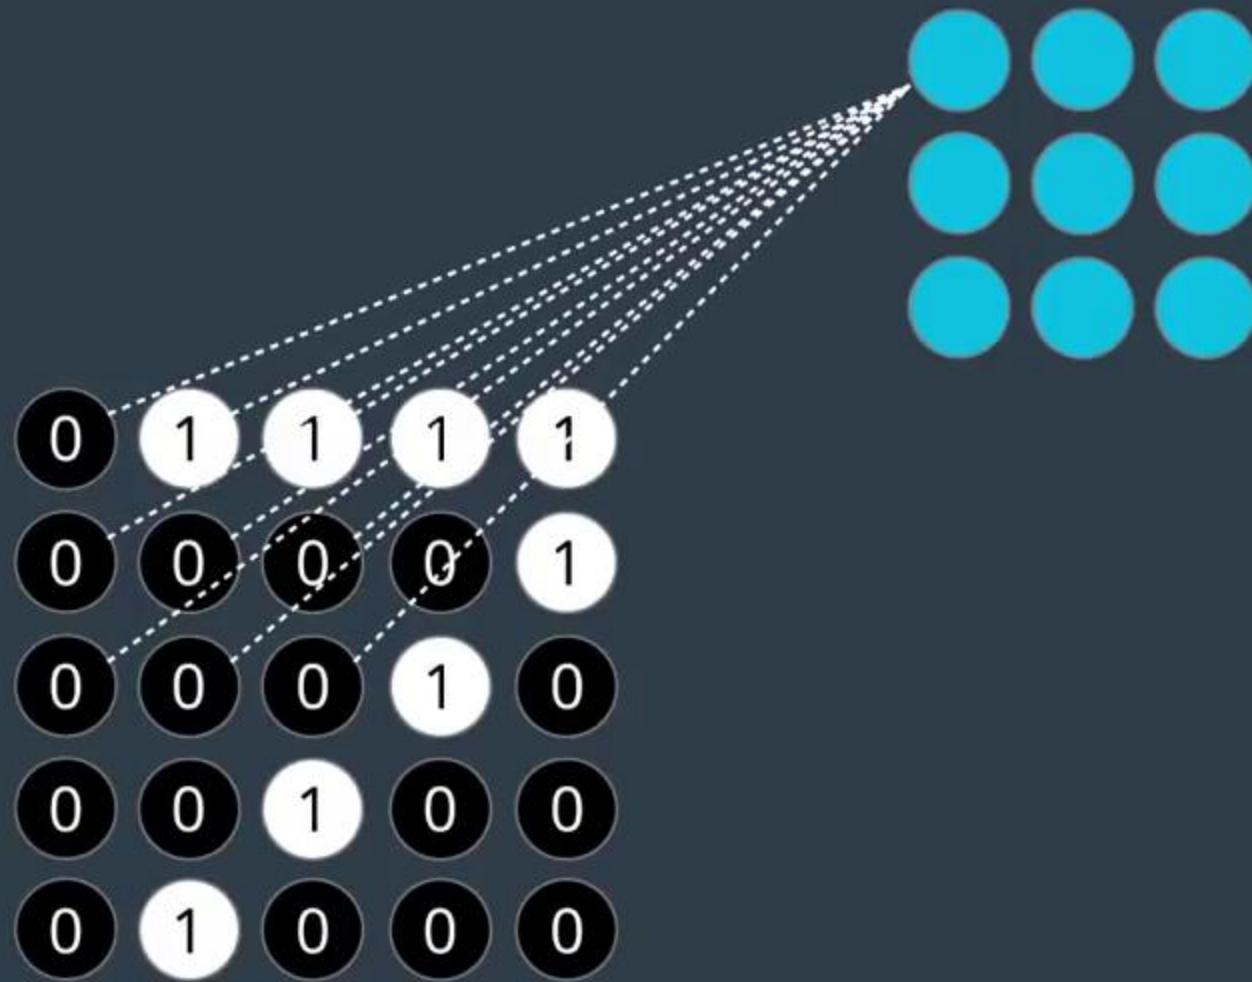
-1	-1	+1
-1	+1	-1
+1	-1	-1

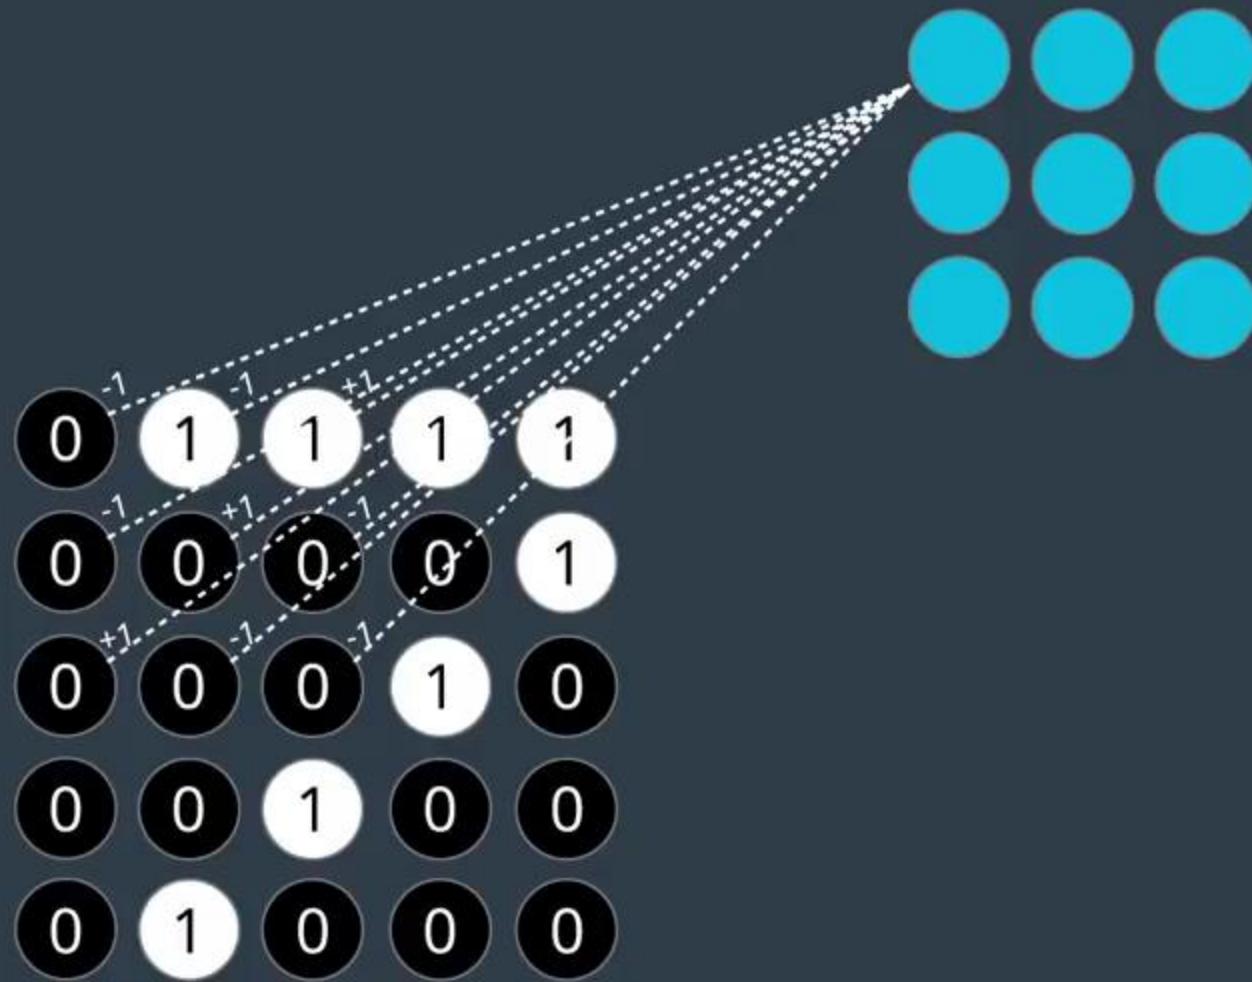


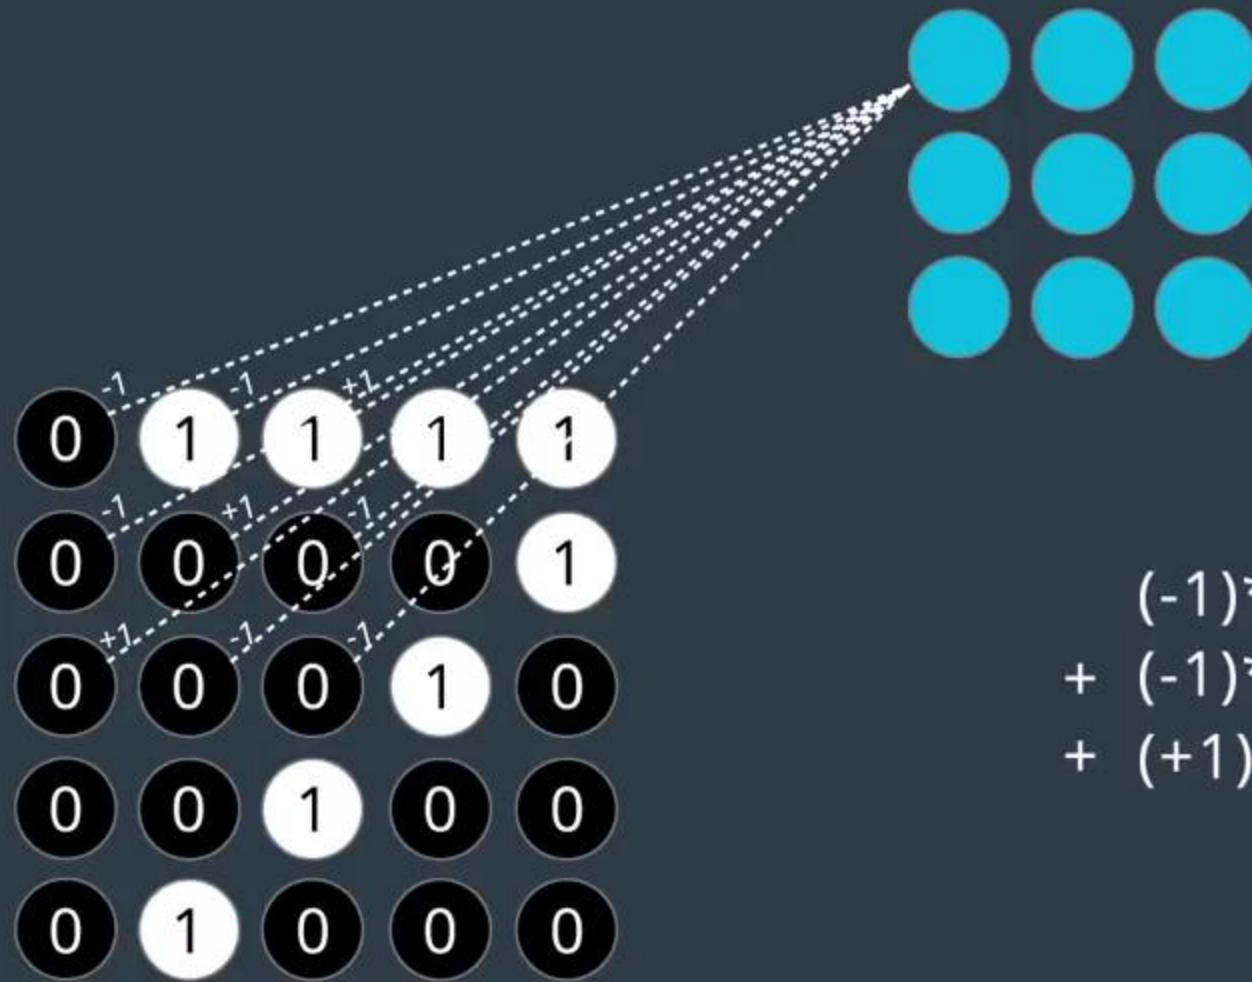
-1	-1	+1
-1	+1	-1
+1	-1	-1

Input Layer —————> Hidden Layer —> Output Layer

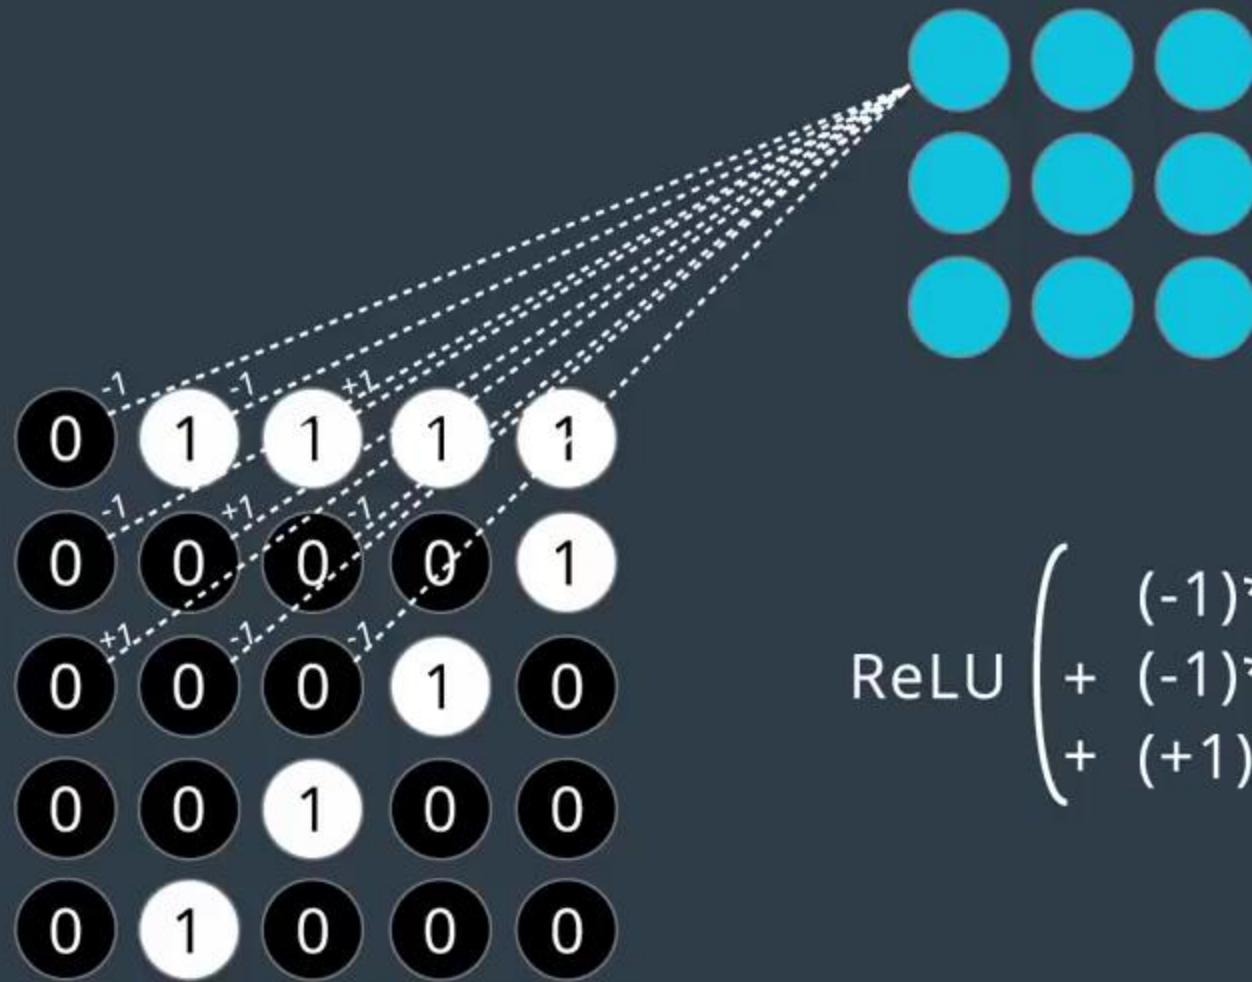




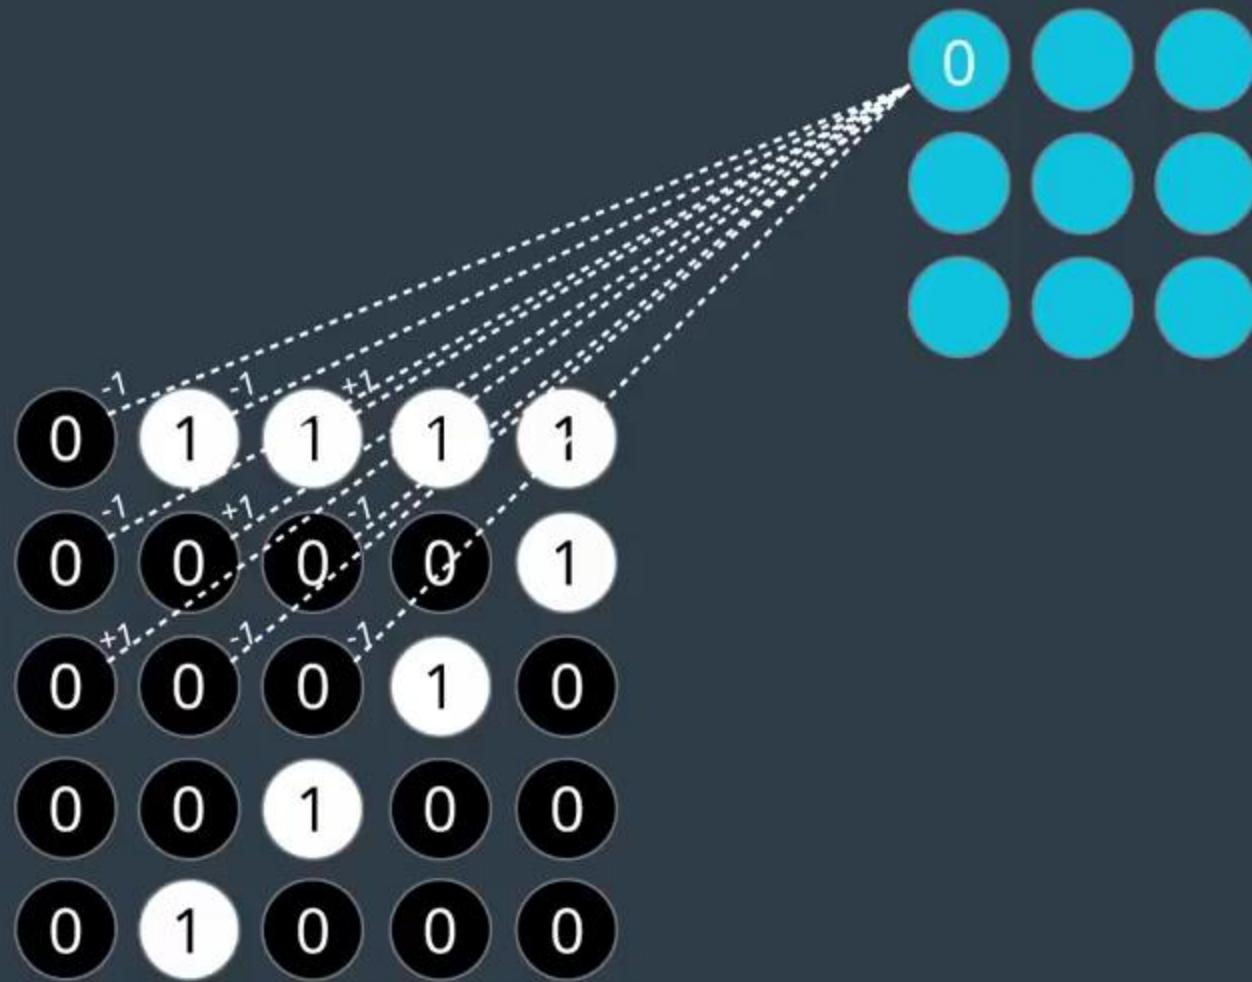


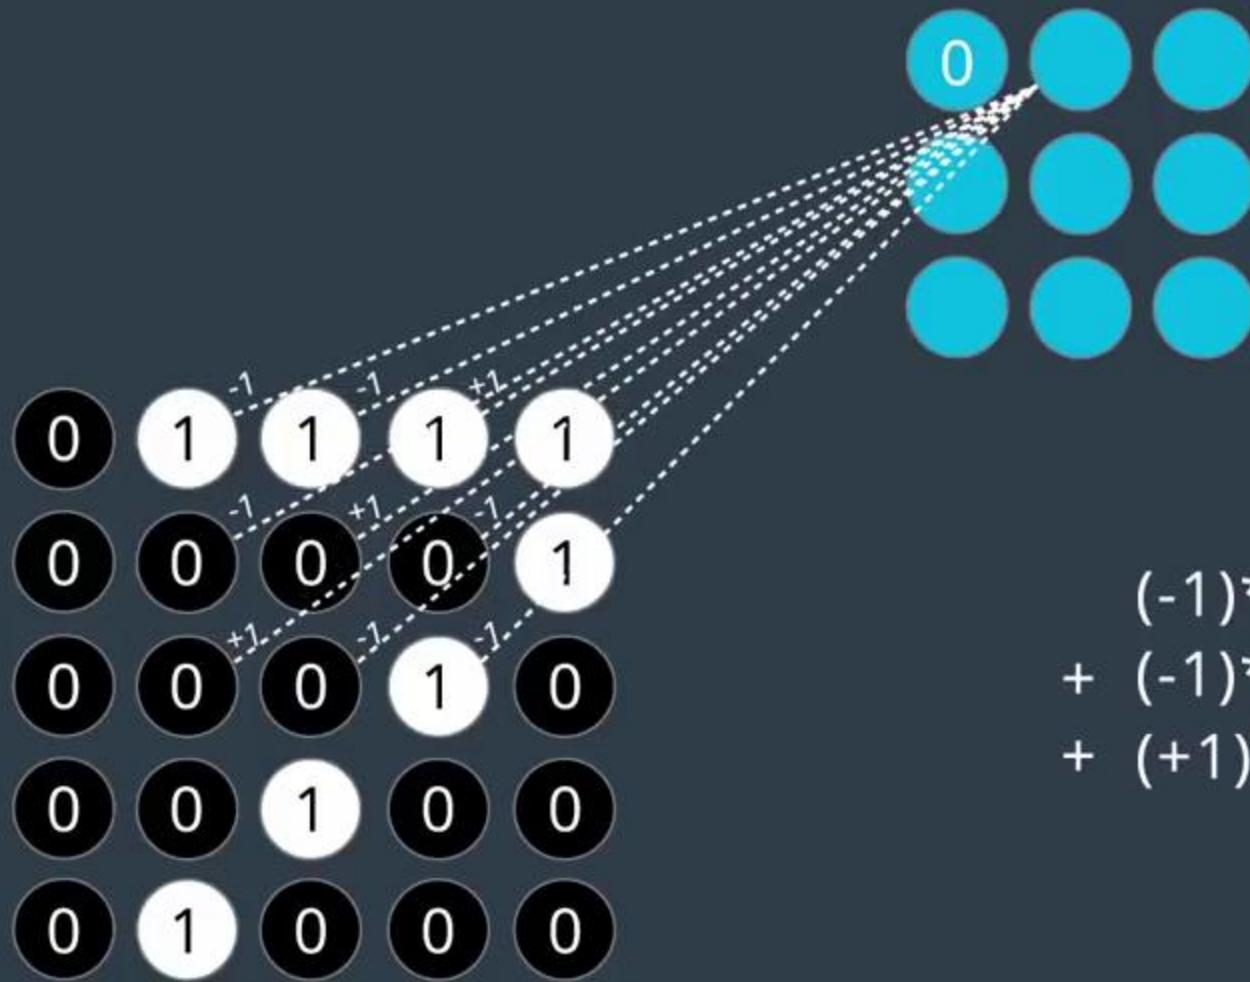


$$\begin{aligned}
 & (-1)^*0 + (-1)^*1 + (+1)^*1 \\
 & + (-1)^*0 + (+1)^*0 + (-1)^*0 = 0 \\
 & + (+1)^*0 + (-1)^*0 + (-1)^*0
 \end{aligned}$$

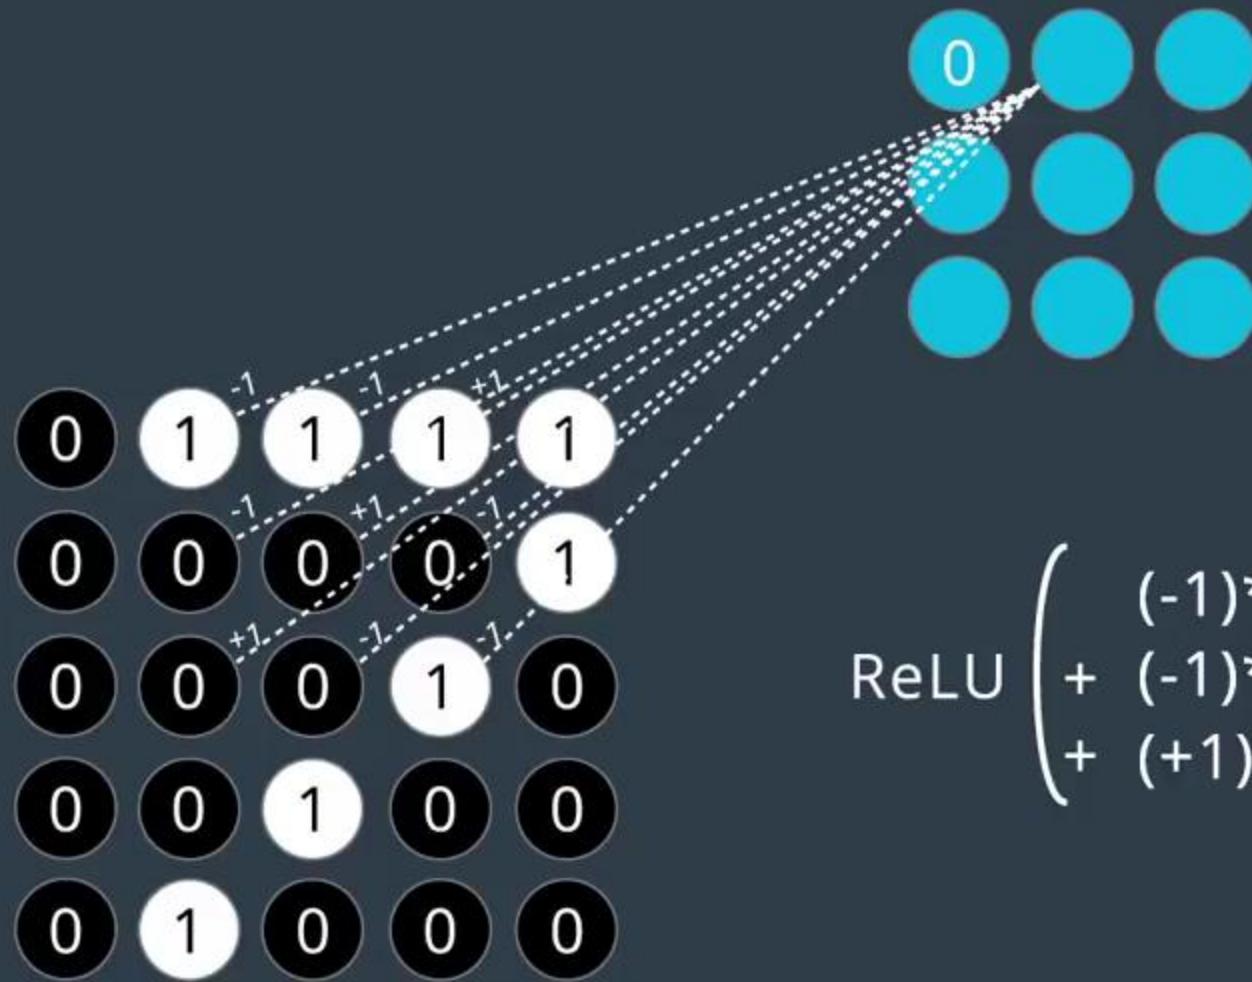


$$\text{ReLU} \begin{pmatrix} (-1)*0 + (-1)*1 + (+1)*1 \\ + (-1)*0 + (+1)*0 + (-1)*0 \\ + (+1)*0 + (-1)*0 + (-1)*0 \end{pmatrix} = 0$$



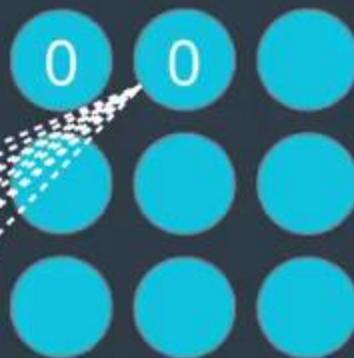


$$\begin{aligned}
 & (-1) * 0 + (-1) * 1 + (+1) * 1 \\
 & + (-1) * 0 + (+1) * 0 + (-1) * 0 = -2 \\
 & + (+1) * 0 + (-1) * 0 + (-1) * 0
 \end{aligned}$$



$$\text{ReLU} \begin{pmatrix} (-1)*0+(-1)*1+(+1)*1 \\ + (-1)*0+(+1)*0+(-1)*0 \\ + (+1)*0+(-1)*0+(-1)*0 \end{pmatrix} = 0$$

0	1	1	1	1
0	0	0	0	1
0	0	0	1	0
0	0	1	0	0
0	1	0	0	0



Filter

-1	-1	+1
-1	+1	-1
+1	-1	-1

0	1	1	1	1
0	0	0	0	1
0	0	0	1	0
0	0	1	0	0
0	1	0	0	0

Input Layer

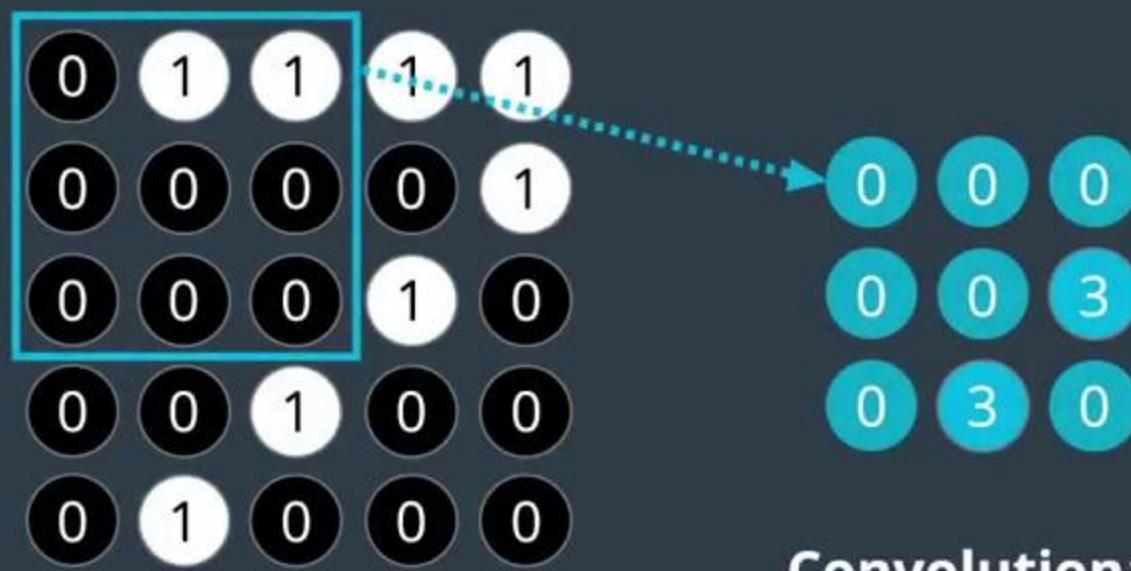
0	0	0
0	0	3
0	3	0

Convolutional Layer

$$\text{ReLU} \left(\text{SUM} \left(\begin{array}{cccc} * & & * & * \\ * & & * & * \\ * & & * & * \\ * & & * & * \end{array} \right) \right) = 0$$

Filter

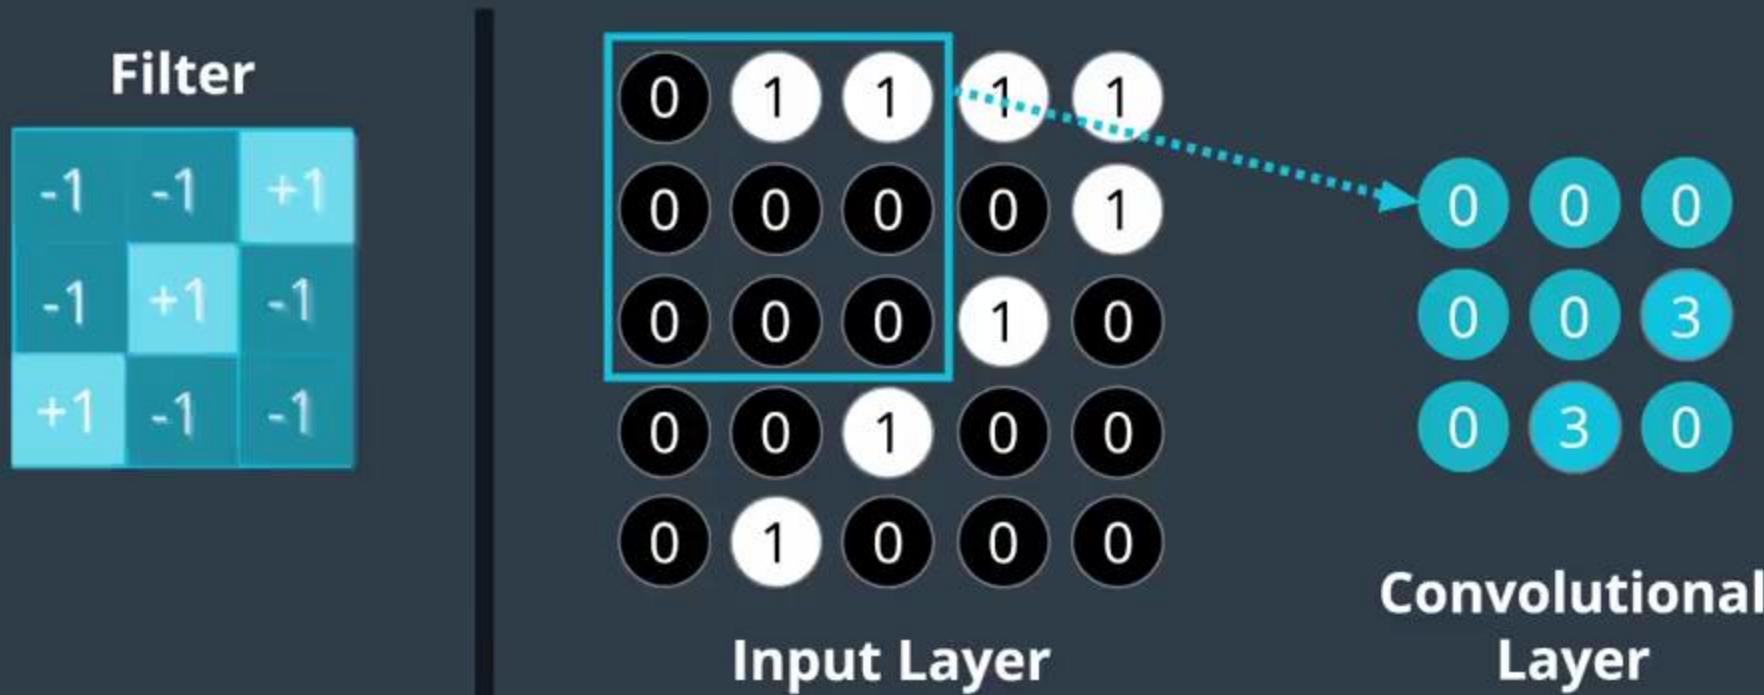
-1	-1	+1
-1	+1	-1
+1	-1	-1



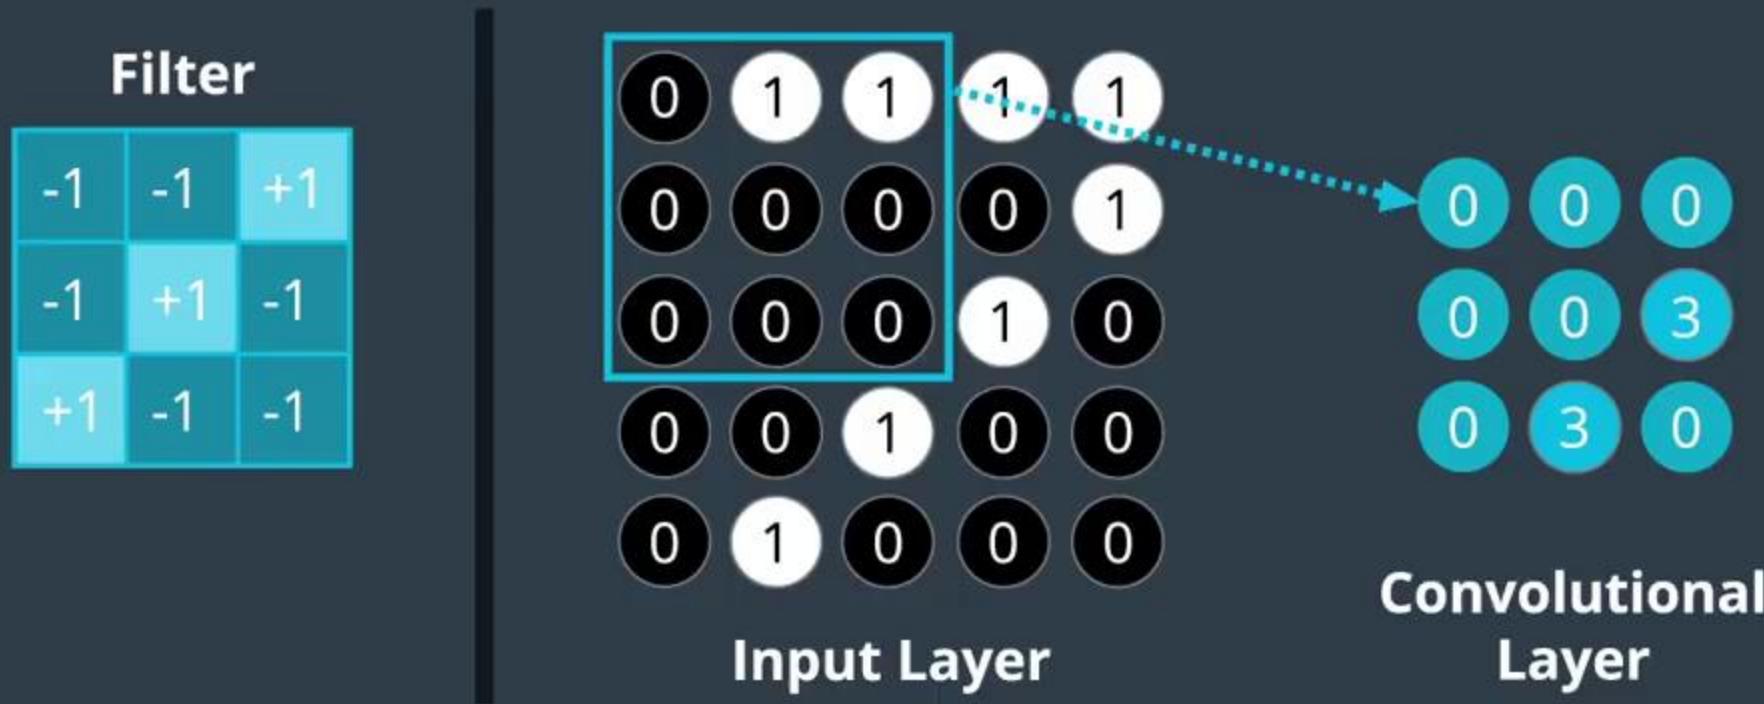
Input Layer

**Convolutional
Layer**

$$\text{ReLU} \left(\text{SUM} \left(\begin{matrix} * & * & * \\ * & * & * \\ * & * & * \end{matrix} \right) \right) = 0$$



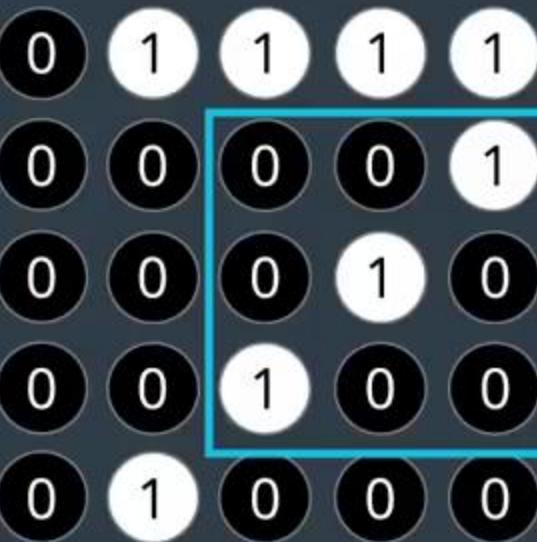
$$\text{ReLU} \left(\text{SUM} \left(\begin{array}{ccc} *0 & *1 & *1 \\ *0 & *0 & *0 \\ *0 & *0 & *0 \end{array} \right) \right) = 0$$



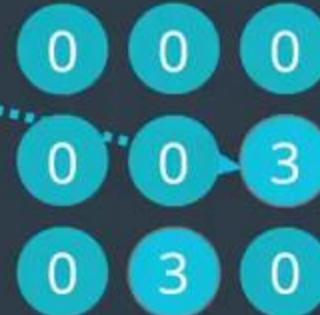
$$\text{ReLU} \left(\text{SUM} \left(\begin{array}{ccc} -1 * 0 & -1 * 1 & +1 * 1 \\ -1 * 0 & +1 * 0 & -1 * 0 \\ +1 * 0 & -1 * 0 & -1 * 0 \end{array} \right) \right) = 0$$

Filter

-1	-1	+1
-1	+1	-1
+1	-1	-1



Input Layer

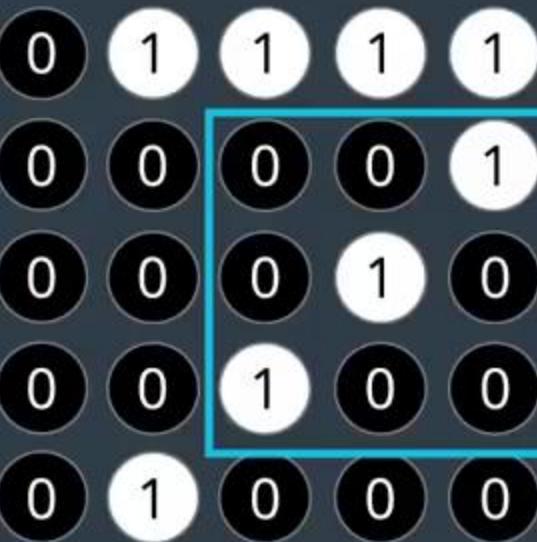


**Convolutional
Layer**

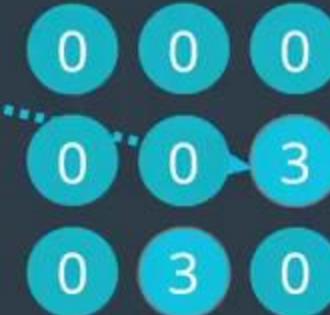
$$\text{ReLU} \left(\text{SUM} \left(\begin{array}{ccc} -1 * 0 & -1 * 0 & +1 * 1 \\ -1 * 0 & +1 * 1 & -1 * 0 \\ +1 * 1 & -1 * 0 & -1 * 0 \end{array} \right) \right) = 3$$

Filter

-1	-1	+1
-1	+1	-1
+1	-1	-1



Input Layer

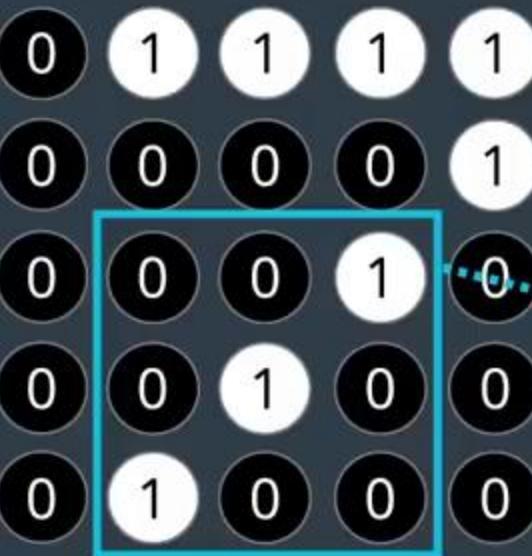


**Convolutional
Layer**

$$\text{ReLU} \left(\text{SUM} \left(\begin{array}{ccc} -1 * 0 & -1 * 0 & +1 * 1 \\ -1 * 0 & +1 * 1 & -1 * 0 \\ +1 * 1 & -1 * 0 & -1 * 0 \end{array} \right) \right) = 3$$

Filter

-1	-1	+1
-1	+1	-1
+1	-1	-1



Input Layer



**Convolutional
Layer**

$$\text{ReLU} \left(\text{SUM} \left(\begin{array}{ccc} -1 * 0 & -1 * 0 & +1 * 1 \\ -1 * 0 & +1 * 1 & -1 * 0 \\ +1 * 1 & -1 * 0 & -1 * 0 \end{array} \right) \right) = 3$$

Filter



Input Layer

Convolutional Layer

$$\text{ReLU}(\text{SUM}\left(\begin{array}{ccc} -1 * 0 & -1 * 0 & +1 * 1 \\ -1 * 0 & +1 * 1 & -1 * 0 \\ +1 * 1 & -1 * 0 & -1 * 0 \end{array}\right)) = 3$$



HEIGHT

WIDTH





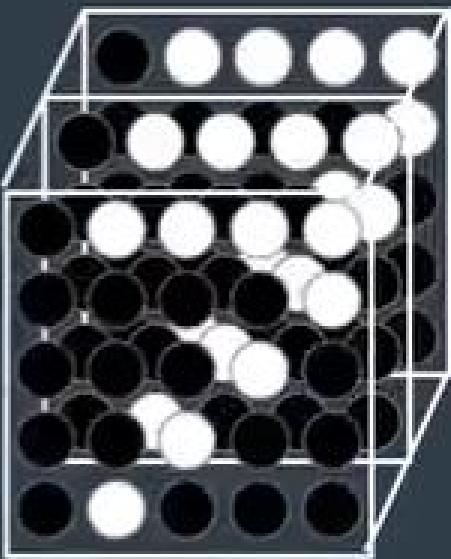


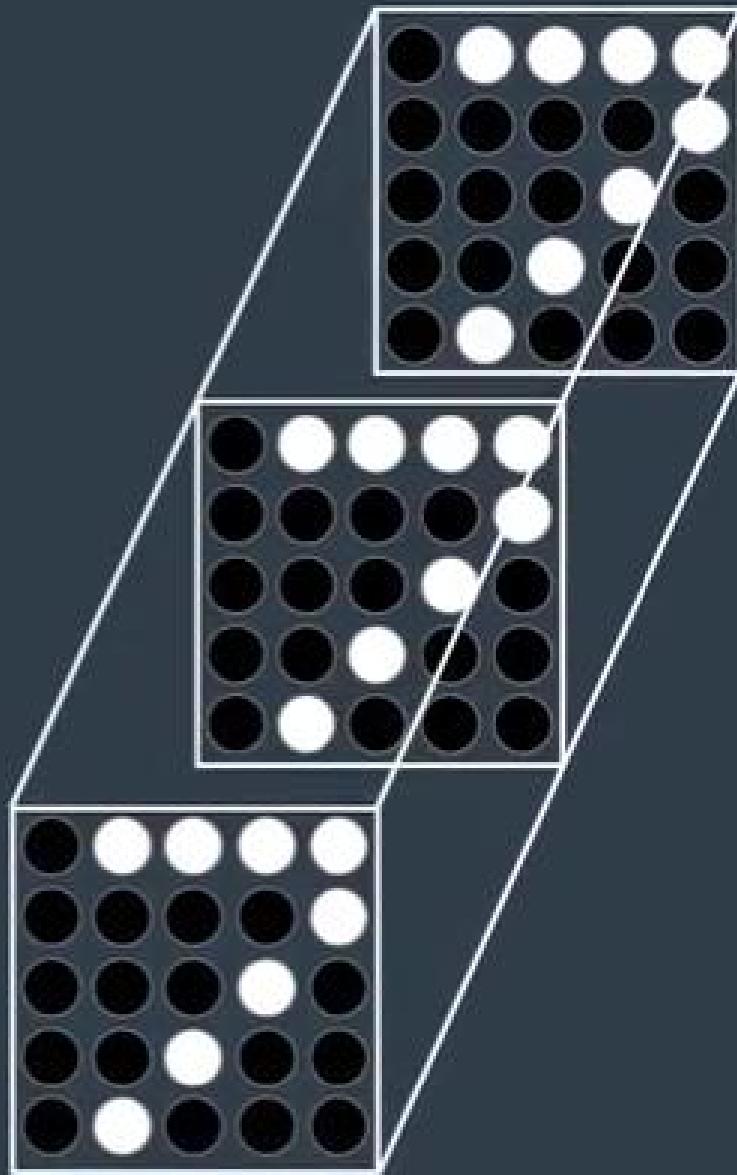


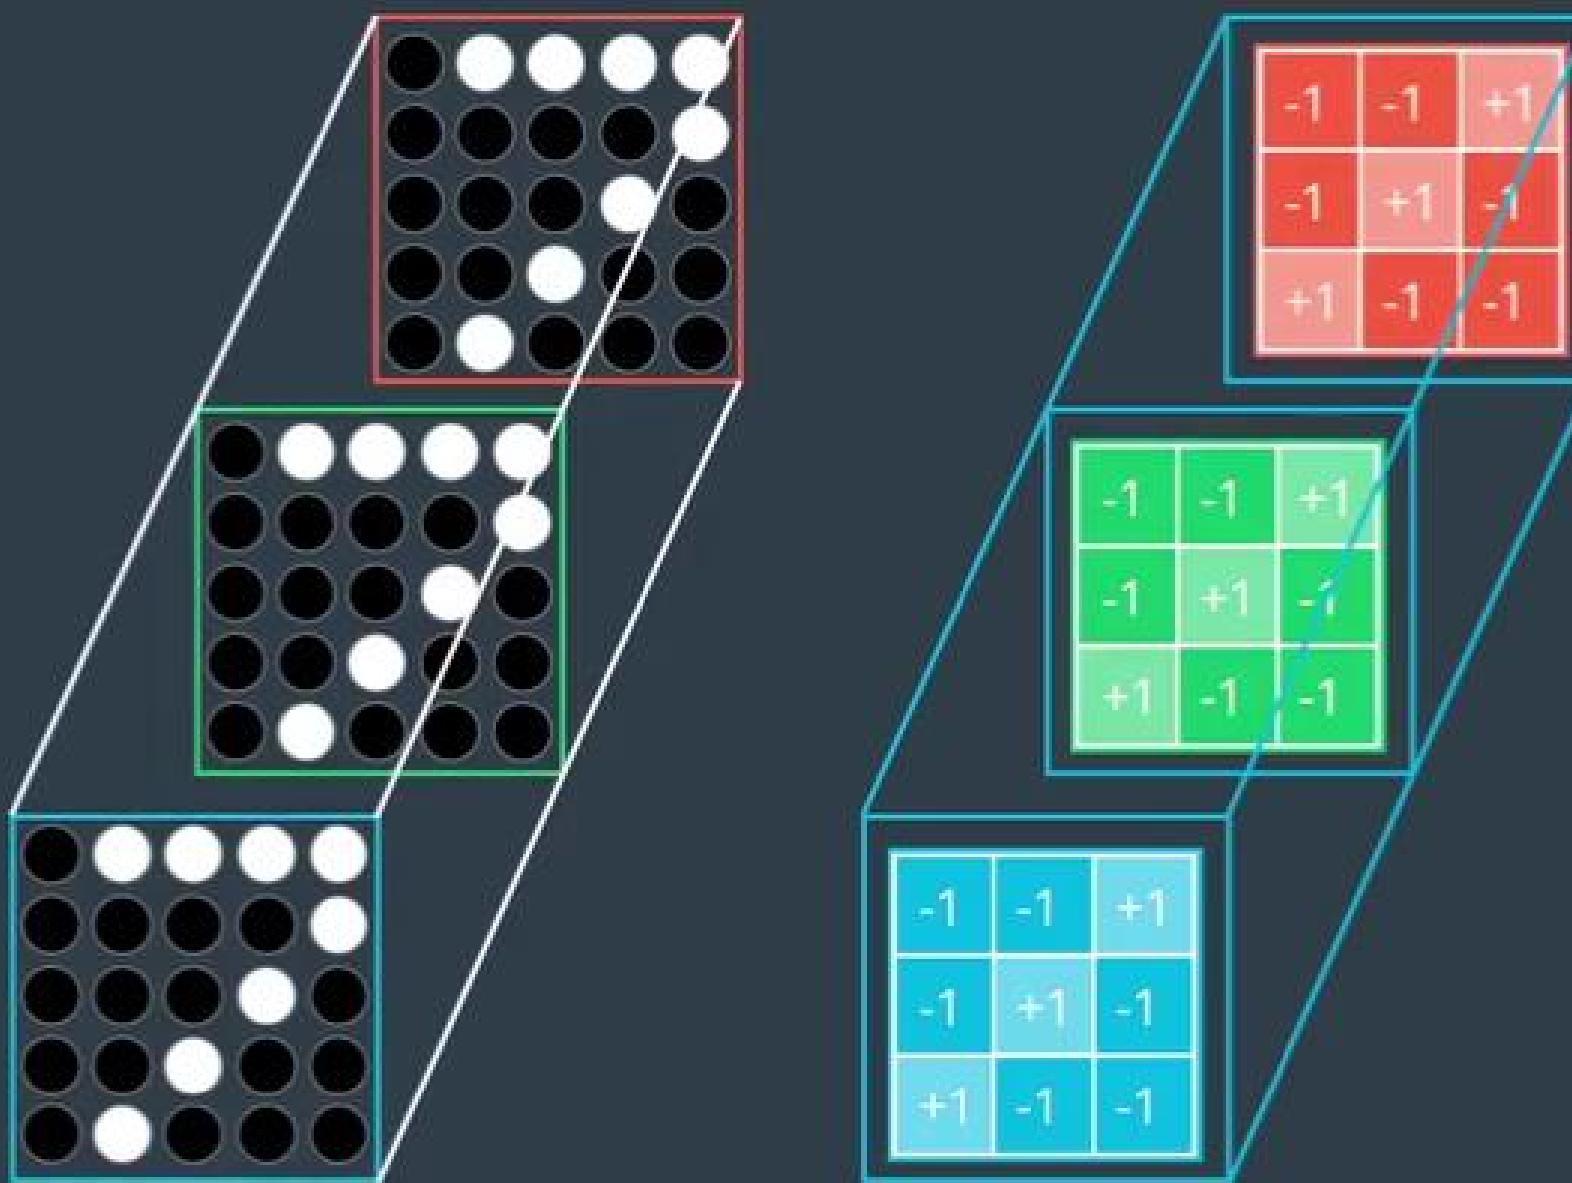


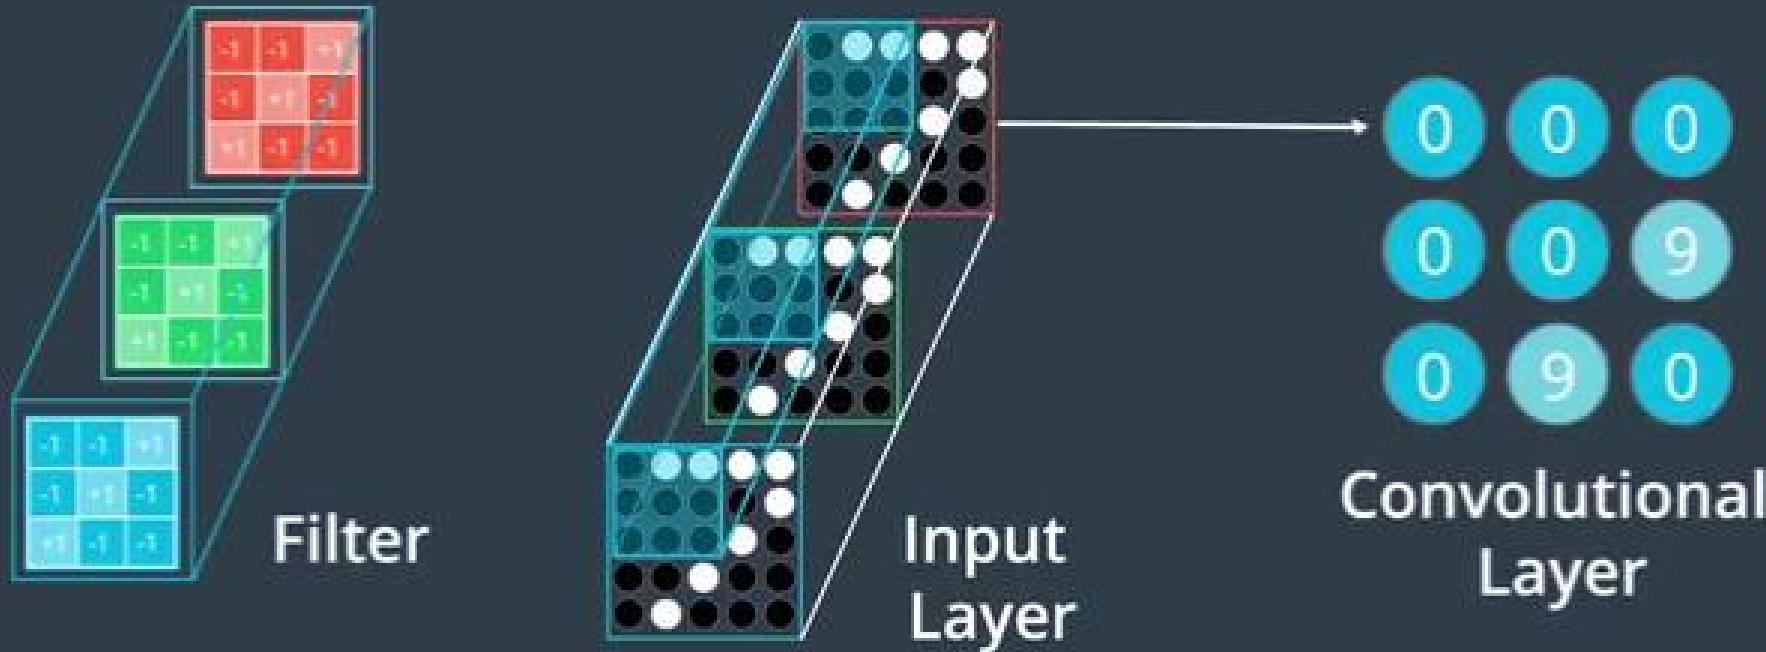


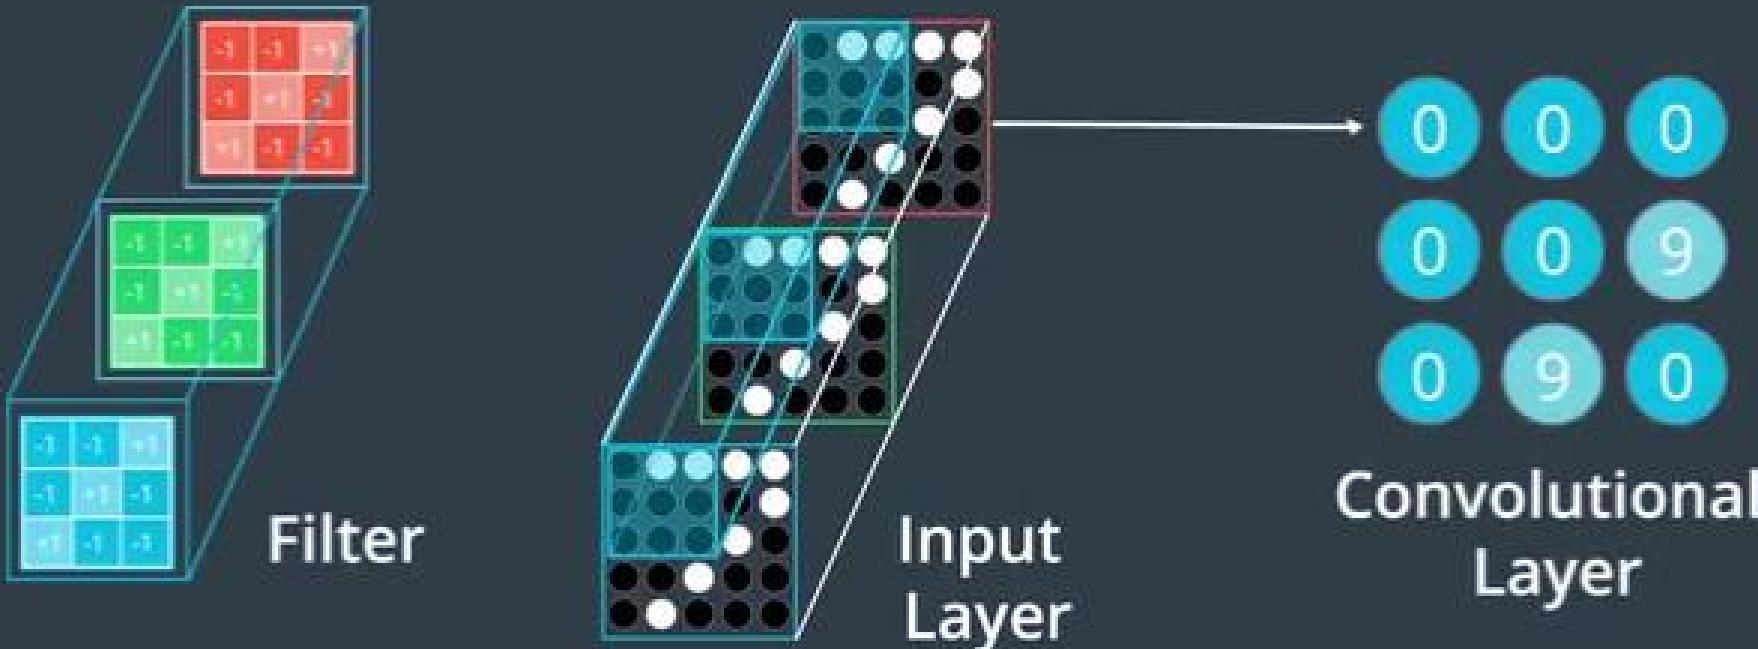
↔
DEPTH





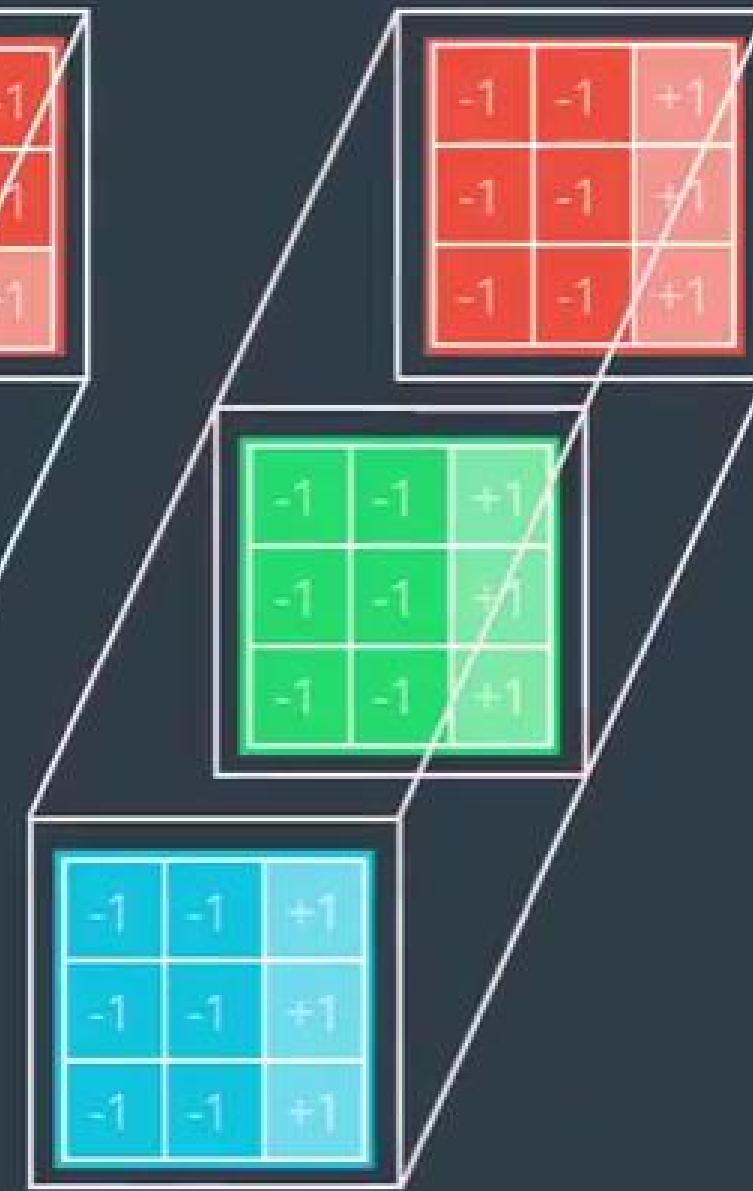






Convolutional Layer

$$\text{ReLU} \left(\text{SUM} \left(\begin{array}{cccccccc} \text{cyan} \times \bullet & \text{cyan} \times \circ & \text{light blue} \times \bullet & \text{green} \times \bullet & \text{green} \times \circ & \text{light green} \times \bullet & \text{red} \times \bullet & \text{red} \times \circ & \text{pink} \times \bullet \\ \text{cyan} \times \bullet & \text{light blue} \times \bullet & \text{cyan} \times \bullet & \text{green} \times \bullet & \text{light green} \times \bullet & \text{green} \times \bullet & \text{red} \times \bullet & \text{pink} \times \bullet & \text{red} \times \bullet \\ \text{light blue} \times \bullet & \text{cyan} \times \bullet & \text{cyan} \times \bullet & \text{green} \times \bullet & \text{light green} \times \bullet & \text{green} \times \bullet & \text{pink} \times \bullet & \text{red} \times \bullet & \text{red} \times \bullet \end{array} \right) \right)$$



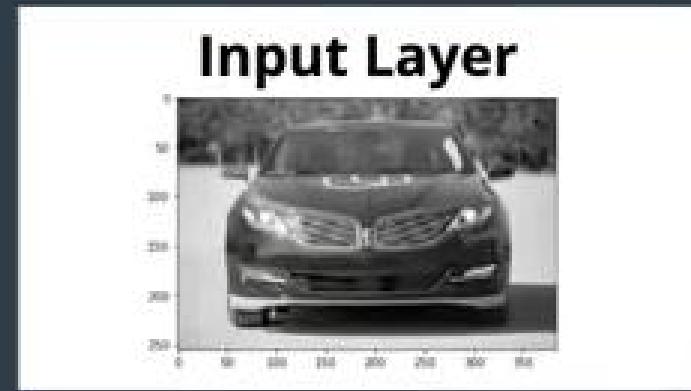
Filter 1

Filter 2

Filter 3

Input Layer





Filter 1

-1	-1	+1	+1
-1	-1	+1	+1
-1	-1	+1	+1
-1	-1	+1	+1

Filter 2

+1	+1	-1	-1
+1	+1	-1	-1
+1	+1	-1	-1
+1	+1	-1	-1

Filter 3

-1	-1	-1	-1
-1	-1	-1	-1
+1	+1	+1	+1
+1	+1	+1	+1

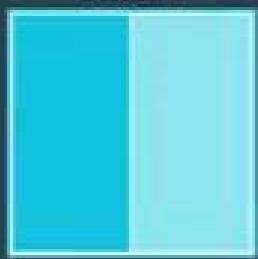
Filter 4

+1	+1	+1	+1
+1	+1	+1	+1
-1	-1	-1	-1
-1	-1	-1	-1

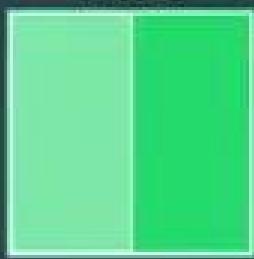
Input Layer



Filter 1



Filter 2

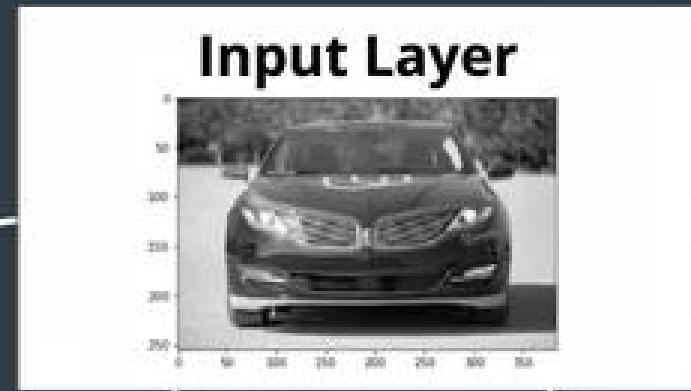


Filter 3

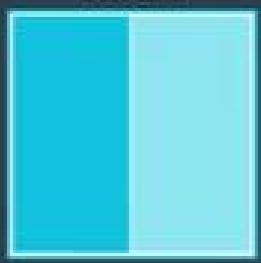


Filter 4





Filter 1



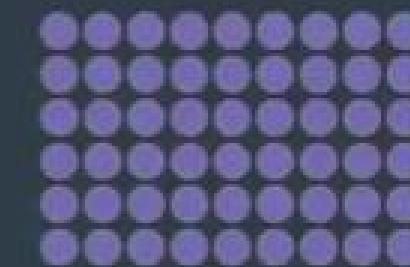
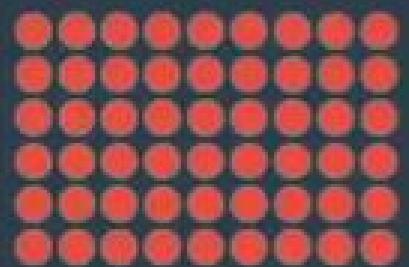
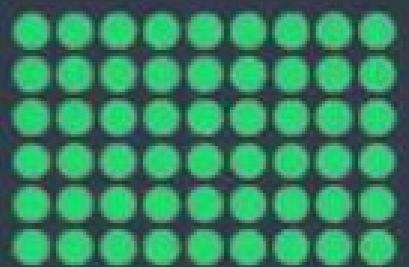
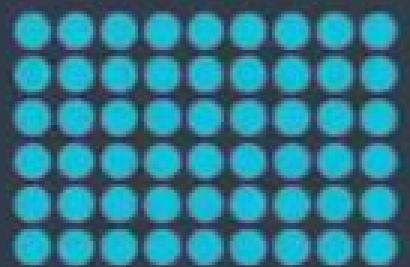
Filter 2

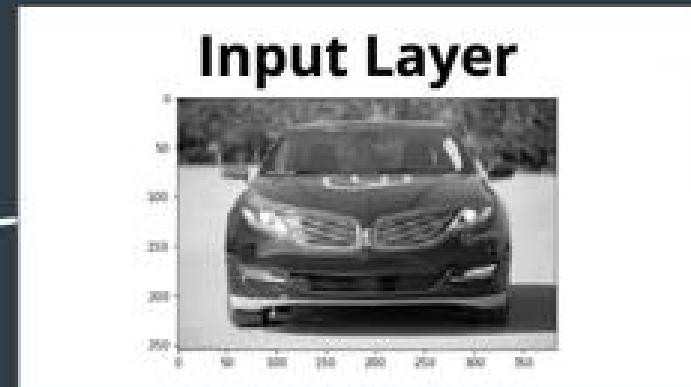


Filter 3

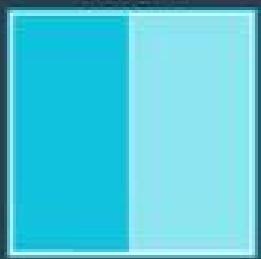


Filter 4

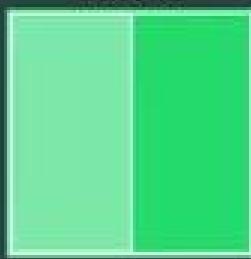




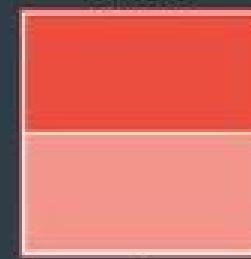
Filter 1



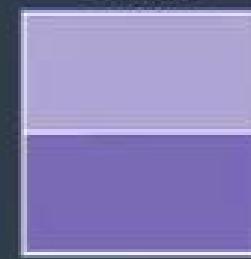
Filter 2



Filter 3



Filter 4

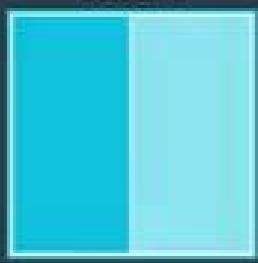


Convolutional Layer

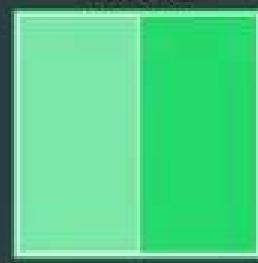
Input Layer



Filter 1



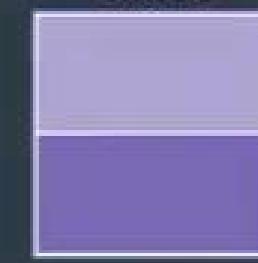
Filter 2



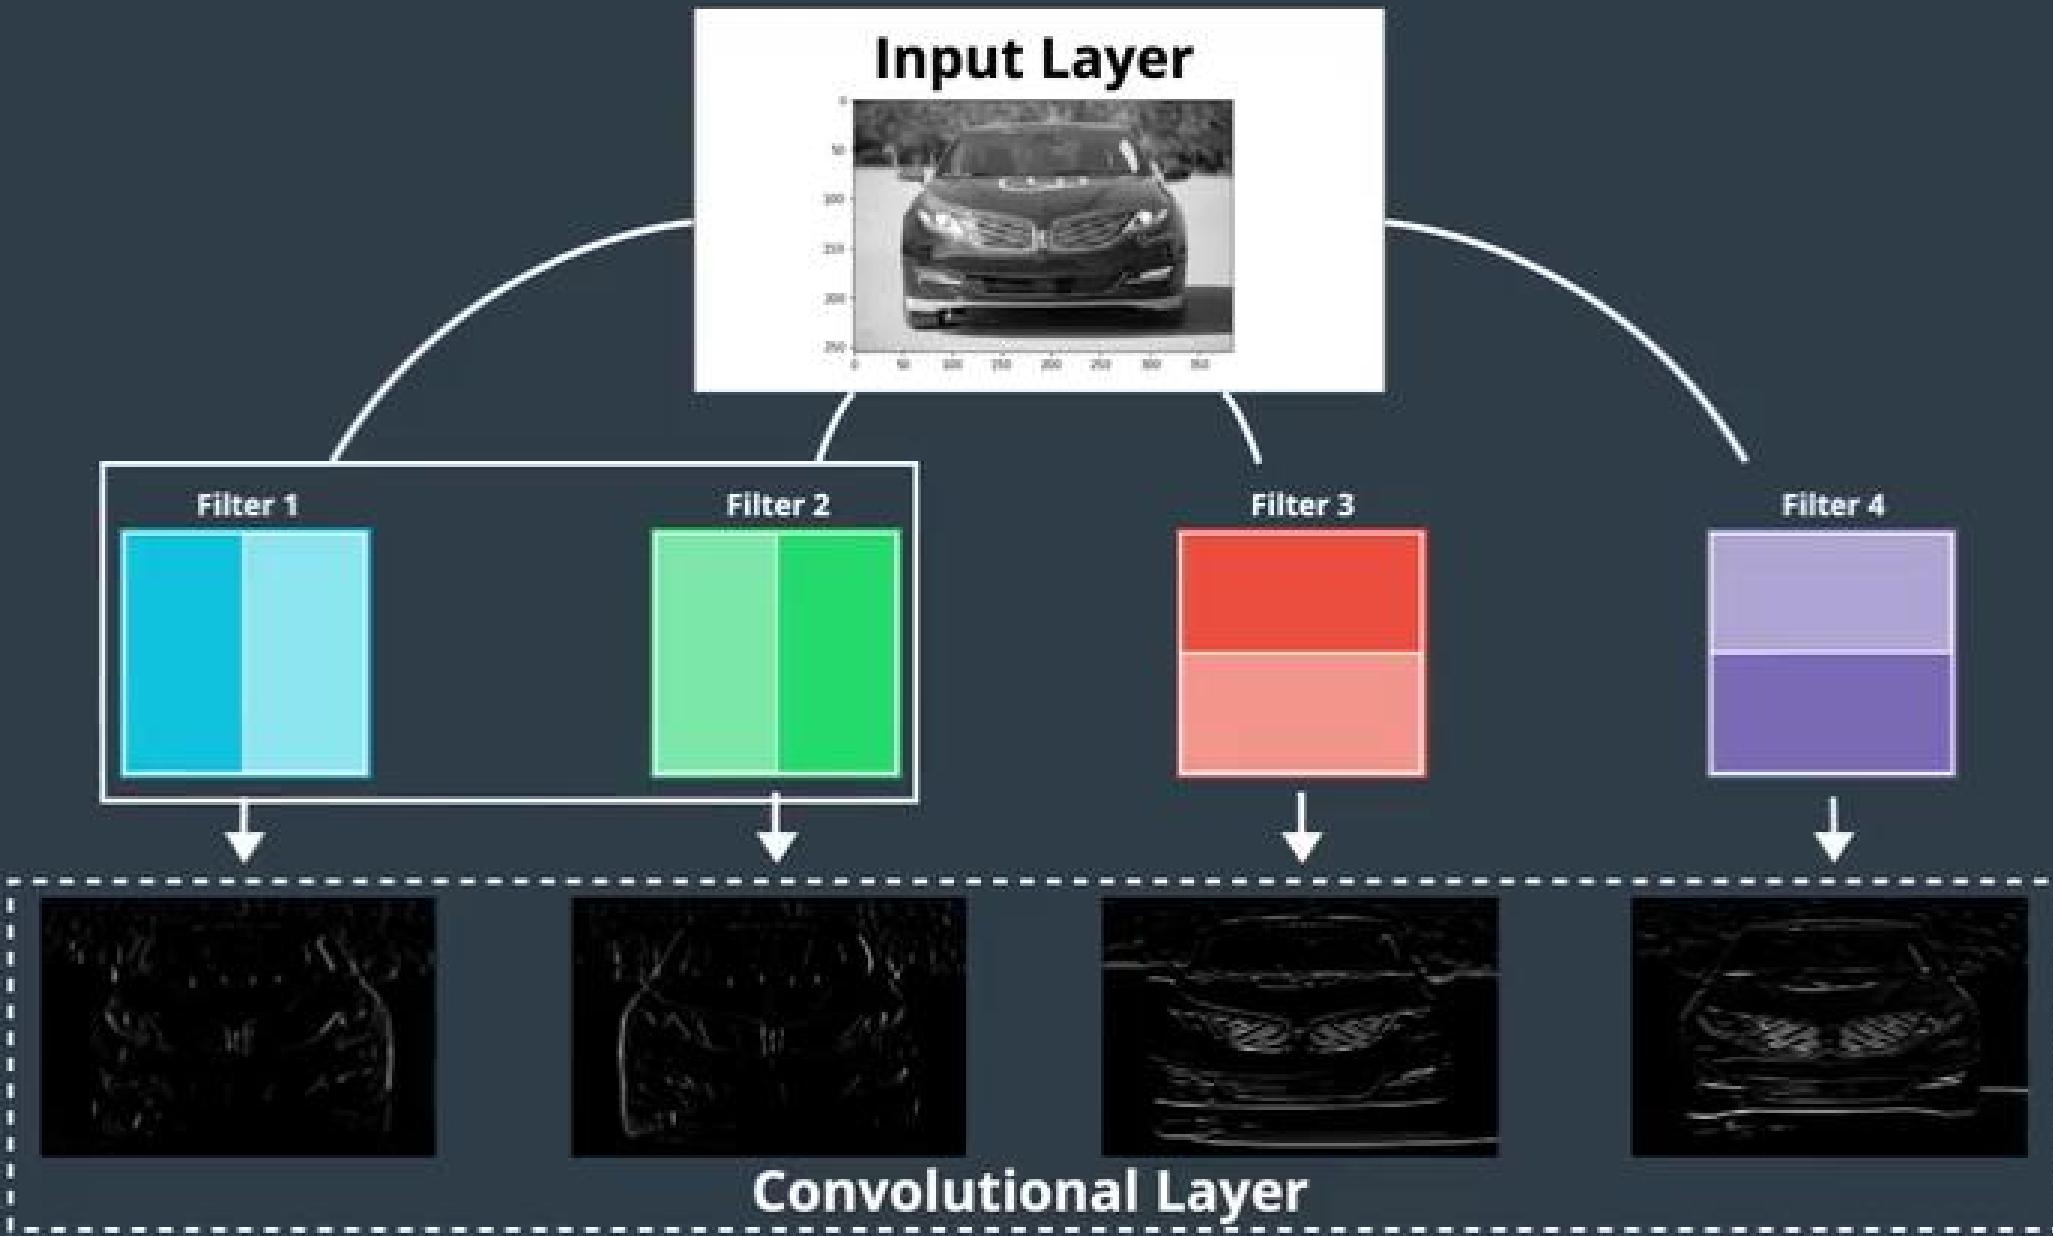
Filter 3



Filter 4



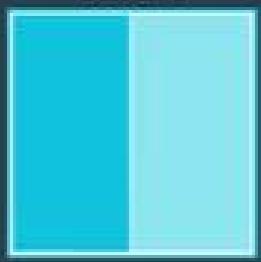
Convolutional Layer



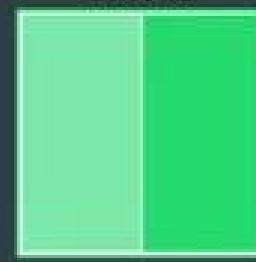
Input Layer



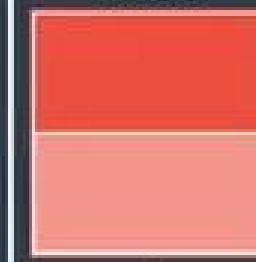
Filter 1



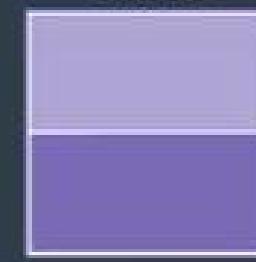
Filter 2



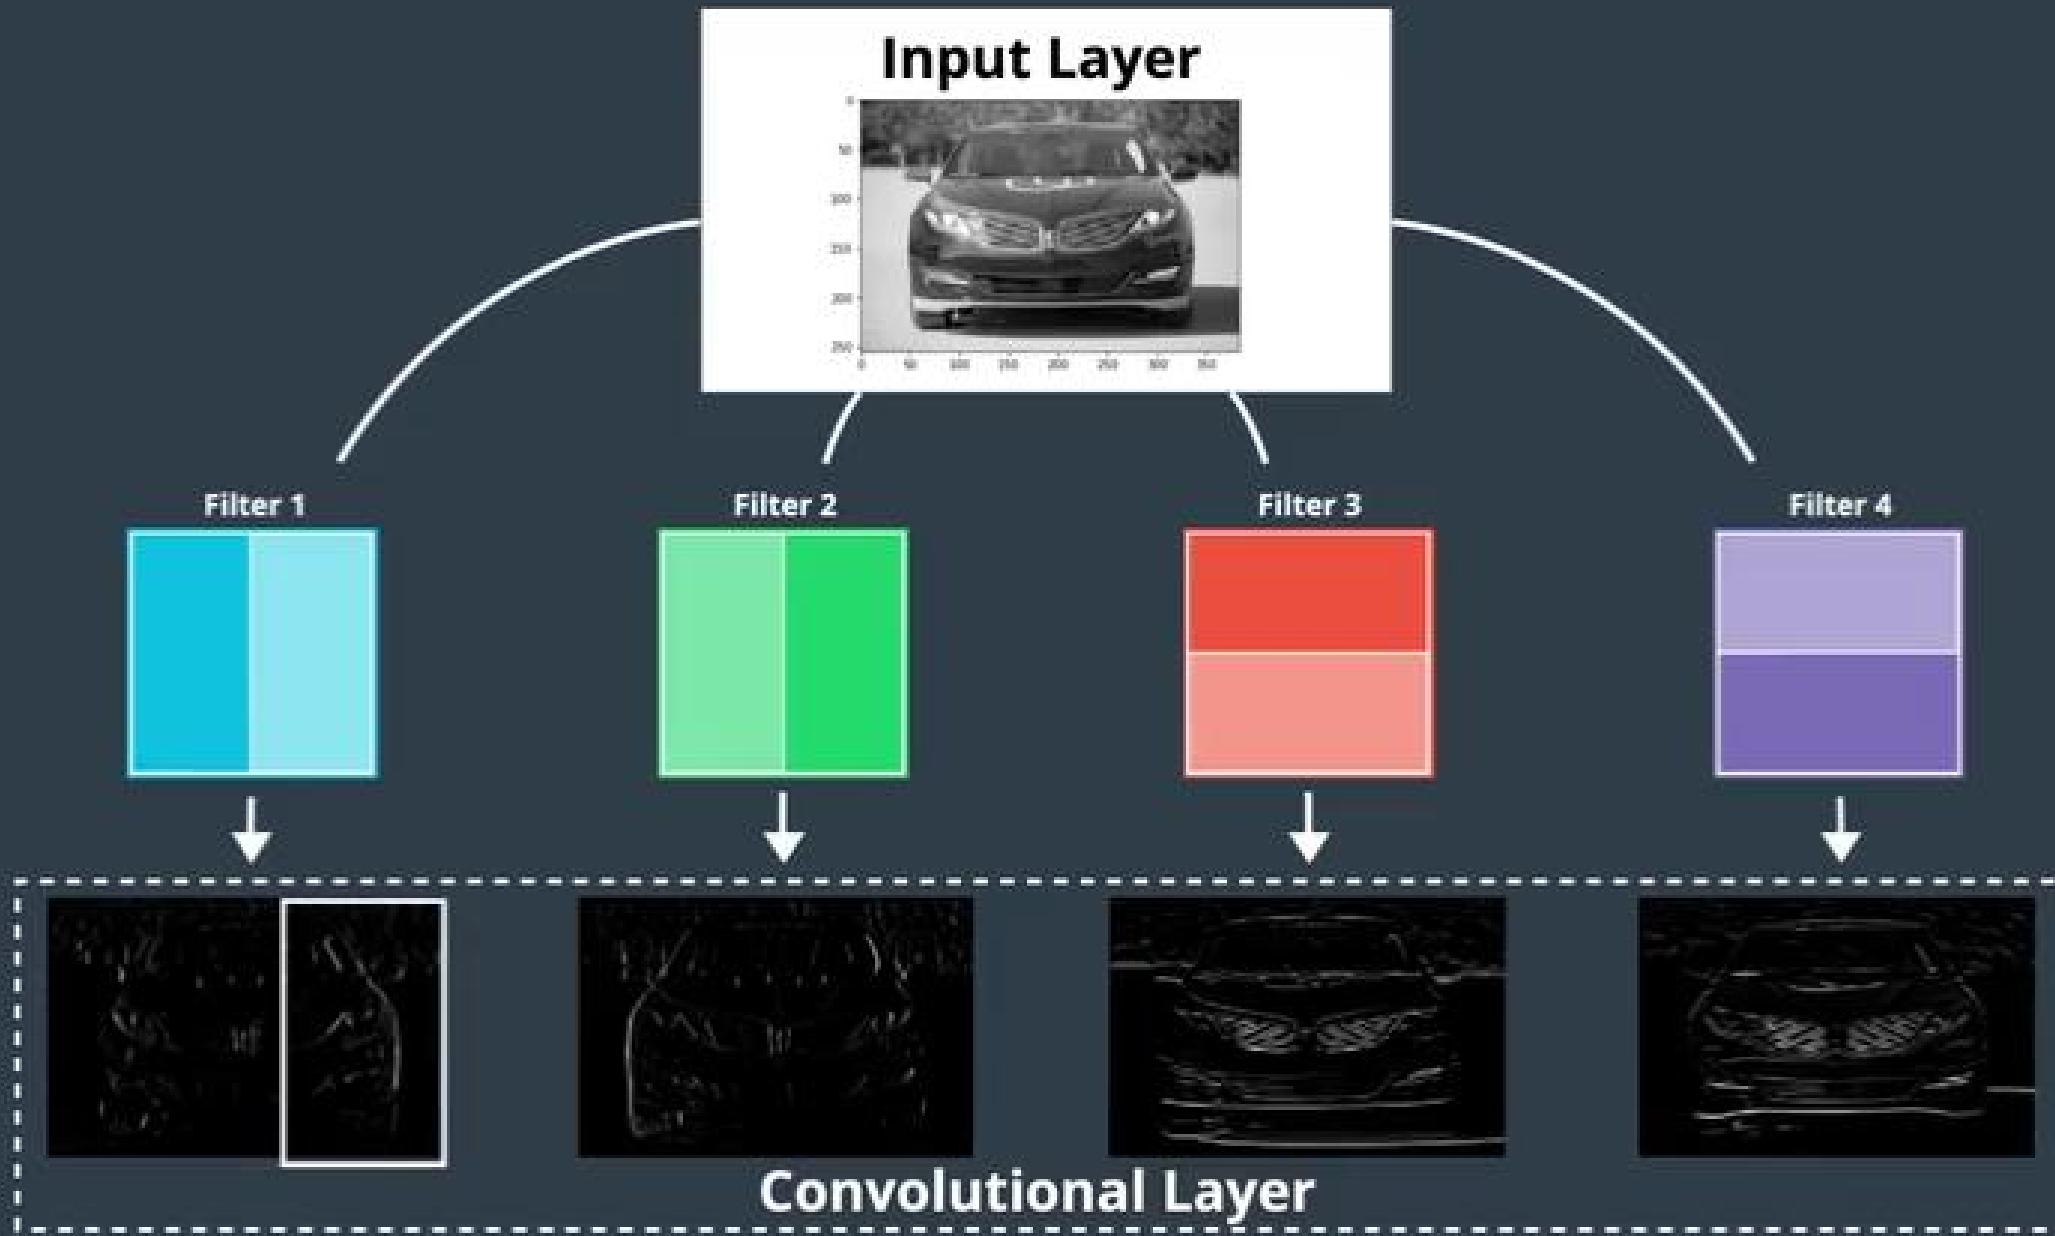
Filter 3



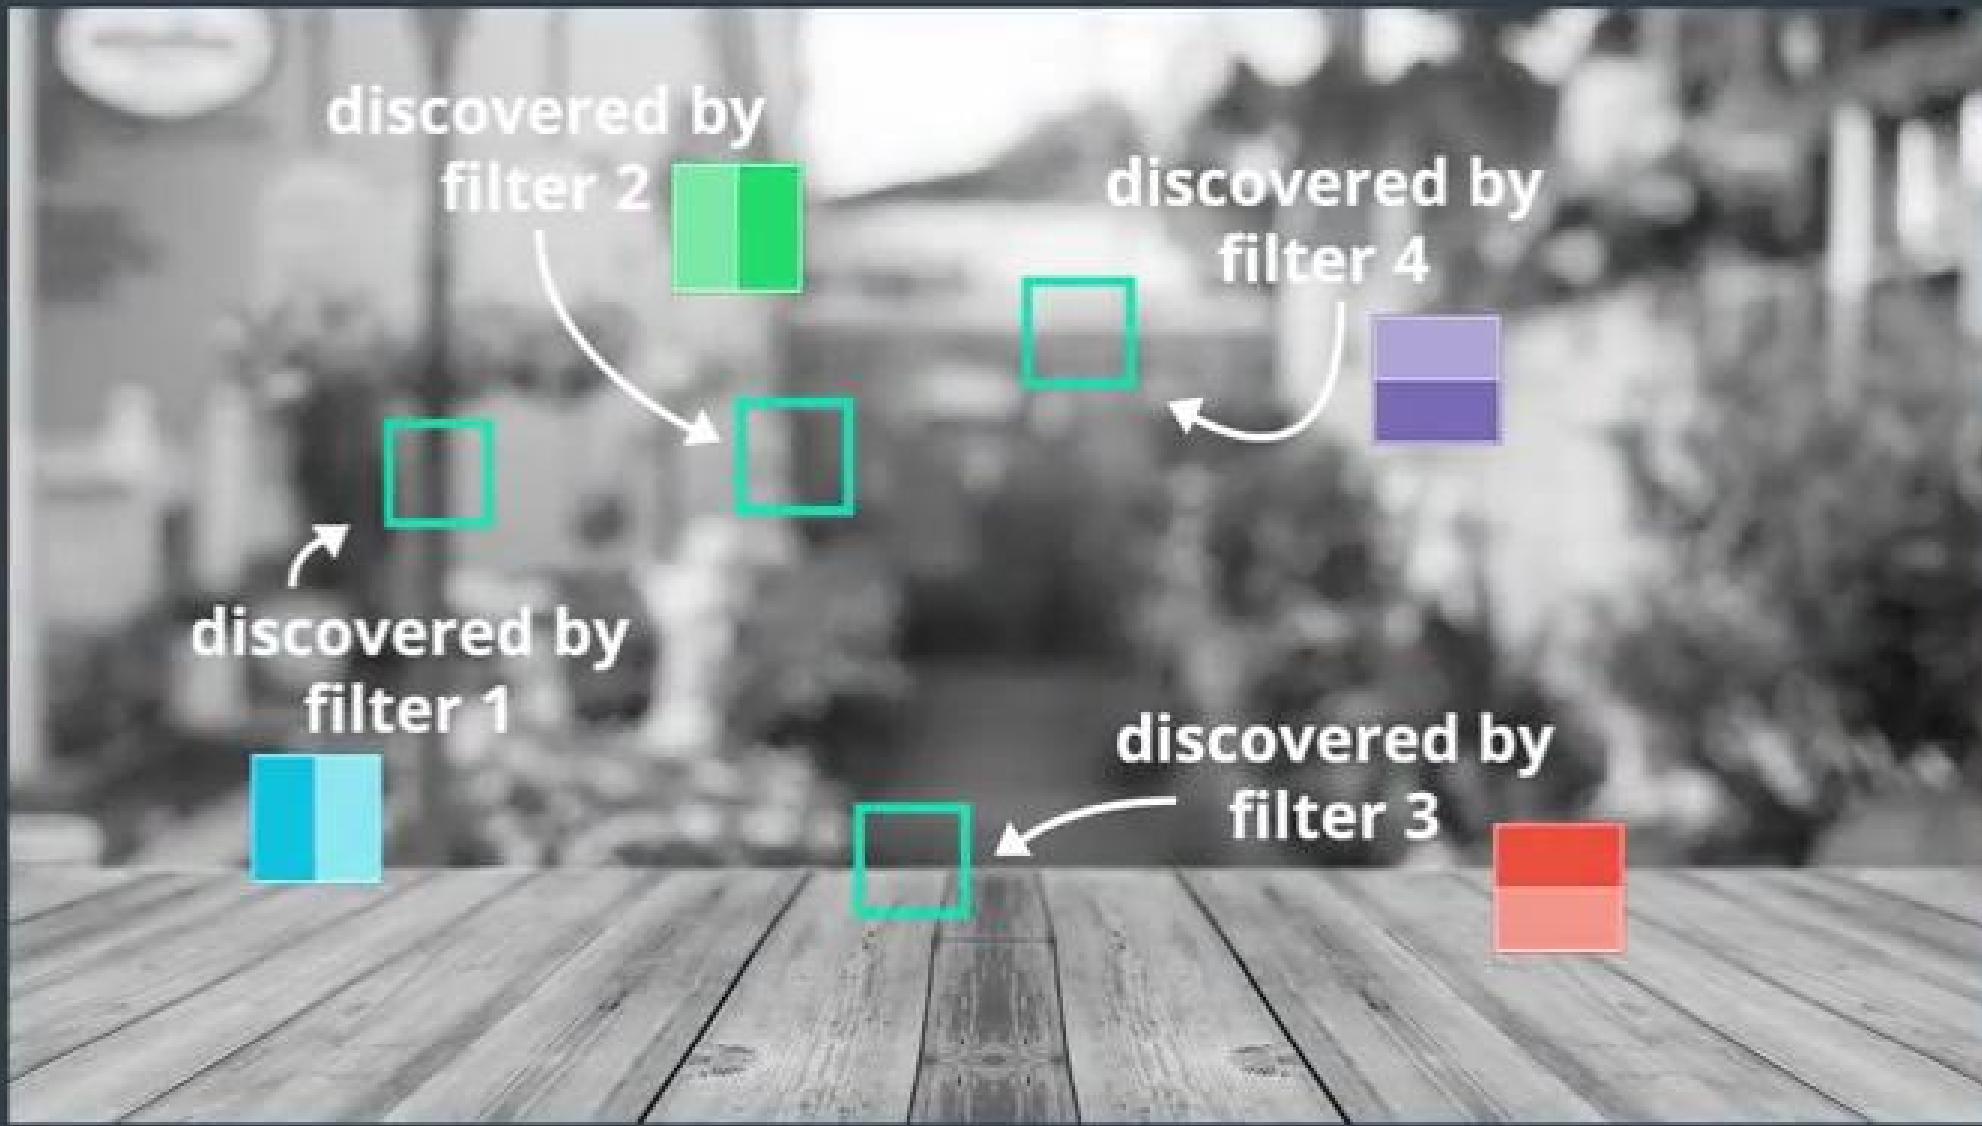
Filter 4

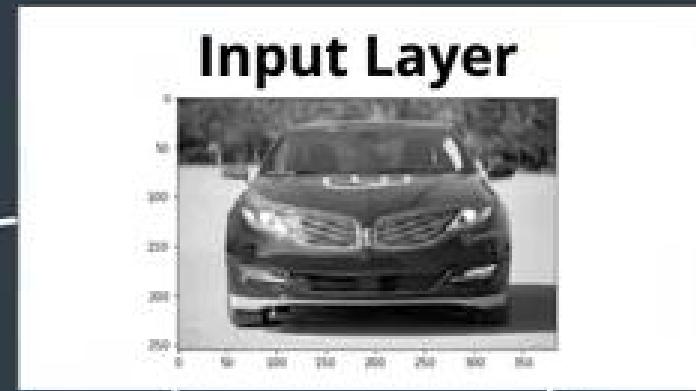


Convolutional Layer

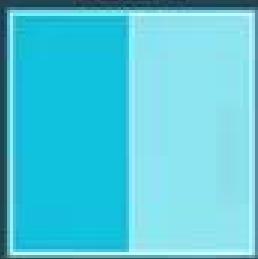




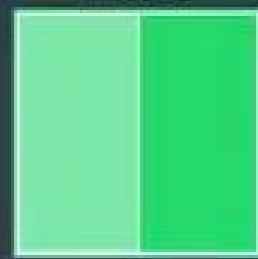




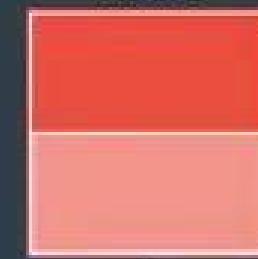
Filter 1



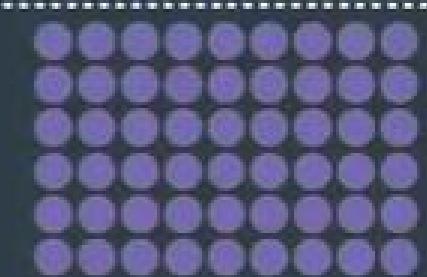
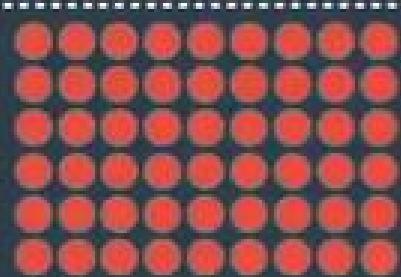
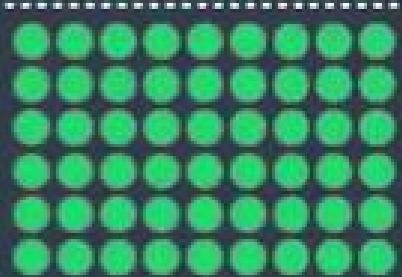
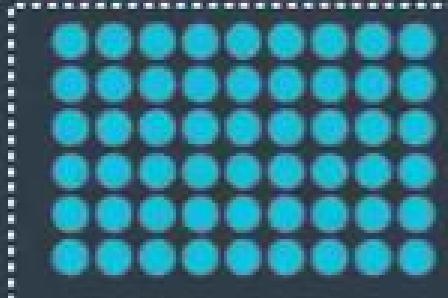
Filter 2



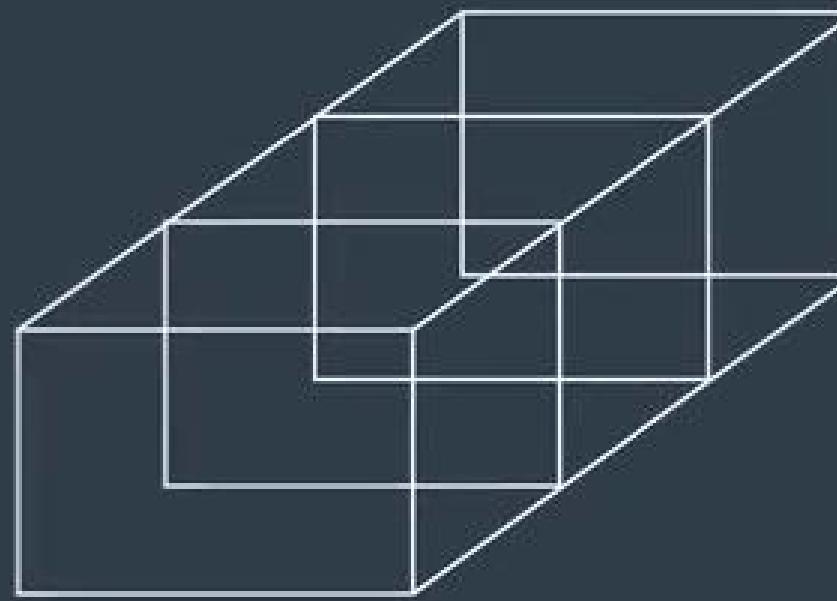
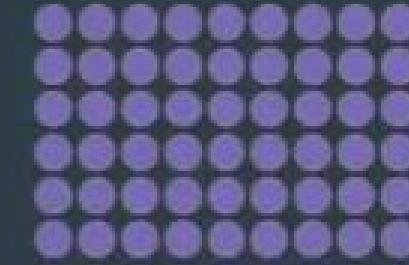
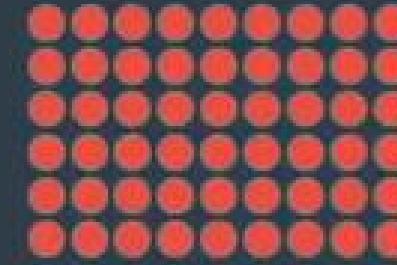
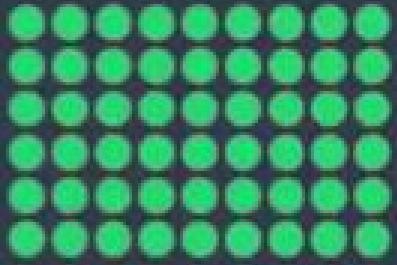
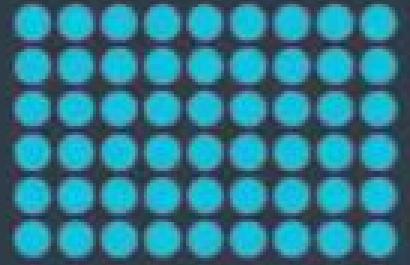
Filter 3

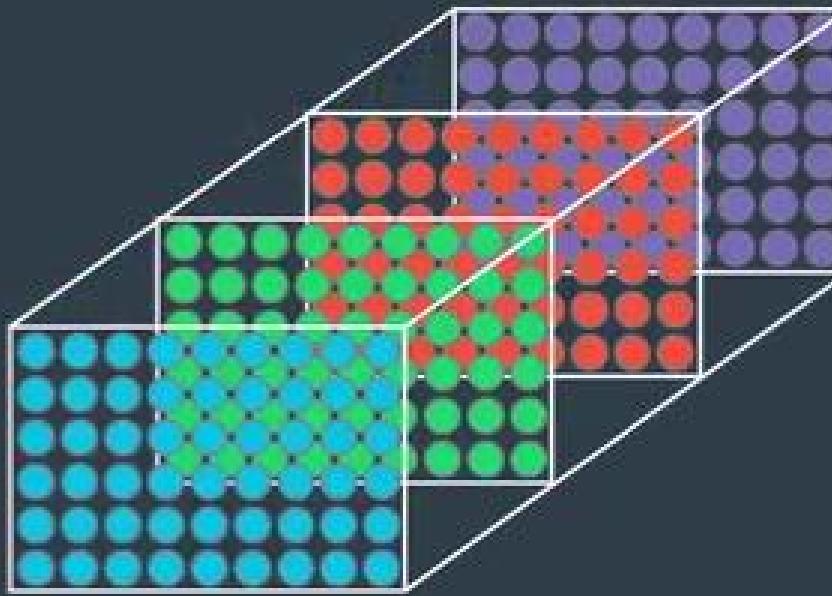


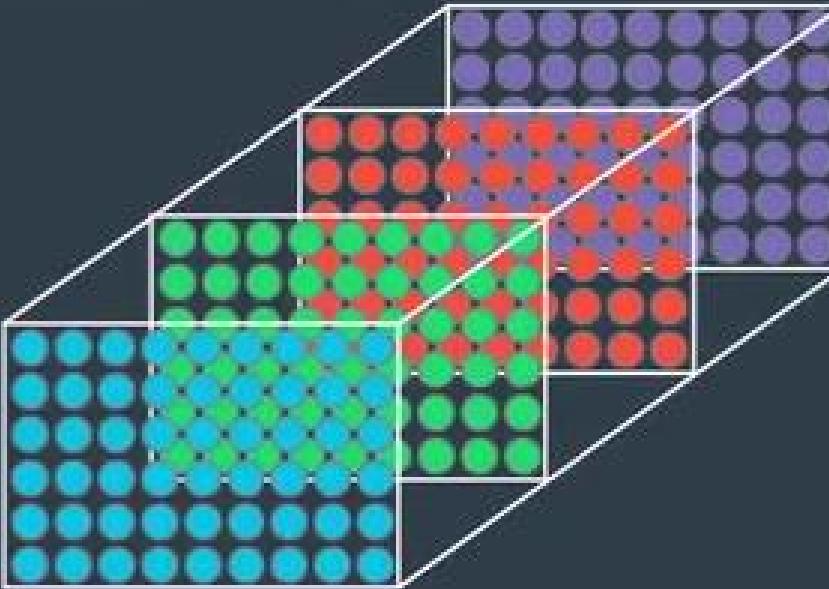
Filter 4



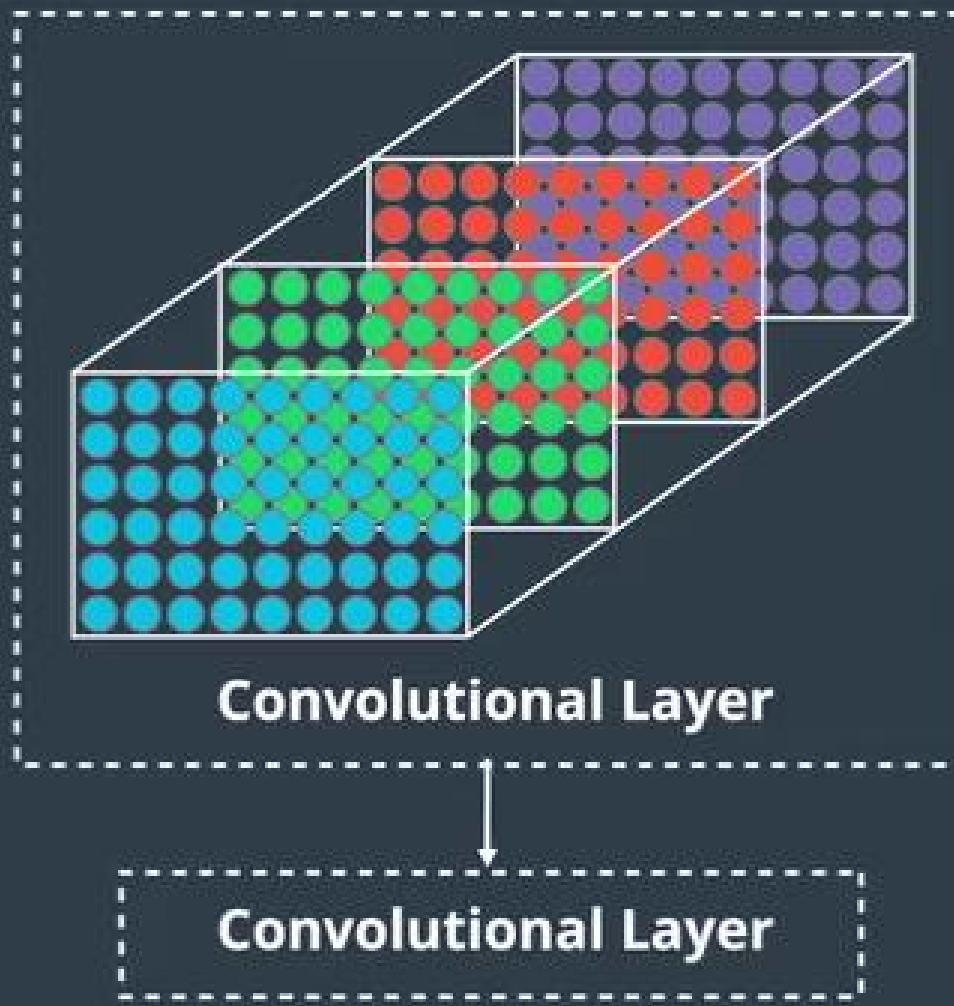
Convolutional Layer

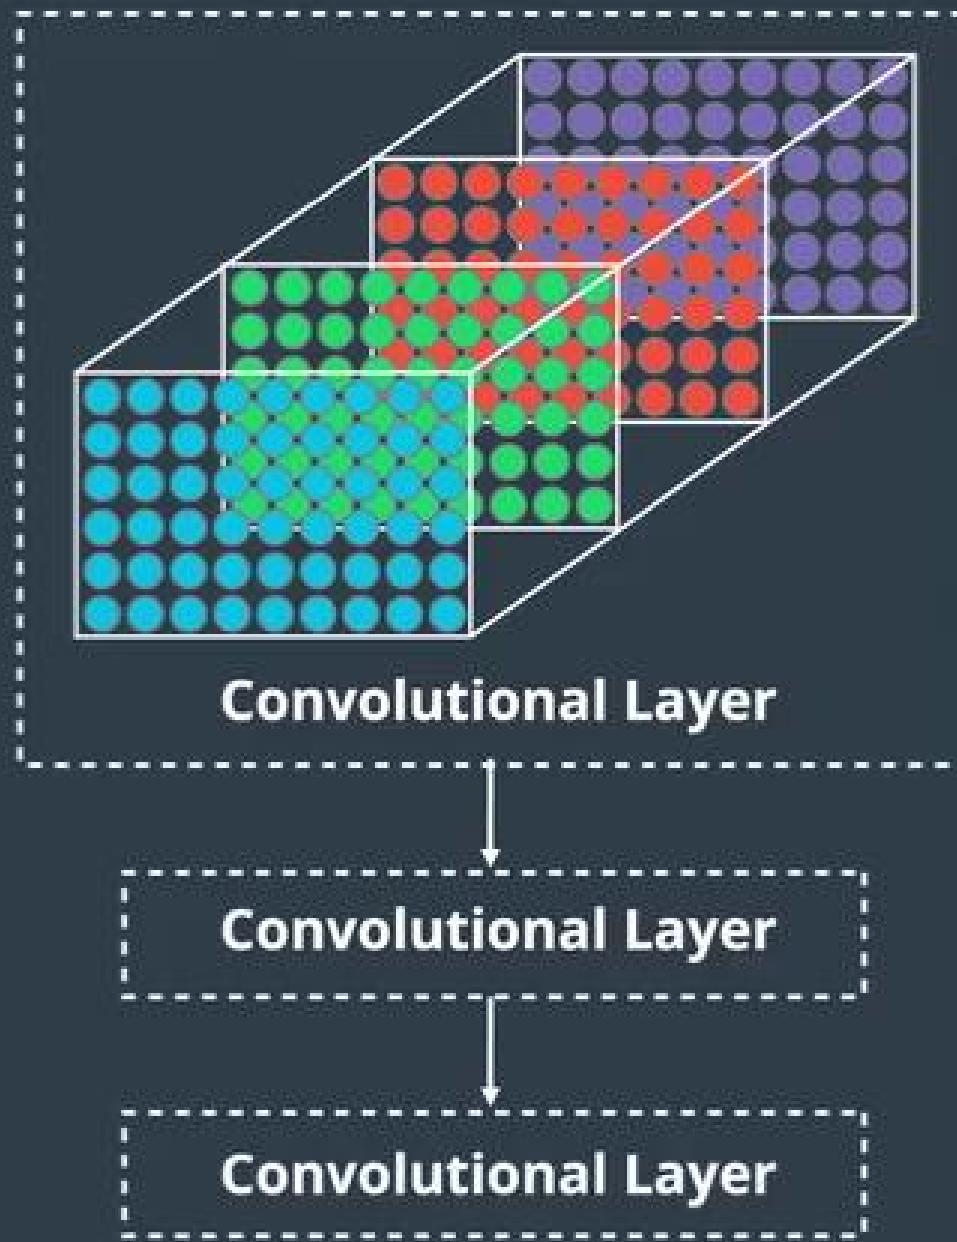




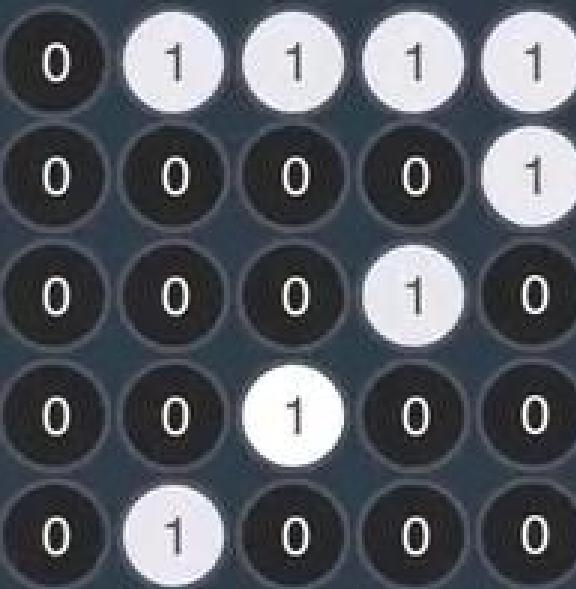
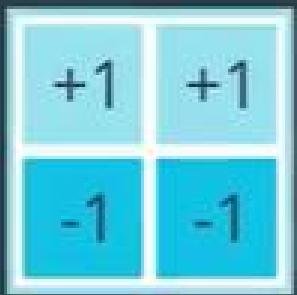


Convolutional Layer





Filter

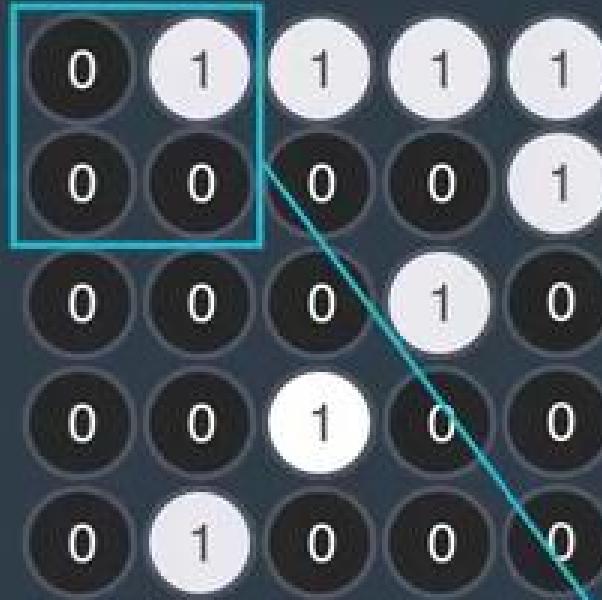


Input Layer

Convolutional
Layer

Filter

$$\begin{bmatrix} +1 & +1 \\ -1 & -1 \end{bmatrix}$$



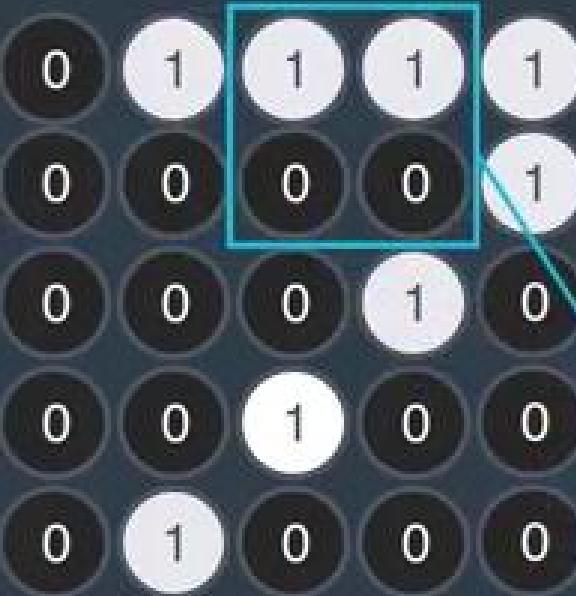
Input Layer

Convolutional
Layer

1

Filter

$$\begin{bmatrix} +1 & +1 \\ -1 & -1 \end{bmatrix}$$



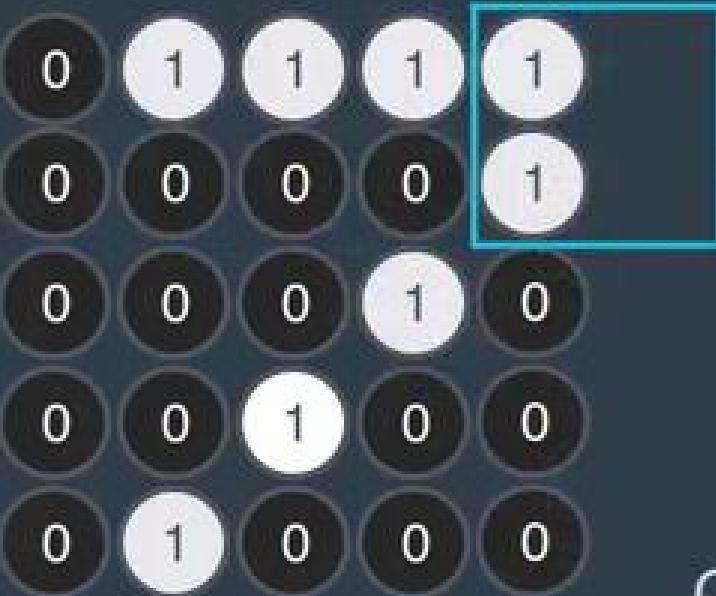
Input Layer

Convolutional
Layer

1 2

Filter

$$\begin{bmatrix} +1 & +1 \\ -1 & -1 \end{bmatrix}$$

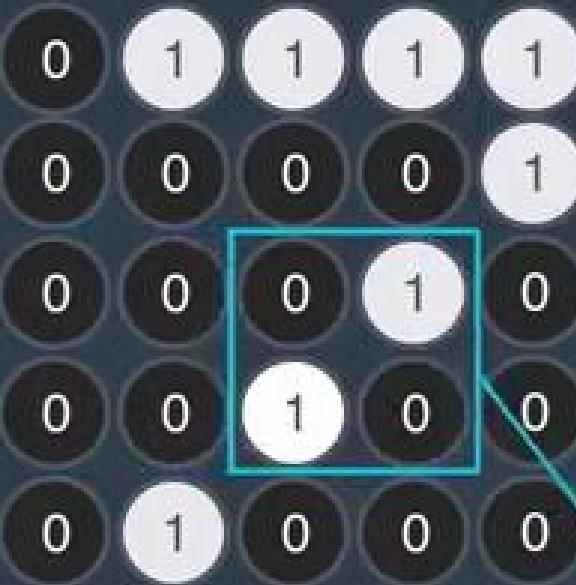
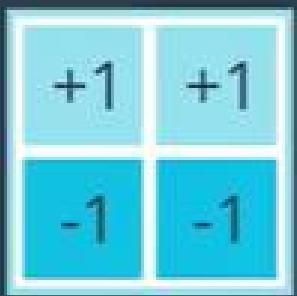


Input Layer

Convolutional
Layer

$$\begin{bmatrix} 1 & 2 \end{bmatrix}$$

Filter

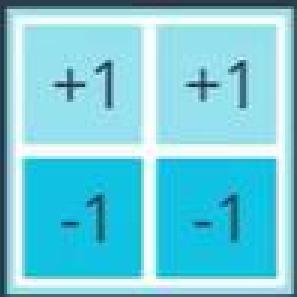


Input Layer

Convolutional
Layer



Filter

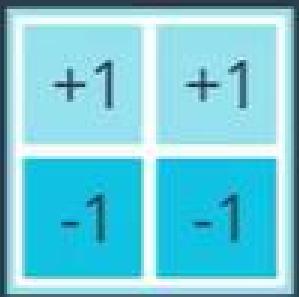


Input Layer

Convolutional
Layer

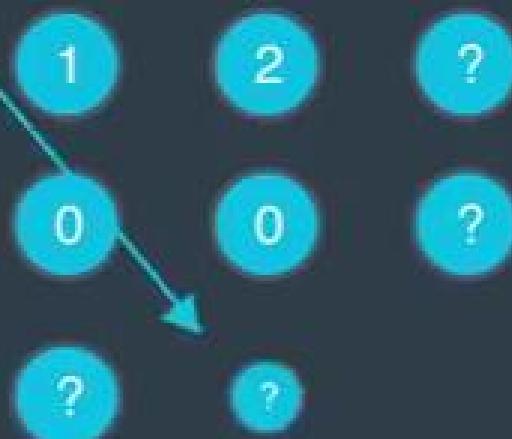


Filter



Input Layer

Convolutional Layer



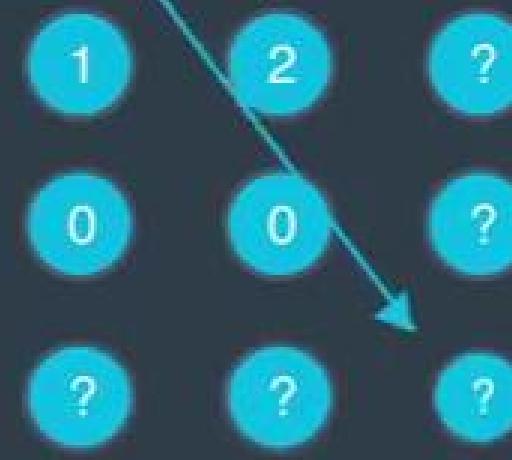
Filter

$$\begin{bmatrix} +1 & +1 \\ -1 & -1 \end{bmatrix}$$

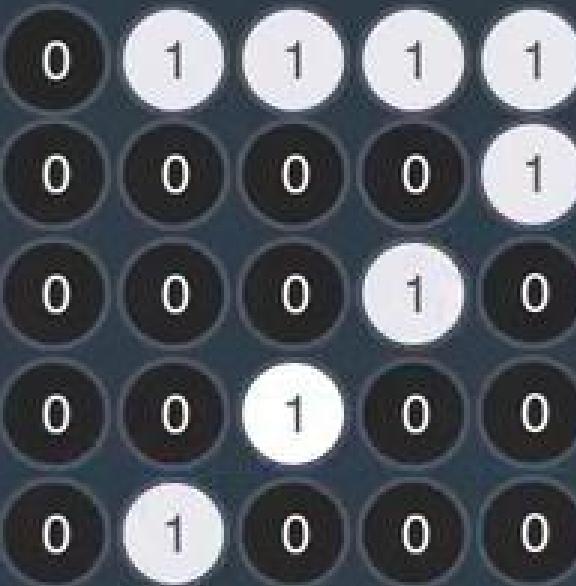
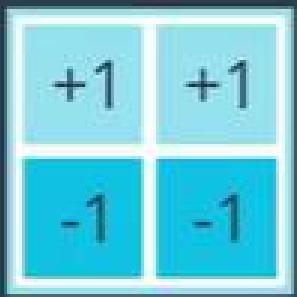


Input Layer

Convolutional
Layer

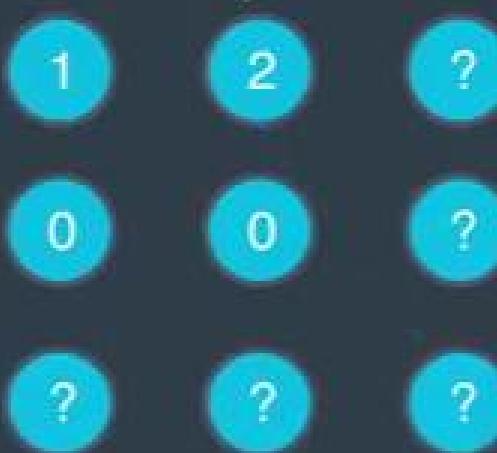


Filter

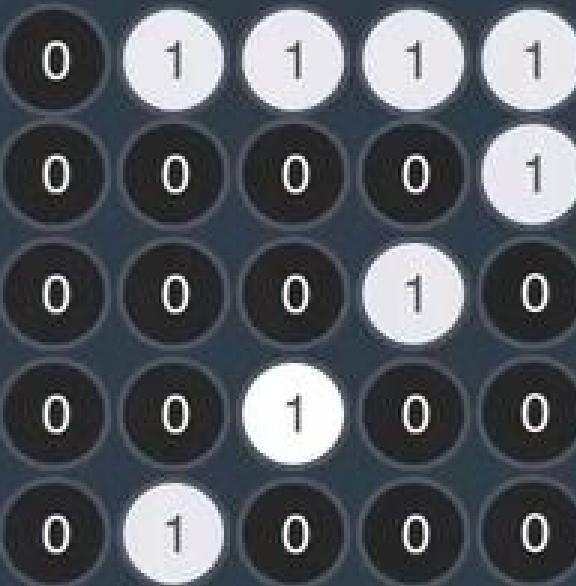
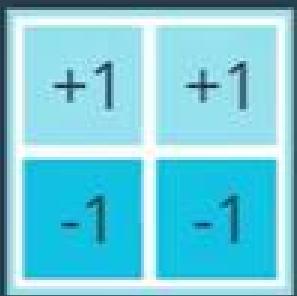


Input Layer

Convolutional
Layer

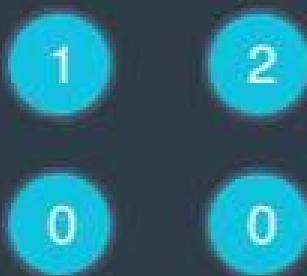


Filter

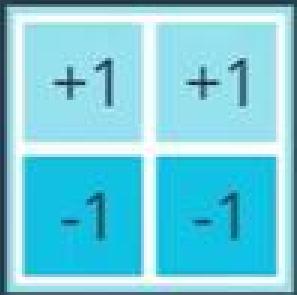


Input Layer

Convolutional
Layer

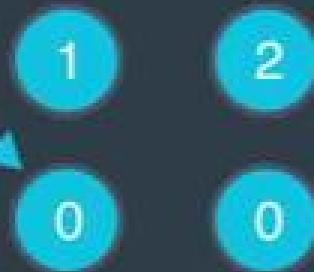


Filter

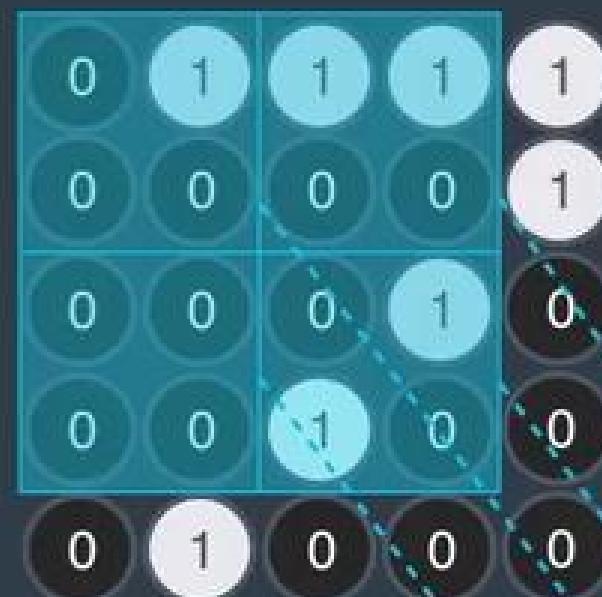
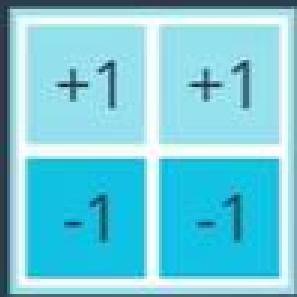


Input Layer

Convolutional
Layer



Filter



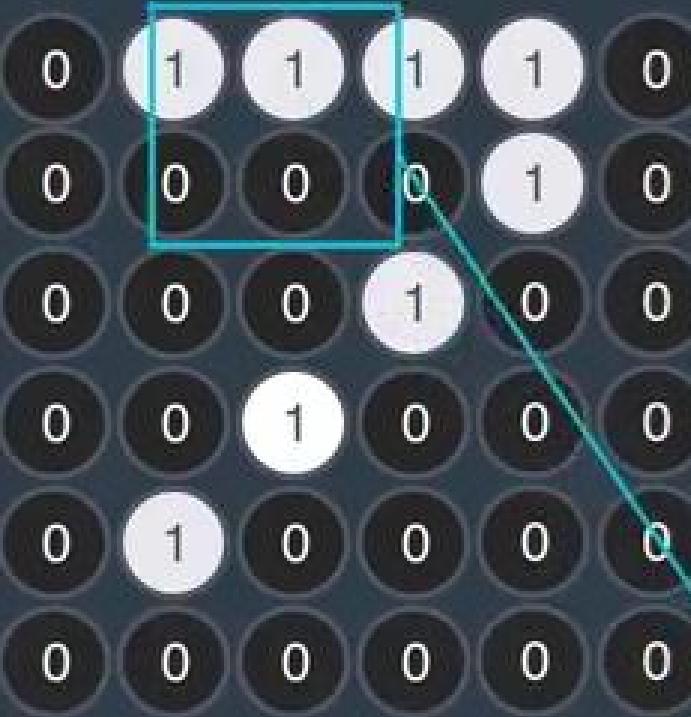
Input Layer

Convolutional
Layer



Filter

$$\begin{bmatrix} +1 & +1 \\ -1 & -1 \end{bmatrix}$$



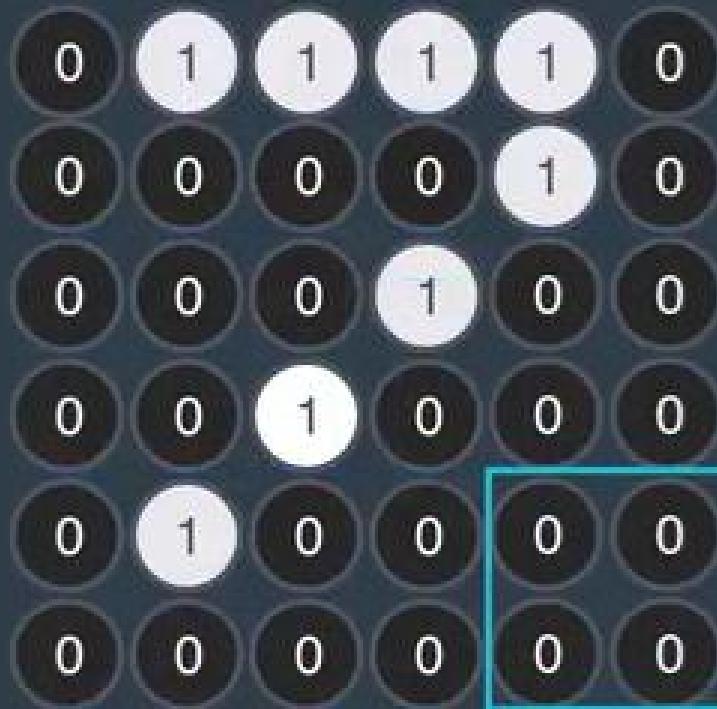
Input Layer

Convolutional
Layer

1

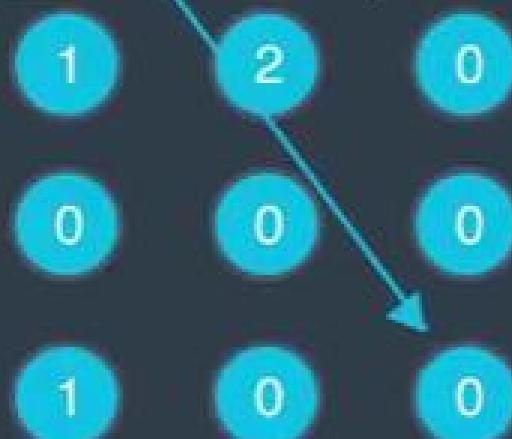
Filter

$$\begin{bmatrix} +1 & +1 \\ -1 & -1 \end{bmatrix}$$



Input Layer

Convolutional
Layer



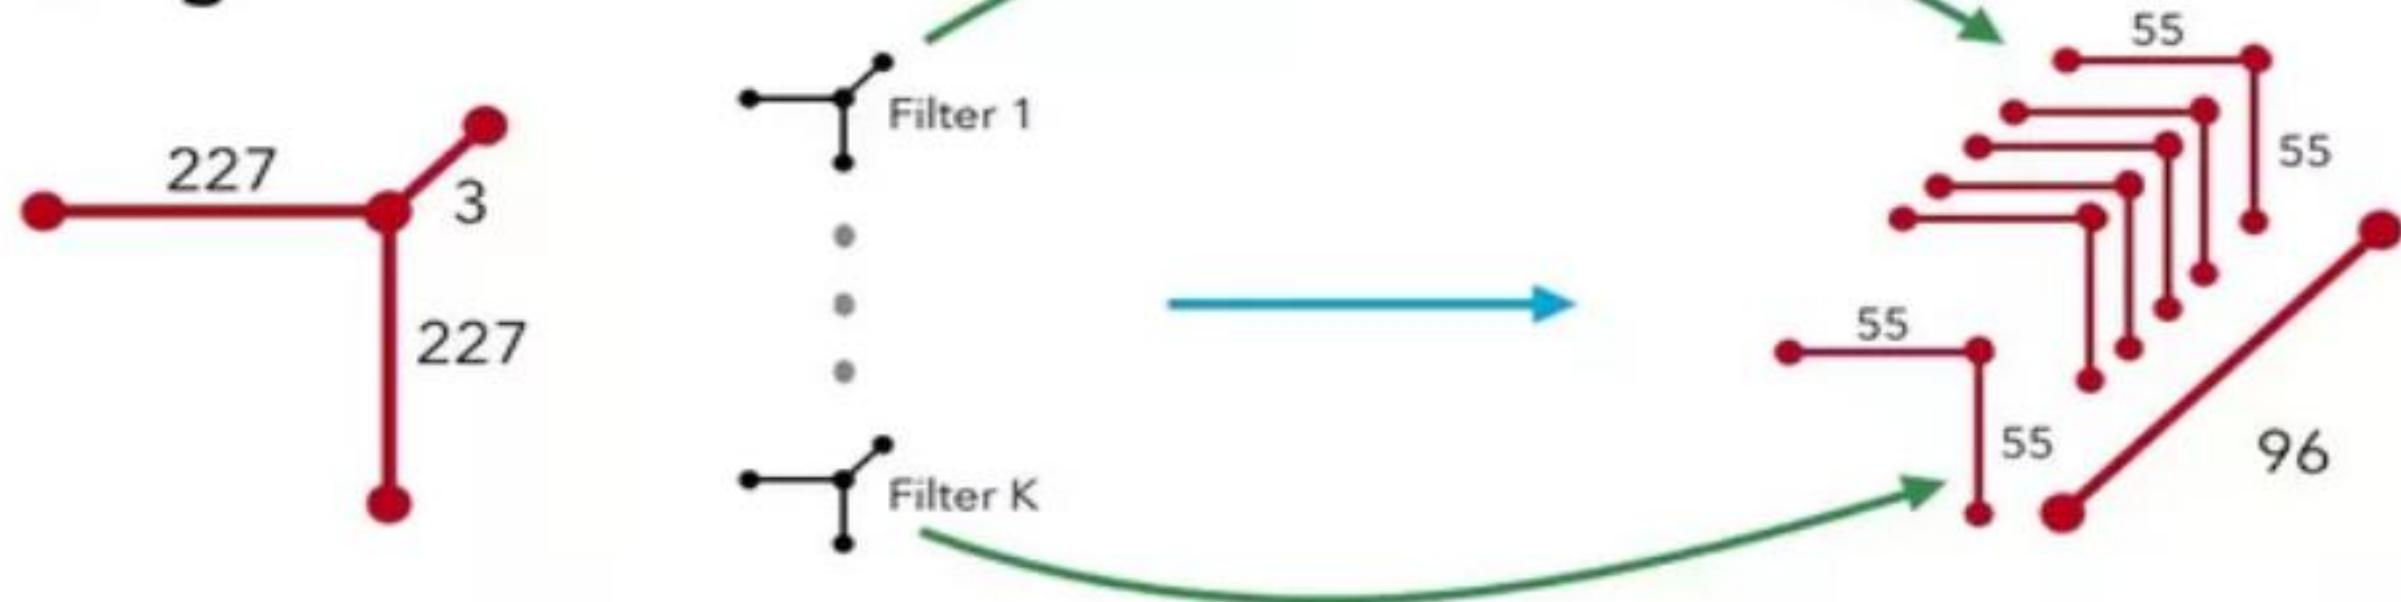
INPUT

CONV

CONV

CONV

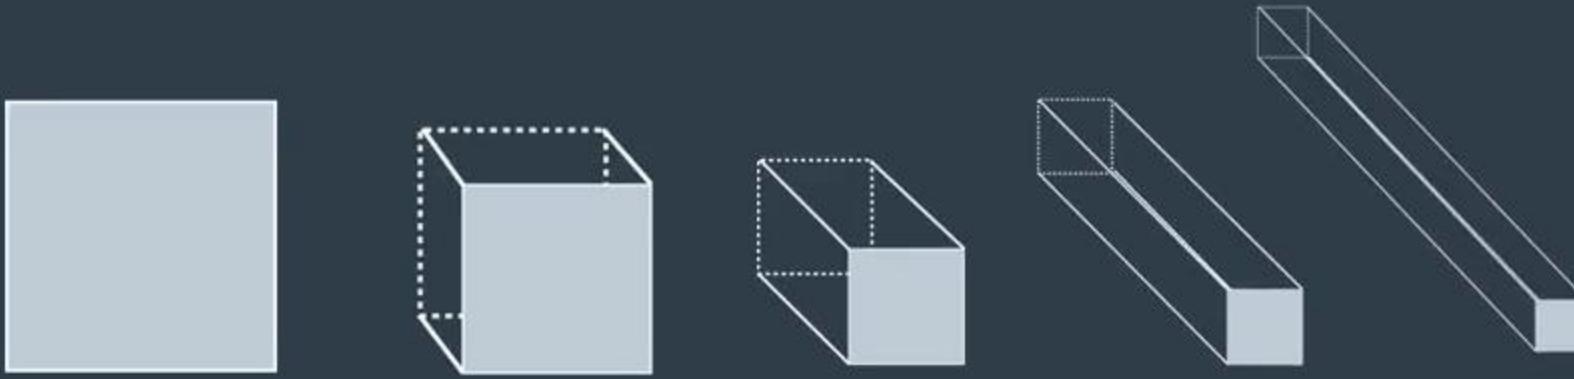
Parameter Sharing



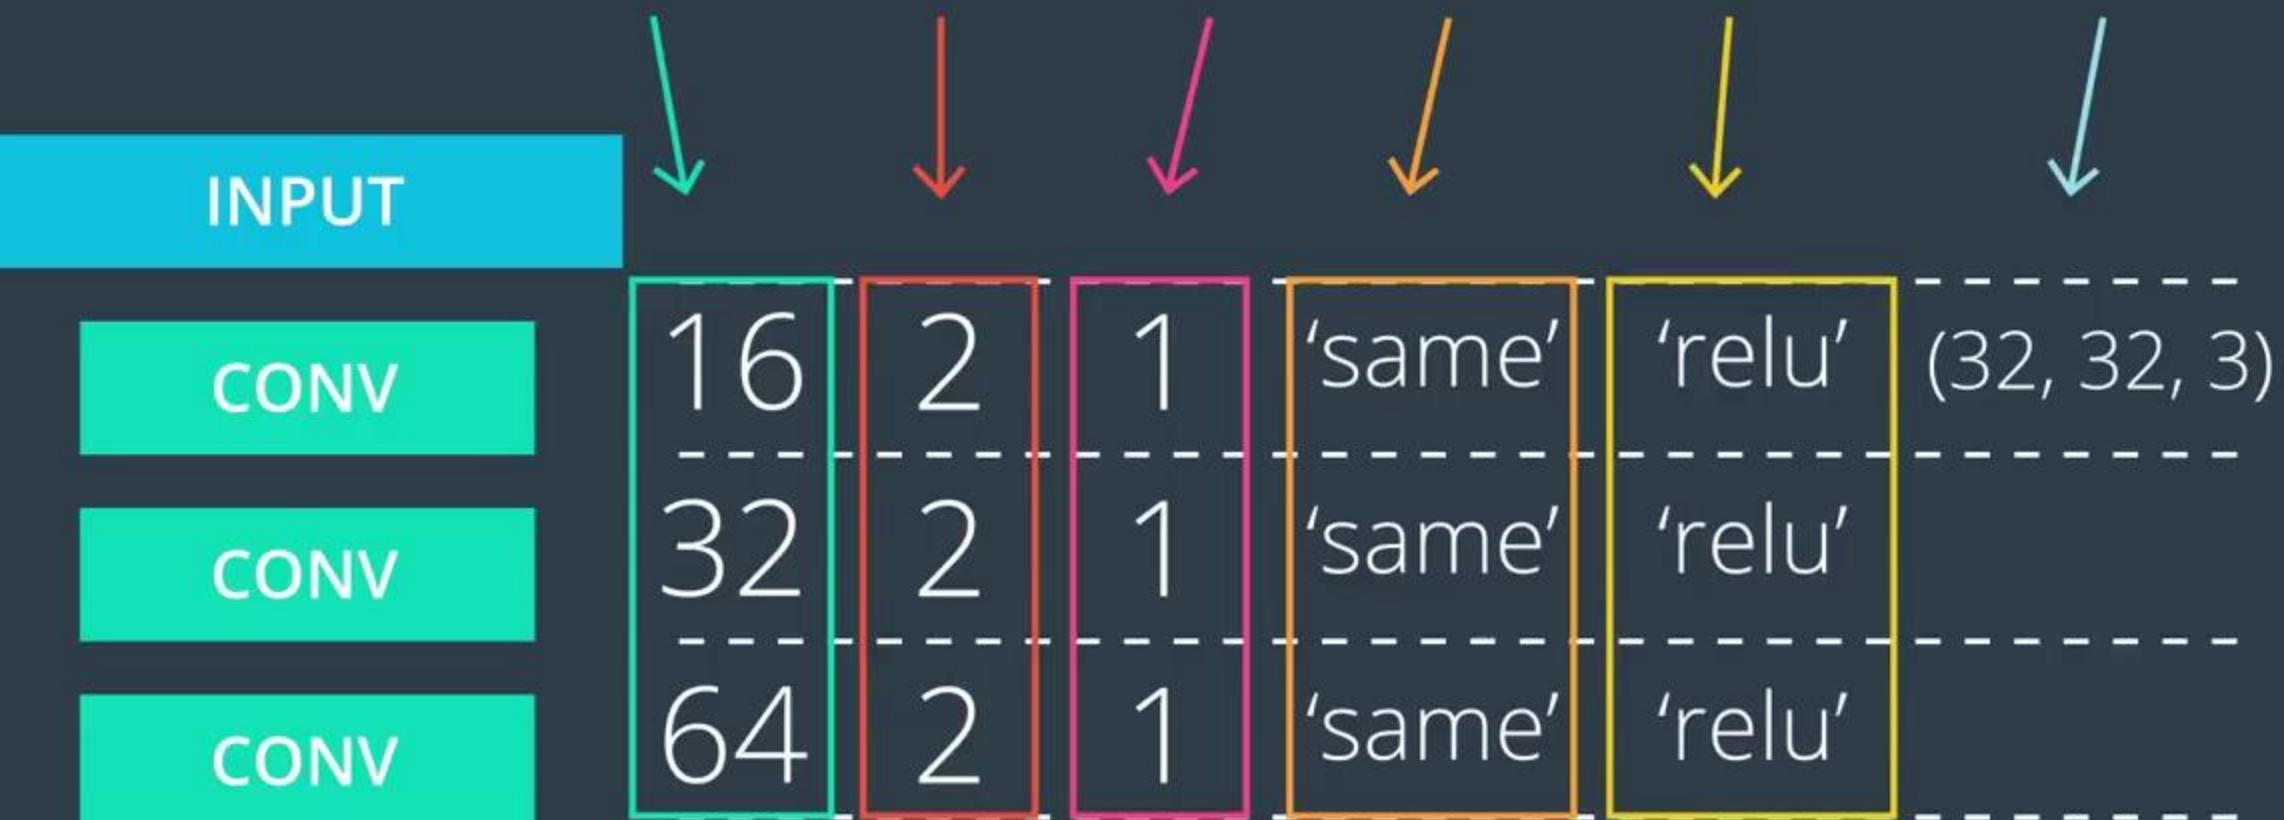
Each filter convolves on the input image and creates a 55×55 2D matrix. The K filters together create the output of the first conv layer which is $55 \times 55 \times K$. Each of the 55×55 neurons for each depth slice of the output from the first conv layer use the same weights ($11 \times 11 \times 3$). The intuition here is that each filter represents a feature and it makes sense to apply the same filter for the entire image using the same weights.

A CNN has multiple layers. Weight sharing happens across the receptive field of the neurons(filters) in a particular layer. Weights are the numbers within each filter. So essentially we are trying to learn a filter. These filters act on a certain receptive field/ small section of the image. When the filter moves through the image, the filter does not change. The idea being, if an edge is important to learn in a particular part of an image, it is important in other parts of the image too.

Convolution Layers increase the image depth and decrease image Width and Height it also converts Spatial Values to Feature Values



`Convolution2D(filters, kernel_size, strides, padding, activation, input_shape)`



```
from keras.models import Sequential
from keras.layers import Convolution2D

model = Sequential()
model.add(Convolution2D(filters=16, kernel_size=2, padding='same', activation='relu',
                       input_shape=(32, 32, 3)))
model.add(Convolution2D(filters=32, kernel_size=2, padding='same', activation='relu'))
model.add(Convolution2D(filters=64, kernel_size=2, padding='same', activation='relu'))

model.summary()
```

Using TensorFlow backend.

Layer (type)	Output Shape	Param #
<hr/>		
conv2d_1 (Conv2D)	(None, 32, 32, 16)	208
conv2d_2 (Conv2D)	(None, 32, 32, 32)	2080
conv2d_3 (Conv2D)	(None, 32, 32, 64)	8256
<hr/>		
Total params:	10,544.0	
Trainable params:	10,544	
Non-trainable params:	0.0	

Define a convolutional layer in PyTorch

```
self.conv1 = nn.Conv2d(3, 16,  
                    kernel_size, stride = 1, padding = 0)
```

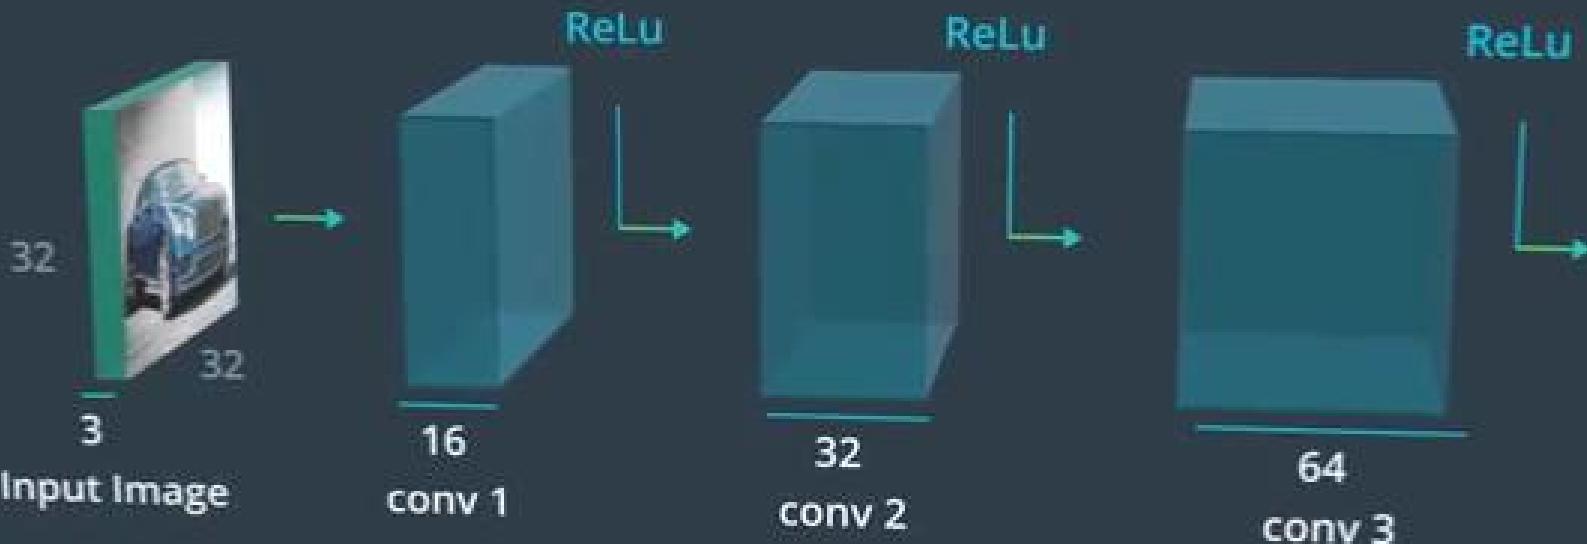


3



Input Image

convolutional
layer



```
self.conv1 = nn.Conv2d(3, 16, 3, padding = 1)
self.conv2 = nn.Conv2d(16, 32, 3, padding = 1)
self.conv3= nn.Conv2d(32, 64, 3, padding = 1)
```



INPUT

CONV

POOL

CONV

POOL

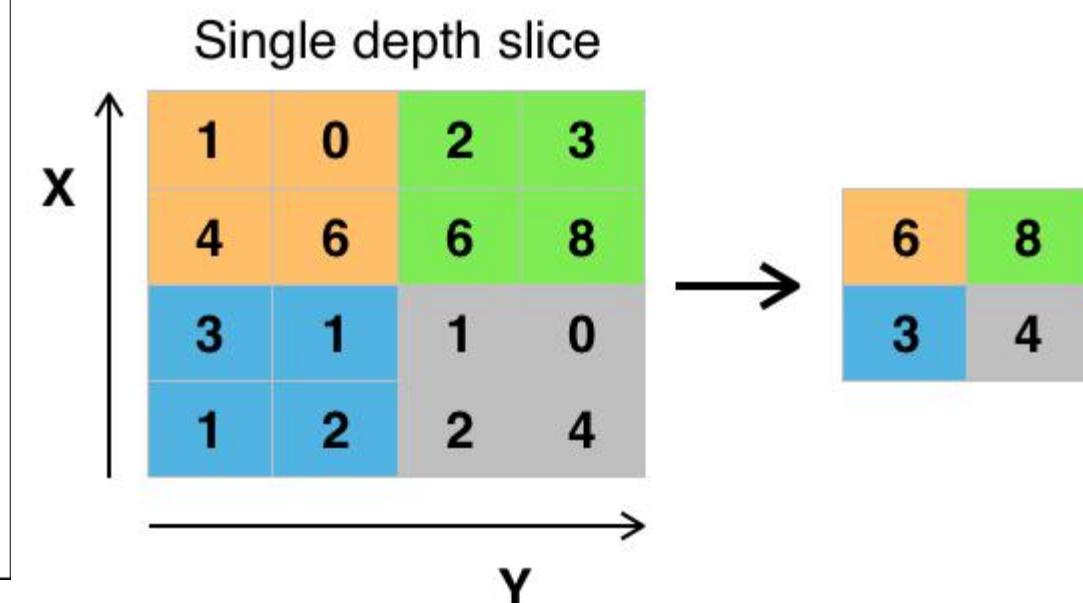
CONV

POOL

POOL layer:

Pool Layer performs a function to reduce the spatial dimensions of the input, and the computational complexity of our model. And it also controls overfitting. It operates independently on every depth slice of the input. There are different functions such as Max pooling, average pooling, or L2-norm pooling. However, Max pooling is the most used type of pooling which only takes the most important part (the value of the brightest pixel) of the input volume.

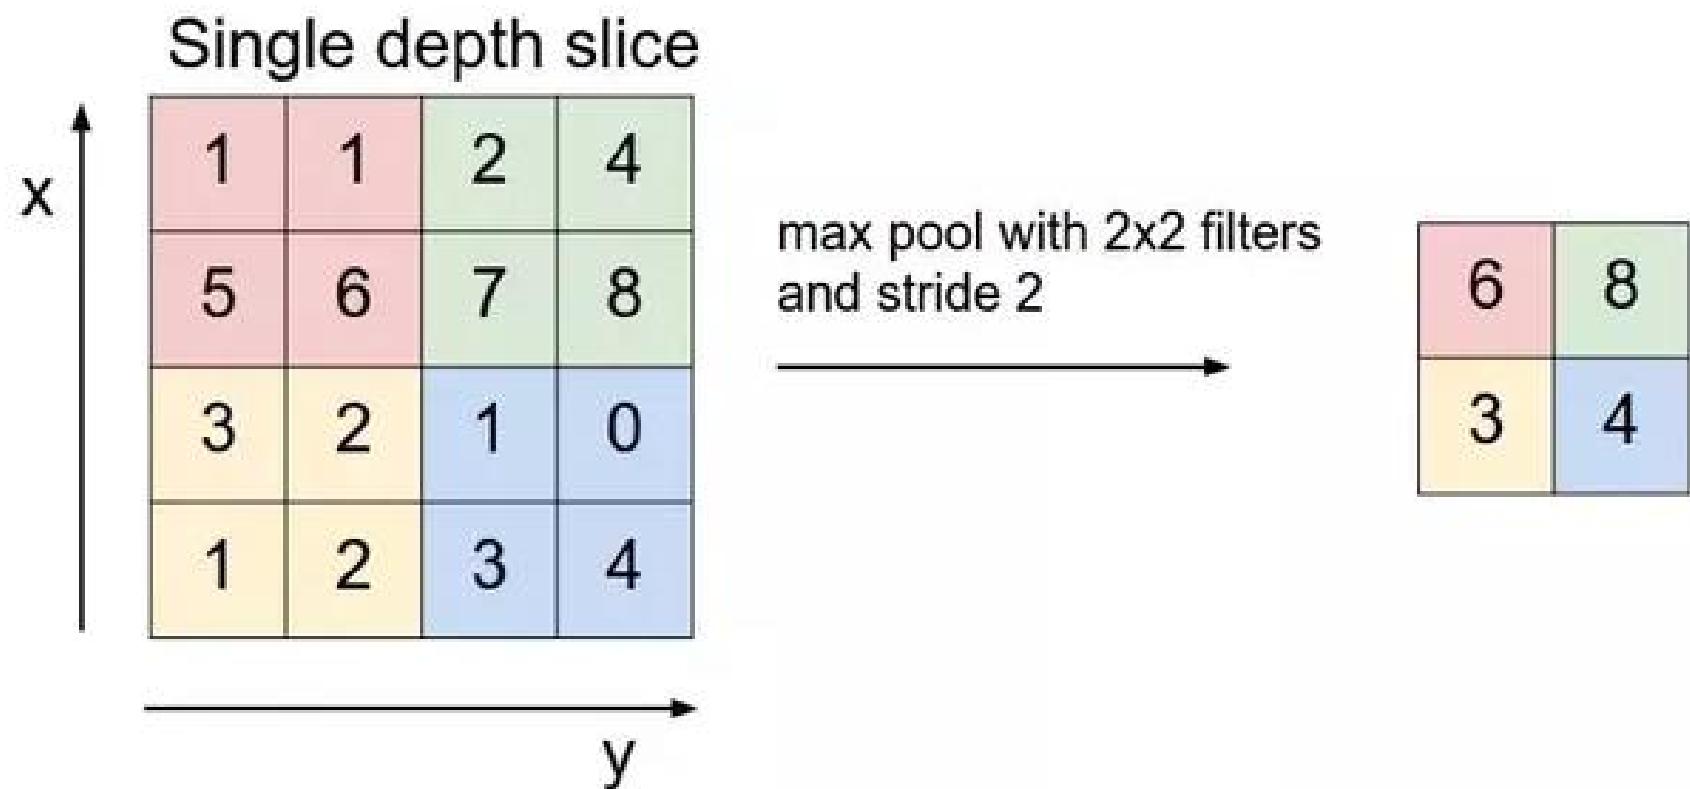
Example of a Max pooling with 2x2 filter and stride = 2. So, for each of the windows, max pooling takes the max value of the 4 pixels.



Pooling Layer

Pooling layers section would reduce the number of parameters when the images are too large. Spatial pooling also called subsampling or downsampling which reduces the dimensionality of each map but retains the important information. Spatial pooling can be of different types:

- Max Pooling
- Average Pooling
- Sum Pooling



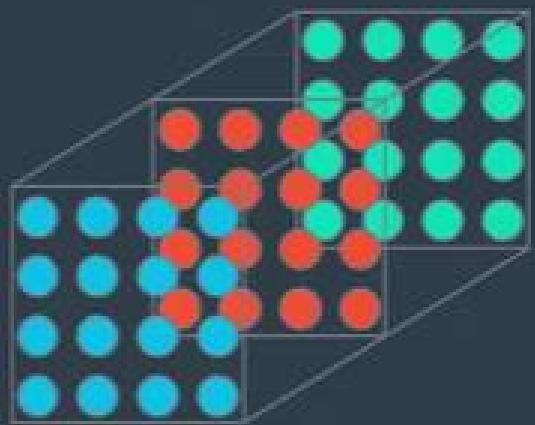
- Pool layer doesn't have parameters (the weights and biases of the neurons), and no zero padding, but it has two hyperparameters: Filter (F) and Stride (S). More generally, having the input $W_1 \times H_1 \times D_1$, the pooling layer produces a volume of size $W_2 \times H_2 \times D_2$ where:

- $W_2 = (W_1 - F)/S + 1$
- $H_2 = (H_1 - F)/S + 1$
- $D_2 = D_1$

A common form of a Max pooling is filters of size 2×2 applied with a stride of 2. The Pooling sizes with larger filters are too destructive and they usually lead to worse performance.

Many people don't like using a pooling layer because it throws away information and they replace it by a Conv layer with increased stride once in a while.

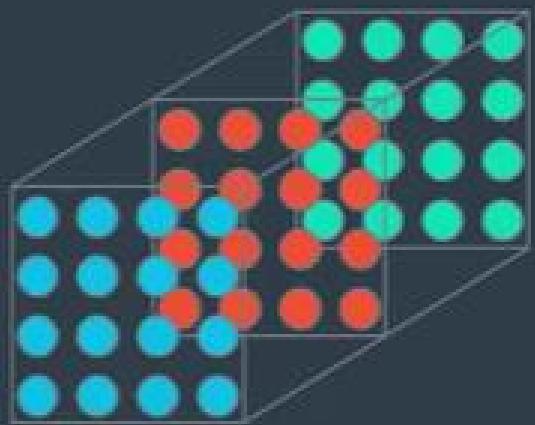
Getting rid of pooling. Many people dislike the pooling operation and think that we can get away without it. For example, [Striving for Simplicity: The All Convolutional Net](#) proposes to discard the pooling layer in favor of architecture that only consists of repeated CONV layers. To reduce the size of the representation they suggest using larger stride in CONV layer once in a while. Discarding pooling layers has also been found to be important in training good generative models, such as variational autoencoders (VAEs) or generative adversarial networks (GANs). It seems likely that future architectures will feature very few to no pooling layers.



Convolutional Layer

1	9	6	4
5	4	7	8
5	1	2	9
6	7	6	0
9	1	7	4
5	6	3	0
1	2	5	4
0	8	9	0
7	6	9	1
5	2	0	4
5	8	3	9
0	2	2	1

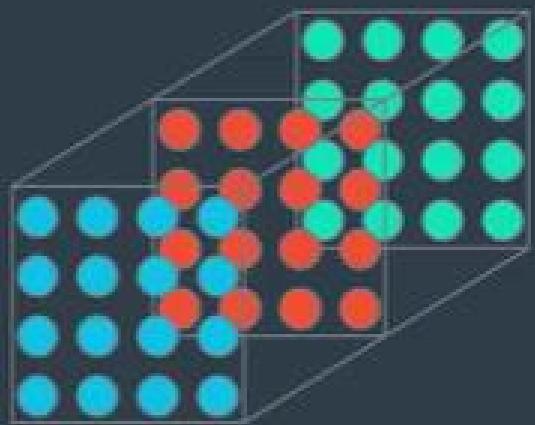
Max Pooling Layer
WINDOW SIZE: 2X2
STRIDE: 2



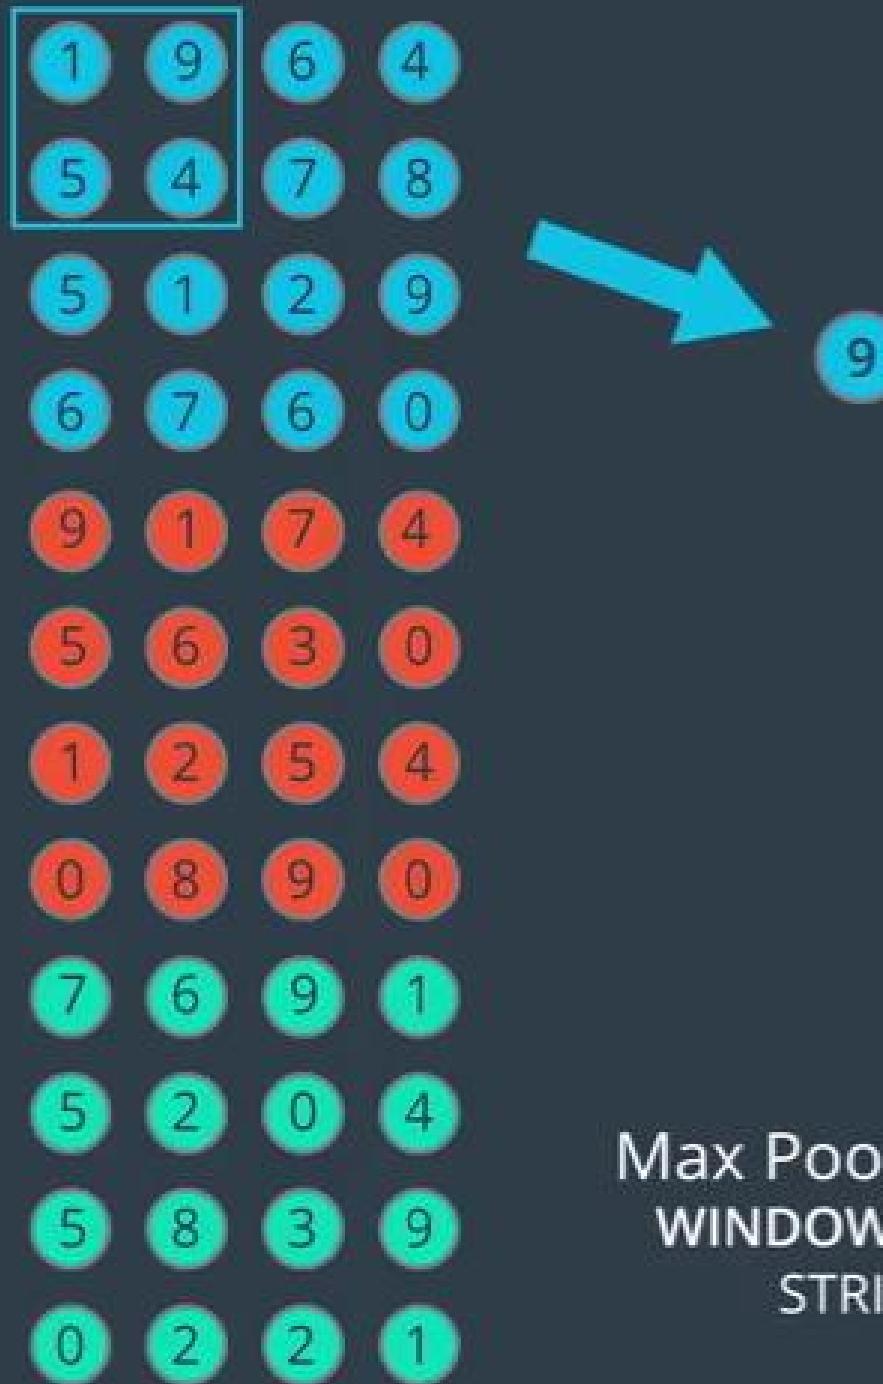
Convolutional Layer

1	9	6	4
5	4	7	8
5	1	2	9
6	7	6	0
9	1	7	4
5	6	3	0
1	2	5	4
0	8	9	0
7	6	9	1
5	2	0	4
5	8	3	9
0	2	2	1

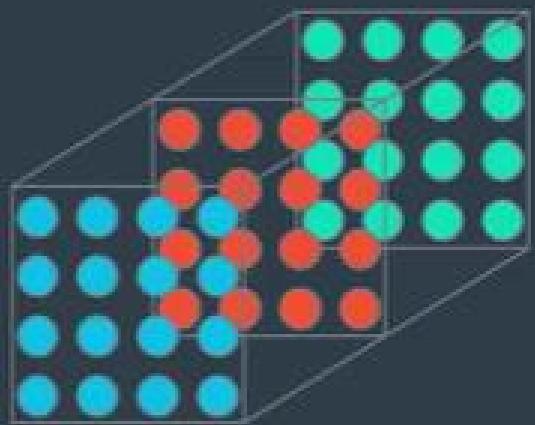
Max Pooling Layer
WINDOW SIZE: 2x2
STRIDE: 2



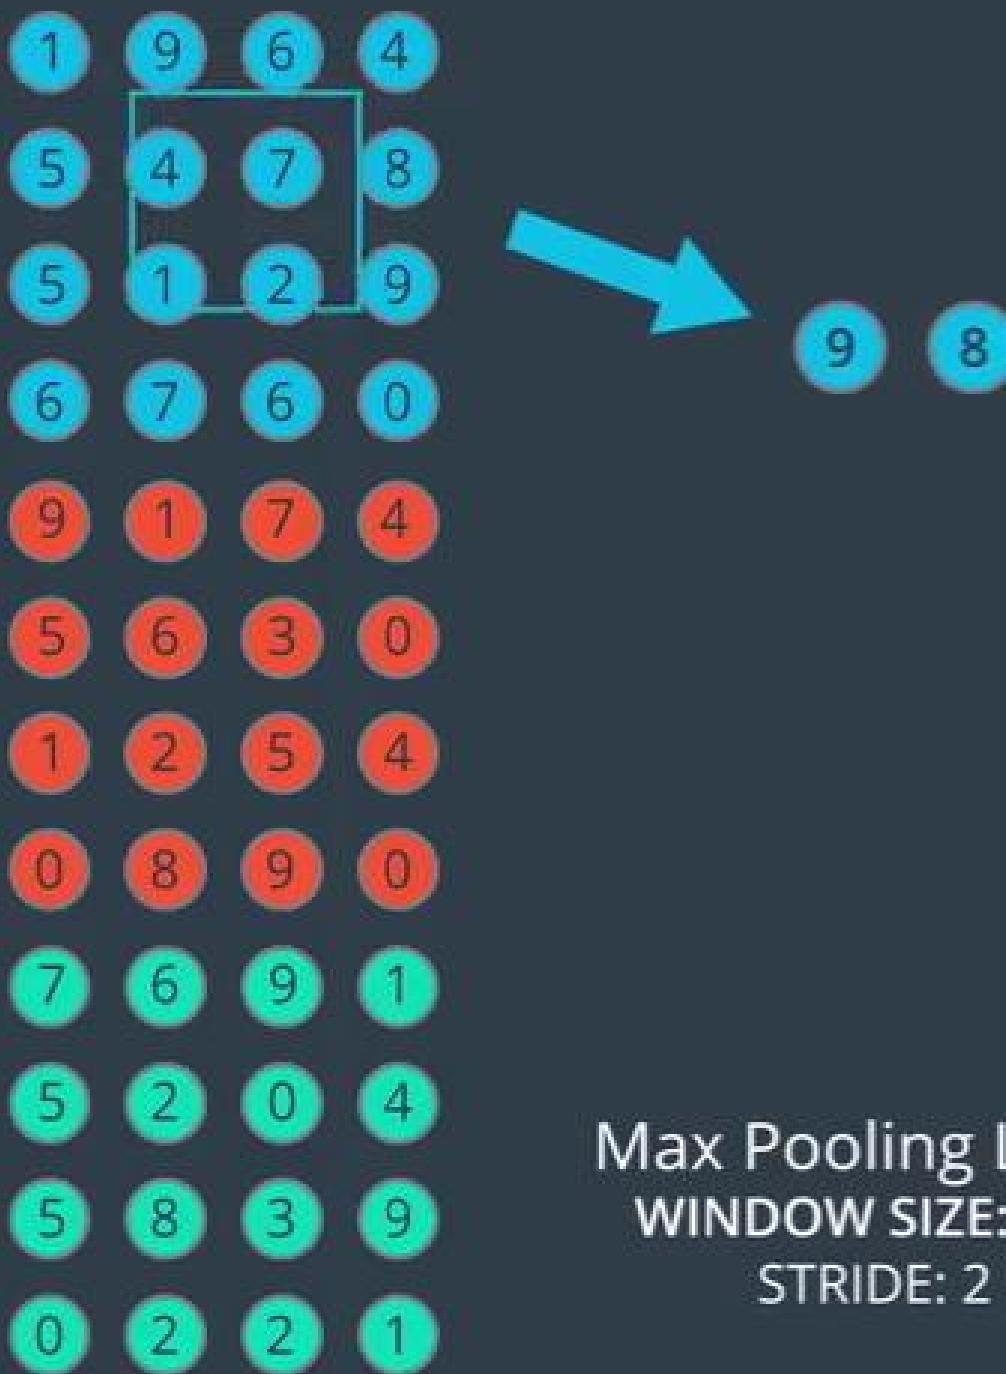
Convolutional Layer

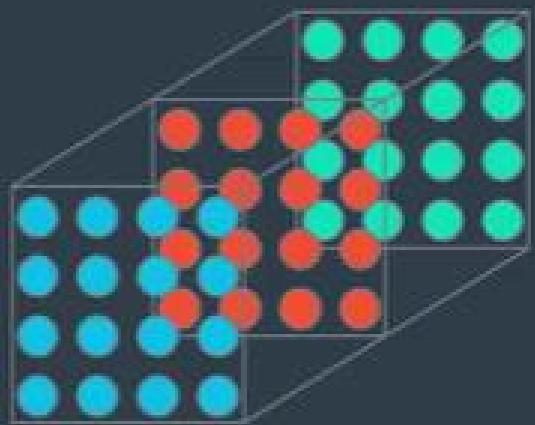


Max Pooling Layer
WINDOW SIZE: 2X2
STRIDE: 2

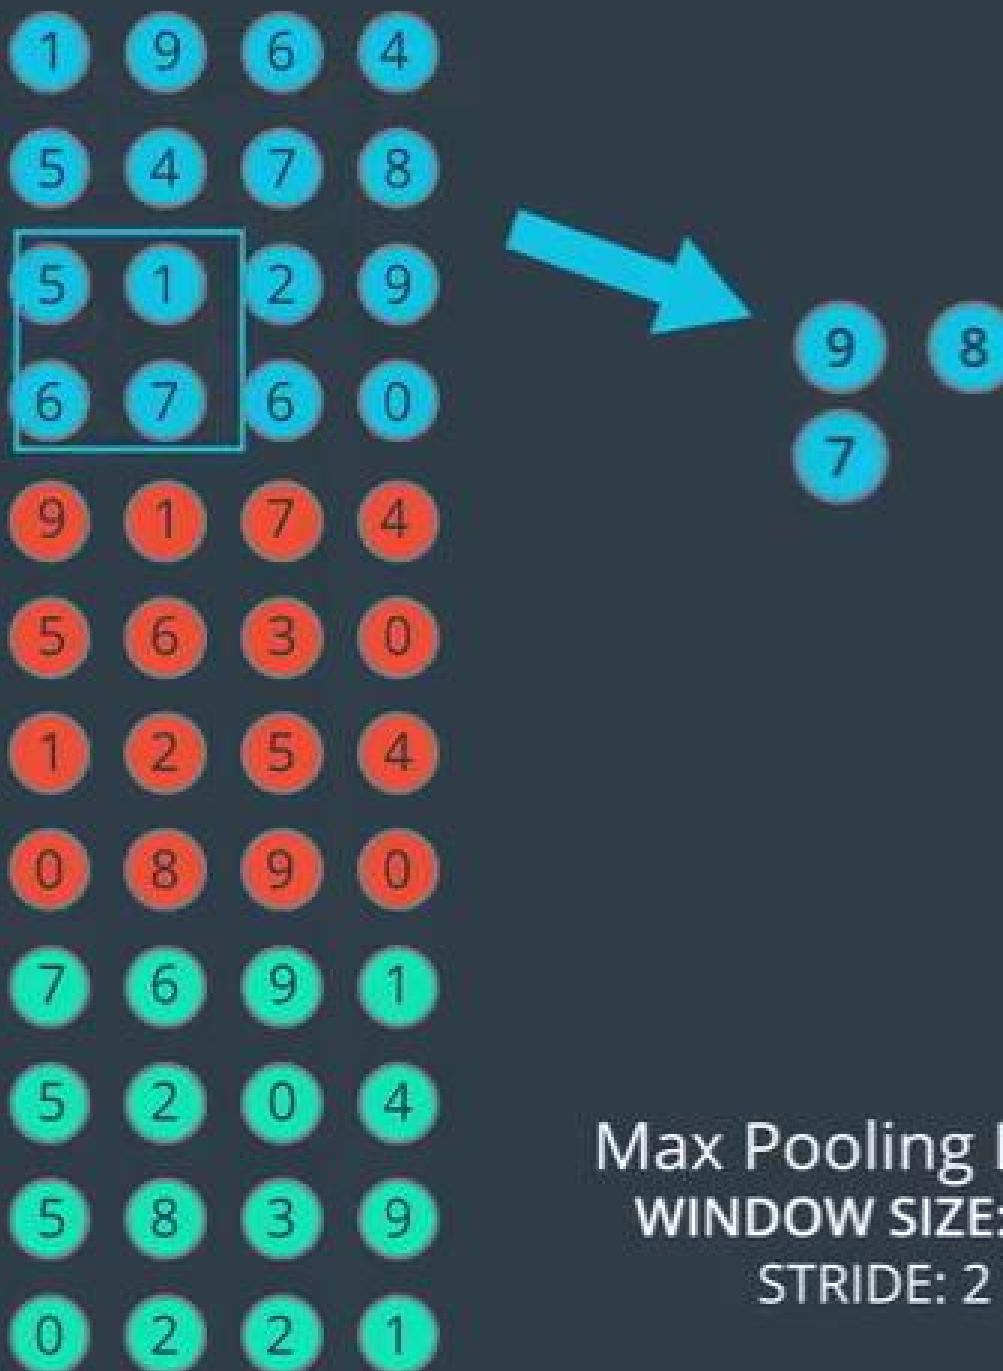


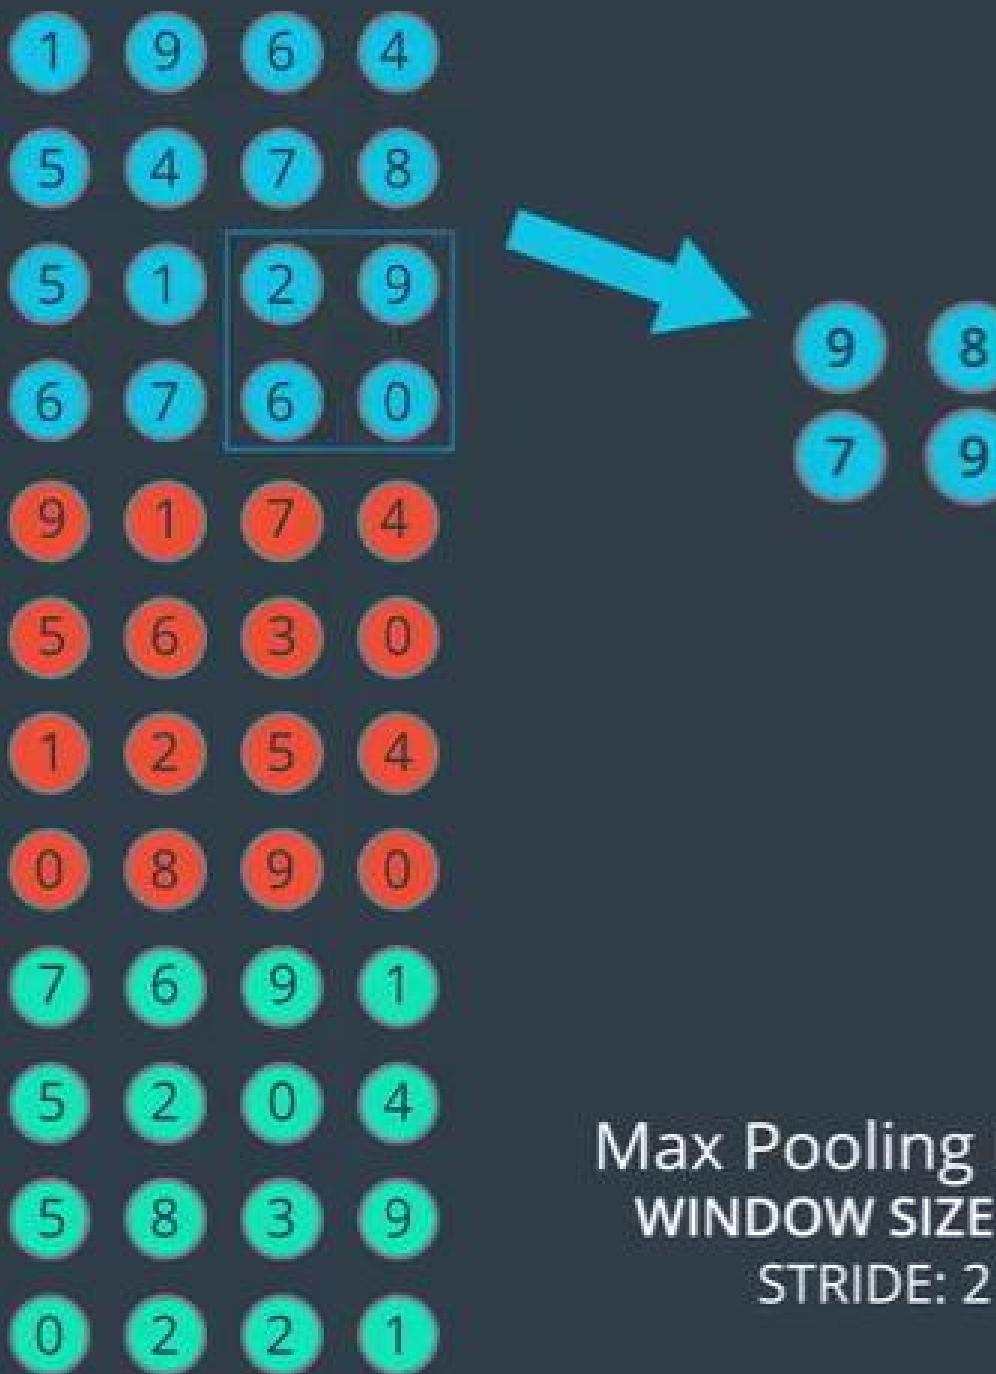
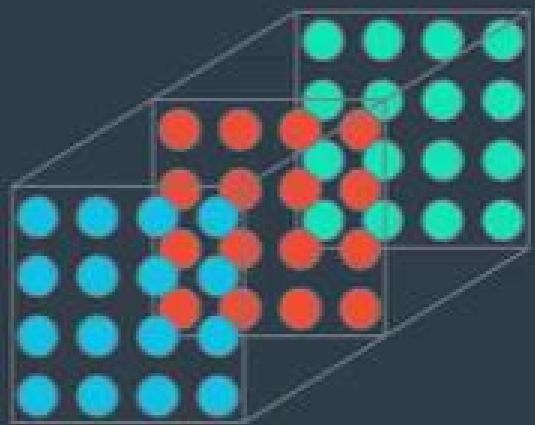
Convolutional Layer

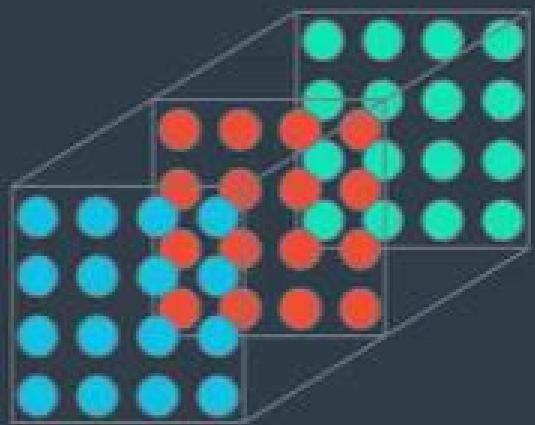




Convolutional Layer

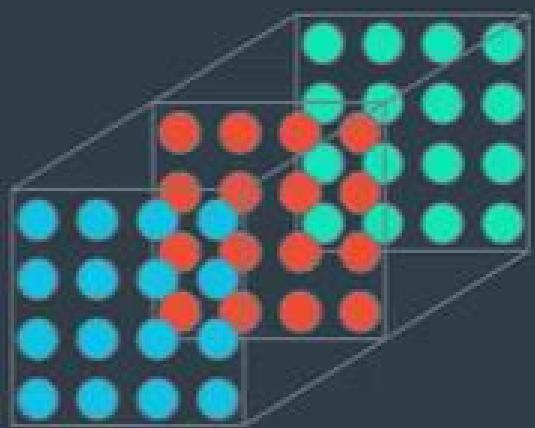




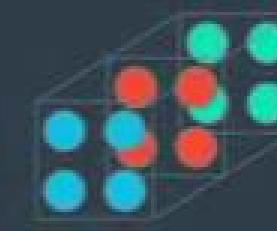


Convolutional Layer



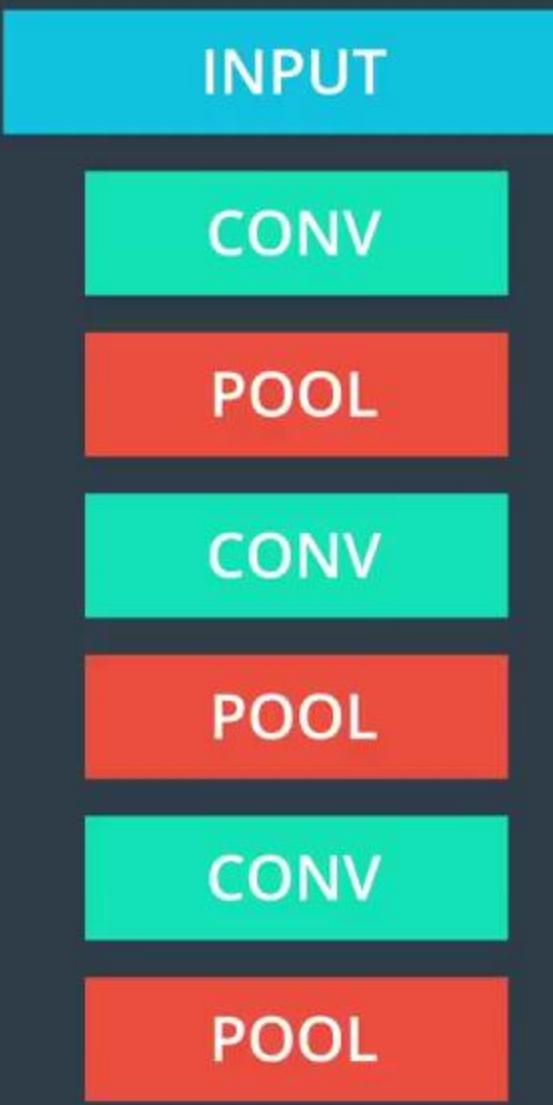


Convolutional Layer



Max Pooling Layer
WINDOW SIZE: 2X2
STRIDE: 2

MaxPooling2D(pool_size, stride, padding)



(default ok)

```
from keras.models import Sequential
from keras.layers import Convolution2D, MaxPooling2D

model = Sequential()
model.add(Convolution2D(filters=16, kernel_size=2, padding='same', activation='relu',
                       input_shape=(32, 32, 3)))
model.add(MaxPooling2D(pool_size=2))
model.add(Convolution2D(filters=32, kernel_size=2, padding='same', activation='relu'))
model.add(MaxPooling2D(pool_size=2))
model.add(Convolution2D(filters=64, kernel_size=2, padding='same', activation='relu'))
model.add(MaxPooling2D(pool_size=2))

model.summary()
```

Layer (type)	Output Shape	Param #
conv2d_4 (Conv2D)	(None, 32, 32, 16)	208
max_pooling2d_1 (MaxPooling2D)	(None, 16, 16, 16)	0
conv2d_5 (Conv2D)	(None, 16, 16, 32)	2080
max_pooling2d_2 (MaxPooling2D)	(None, 8, 8, 32)	0
conv2d_6 (Conv2D)	(None, 8, 8, 64)	8256
max_pooling2d_3 (MaxPooling2D)	(None, 4, 4, 64)	0
<hr/>		
Total params:	10,544.0	
Trainable params:	10,544.0	
Non-trainable params:	0.0	

```
model = tf.keras.models.Sequential([
    tf.keras.layers.Conv2D(64, (3,3), activation='relu',
                         input_shape=(28, 28, 1)),
    tf.keras.layers.MaxPooling2D(2, 2),
    tf.keras.layers.Conv2D(64, (3,3), activation='relu'),
    tf.keras.layers.MaxPooling2D(2,2),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(128, activation='relu'),
    tf.keras.layers.Dense(10, activation='softmax')
])
```

Layer (type)	Output Shape	Param #
conv2d_12 (Conv2D)	(None, 26, 26, 64)	640
max_pooling2d_12 (MaxPooling)	(None, 13, 13, 64)	0
conv2d_13 (Conv2D)	(None, 11, 11, 64)	36928
max_pooling2d_13 (MaxPooling)	(None, 5, 5, 64)	0
flatten_5 (Flatten)	(None, 1600)	0
dense_10 (Dense)	(None, 128)	204928
dense_11 (Dense)	(None, 10)	1290





Input Image



Input Image



CNN



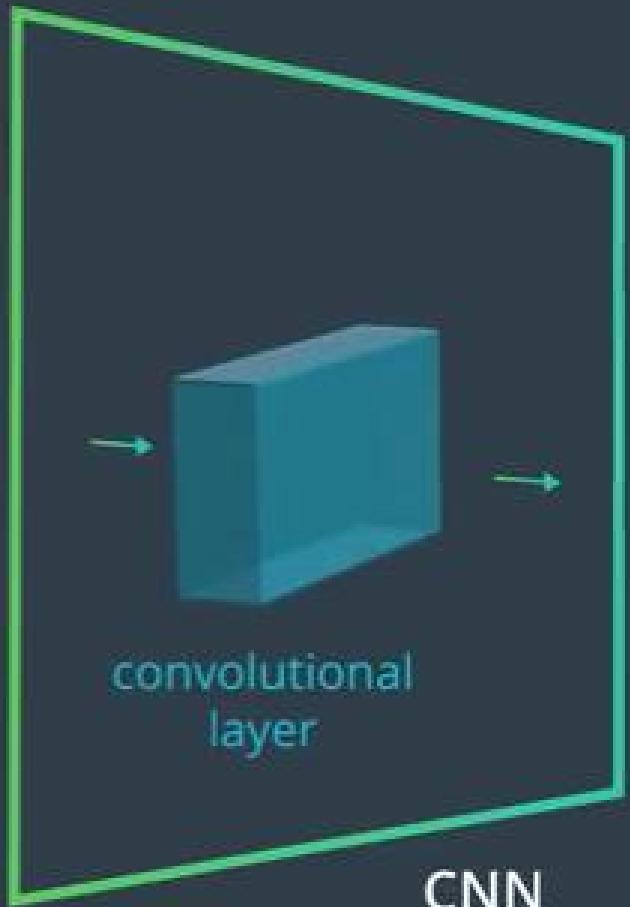
Input Image

convolutional
layer

CNN



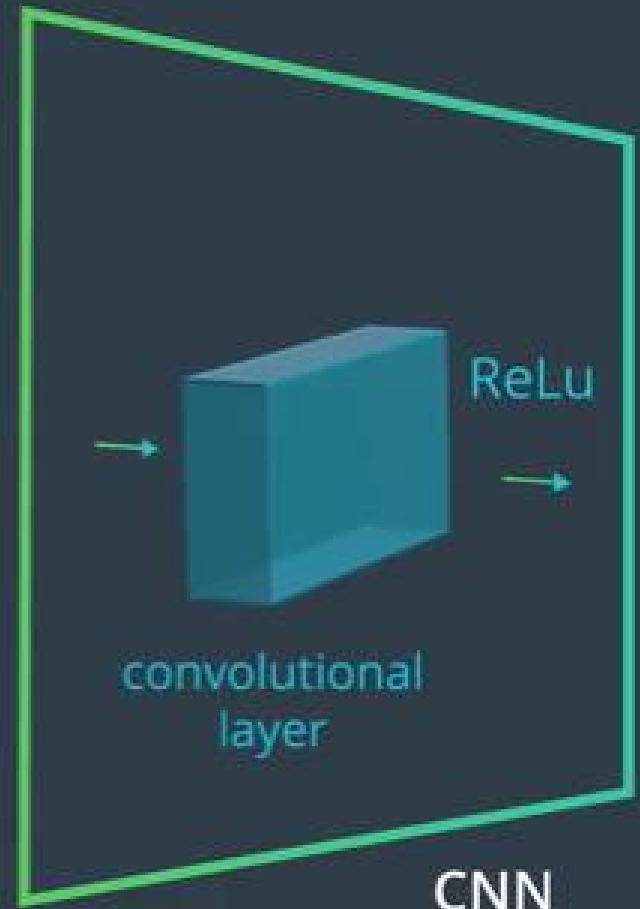
Input Image



CNN

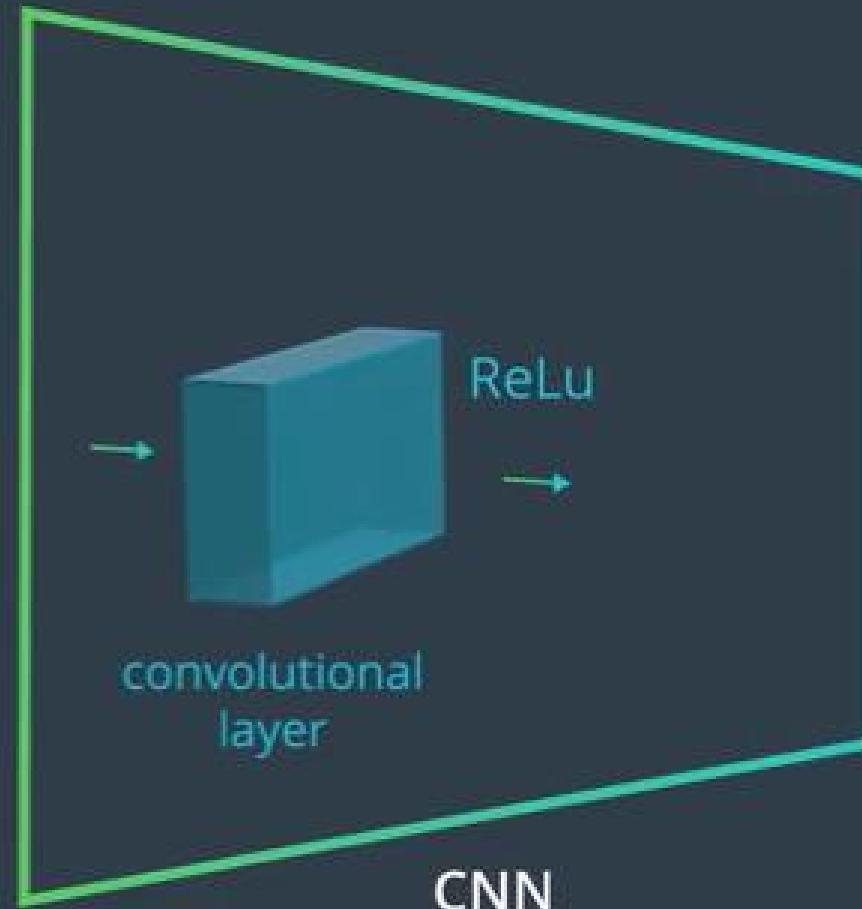


Input Image



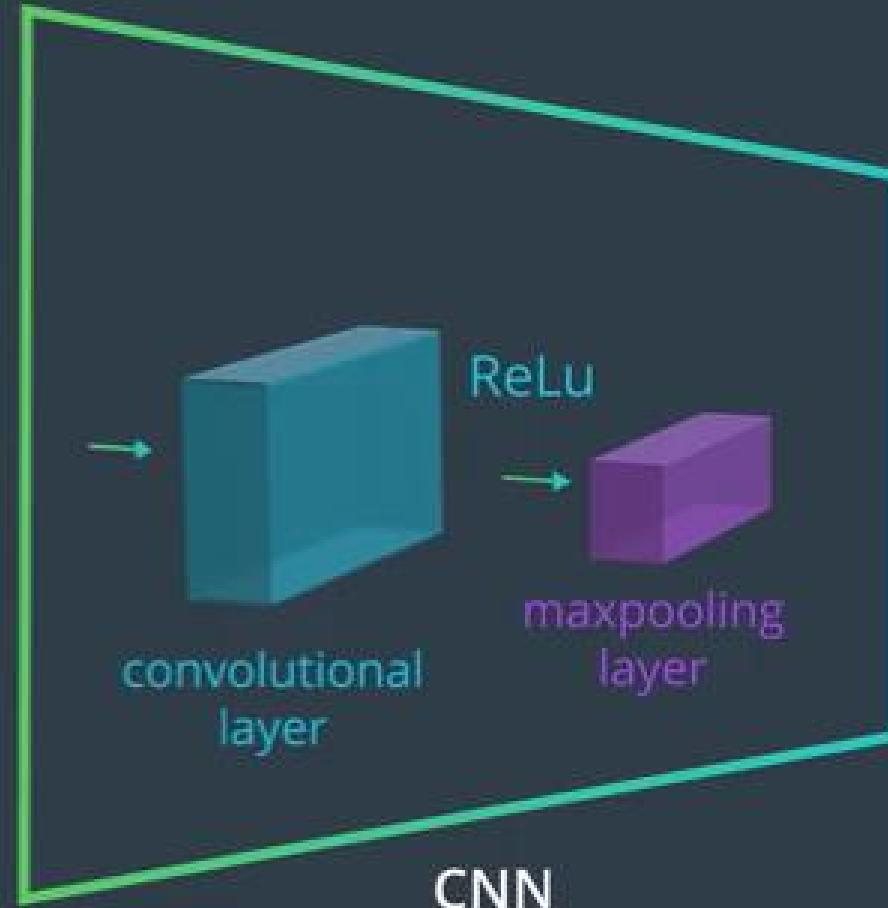


Input Image



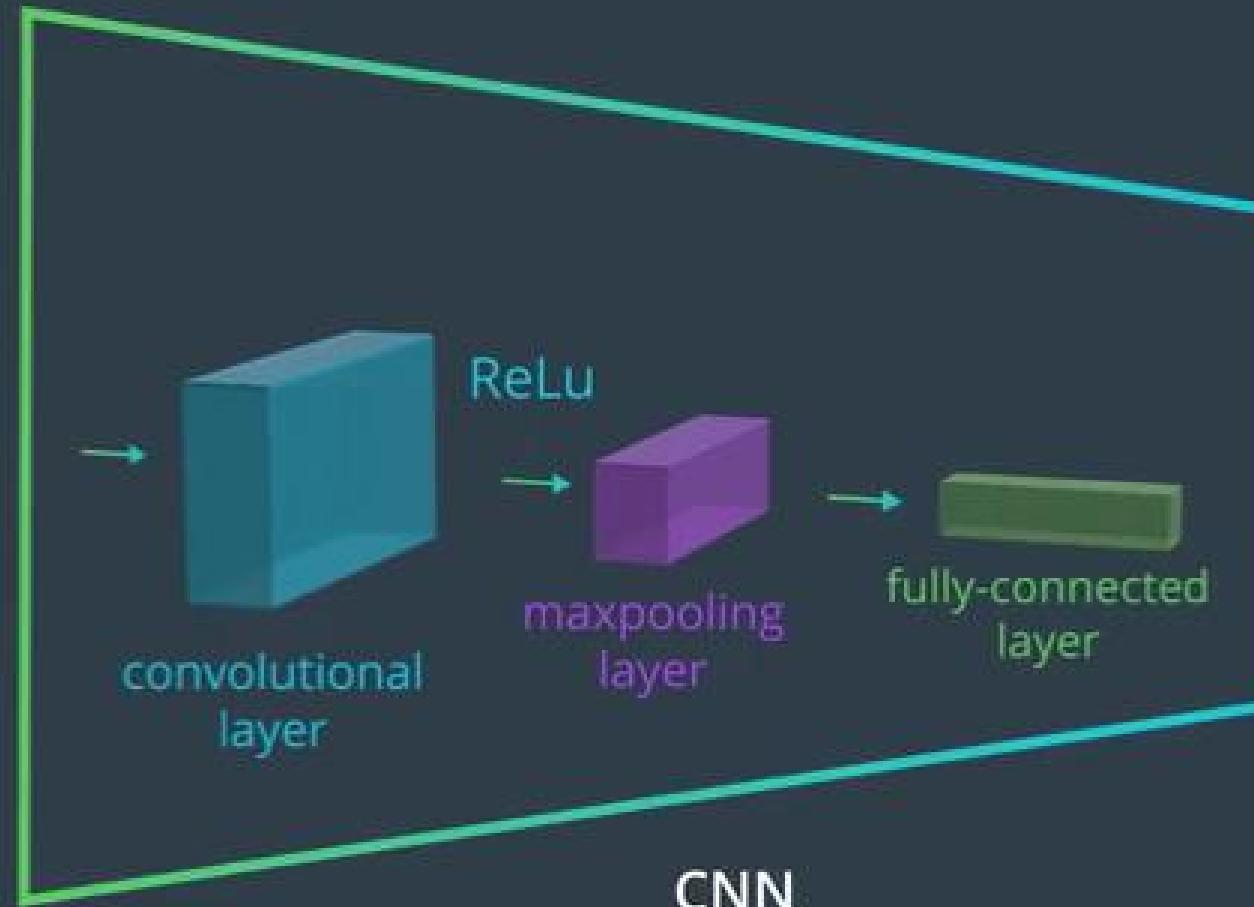


Input Image





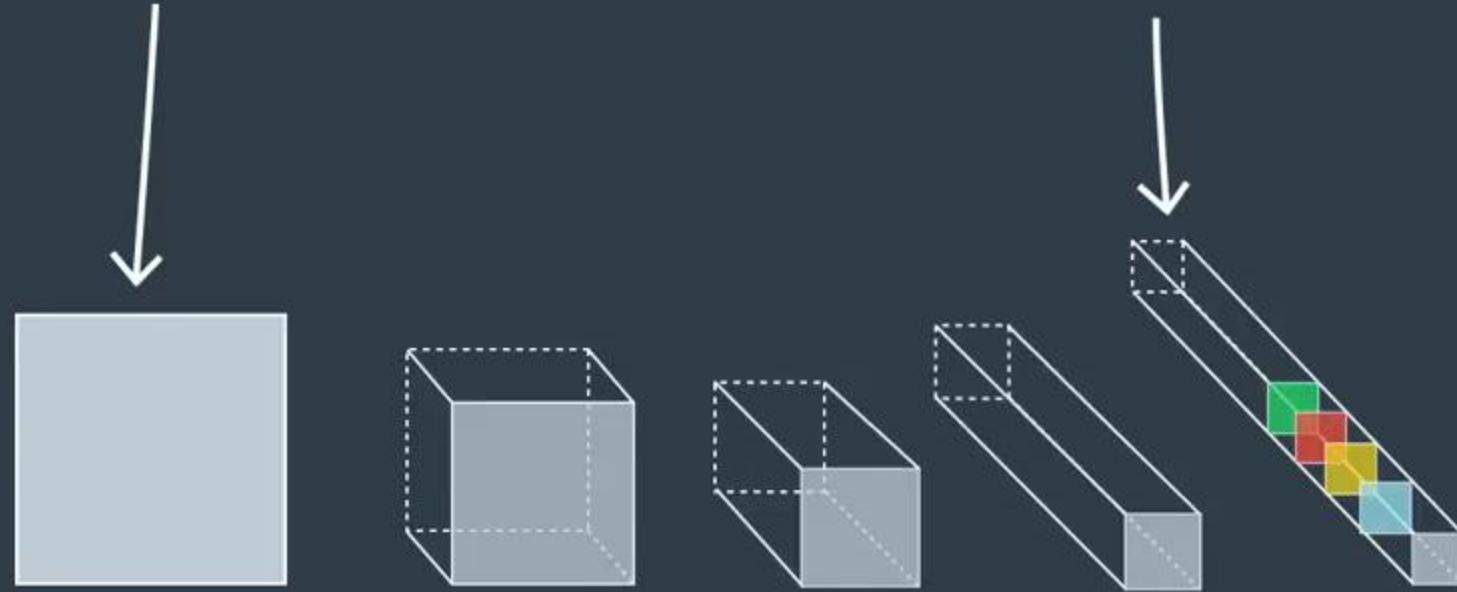
Input Image



CAR
Predicted
Class

Contains a ton
of spatial information

Contains no
spatial information



Wheels? No
Eyes? Yes
Legs? Yes
Tails? Yes



Fully Connected
Layer(s)

Wheels? Yes

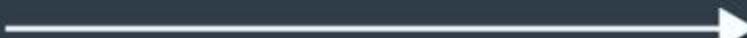
Eyes? No

Legs? No

Tails? No



$\text{prob(car)} = .99$
 $\text{prob(dog)} = .01$

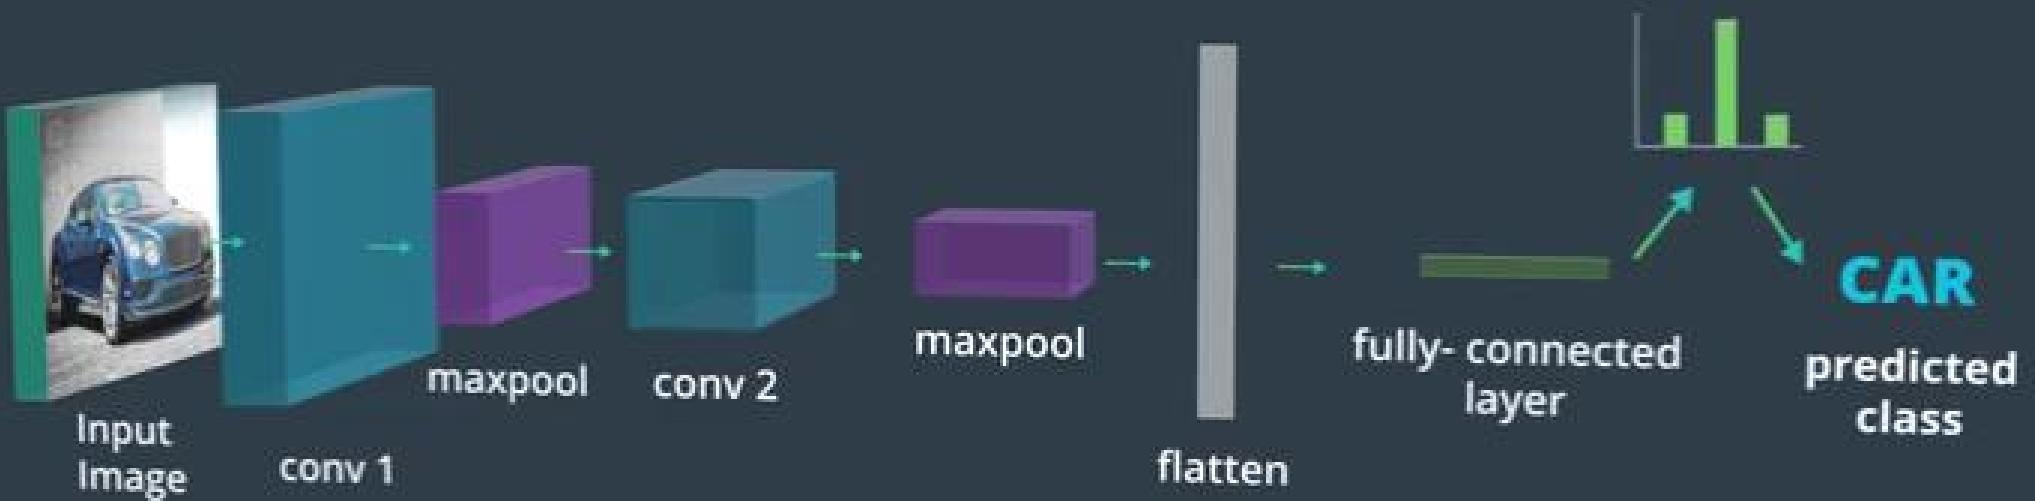


Fully Connected
Layer(s)

Wheels? No
Eyes? Yes
Legs? Yes
Tails? Yes



$\text{prob(car)} = .01$
 $\text{prob(dog)} = .99$



extracting more and more complex features

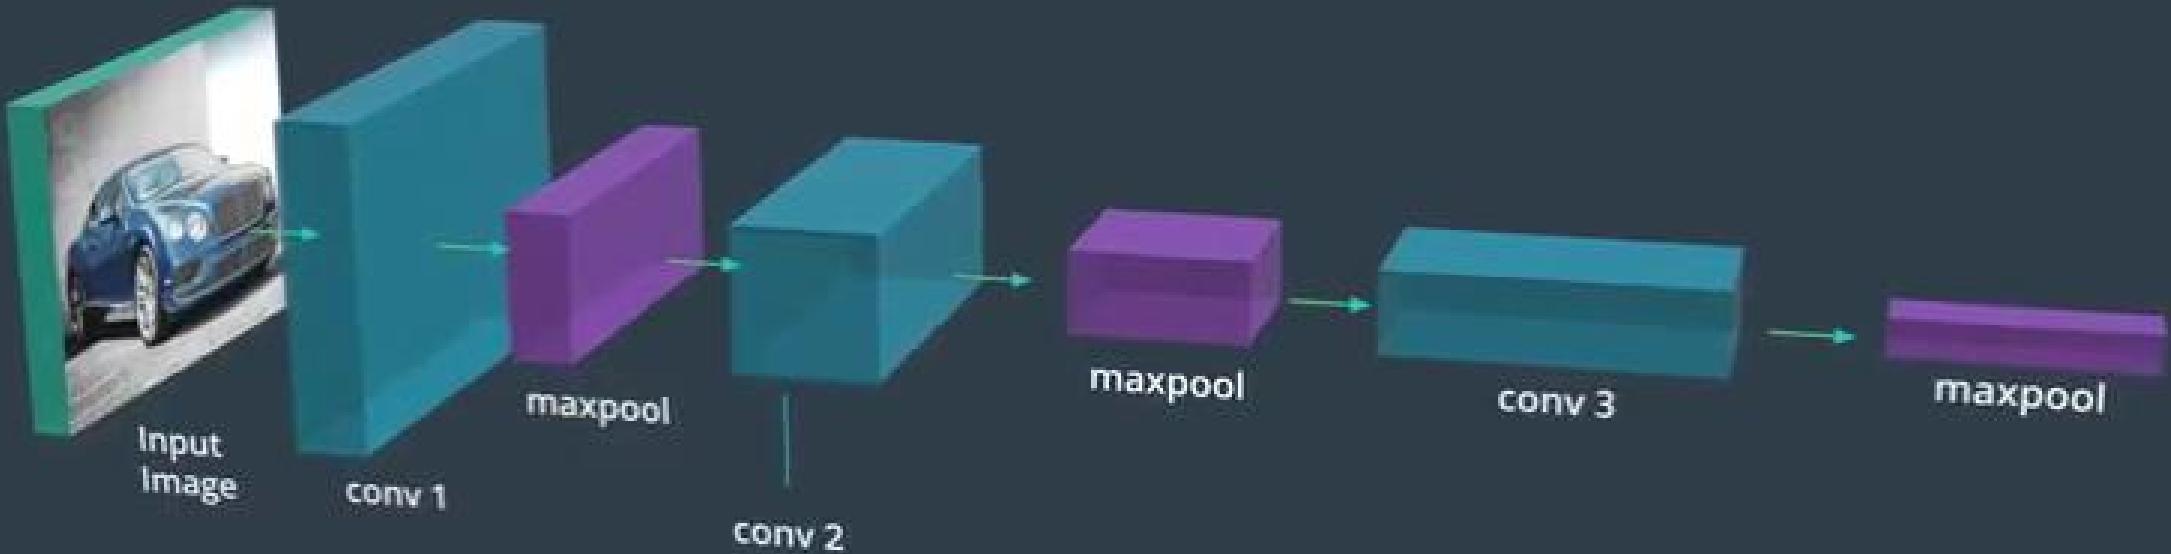
learns to distill the **content** of the image

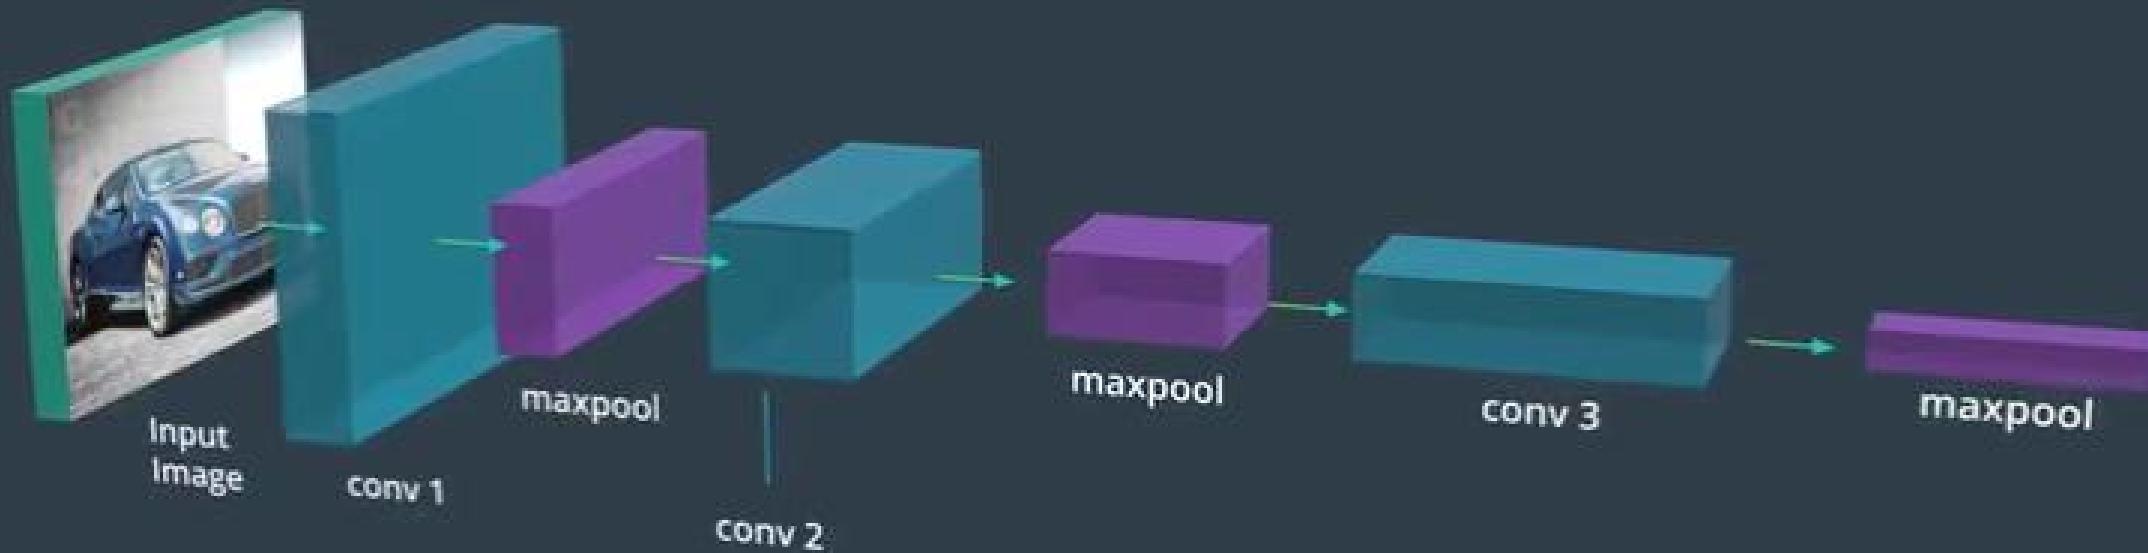
```
from keras.models import Sequential
from keras.layers import Convolution2D, MaxPooling2D, Flatten, Dense

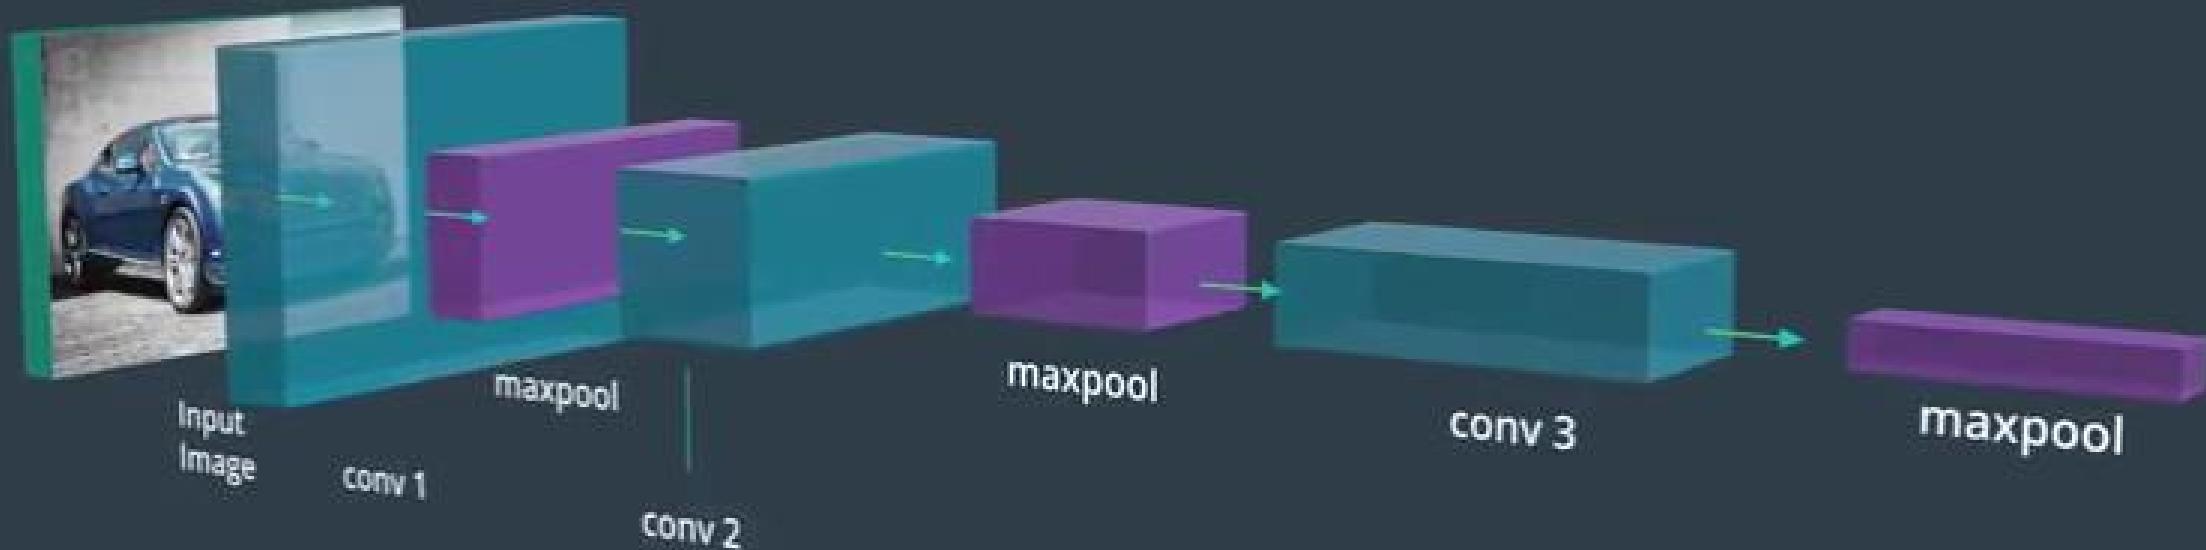
model = Sequential()
model.add(Convolution2D(filters=16, kernel_size=2, padding='same', activation='relu',
                       input_shape=(32, 32, 3)))
model.add(MaxPooling2D(pool_size=2))
model.add(Convolution2D(filters=32, kernel_size=2, padding='same', activation='relu'))
model.add(MaxPooling2D(pool_size=2))
model.add(Convolution2D(filters=64, kernel_size=2, padding='same', activation='relu'))
model.add(MaxPooling2D(pool_size=2))
model.add(Flatten())
model.add(Dense(500, activation='relu'))
model.add(Dense(10, activation='softmax'))

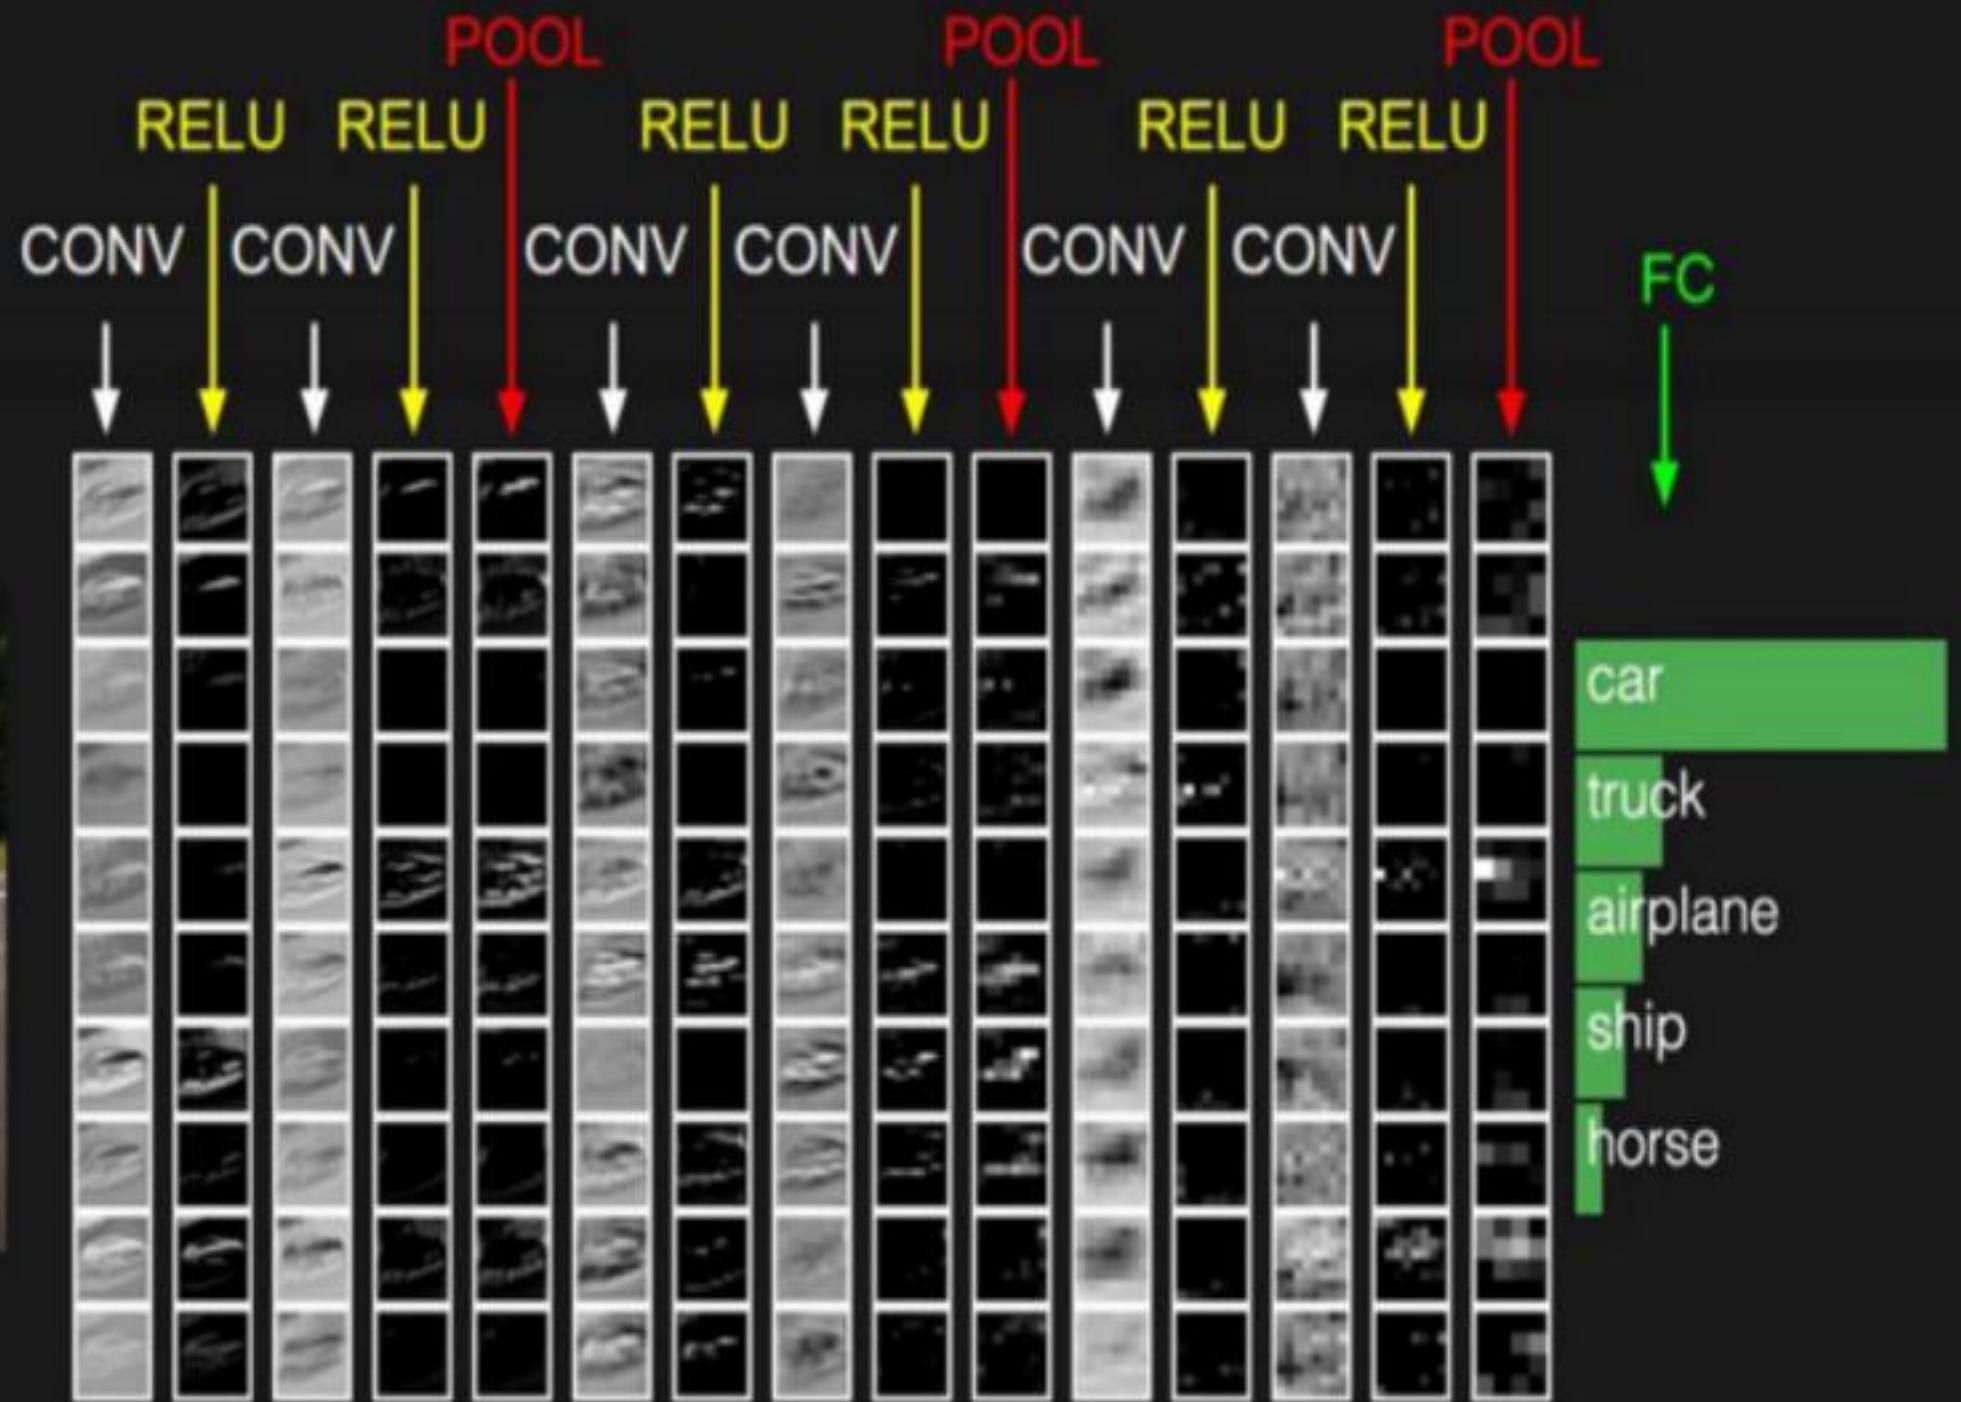
model.summary()
```

Layer (type)	Output Shape	Param #
conv2d_7 (Conv2D)	(None, 32, 32, 16)	208
max_pooling2d_4 (MaxPooling2D)	(None, 16, 16, 16)	0
conv2d_8 (Conv2D)	(None, 16, 16, 32)	2080
max_pooling2d_5 (MaxPooling2D)	(None, 8, 8, 32)	0
conv2d_9 (Conv2D)	(None, 8, 8, 64)	8256
max_pooling2d_6 (MaxPooling2D)	(None, 4, 4, 64)	0
flatten_1 (Flatten)	(None, 1024)	0
dense_1 (Dense)	(None, 500)	512500
dense_2 (Dense)	(None, 10)	5010
Total params: 528,054.0		
Trainable params: 528,054.0		
Non-trainable params: 0.0		









Draw your number here



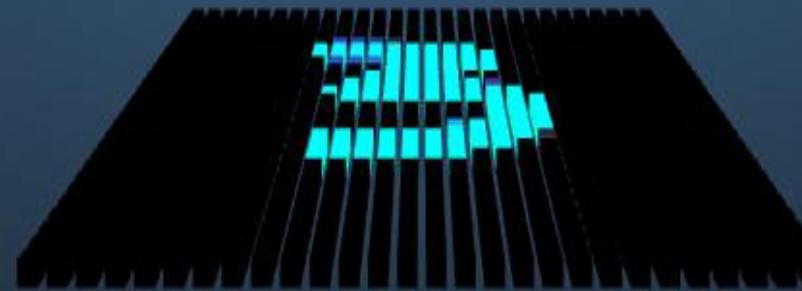
Downsampled drawing: 

First guess: 

Second guess: 

[Click Here](#)

0 1 2 3 4 5 6 7 8 9



Draw your number here



Downsampled drawing:



First guess:



Second guess:



Layer visibility

Input layer



Show

Convolution layer 1

Show

Downsampling layer 1

Show

Convolution layer 2

Show

Downsampling layer 2

Show

Fully-connected layer 1

Show

Fully-connected layer 2

Show

Output layer

Show

Click Here

0123456789

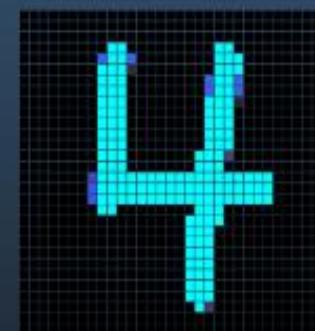
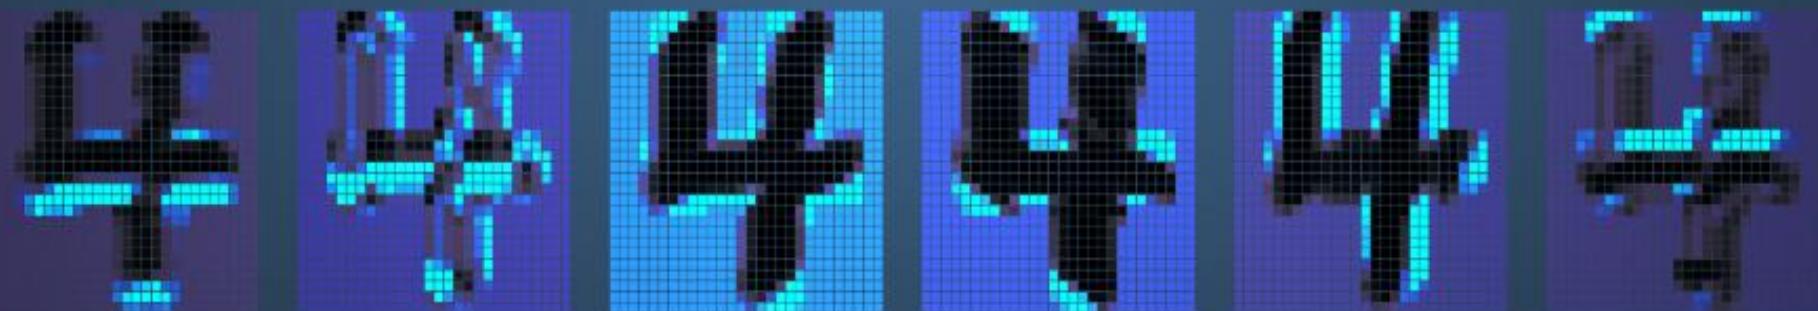
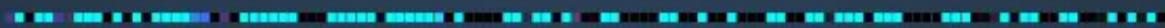


Image Preprocessing and image augmentation



Image Scaling



32x32



32x32



32x32

Invariant representation



Scale Invariance



Rotation Invariance



Translation Invariance



Data Augmentation





Rotation invariance



Translation invariance

```
from tensorflow.keras.preprocessing.image  
import ImageDataGenerator  
  
train_datagen = ImageDataGenerator(rescale=1./255)  
  
train_generator = train_datagen.flow_from_directory(  
    train_dir,  
    target_size=(300, 300),  
    batch_size=128,  
    class_mode='binary')
```

5. Create and configure augmented image generator

```
from keras.preprocessing.image import ImageDataGenerator  
  
# create and configure augmented image generator  
datagen = ImageDataGenerator(  
    width_shift_range=0.1, # randomly shift images horizontally (10% of total width)  
    height_shift_range=0.1, # randomly shift images vertically (10% of total height)  
    horizontal_flip=True) # randomly flip images horizontally  
  
# fit augmented image generator on data  
datagen.fit(x_train)
```

6. Visualize original and augmented images

```
import matplotlib.pyplot as plt

# take subset of training data
x_train_subset = x_train[:12]

# visualize subset of training data
fig = plt.figure(figsize=(20,2))
for i in range(0, len(x_train_subset)):
    ax = fig.add_subplot(1, 12, i+1)
    ax.imshow(x_train_subset[i])
fig.suptitle('Subset of Original Training Images', fontsize=20)

# visualize augmented images
fig = plt.figure(figsize=(20,2))
for x_batch in datagen.flow(x_train_subset, batch_size=12):
    for i in range(0, 12):
        ax = fig.add_subplot(1, 12, i+1)
        ax.imshow(x_batch[i])
    fig.suptitle('Augmented Images', fontsize=20)
    plt.show()
    break;
```



7. Define the model architecture

```
from keras.models import Sequential
from keras.layers import Convolution2D, MaxPooling2D, Flatten, Dense, Dropout

model = Sequential()
model.add(Convolution2D(filters=16, kernel_size=2, padding='same', activation='relu',
                       input_shape=(32, 32, 3)))
model.add(MaxPooling2D(pool_size=2))
model.add(Convolution2D(filters=32, kernel_size=2, padding='same', activation='relu'))
model.add(MaxPooling2D(pool_size=2))
model.add(Convolution2D(filters=64, kernel_size=2, padding='same', activation='relu'))
model.add(MaxPooling2D(pool_size=2))
model.add(Dropout(0.3))
model.add(Flatten())
model.add(Dense(500, activation='relu'))
model.add(Dropout(0.4))
model.add(Dense(10, activation='softmax'))

model.summary()
```

8. Compile the model

```
# compile the model  
model.compile(loss='categorical_crossentropy', optimizer='rmsprop',  
               metrics=['accuracy'])
```

9. Train the model

10. Load the model with the best validation accuracy

```
# load the weights that yielded the best validation accuracy  
model.load_weights('aug_model.weights.best.hdf5')
```

11. Calculate classification accuracy on test set

```
# evaluate and print test accuracy
score = model.evaluate(x_test, y_test, verbose=0)
print('\n', 'Test accuracy:', score[1])
```

Test accuracy 0.6825

Train your Own Data

```
from keras.preprocessing.image import ImageDataGenerator  
  
train_datagen = ImageDataGenerator(rescale = 1./255,  
                                     shear_range = 0.2,  
                                     zoom_range = 0.2,  
                                     horizontal_flip = True)  
  
test_datagen = ImageDataGenerator(rescale = 1./255)  
  
training_set = train_datagen.flow_from_directory('dataset/training_set',  
                                                target_size = (64, 64),  
                                                batch_size = 100,  
                                                class_mode = 'binary')  
  
test_set = test_datagen.flow_from_directory('dataset/test_set',  
                                            target_size = (64, 64),  
                                            batch_size = 100,  
                                            class_mode = 'binary')  
  
classifier.fit_generator(training_set,  
                         samples_per_epoch = 8000,  
                         nb_epoch = 25,]  
                         validation_data = test_set,  
                         nb_val_samples = 2000)]
```

Train your Own Data

`samples_per_epoch` is usually setted as:

```
samples_per_epoch=train_generator.nb_samples
```

This way you are ensuring that every epoch you are seeing a number of samples equal to the size of your training set. This means that you are seeing all of your training samples at every epoch.

nb_epoch is pretty much up to you. It determines how many times you iterate over a number defined by `samples_per_epoch`.

To give you an example, in your code right now your model is 'seeing' `(nb_epoch * samples_per_epoch)` images, which in this case are 65000 images.

nb_val_samples determines over how many validation samples your model is evaluated after finishing every epoch. It is up to you as well. The usual thing is to set:

```
nb_val_samples=validation_generator.nb_samples
```

In order to evaluate your model on the full validation set.

batch_size determines how many images are feeded **at the same time** to your gpu (or cpu). Rule of thumb is to set the largest `batch_size` that your gpu's memory allows. Ideal batch_size is active area of research nowadays, but usually a bigger batch_size will work better.

```
classifier.fit_generator(training_set,  
                         samples_per_epoch = 8000,  
                         nb_epoch = 25,]  
                         validation_data = test_set,  
                         nb_val_samples = 2000)]
```

MNIST using CNN

```
from __future__ import print_function
import keras
from keras.datasets import mnist
from keras.models import Sequential
from keras.layers import Dense, Dropout, Flatten
from keras.layers import Conv2D, MaxPooling2D
from keras import backend as K

batch_size = 128
num_classes = 10
epochs = 12

# input image dimensions
img_rows, img_cols = 28, 28

# the data, split between train and test sets
(x_train, y_train), (x_test, y_test) = mnist.load_data()

x_train = x_train.astype('float32')
x_test = x_test.astype('float32')
x_train /= 255
x_test /= 255

# convert class vectors to binary class matrices
y_train = keras.utils.to_categorical(y_train, num_classes)
y_test = keras.utils.to_categorical(y_test, num_classes)
```

MNIST using CNN

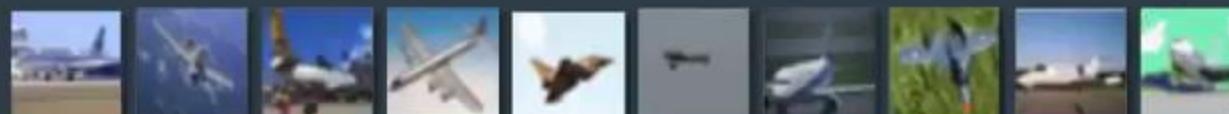
```
model = Sequential()
model.add(Conv2D(32, kernel_size=(3, 3),
                activation='relu',
                input_shape=input_shape))
model.add(Conv2D(64, (3, 3), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))
model.add(Flatten())
model.add(Dense(128, activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(num_classes, activation='softmax'))

model.compile(loss=keras.losses.categorical_crossentropy,
              optimizer=keras.optimizers.Adadelta(),
              metrics=['accuracy'])

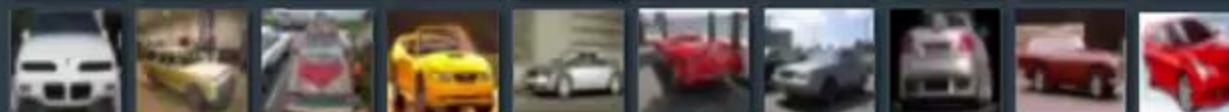
model.fit(x_train, y_train,
          batch_size=batch_size,
          epochs=epochs,
          verbose=1,
          validation_data=(x_test, y_test))
score = model.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])
```

CIFAR10-CNN

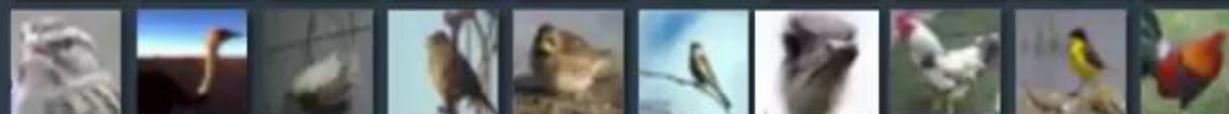
Airplane



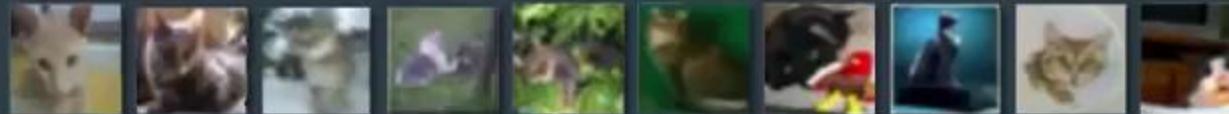
Automobile



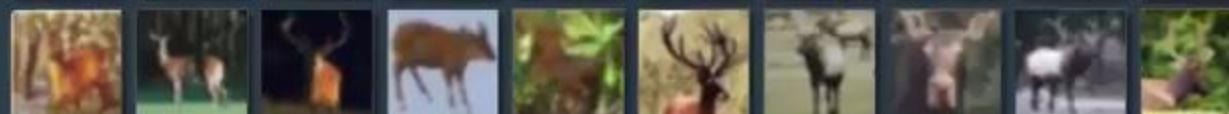
Bird



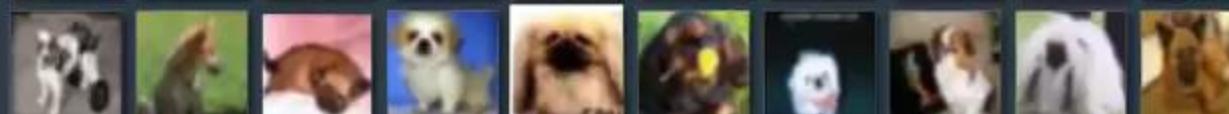
Cat



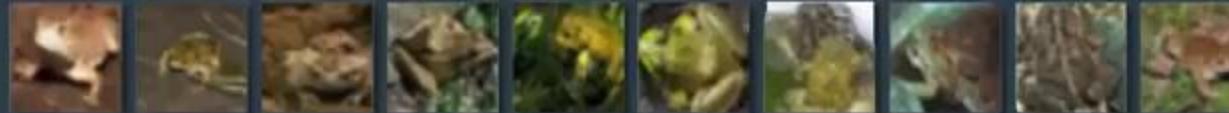
Deer



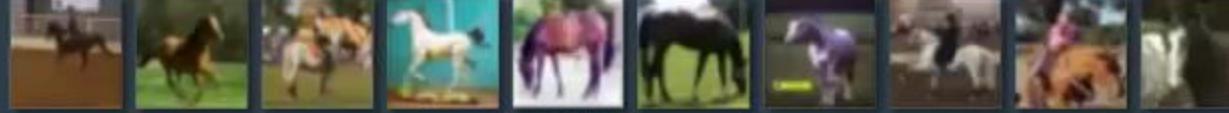
Dog



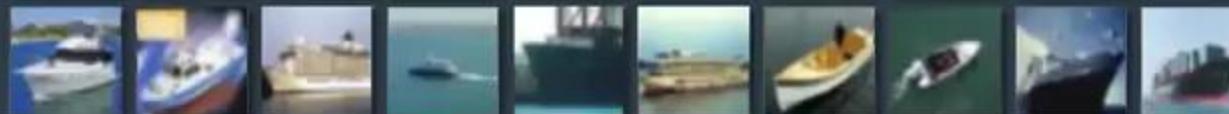
Frog



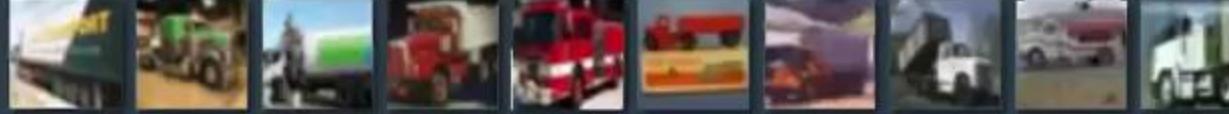
Horse



Ship



Truck



1. Load CIFAR -10 Database

```
import keras
from keras.datasets import cifar10

# load the pre-shuffled train and test data
(x_train, y_train), (x_test, y_test) = cifar10.load_data ()
```

Using TensorFlow backend.

2. Visualize the first 24 training images

```
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline

fig = plt.figure(figsize=(20,5))
for i in range (36):
    ax = fig.add_subplot(3, 12, i + 1, xticks=[], yticks[])
    ax.imshow(np.squeeze(x_train[i]))
```



3. Rescale the images by dividing every pixel in every image by 255

```
# rescale [0,255] --> [0,1]
x_train = x_train.astype('float32')/255
x_tests = x_test.astype('float32')/255
```

4. Break dataset into Training, Testing, and Validation sets

```
from keras.utils import np_utils

# one-hot encode the labels
num_classes = len(np.unique(y_train))
y_train = keras.utils.to_categorical(y_train, num_classes)
y_test = keras.utils.to_categorical(y_test, num_classes)

# break training set into training and validation sets
(x_train, x_valid) = x_train[5000:], x_train[:5000]
(y_train, y_valid) = y_train[5000:], y_train[:5000]

# print shape of training sets
print('x_train shape:', x_train.shape)

# print number of training, validation, and test images
print(x_train.shape[0], 'train samples')
print(x_test.shape[0], 'test samples')
print(x_valid.shape[0], 'validation samples')
```

```
x_train shape: (45000, 32, 32, 3)
45000 train samples
10000 test samples
5000 validation samples
```

5. Define the model architecture

```
from keras.models import Sequential
from keras.layers import Dense, Dropout, Flatten

# define the model
model = Sequential()
model.add(Flatten(input_shape = x_train.shape[1:]))
model.add(Dense(1000, activation='relu'))
model.add(Dropout(0.2))
model.add(Dense(512, activation='relu'))
model.add(Dropout(0.2))
model.add(Dense(num_classes, activation='softmax'))
model.summary()
```

Layer (type)	Output Shape	Param #
flatten_1 (Flatten)	(None, 3072)	0
dense_1 (Dense)	(None, 1000)	3073000
dropout_1 (Dropout)	(None, 1000)	0
dense_2 (Dense)	(None, 512)	512512
dropout_2 (Dropout)	(None, 512)	0
dense_3 (Dense)	(None, 10)	5130

Total params: 3,590,642.0
Trainable params: 3,590,642.0
Non-trainable params: 0.0

6. Compile the model

```
# compile the model  
model.compile(loss='categorical_crossentropy', optimizer= 'rmsprop',  
metrics=[ 'accuracy'])
```

7. Train the model

```
from keras.callbacks import ModelCheckpoint

# train the model
checkpointer = ModelCheckpoint(filepath='MLP.weights.best.hdf5', verbose=1,
                                save_best_only=True)
hist = model.fit(x_train, y_train, batch_size=32, epochs=20,
                  validation_data=(x_valid, y_valid), callbacks=[checkpointer],
                  verbose=2, shuffle=True)
```

8. Load the model with the best classification accuracy on the validation set

```
# load the weights that yielded the best validation accuracy  
model.load_weights('MLP.weights.best.hdf5')
```

9. Calculate classification accuracy on test set

```
# evaluate and print test accuracy
score = model.evaluate(x_test, y_test, verbose=0)
print('\n', 'Test accuracy:' , score[1])
```

Test accuracy 0.4

5. Define the model architecture

```
from keras.models import Sequential
from keras.layers import Convolution2D, MaxPooling2D, Flatten, Dense, Dropout

model = Sequential()
model.add(Convolution2D(filters=16, kernel_size=2, padding='same', activation='relu',
                       input_shape=(32, 32, 3)))
model.add(MaxPooling2D(pool_size=2))
model.add(Convolution2D(filters=32, kernel_size=2, padding='same', activation='relu'))
model.add(MaxPooling2D(pool_size=2))
model.add(Convolution2D(filters=64, kernel_size=2, padding='same', activation='relu'))
model.add(MaxPooling2D(pool_size=2))
model.add(Dropout(0.3))
model.add(Flatten())
model.add(Dense(500, activation='relu'))
model.add(Dropout(0.4))
model.add(Dense(10, activation='softmax'))

model.summary()
```

Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 32, 32, 16)	208
max_pooling2d_1 (MaxPooling2D)	(None, 16, 16, 16)	0
conv2d_2 (Conv2D)	(None, 16, 16, 32)	2080
max_pooling2d_2 (MaxPooling2D)	(None, 8, 8, 32)	0
conv2d_3 (Conv2D)	(None, 8, 8, 64)	8256
max_pooling2d_3 (MaxPooling2D)	(None, 4, 4, 64)	0
dropout_1 (Dropout)	(None, 4, 4, 64)	0
flatten_1 (Flatten)	(None, 1024)	0
dense_1 (Dense)	(None, 500)	512500
dropout_2 (Dropout)	(None, 500)	0
dense_2 (Dense)	(None, 10)	5010
Total params: 528,054.0		
Trainable params: 528,054.0		
Non-trainable params: 0.0		

6. Compile the model

```
# compile the model
model.compile(loss='categorical_crossentropy', optimizer= 'rmsprop',
               metrics=[ 'accuracy'])
```

7. Train the model

```
from keras.callbacks import ModelCheckpoint

# train the model
checkpointer = ModelCheckpoint(filepath='MLP.weights.best.hdf5', verbose=1,
                                save_best_only=True)
hist = model.fit(x_train, y_train, batch_size=32, epochs=100,
                  validation_data=(x_valid, y_valid), callbacks=[checkpointer],
                  verbose=2, shuffle=True)
```

8. Load the model with the best classification accuracy on the validation set

```
# load the weights that yielded the best validation accuracy
model.load_weights('model.weights.best.hdf5')
```

9. Calculate classification accuracy on test set

```
# evaluate and print test accuracy
score = model.evaluate(x_test, y_test, verbose=0)
print('\n', 'Test accuracy:' , score[1])
```

Test accuracy 0.6598